

Distributed Systems: Coordination & Agreement



Agenda

- Synchronous vs. asynchronous systems
- Distributed mutual exclusion
 - Mutual exclusion with shared variables (recap)
 - Algorithms for distributed mutual exclusion
- Leader election
 - Problem definition
 - Algorithms

DISTRIBUTED SYSTEM MODELS

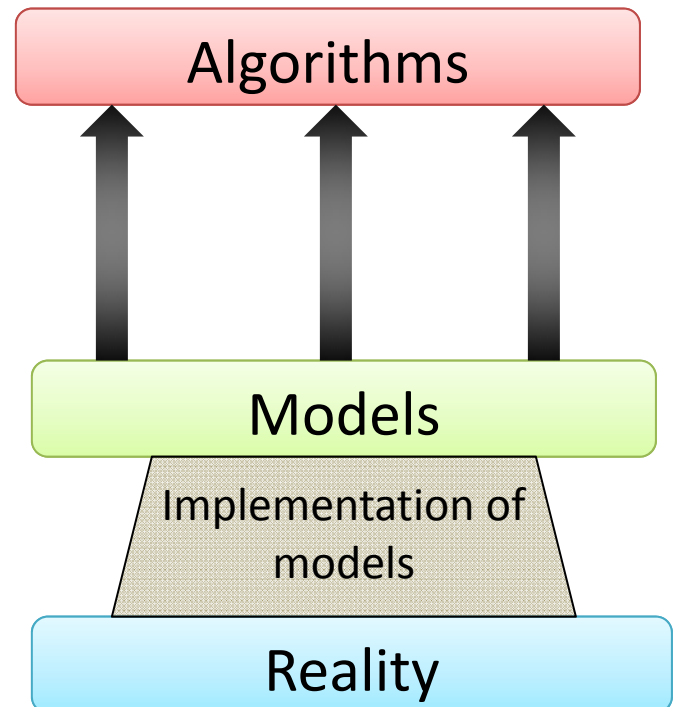
Distributed Systems (H.-A. Jacobsen)

Distributed system model

Of theoretical relevance for designing algorithms

- Model captures all the **assumptions** made about the system
- Including network, clocks, processor, etc.
- Algorithms always **assume** a certain model
- Some algorithms are only possible in **stronger models** (more restrictions)
- Model is **theoretical**: whether its assumptions hold in practice is a different question

"All models are wrong, but some are useful"



Synchronous vs. asynchronous model

Property	Synchronous system model	Asynchronous system model
Clocks	Bound on drift	No bound on drift
Processes	Bound on execution time	No bound on execution time
Network	Bound on latency	No bound on latency

- There are other models.

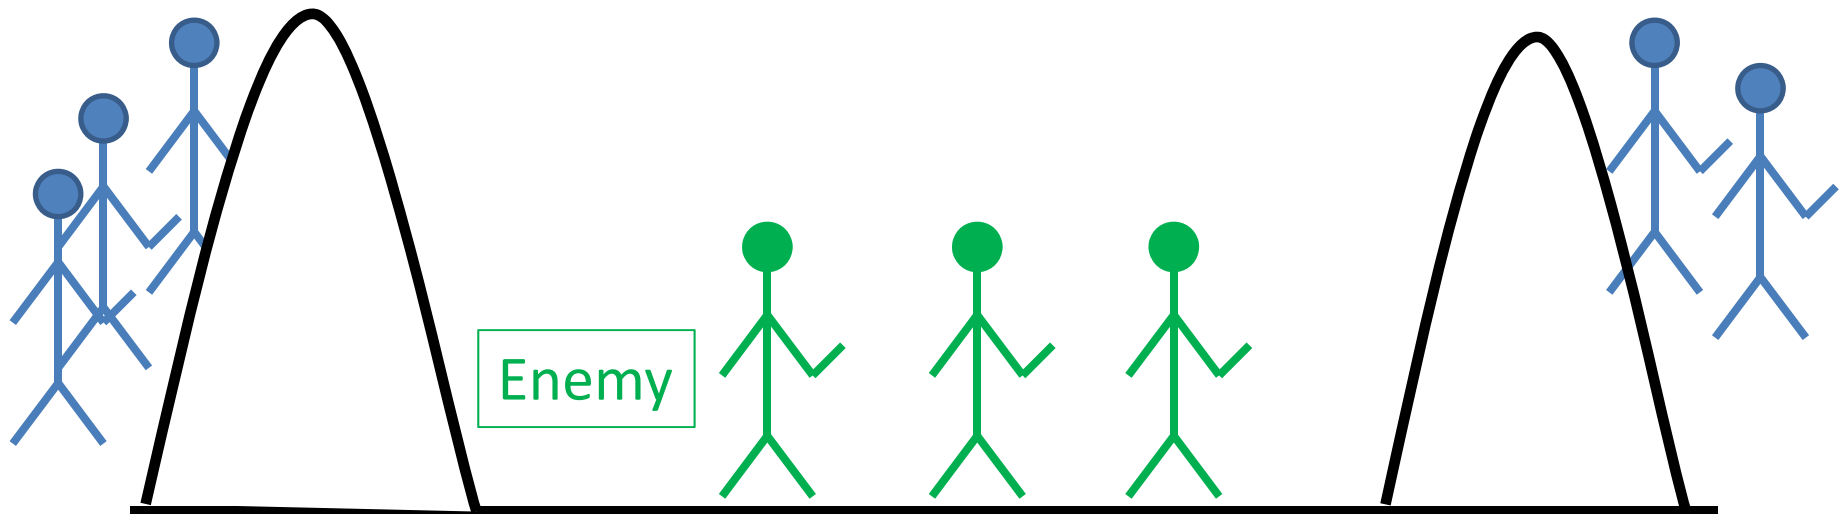
Two General's Problem (Agreement)

Armies need to agree on:

**A thought
experiment**

Army 1

Army 2



**Armies are safe if they
don't attack (or win)
(Safety)**

**Armies need to
coordinate attack to win.
(Liveness)**

Two General's Problem (Agreement)

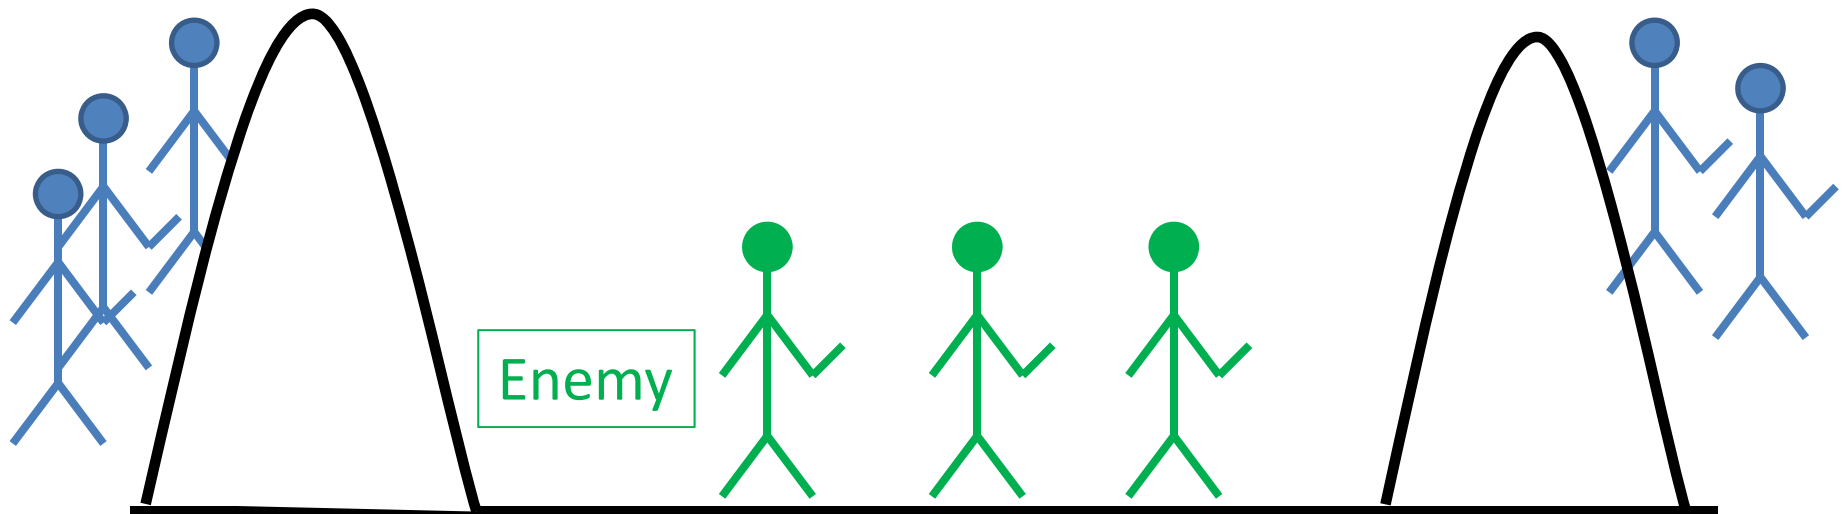
Armies need to agree on:

**A thought
experiment**

*Who leads the attack?
When to attack?*

Army 1

Army 2



**Armies are safe if they
don't attack (or win)
(Safety)**

**Armies need to
coordinate attack to win.
(Liveness)**

Two General's Problem (Agreement)

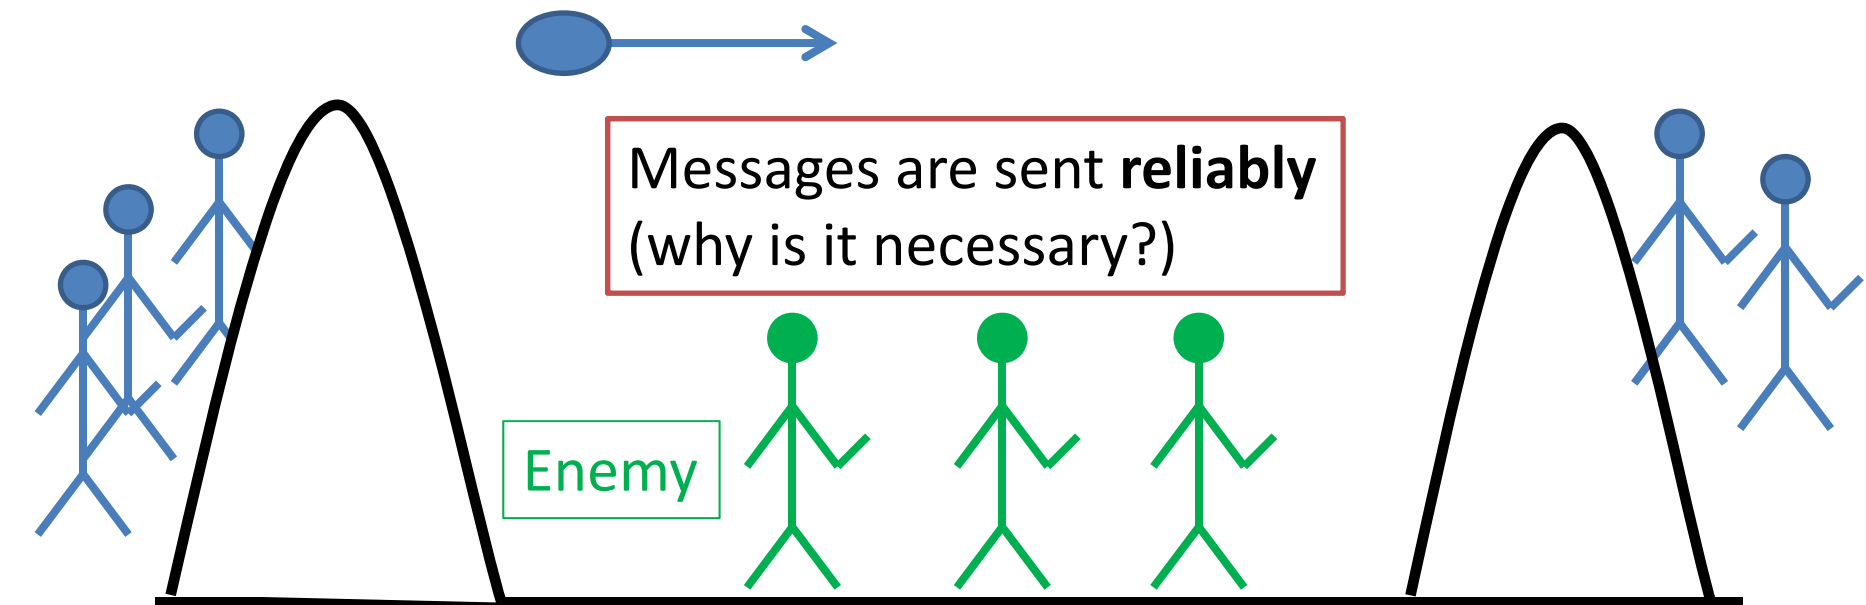
Armies need to agree on:

A thought
experiment

Who leads the attack?
When to attack?

Army 1

Army 2



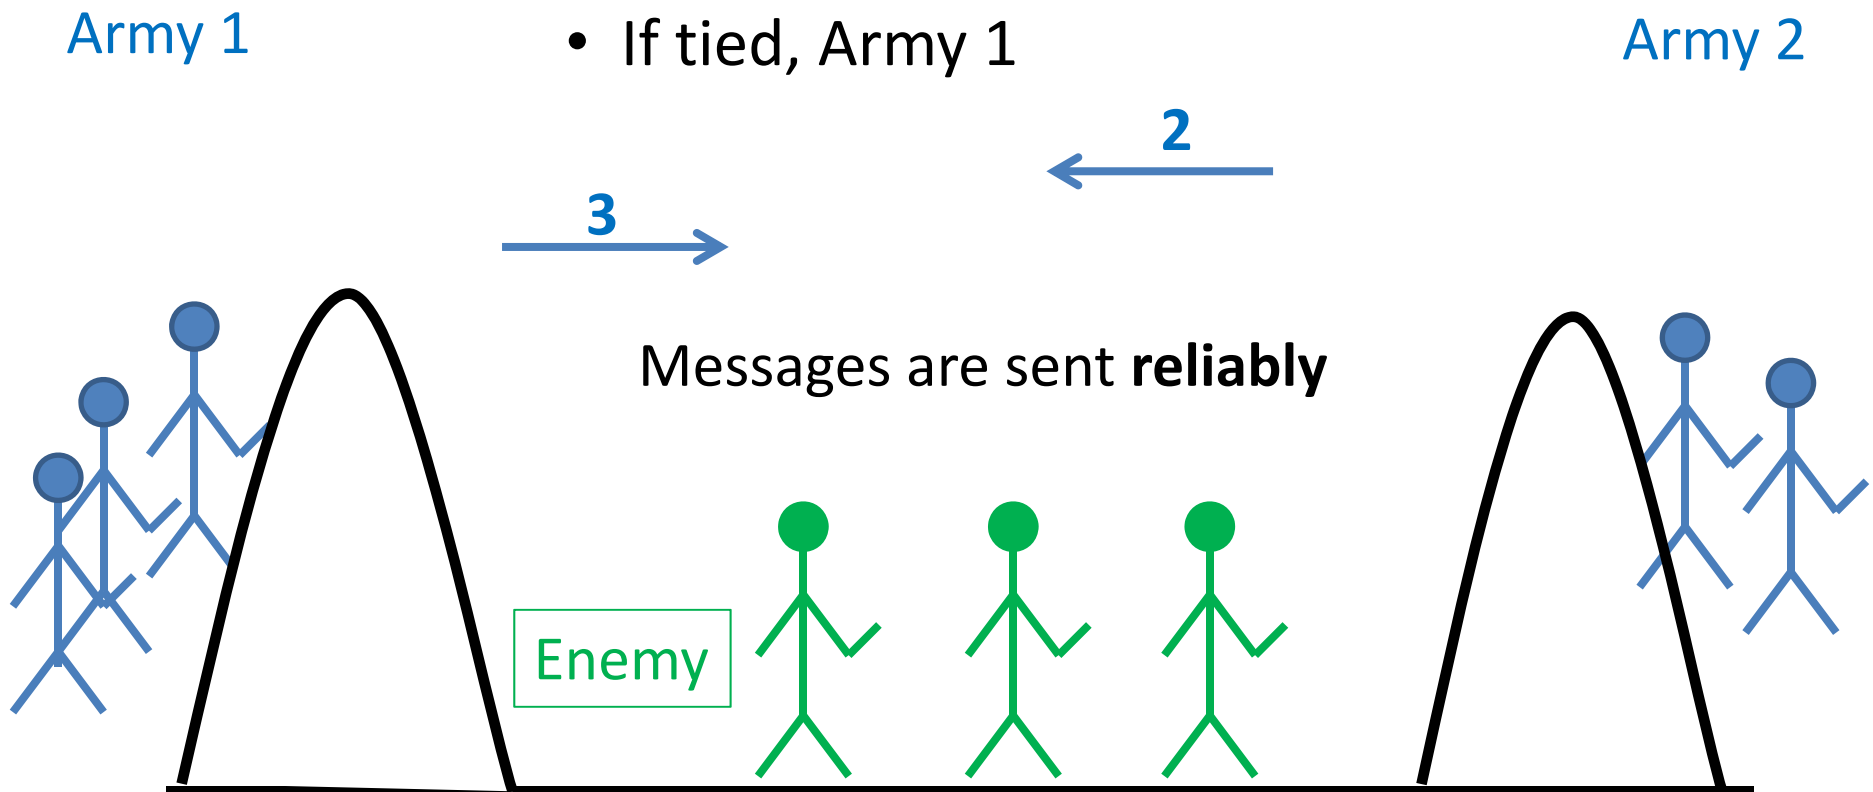
Armies are safe if they
don't attack (or win)
(Safety)

Armies need to
coordinate attack to win.
(Liveness)

Synchronous vs. asynchronous agreement

Who leads the attack?

- Largest army
- If tied, Army 1

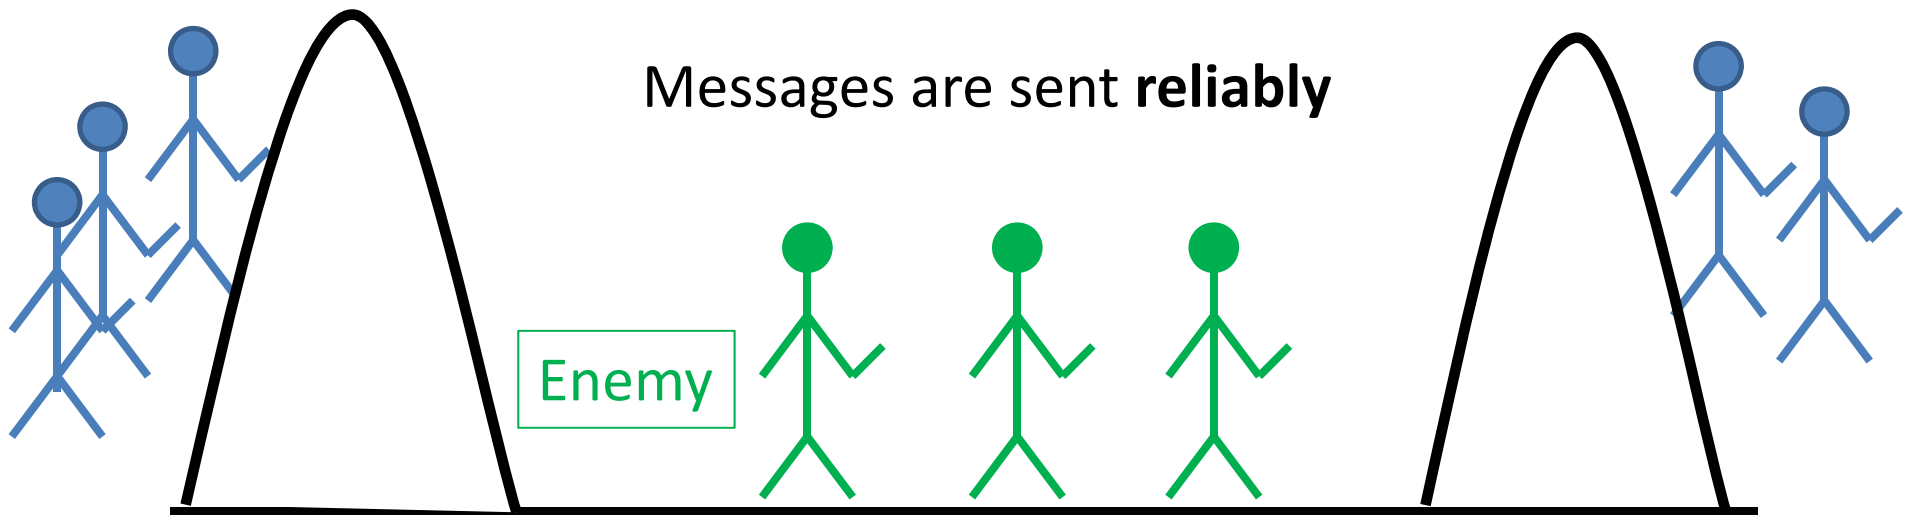


Asynchronous agreement

When to attack? No bound on delivery!

Army 1

Army 2

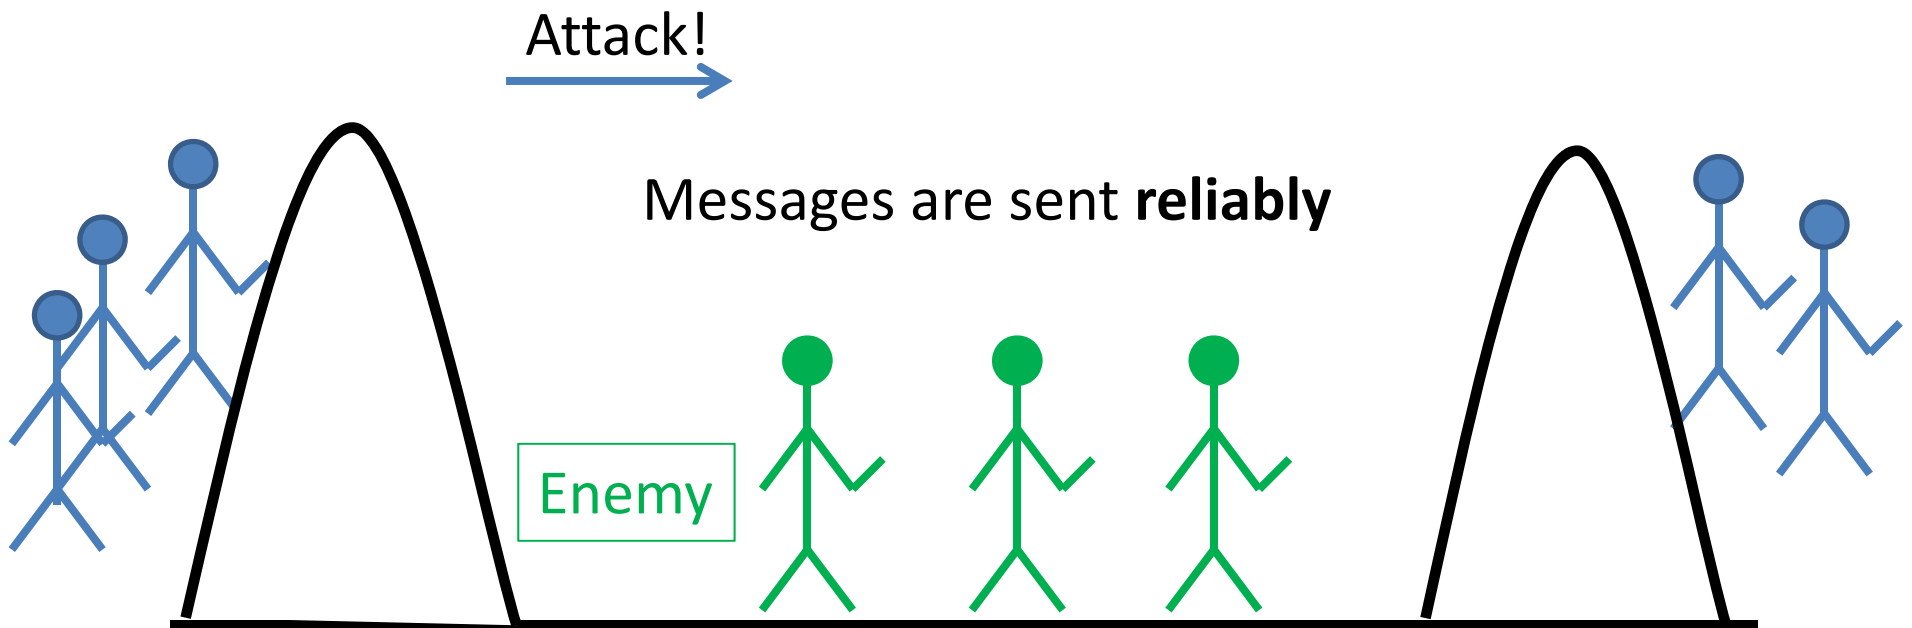


Asynchronous agreement

When to attack? No bound on delivery!

Army 1

Army 2

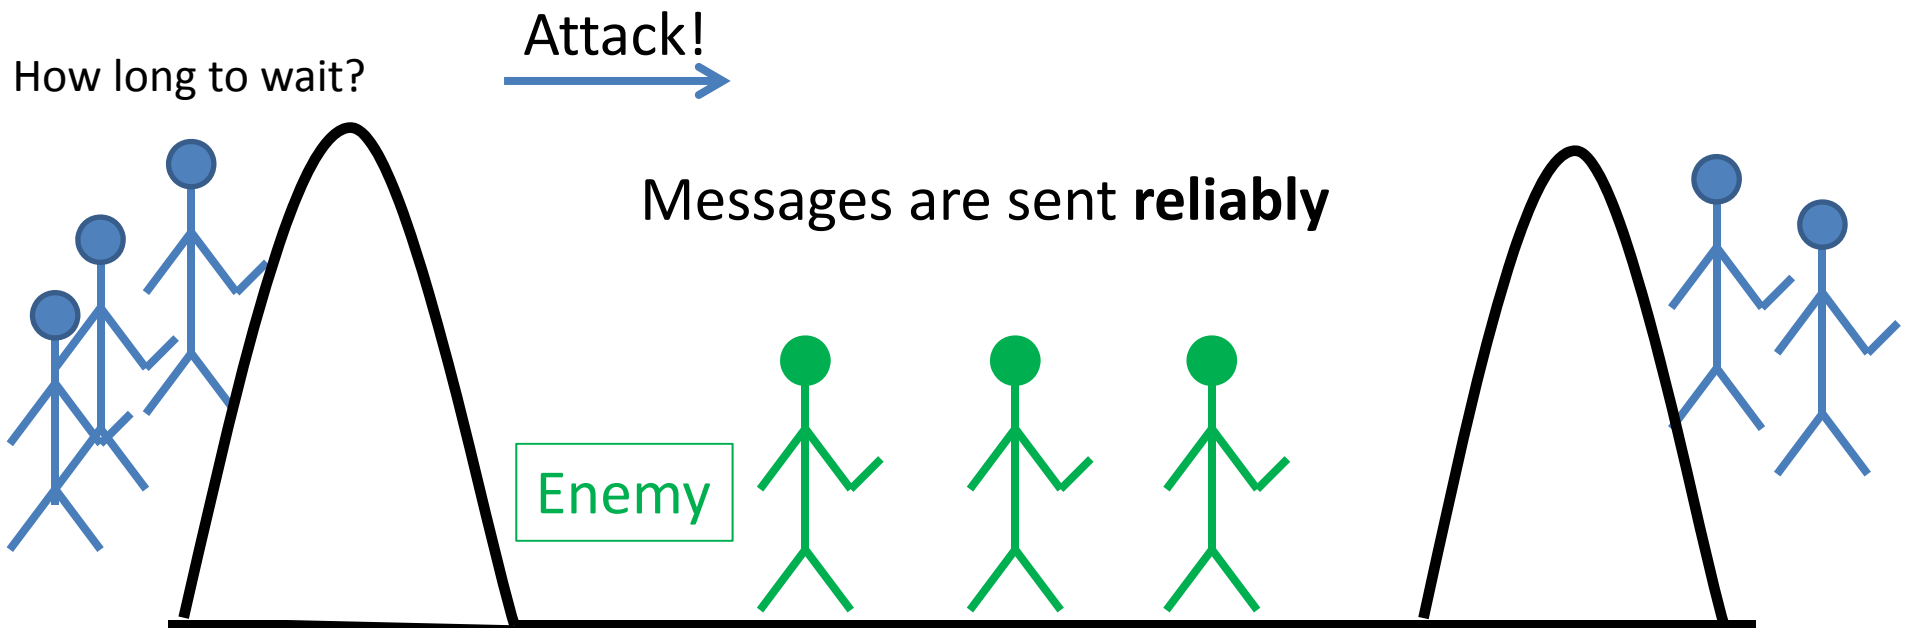


Asynchronous agreement

When to attack? No bound on delivery!

Army 1

Army 2



Synchronous agreement

When to attack?

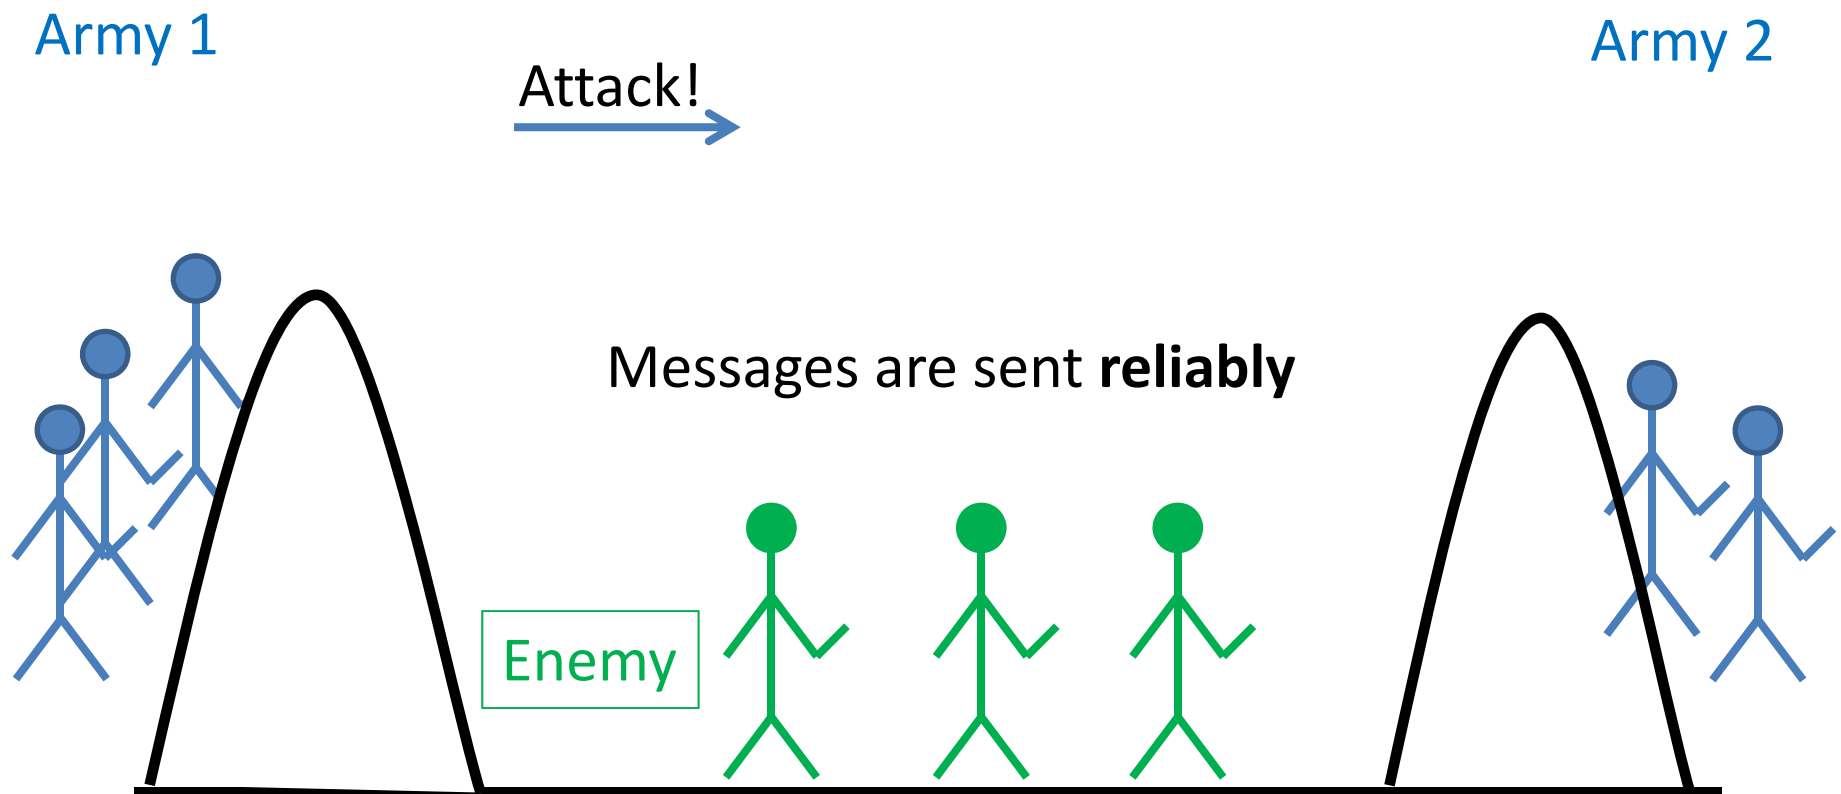
Army 1

Army 2



Synchronous agreement

When to attack?



Synchronous agreement

When to attack?

Message takes at least ***min*** time
and at most ***max*** time to arrive

Army 1

Army 2

Attack!
→



Synchronous agreement

When to attack?

Message takes at least **min** time
and at most **max** time to arrive

Army 1

Army 2

Attack!
→

Waits for min time, then attacks

Messages are sent **reliably**



Synchronous agreement

When to attack?

Message takes at least **min** time
and at most **max** time to arrive

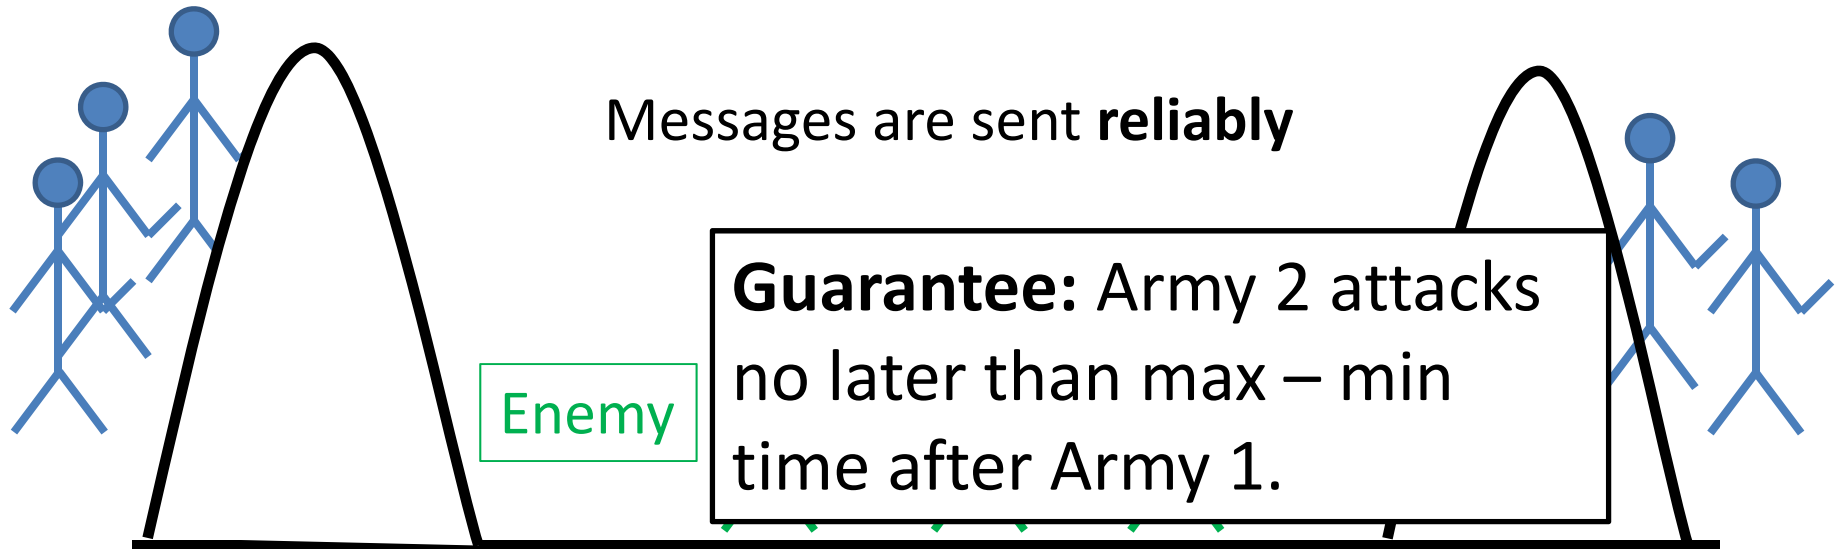
Army 1

Army 2

Attack!
→

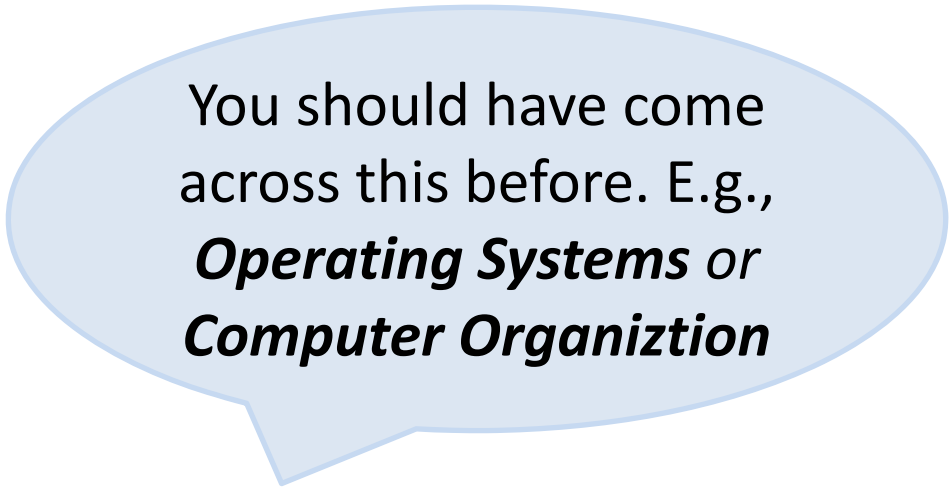
Waits for min time, then attacks

Messages are sent **reliably**



Some takeaways

- Internet and many practical distributed applications are closer to asynchronous than synchronous model
- A solution valid for asynchronous distributed systems is also valid for synchronous ones (synchronous model is a stronger model)
- Many design problems cannot be solved in an asynchronous world (e.g., *when* vs. *who* leads attack)
- Apply **timeouts** and **timing assumptions** to reduce uncertainty and to bring elements of the synchronous model into the picture

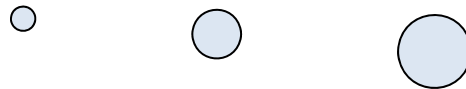


You should have come
across this before. E.g.,
***Operating Systems or
Computer Organization***

RECAP: MUTUAL EXCLUSION

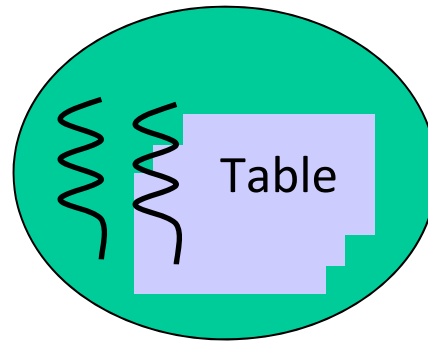
Problem: Access to shared variables

- Imaging a **globally shared variable** counter in a process accessible to multiple threads
 - For example, the **key-value records** managed by a storage server (or more complex data structure)



**Let's dissect
the issue in
detail**

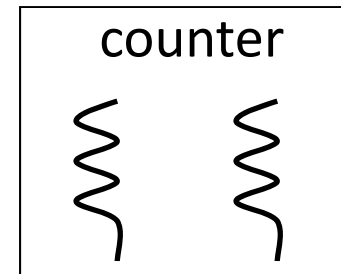
Shared data & synchronization



What may happen if multiple threads concurrently access shared process state (e.g., variables in memory)?

Concurrently manipulating shared data

- Two threads execute concurrently as part of the same process
- Shared variable (e.g., global variable)
 - `counter = 5`
- Thread 1 executes
 - `counter++`
- Thread 2 executes
 - `counter--`
- *What are **all the possible values** of `counter` after Thread 1 and Thread 2 finish executing?*



Machine-level implementation

- Implementation of “counter++”

register₁ = counter

register₁ = register₁ + 1

counter = register₁

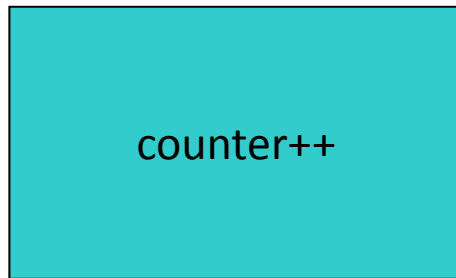
- Implementation of “counter--”

register₂ = counter

register₂ = register₂ - 1

counter = register₂

Possible execution sequences



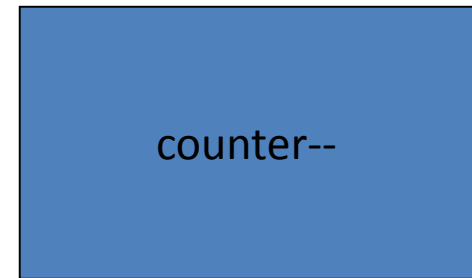
Context Switch



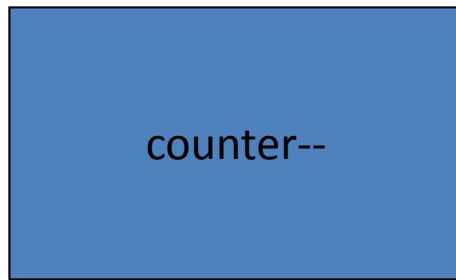
Context Switch



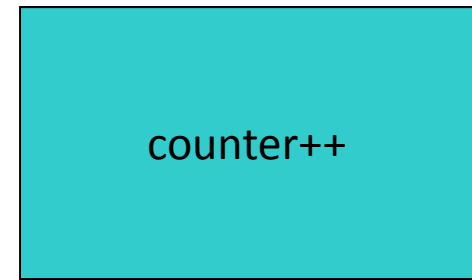
Context Switch



Context Switch



Context Switch



Interleaved execution

- Assume **counter** is 5 and interleaved execution of $\text{counter}++$ (in T_1) and $\text{counter}--$ (in T_2)

$T_1: r_1 = \text{counter}$ ($\text{register}_1 = 5$)

$T_1: r_1 = r_1 + 1$ ($\text{register}_1 = 6$)

$T_2: r_2 = \text{counter}$ ($\text{register}_2 = 5$)

$T_2: r_2 = r_2 - 1$ ($\text{register}_2 = 4$)

$T_1: \text{counter} = r_1$ ($\text{counter} = 6$)

$T_2: \text{counter} = r_2$ ($\text{counter} = 4$)

Context
switch

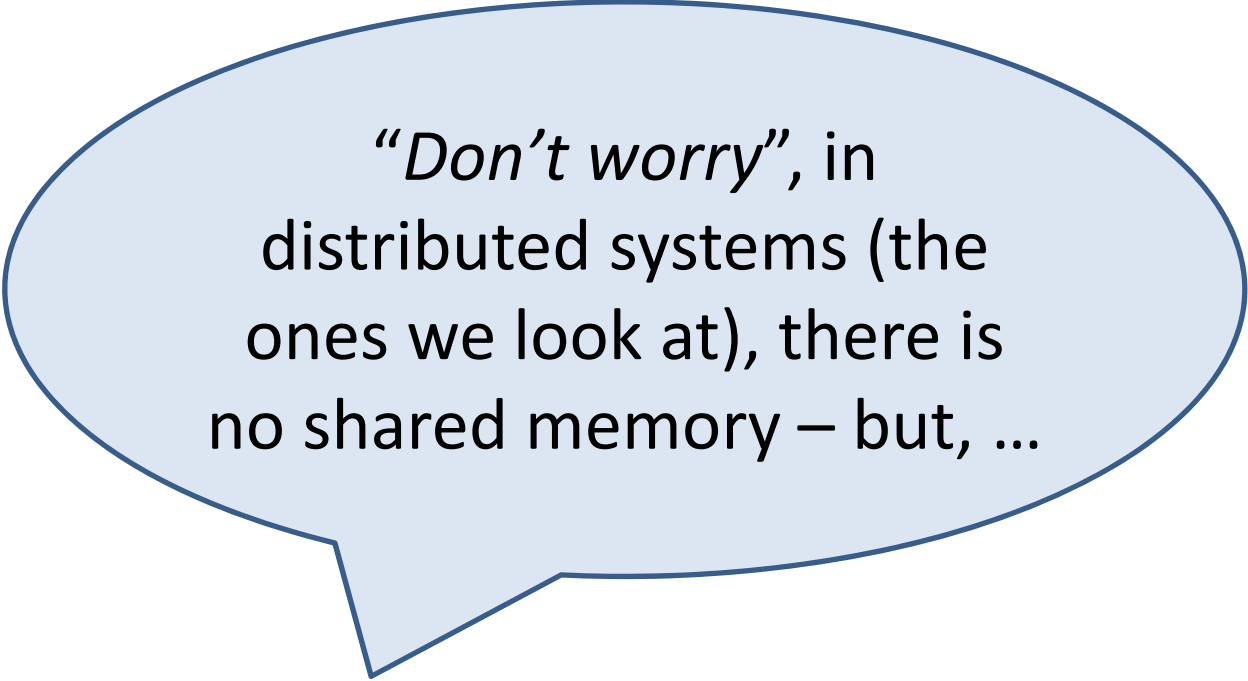
- The value of **counter** may be 4 or 6, whereas the correct result should be 5!

Race condition

- **Race condition**
 - **Several threads manipulate shared data concurrently.** The **final value** of the data **depends upon which thread finishes last.**
- In our example (interleaved execution) of counter++ with counter--
- To prevent race conditions, concurrent processes must be **synchronized!**

The moral of this story

- The statements
counter++;
counter--;
must each be executed *atomically*.
- Atomic operation means an operation that **completes in its entirety without interruption**.
- This is achieved through **synchronization primitives**
- Shared variable accessed in **critical section**, protected by synchronization primitives
- Known as the **critical section problem** or as **mutual exclusion**



“Don’t worry”, in
distributed systems (the
ones we look at), there is
no shared memory – but, ...

DISTRIBUTED MUTUAL EXCLUSION

Distributed mutual exclusion

- In distributed systems, mutual exclusion is more complex due to lack of:
 - Shared memory
 - Timing issues
 - Lack of a global clock
 - Event ordering
- Applications
 - Accessing a shared resource in distributed systems
 - One active Bigtable master to coordinate

Critical section (CS)

problem: No shared memory

- System with n processes
- Processes access shared resources in CS
- **Coordinate access to CS via message passing**
- Application-level protocol for accessing CS
 - Enter_CS() – enter CS, block if necessary
 - ResourceAccess() – access shared resource in CS
 - Exit_CS() – leave CS

Current assumptions

Clearly, not practical, more of theoretical nature

- System is asynchronous
 - No bound on delays, no bound on clock drift, etc.
- Processes do not fail
- Message delivery is reliable
 - Any message sent, is eventually delivered intact and exactly once
- Client processes are well-behaved and spent finite time accessing resources in CS

Mutual exclusion requirements

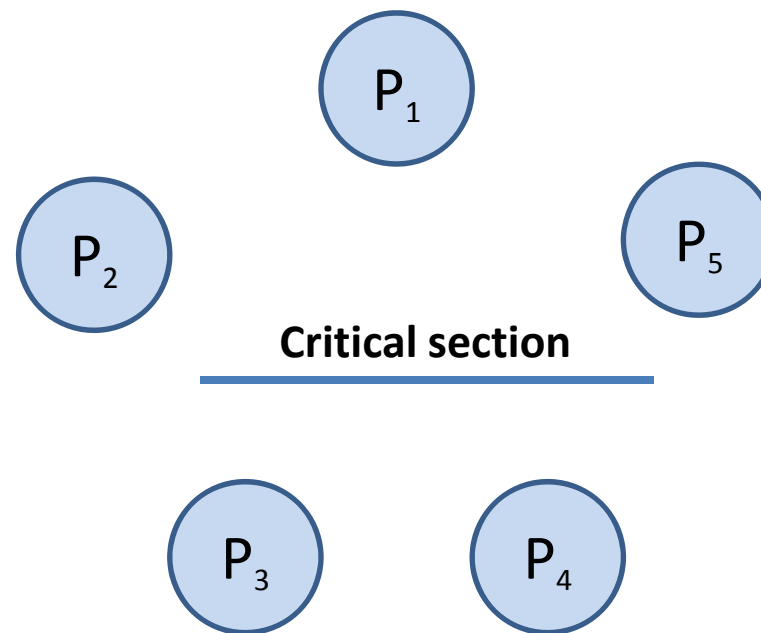
- **Safety** – correctness
 - At most one process in the critical section at a time
- **Liveness** – progress (something good happens)
 - Requests to enter/exit CS eventually succeed
 - No deadlock
- **Fairness** (order & starvation)
 - If one request to enter CS **happened-before** another one, then entry to CS is granted in that **order**
 - Requests are ordered such that no process enters the critical section twice while another waits to enter (i.e., **no starvation**)

Possible performance metrics

- **Bandwidth:** Number of messages sent, received or both
- **Synchronization delay:** Time between one process **exiting** critical section and next one **entering**
- **Client delay:** Delay **at entry** and **exit** (response time)
- We do not measure client access to resources protected by the critical section (**assume finite**)

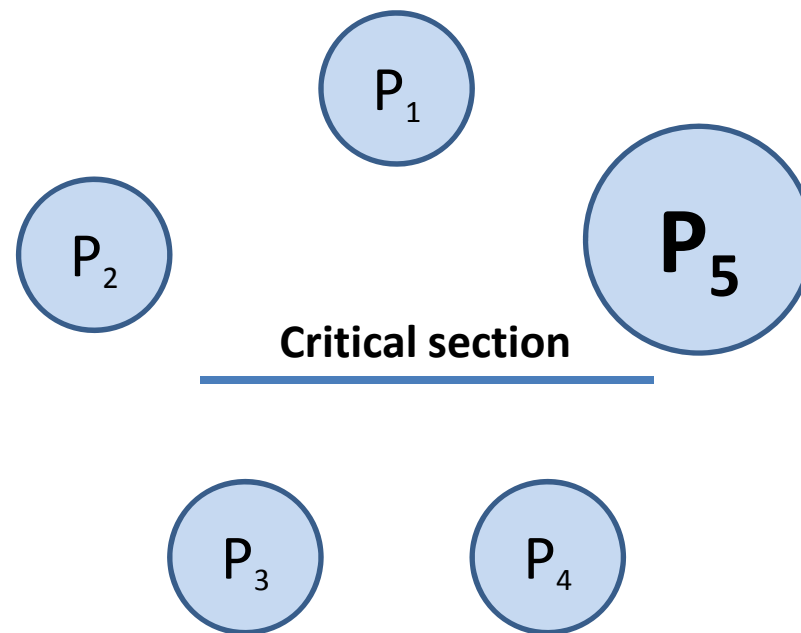
Centralized strategy

1. Elect a leader (details, cf. soon)

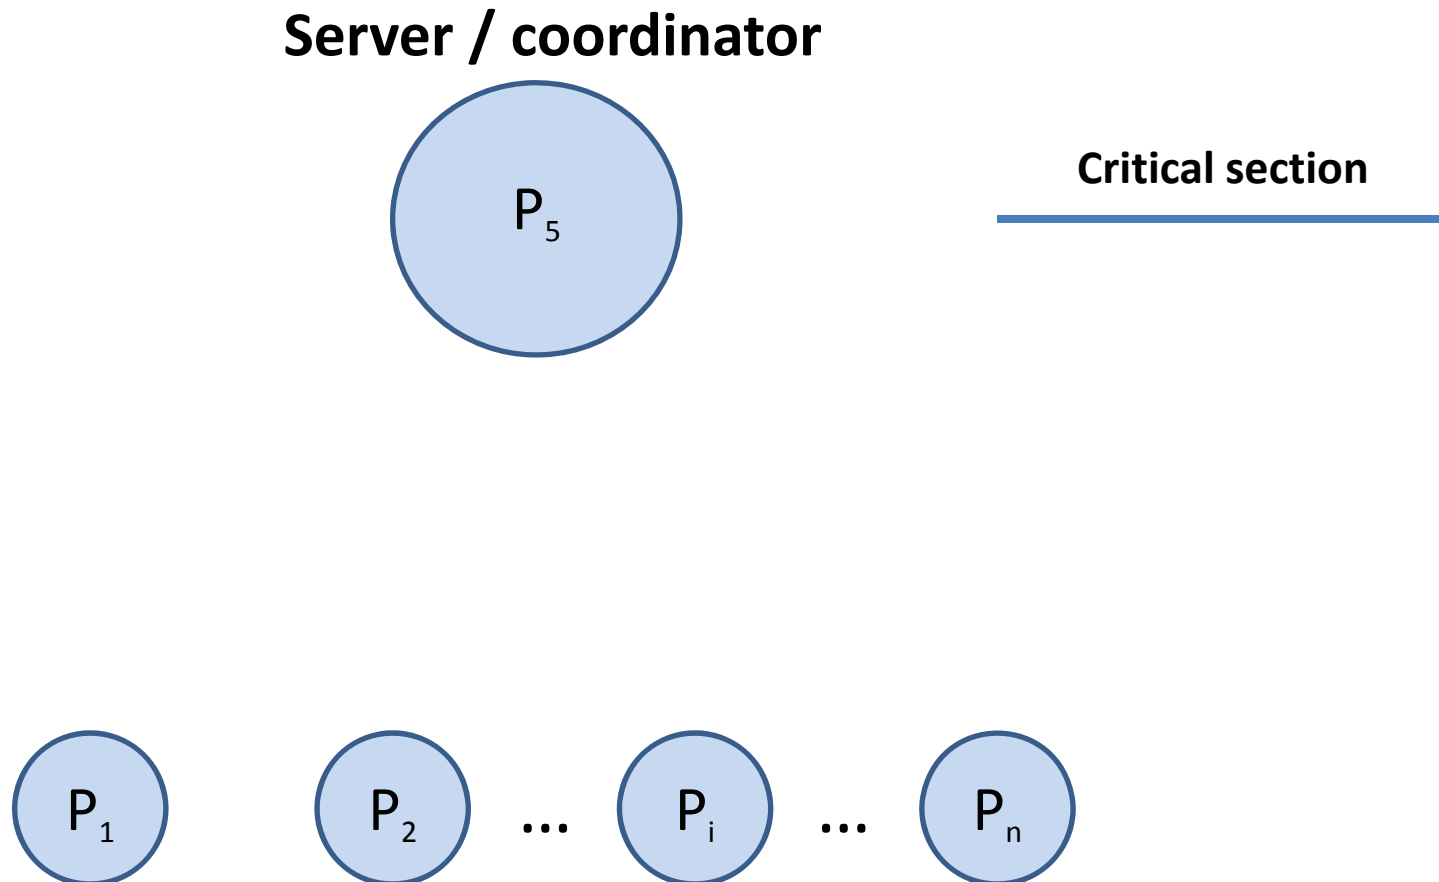


Centralized strategy

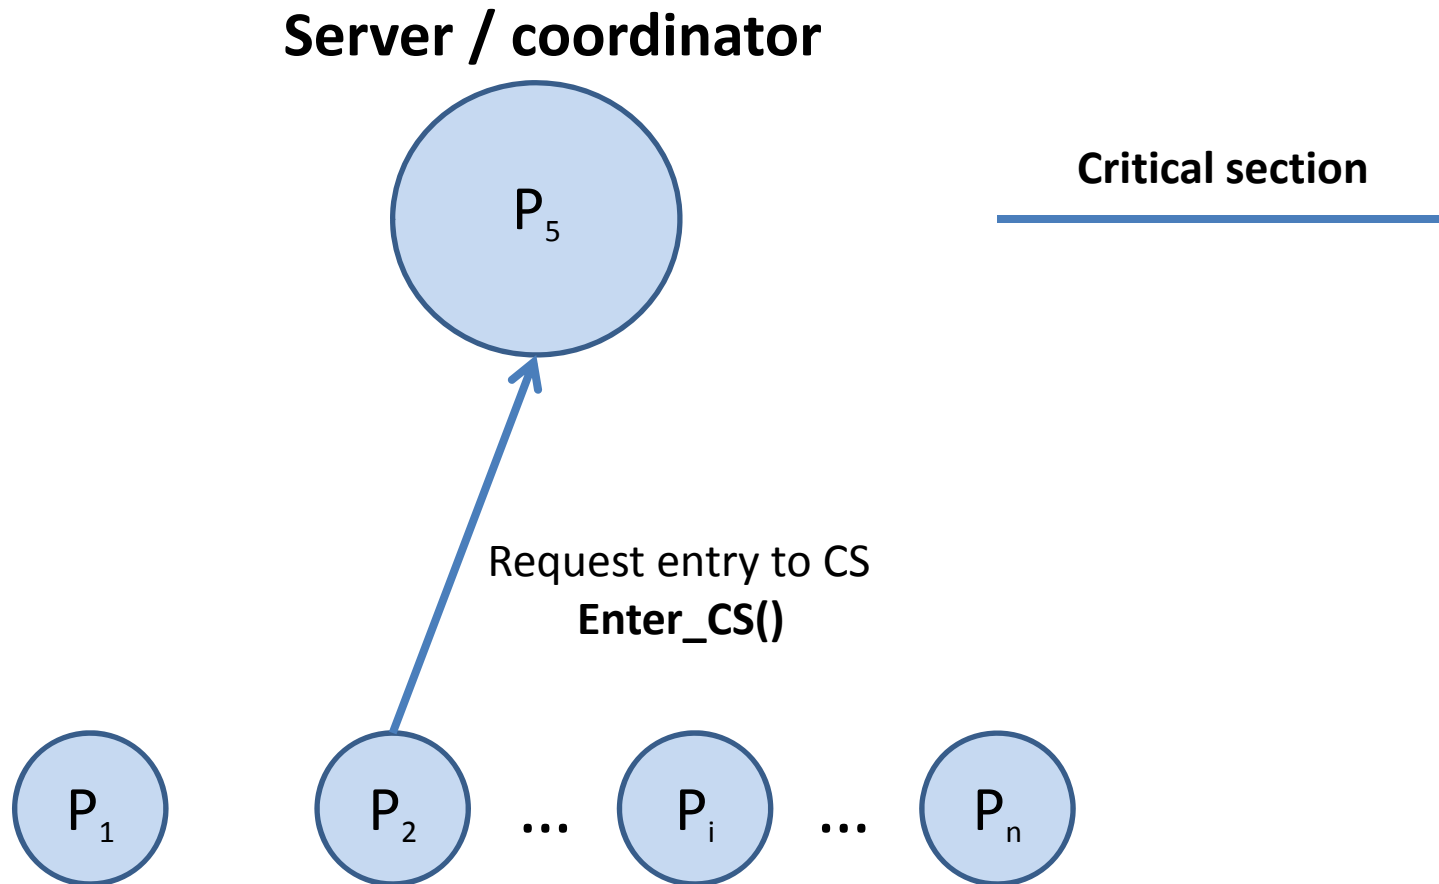
1. Elect a leader (details, cf. soon)



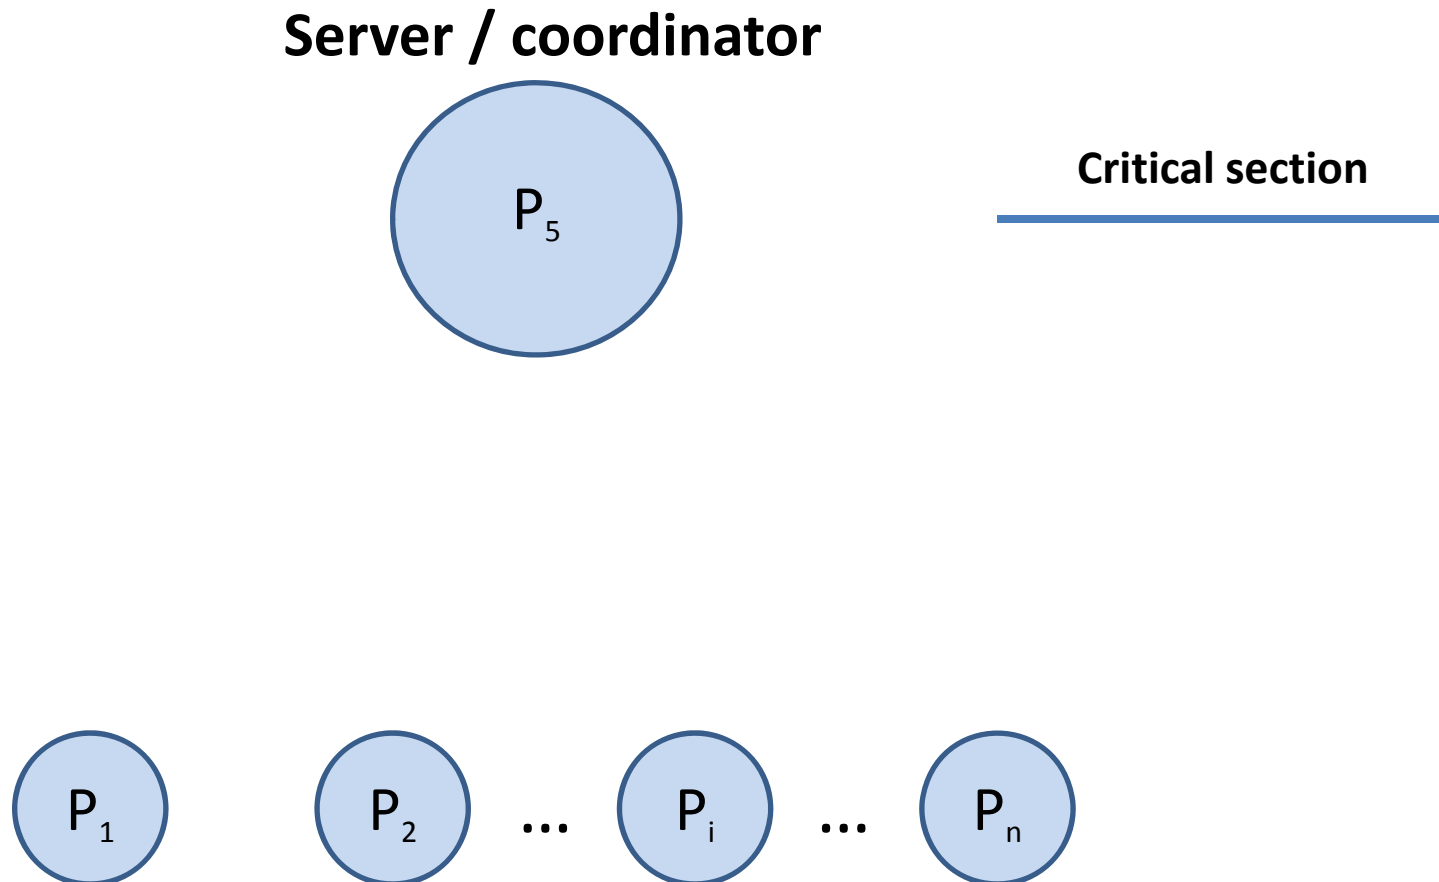
Centralized strategy: Empty CS



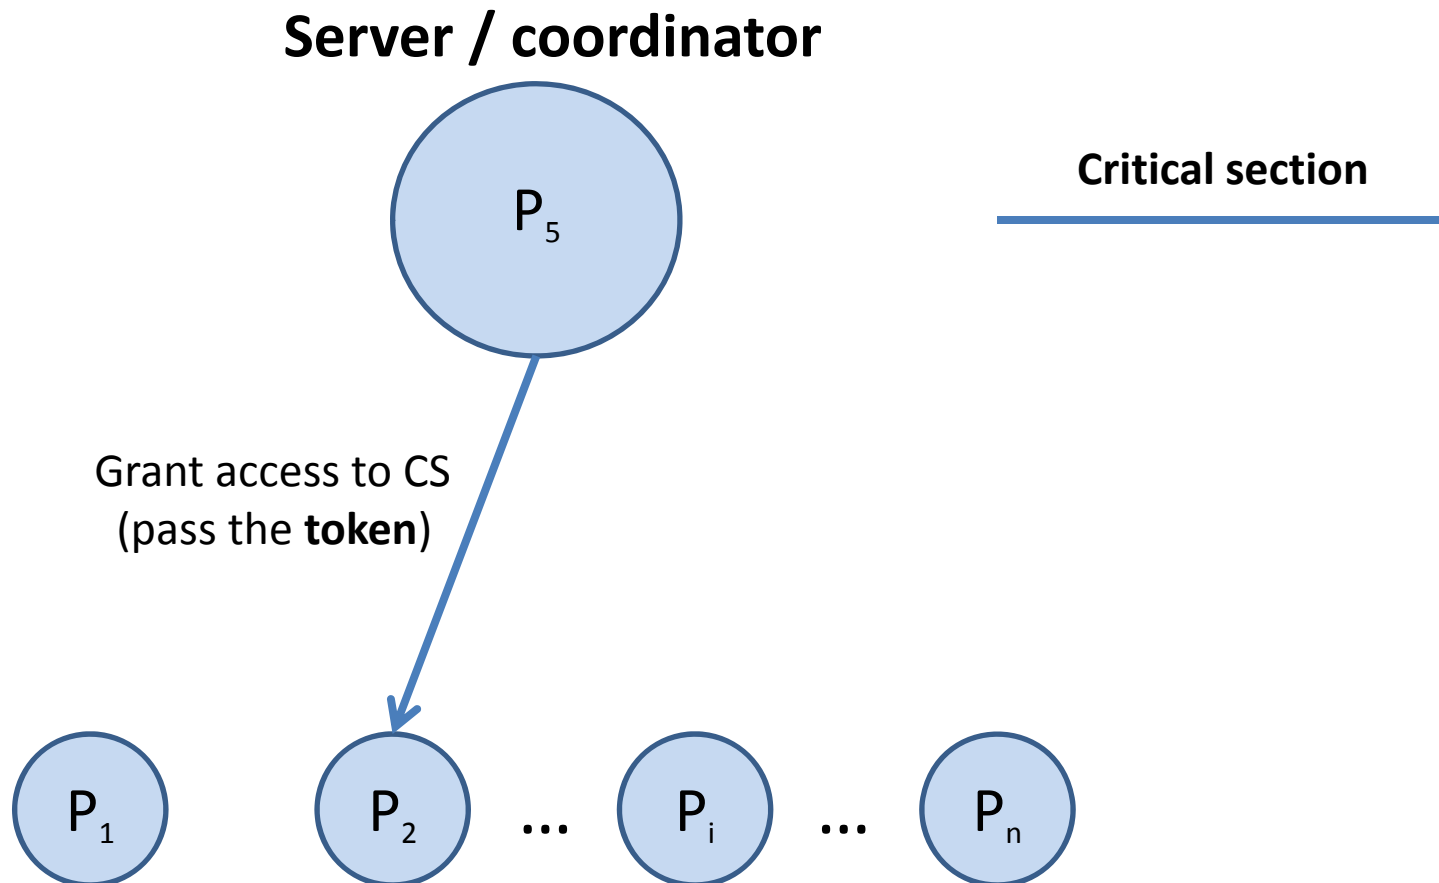
Centralized strategy: Empty CS



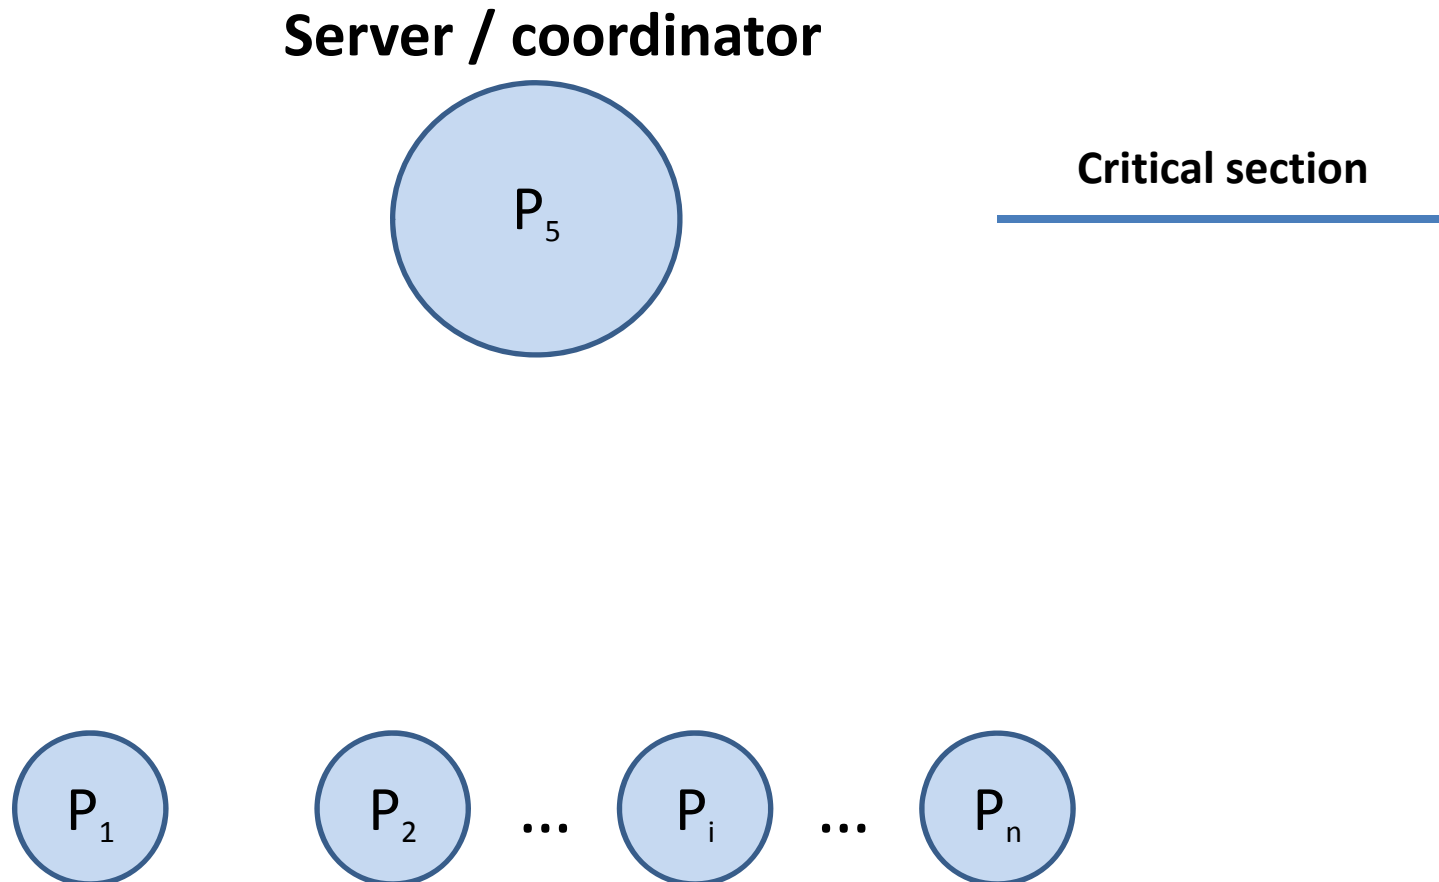
Centralized strategy: Empty CS



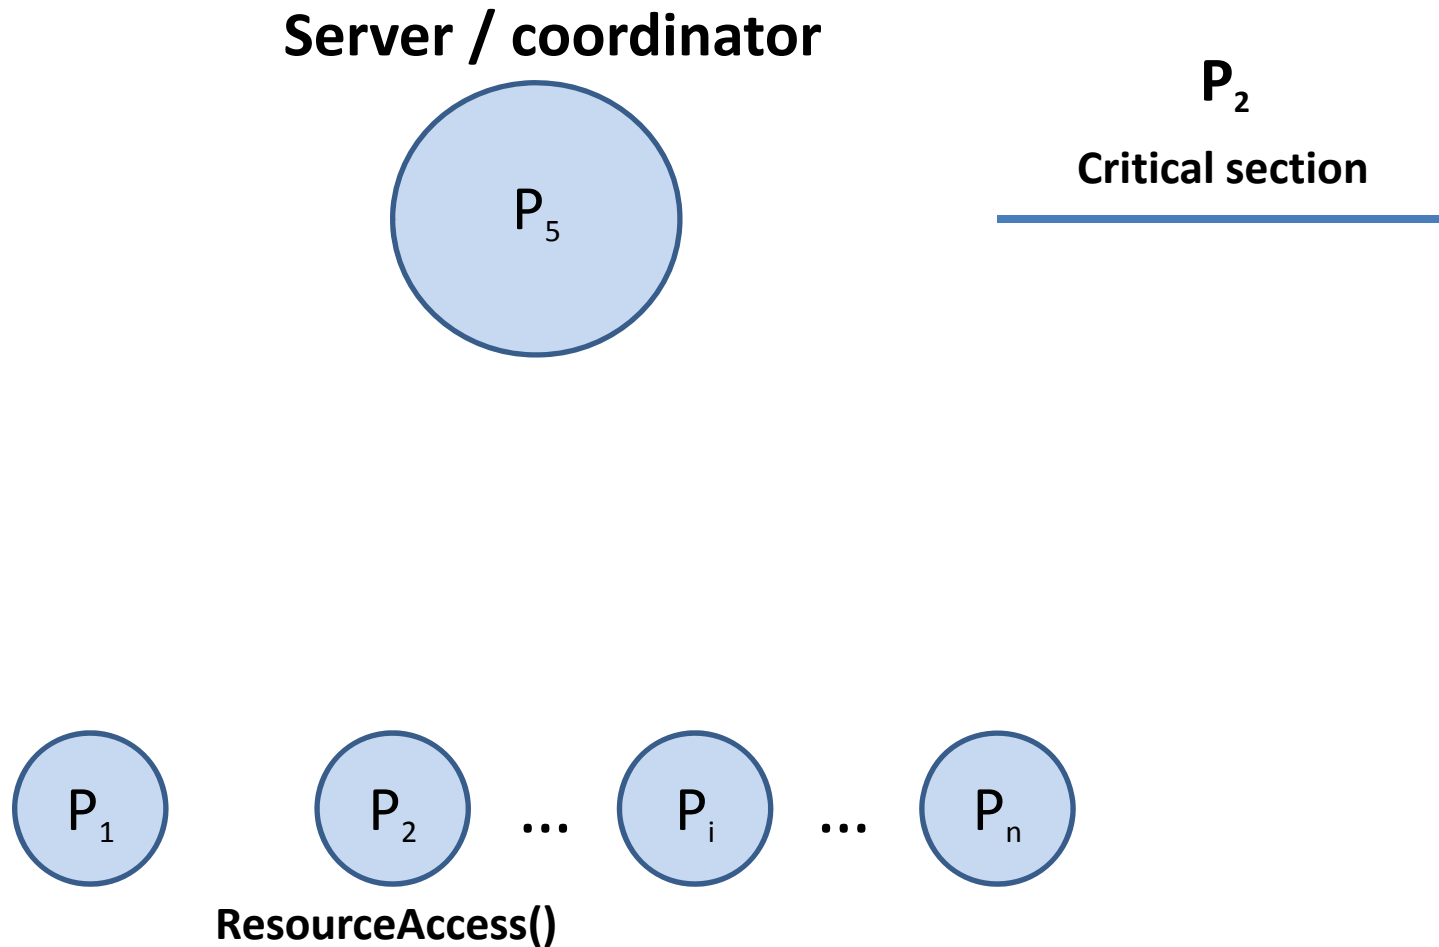
Centralized strategy: Empty CS



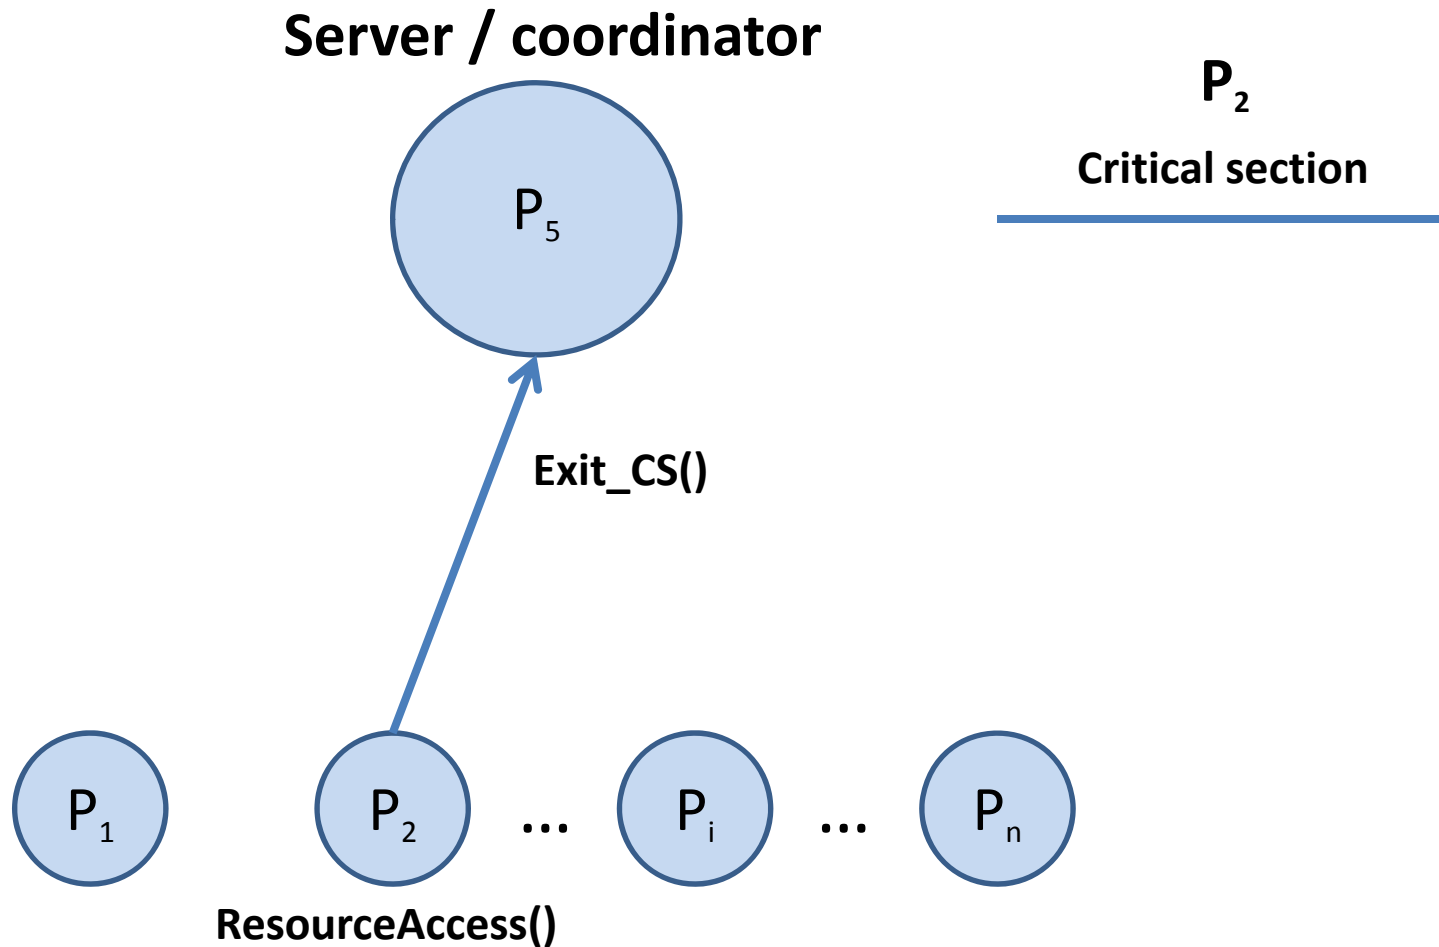
Centralized strategy: Empty CS



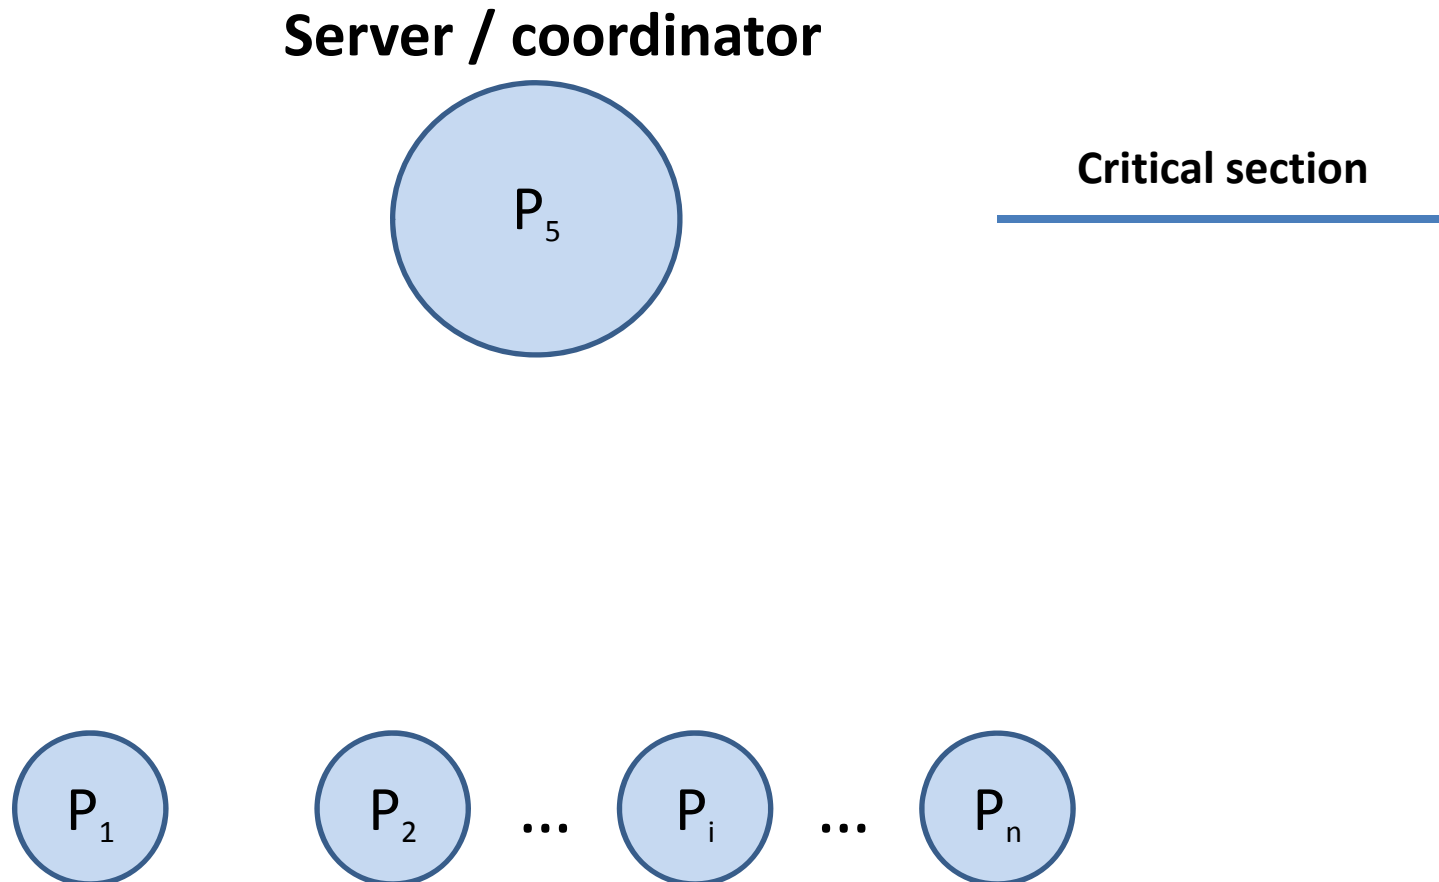
Centralized strategy: Empty CS



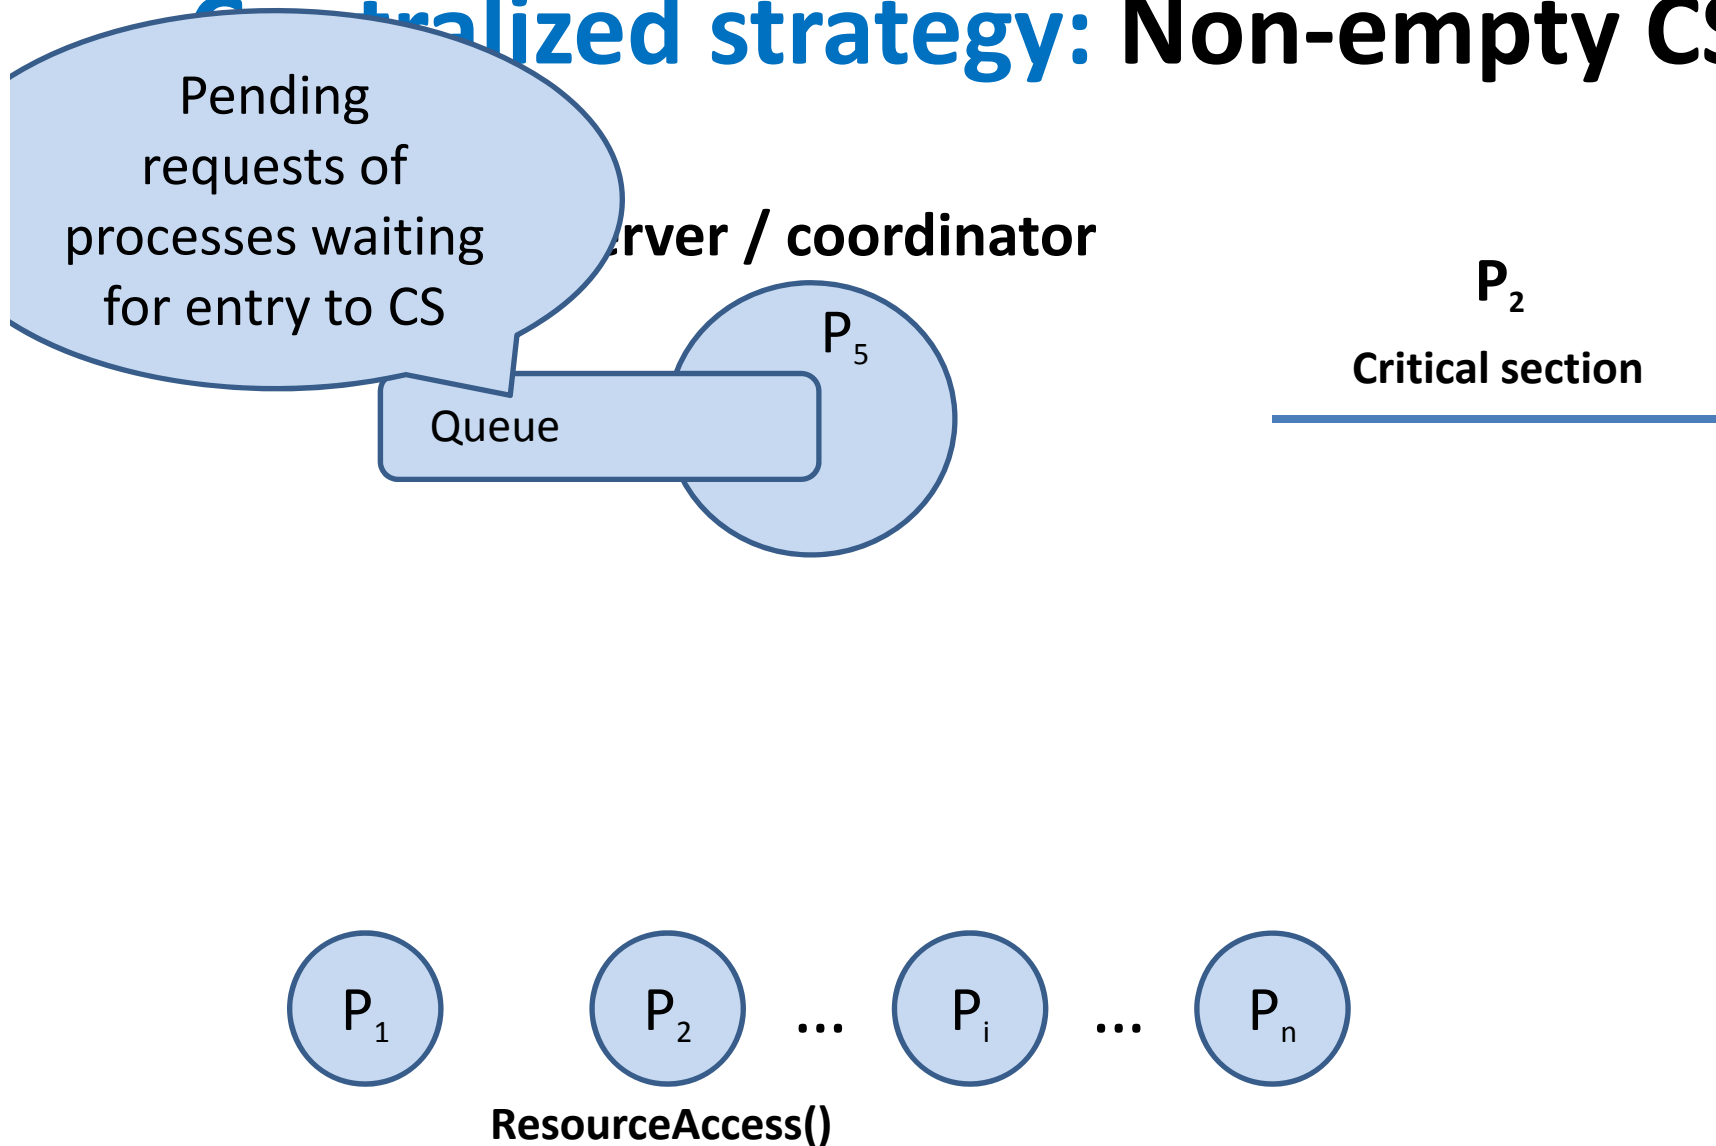
Centralized strategy: Empty CS



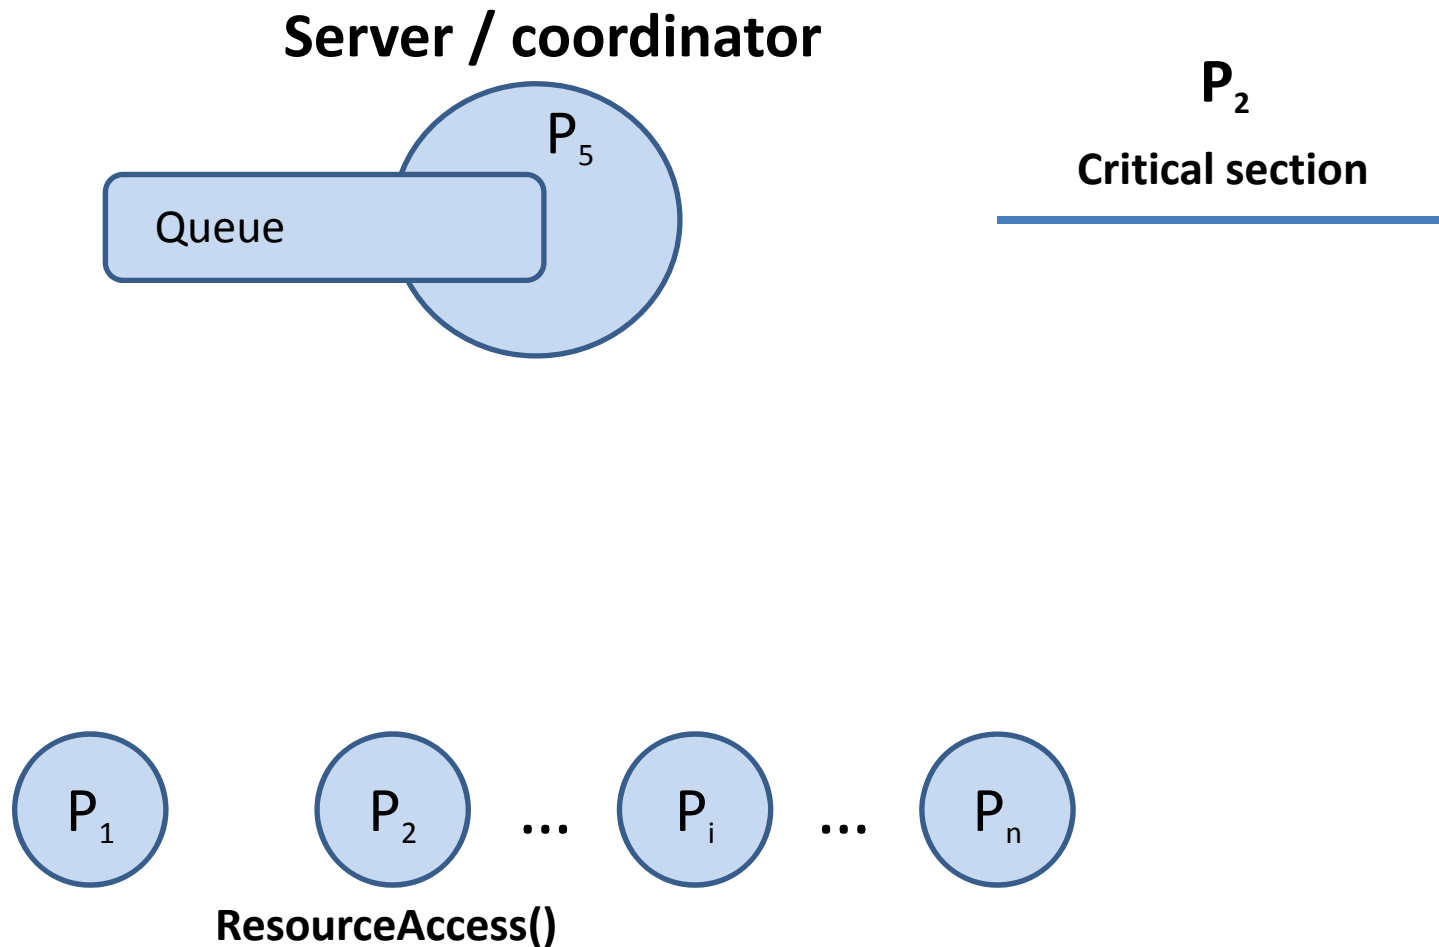
Centralized strategy: Empty CS



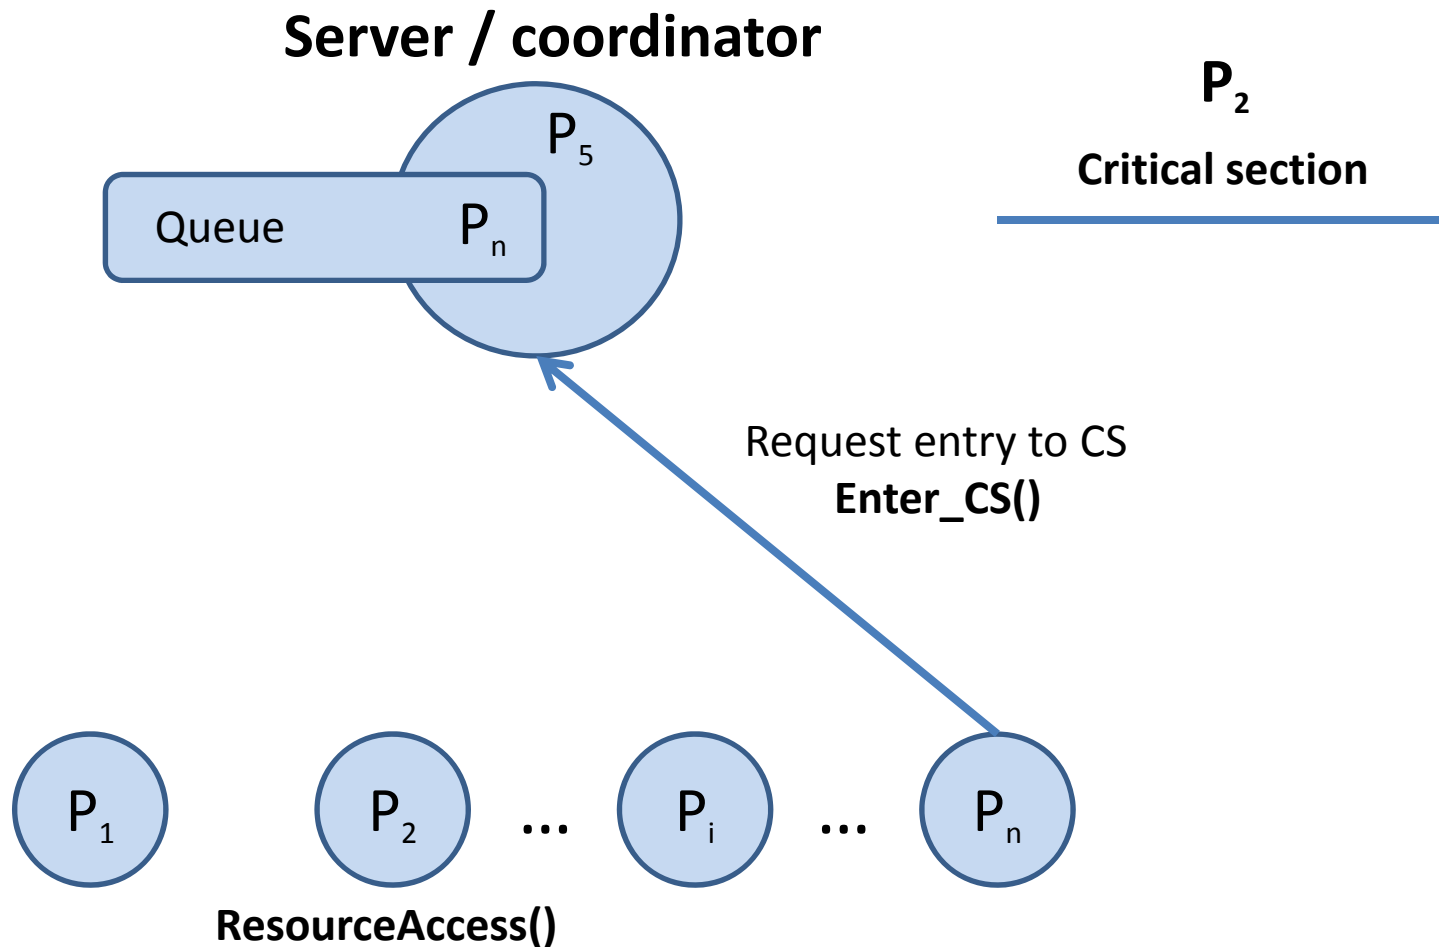
Centralized strategy: Non-empty CS



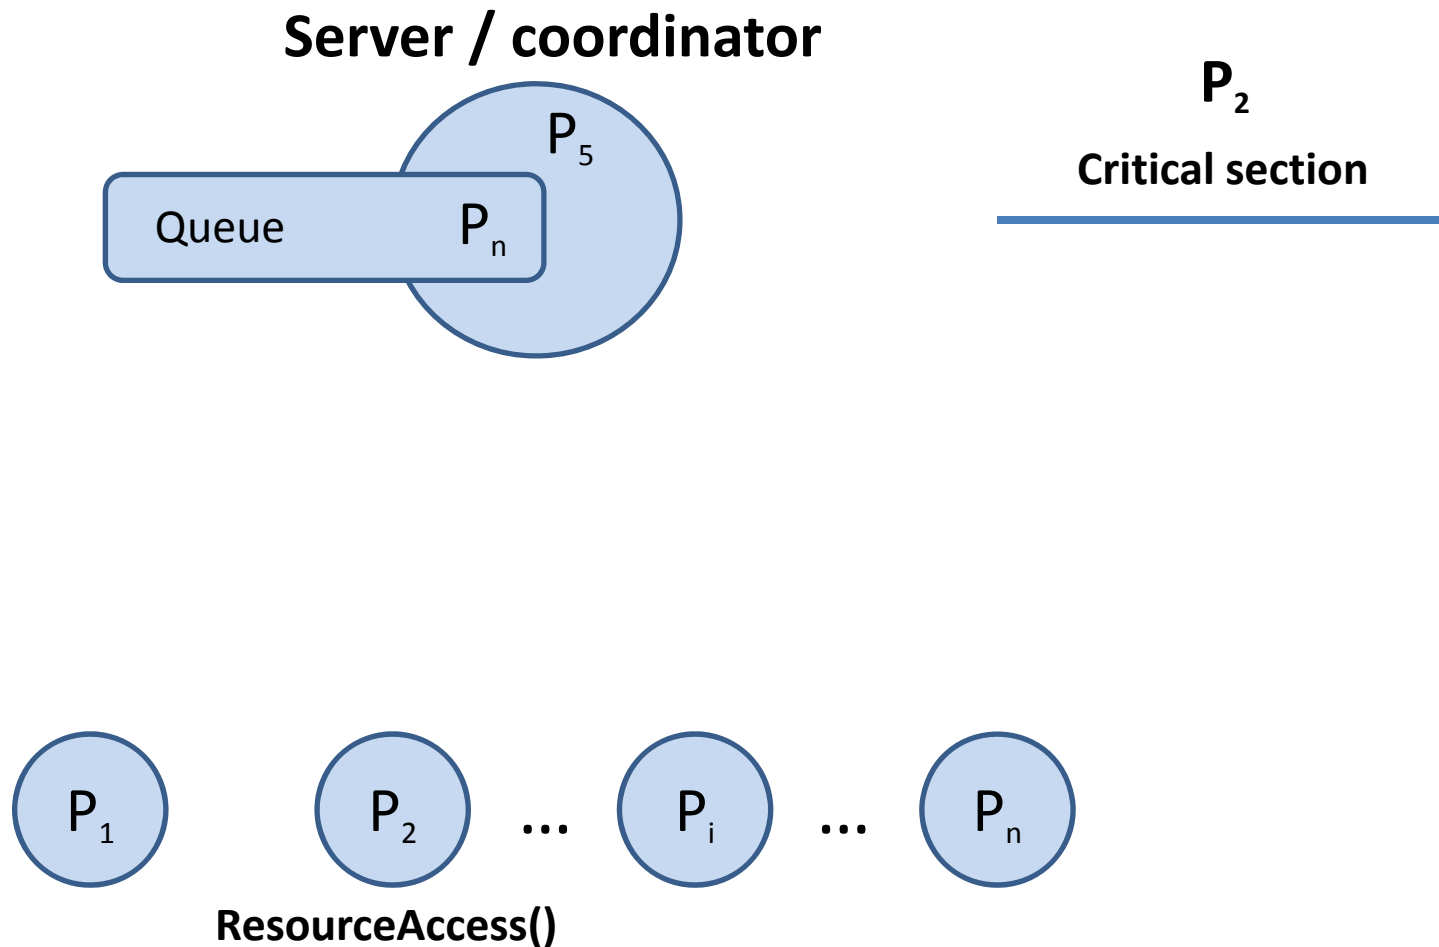
Centralized strategy: Non-empty CS



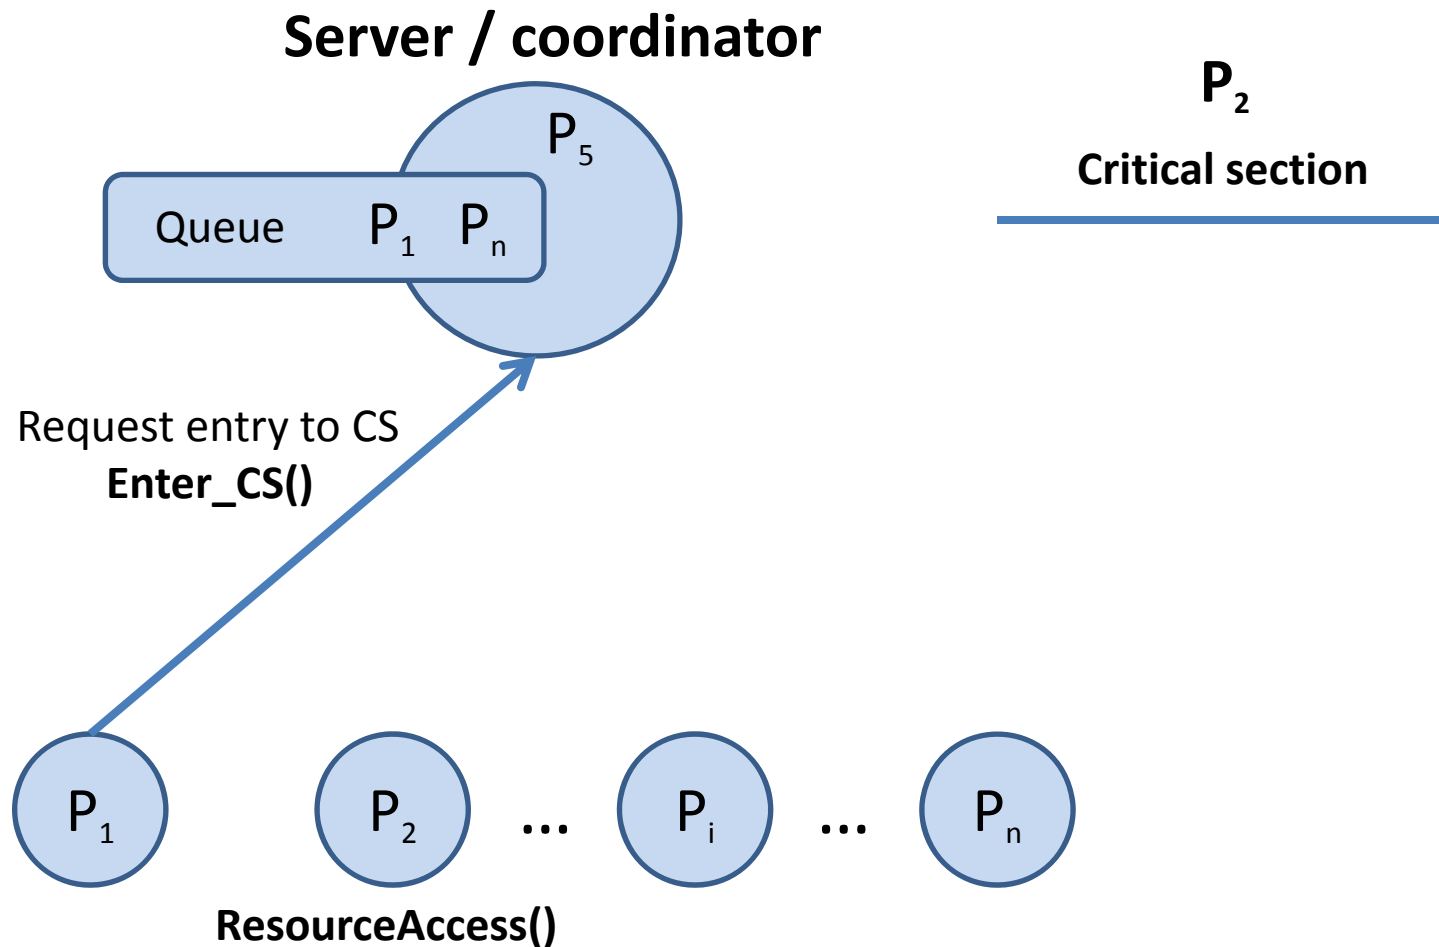
Centralized strategy: Non-empty CS



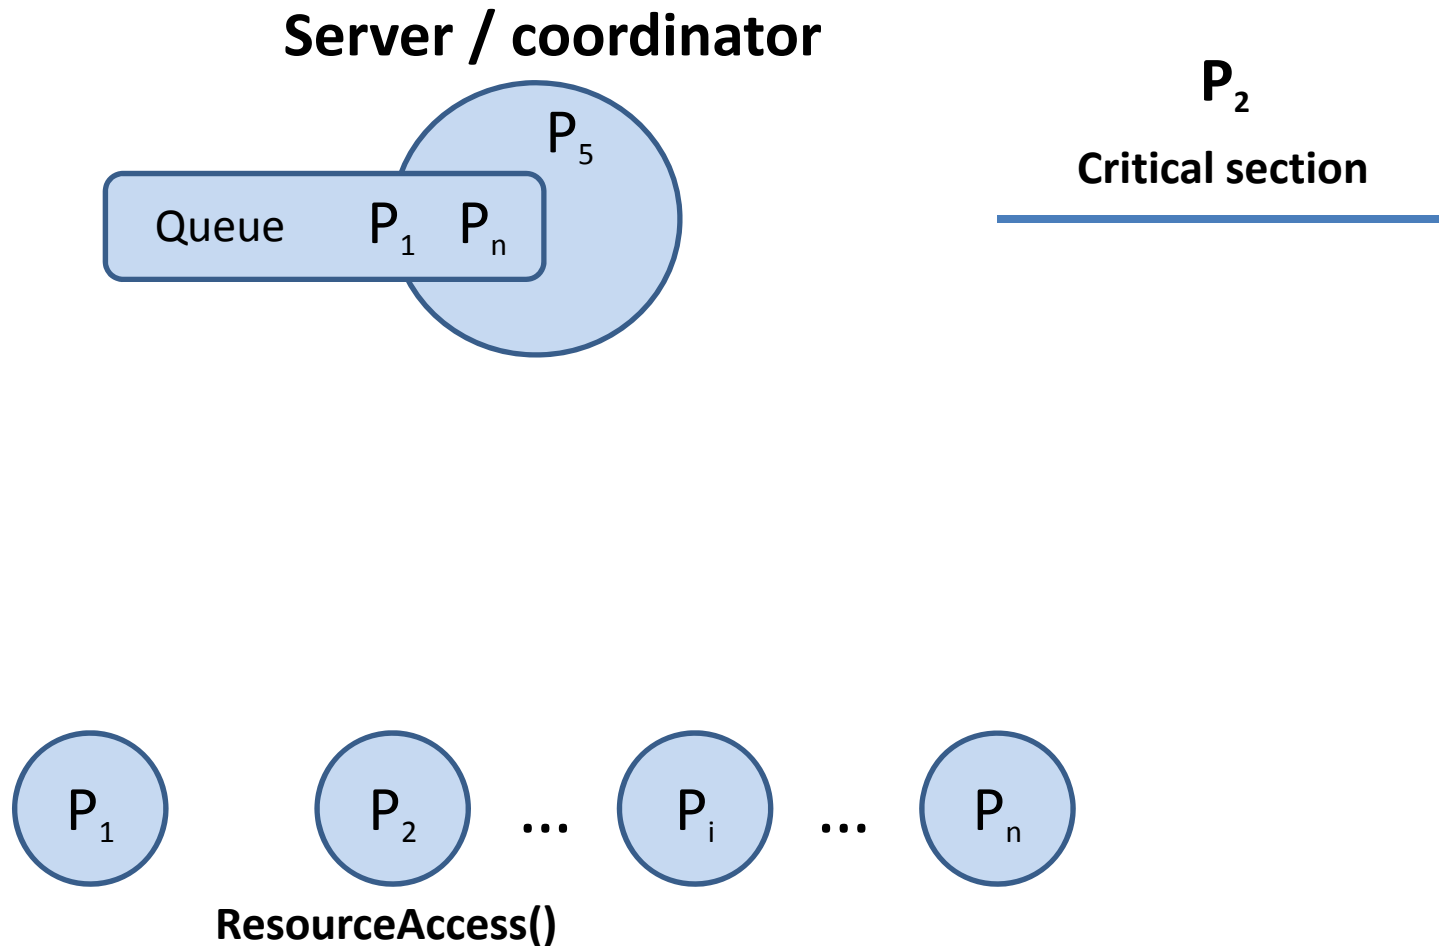
Centralized strategy: Non-empty CS



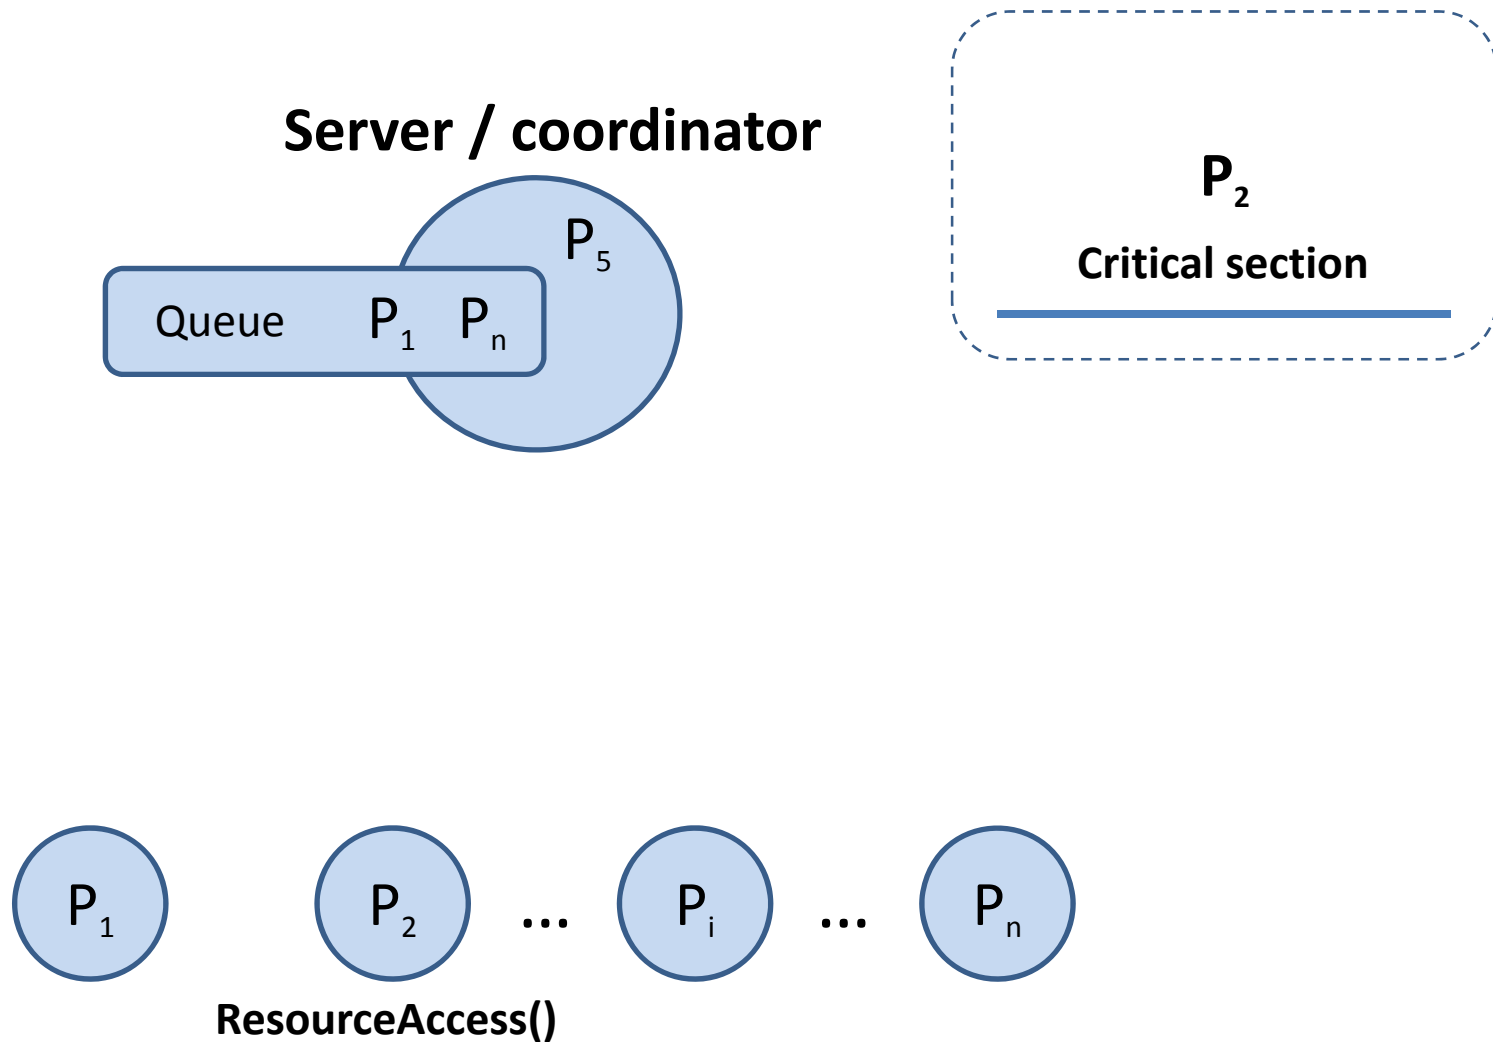
Centralized strategy: Non-empty CS



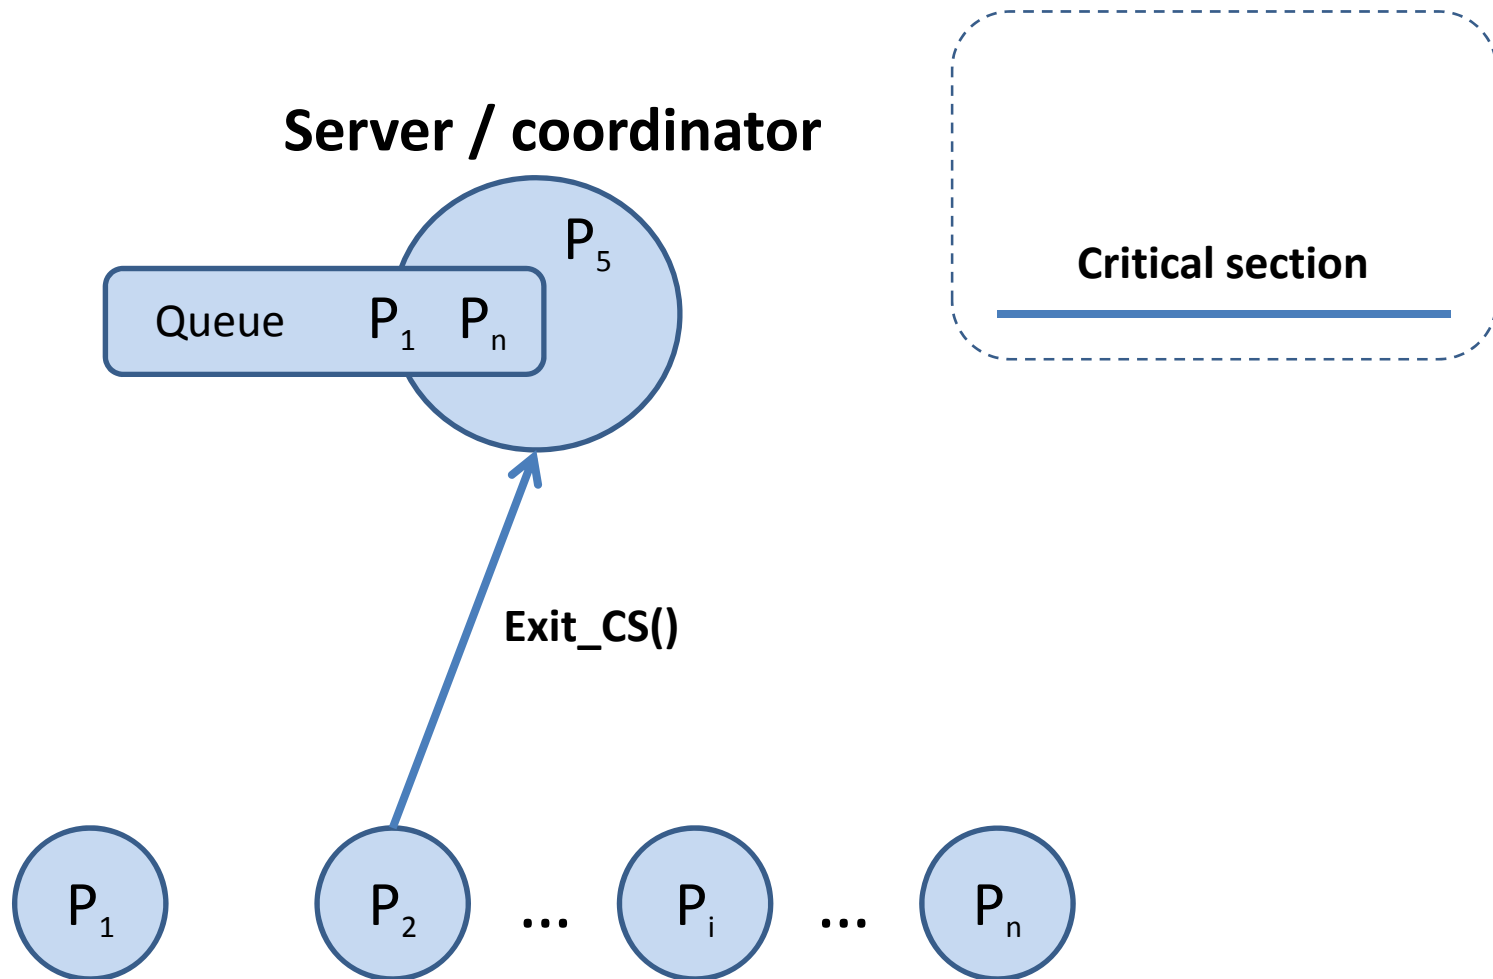
Centralized strategy: Non-empty CS



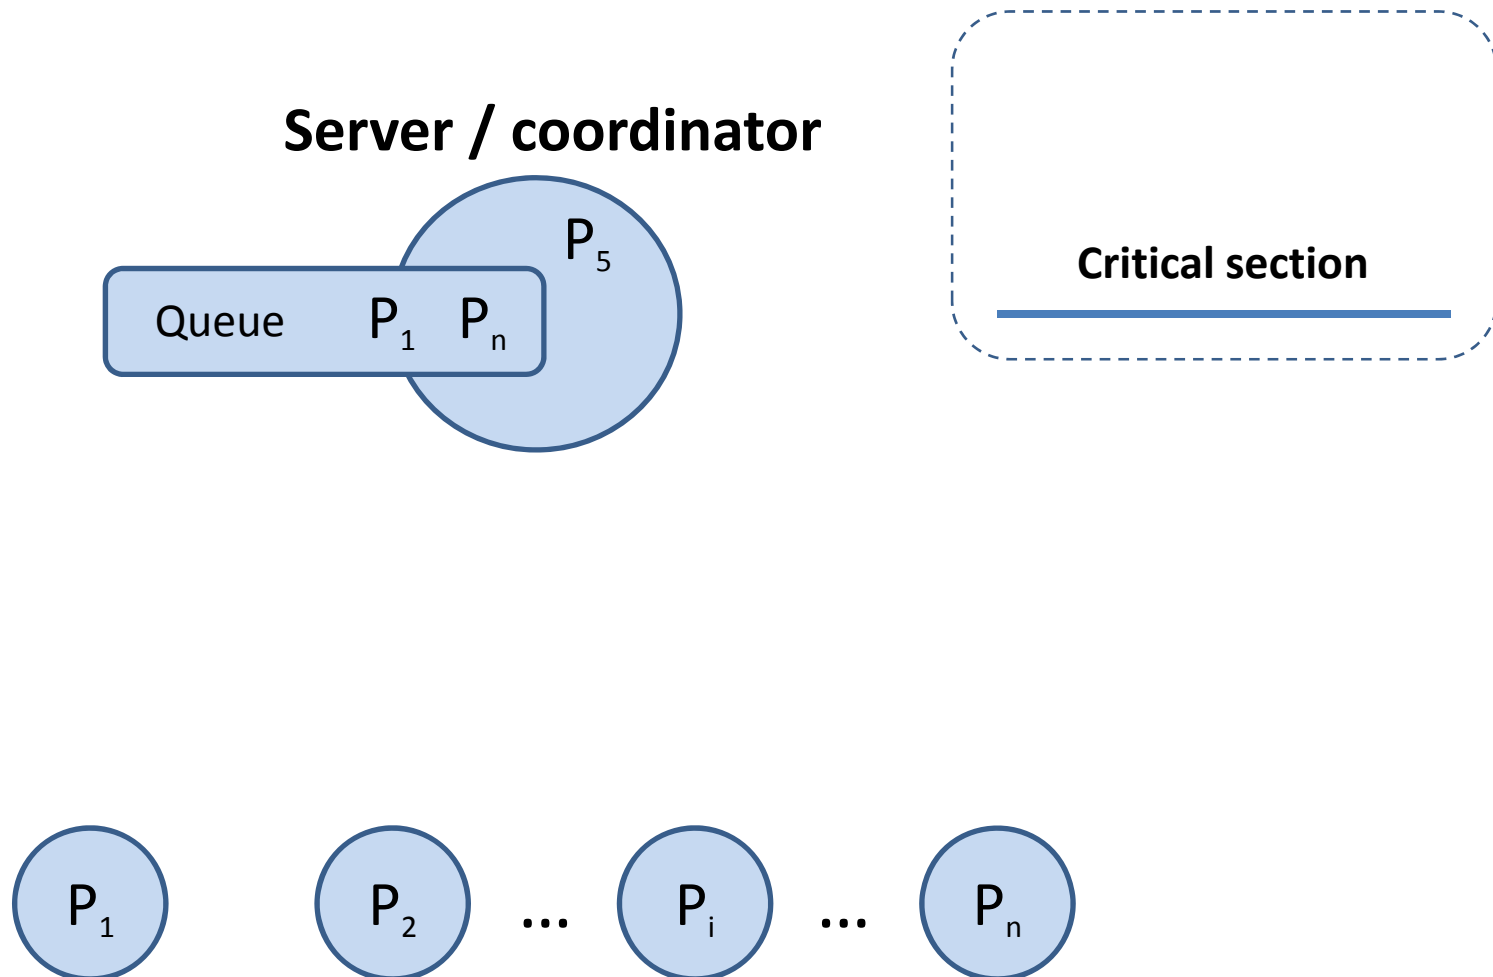
Centralized strategy: Exit CS



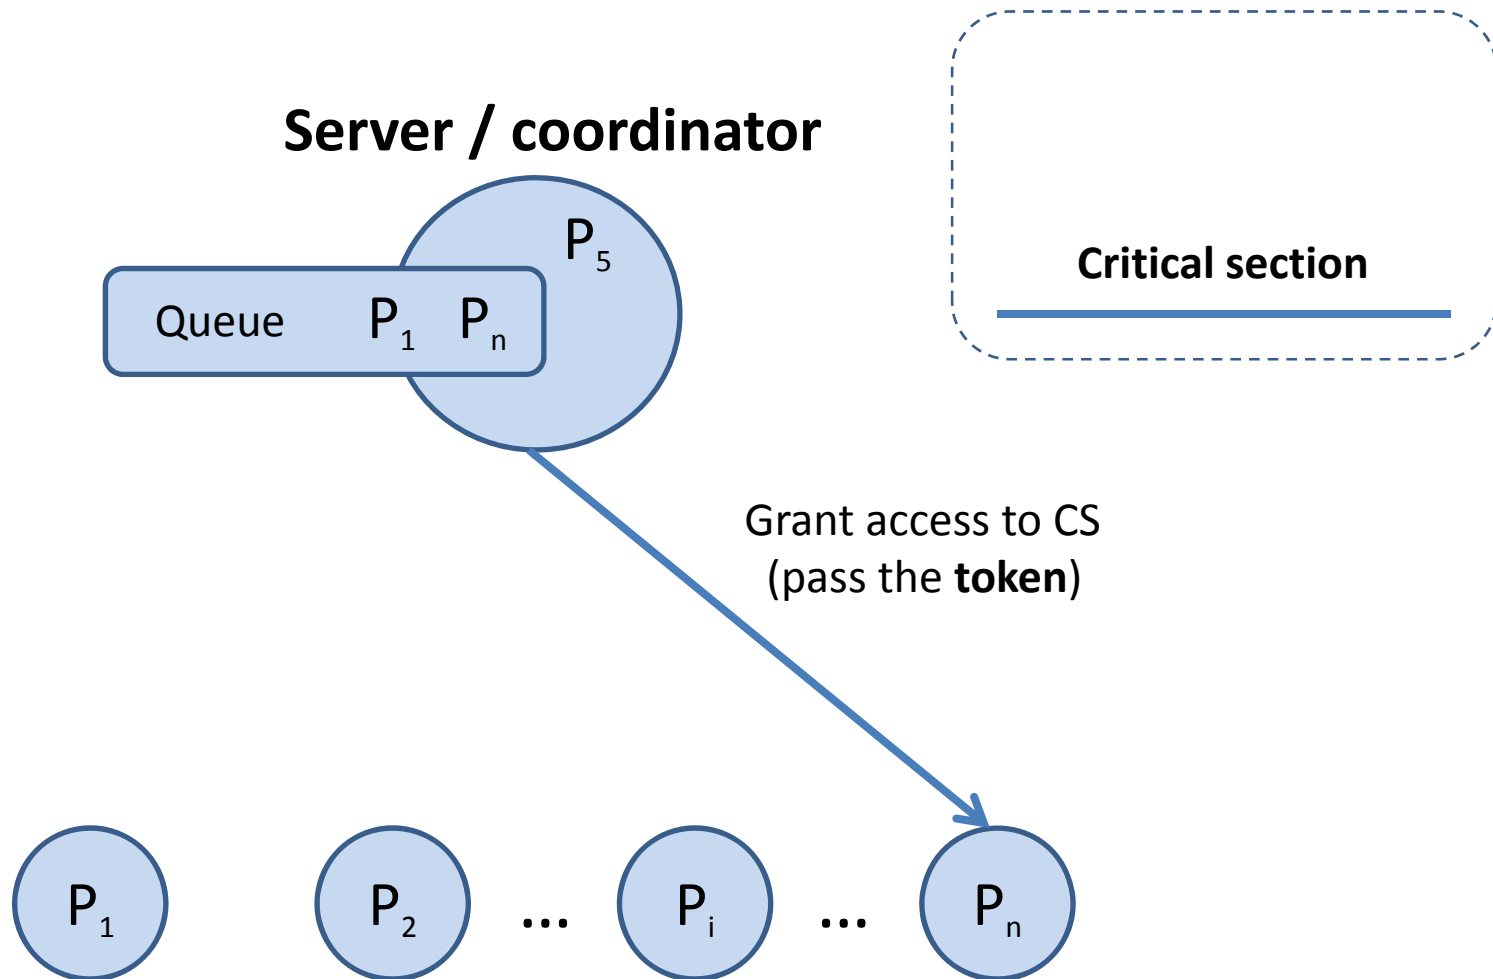
Centralized strategy: Exit CS



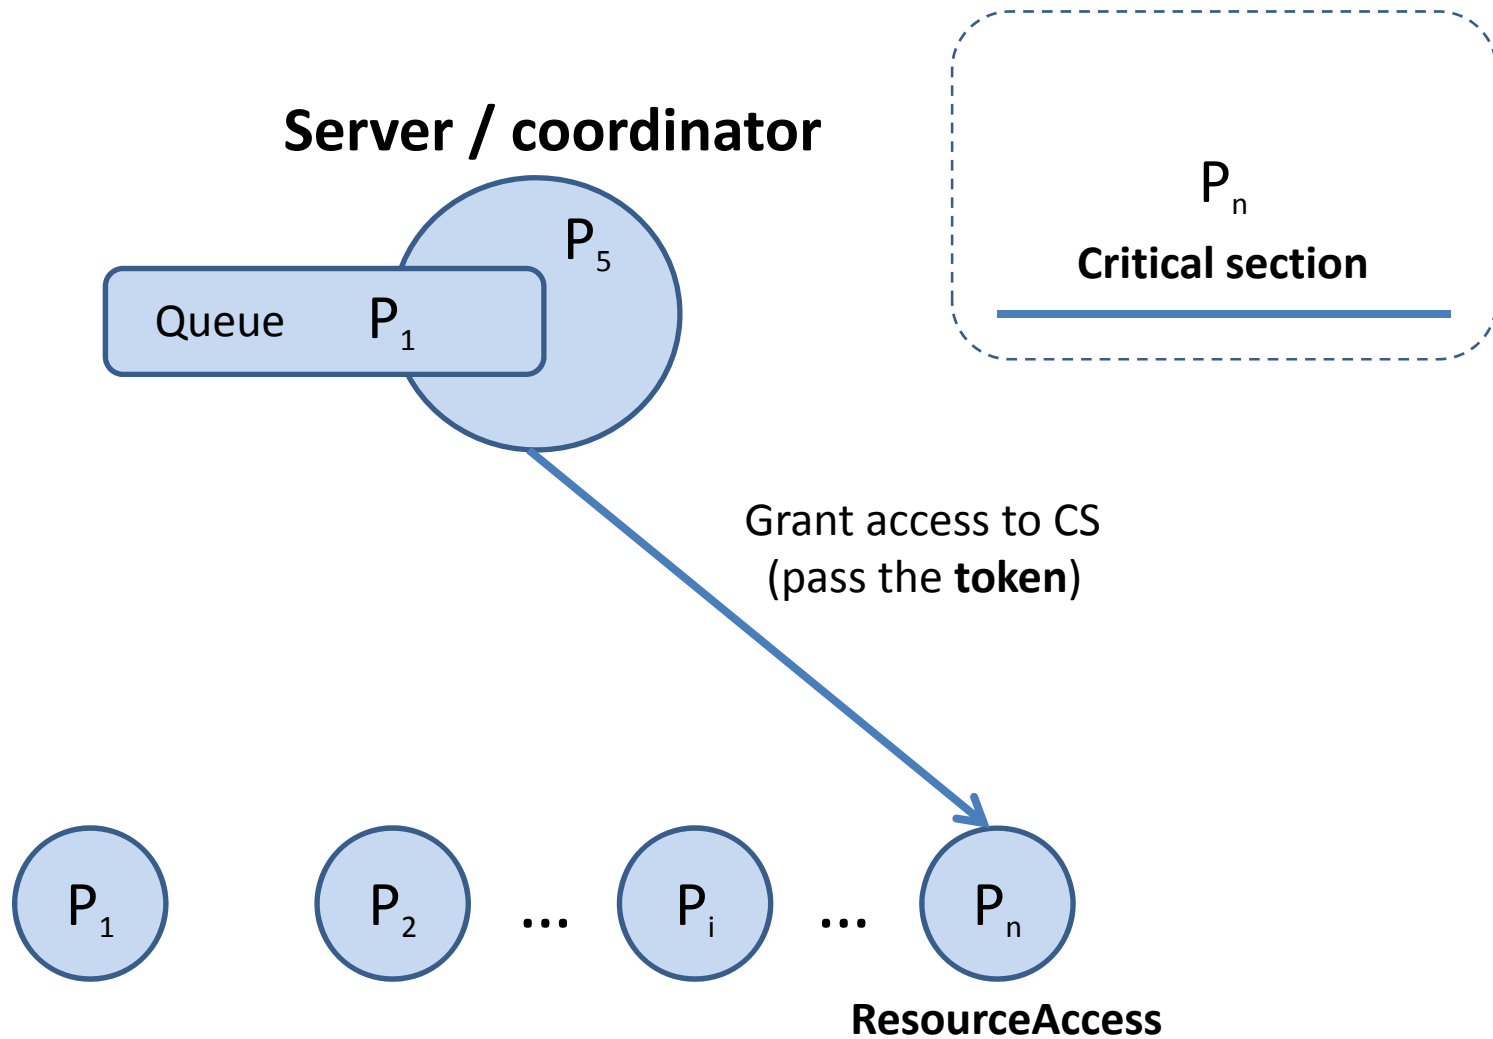
Centralized strategy: Exit CS



Centralized strategy: Exit CS



Centralized strategy: Exit CS



Centralized strategy analysis I

- Meets requirements: Safety, liveness, no starvation
- ***Does solution meet the ordering requirement?***
- Advantages
 - **Simple to implement**
- Disadvantages
 - **Single point of failure**
 - Bottleneck, network congestion, timeout
- Deadlock potential for multiple resources with separate servers

Centralized strategy analysis II

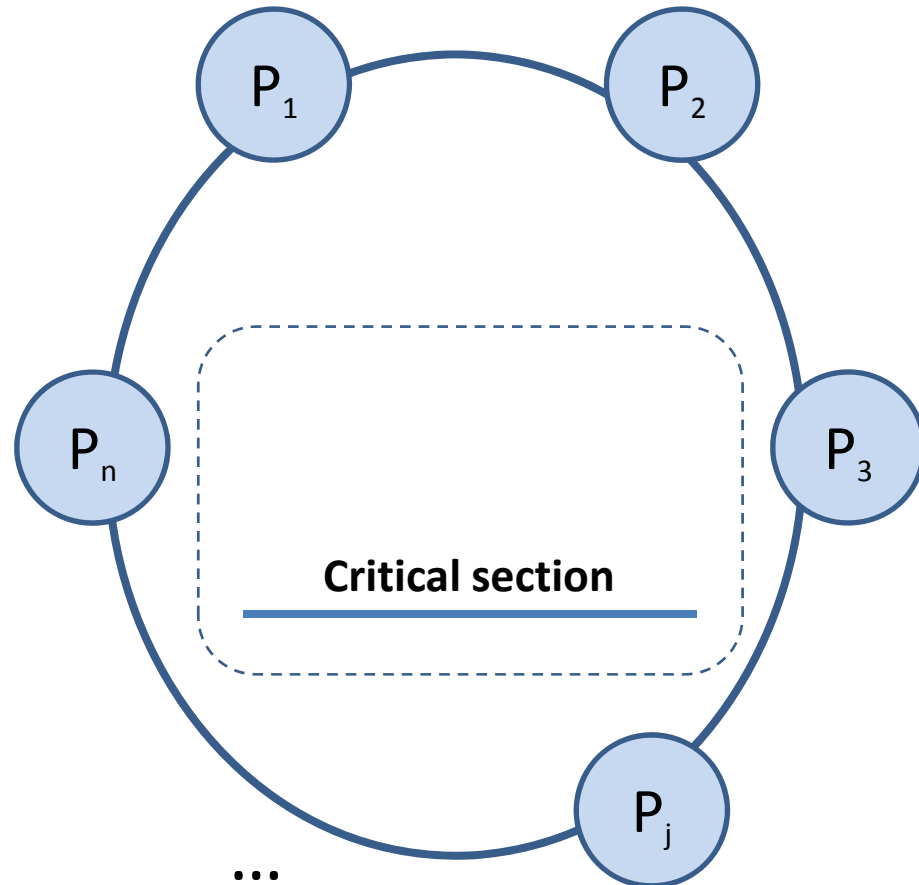
- Enter_CS()
 - **Two messages:** Request & Grant
 - One round of communication (RTT delay)
- Exit_CS()
 - **One message:** Release message
 - No delay for the process in CS

Distributed strategy

- In our distributed strategies, **the same decision made on each node**, independent of the other nodes in the system.
- Selected algorithms
 - Ring-based algorithm
 - **Logical clocks-based algorithm (Lamport, 1976)**
 - Ricart & Agrawala, 1981
 - Maekawa, 1985
 - Many more

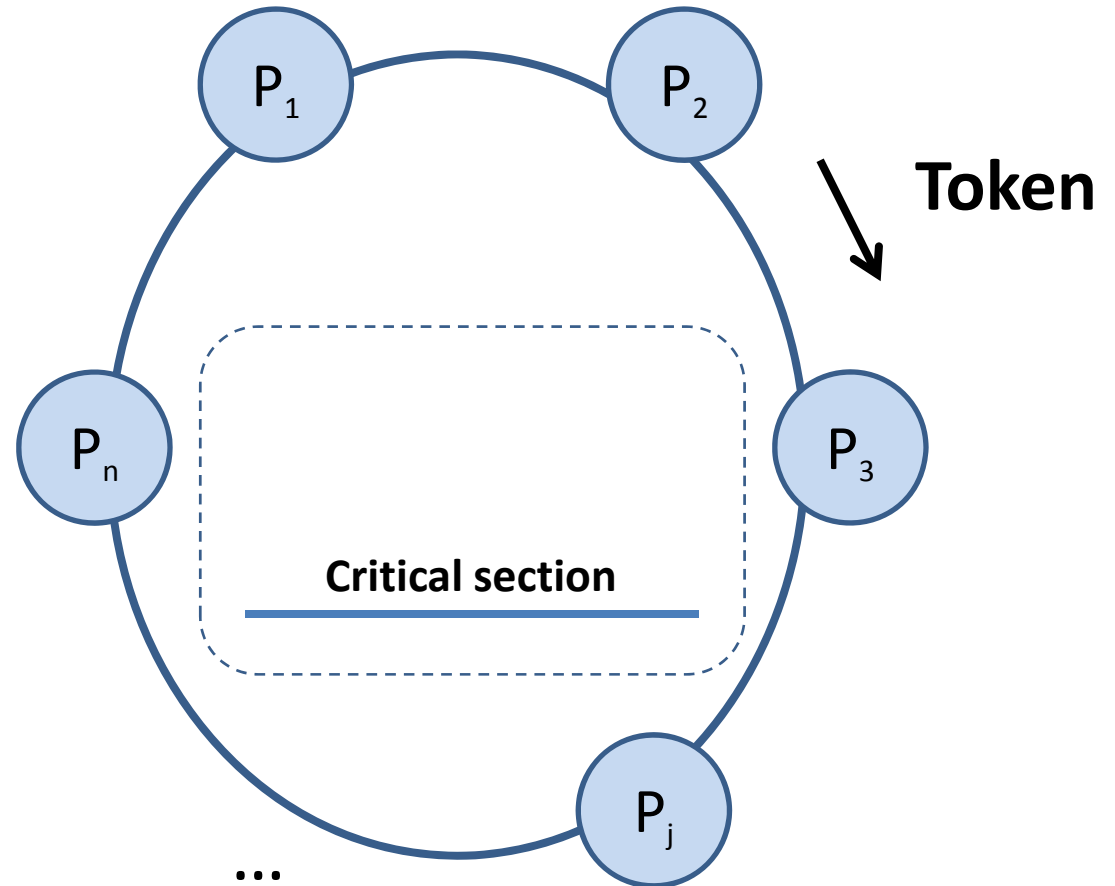
Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology



Ring-based algorithm

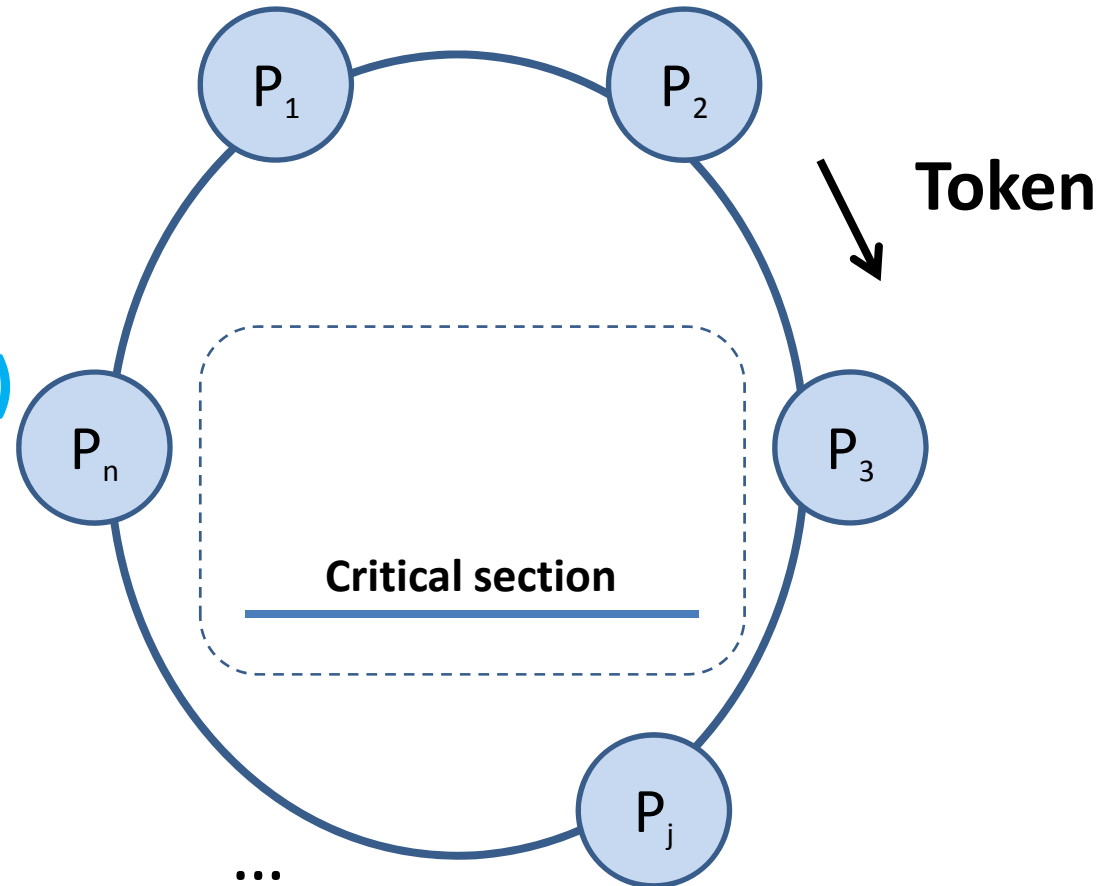
- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

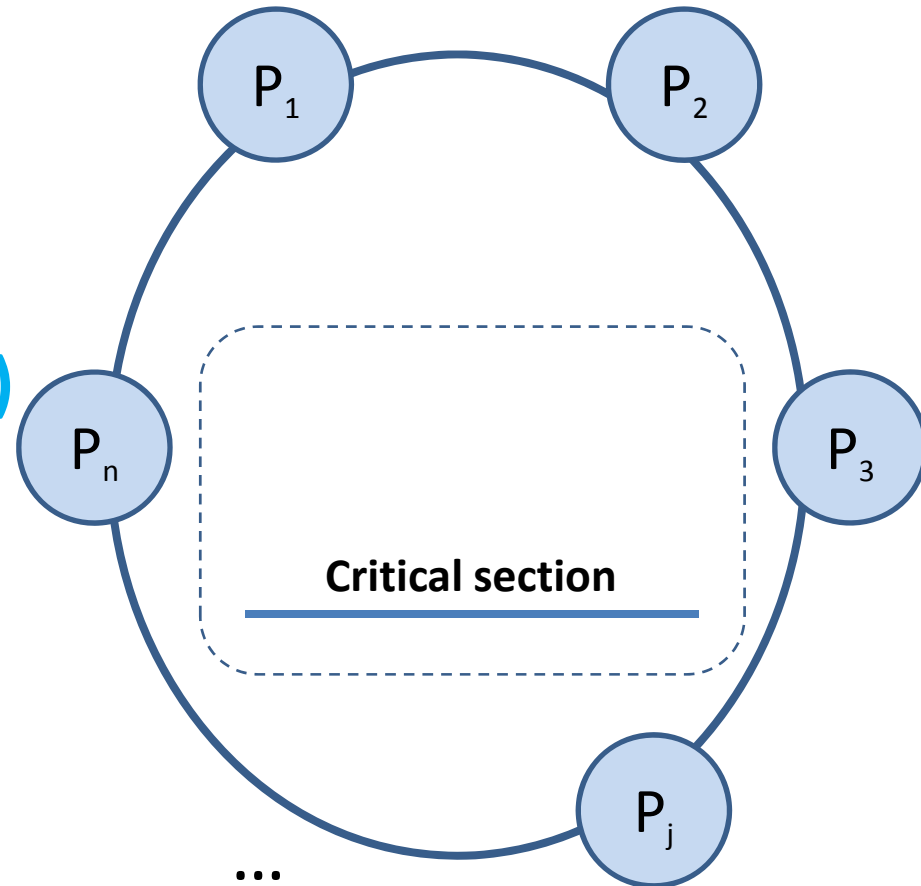
Enter_CS()



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

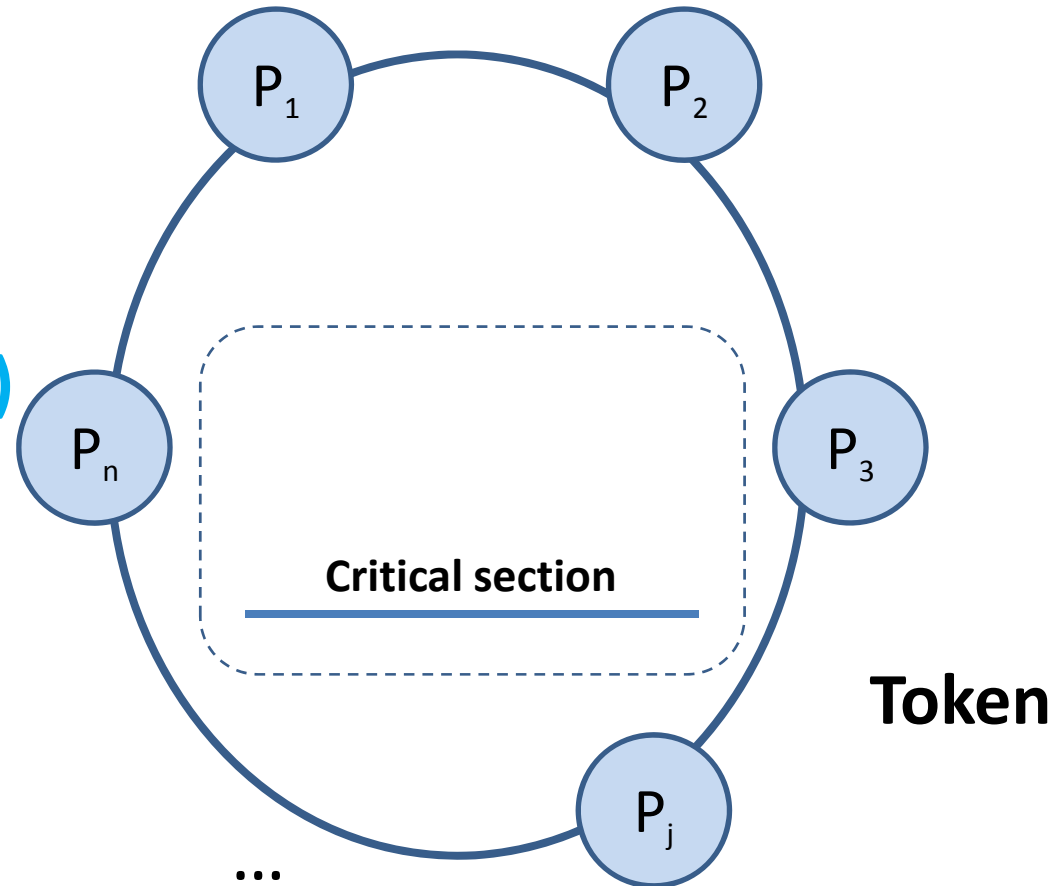
Enter_CS()



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

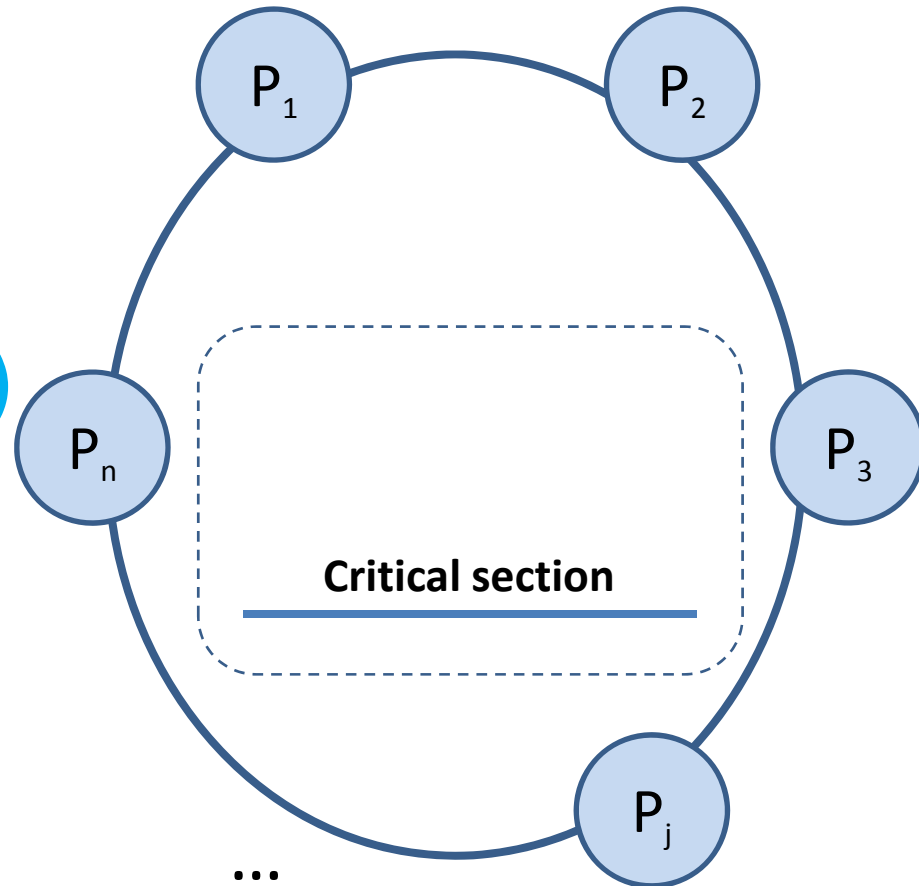
Enter_CS()



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

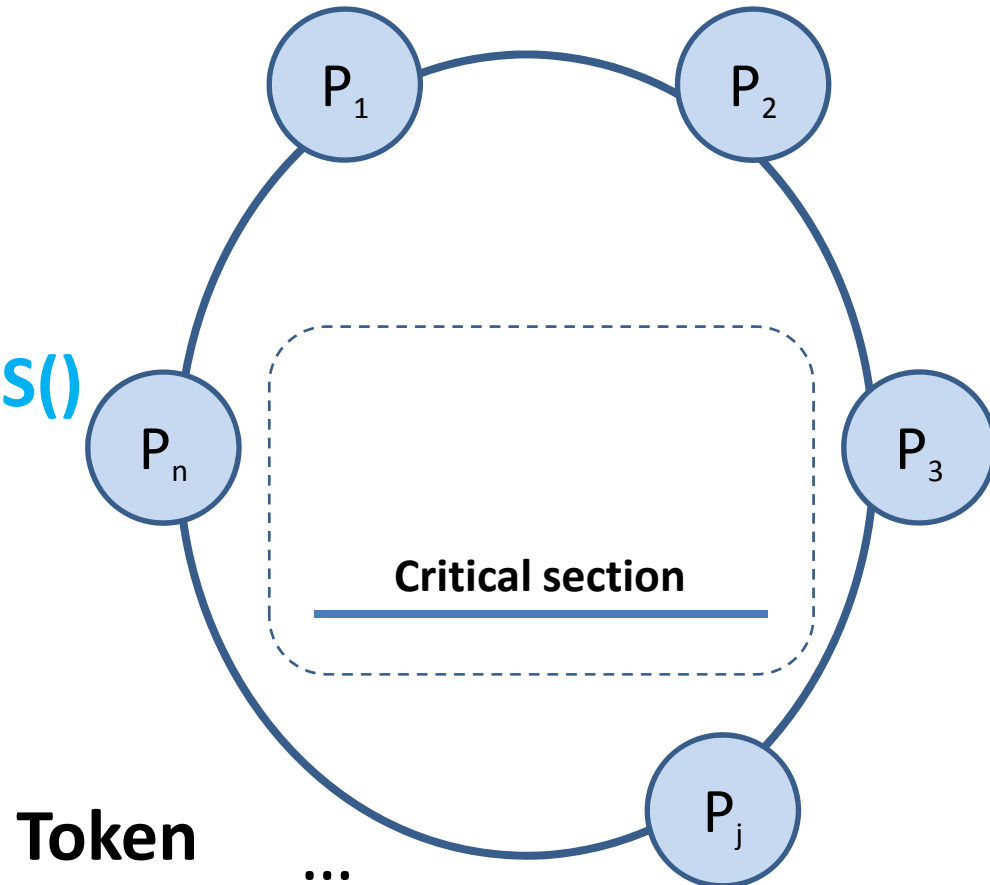
Enter_CS()



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

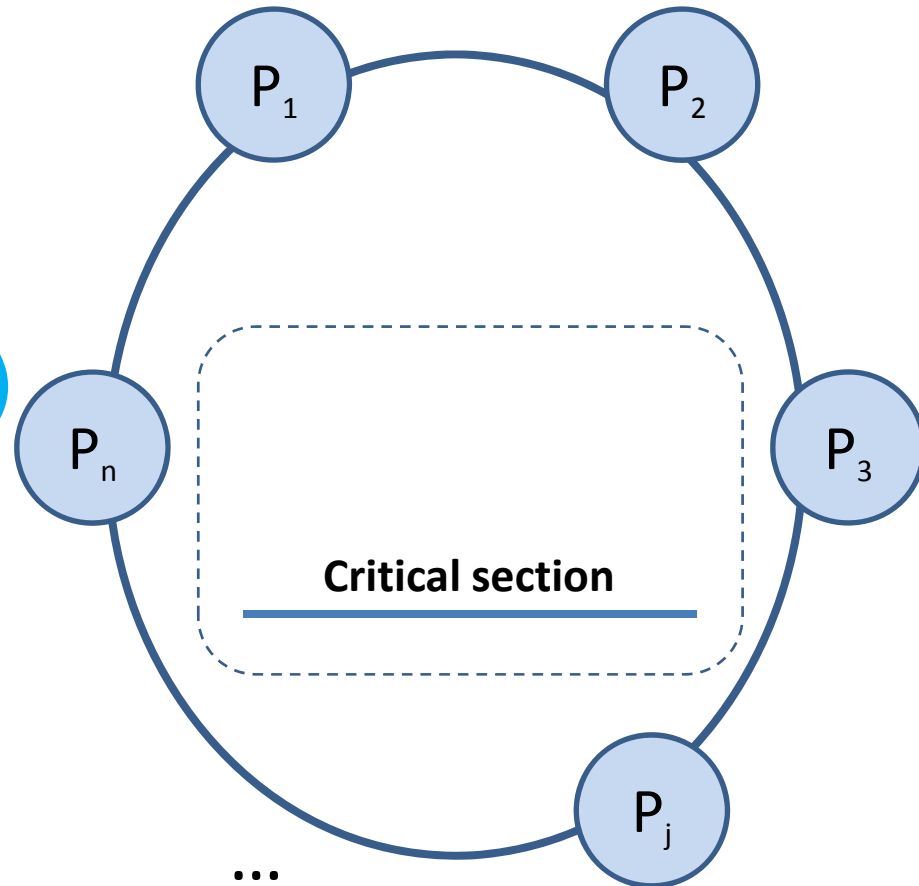
Enter_CS()



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

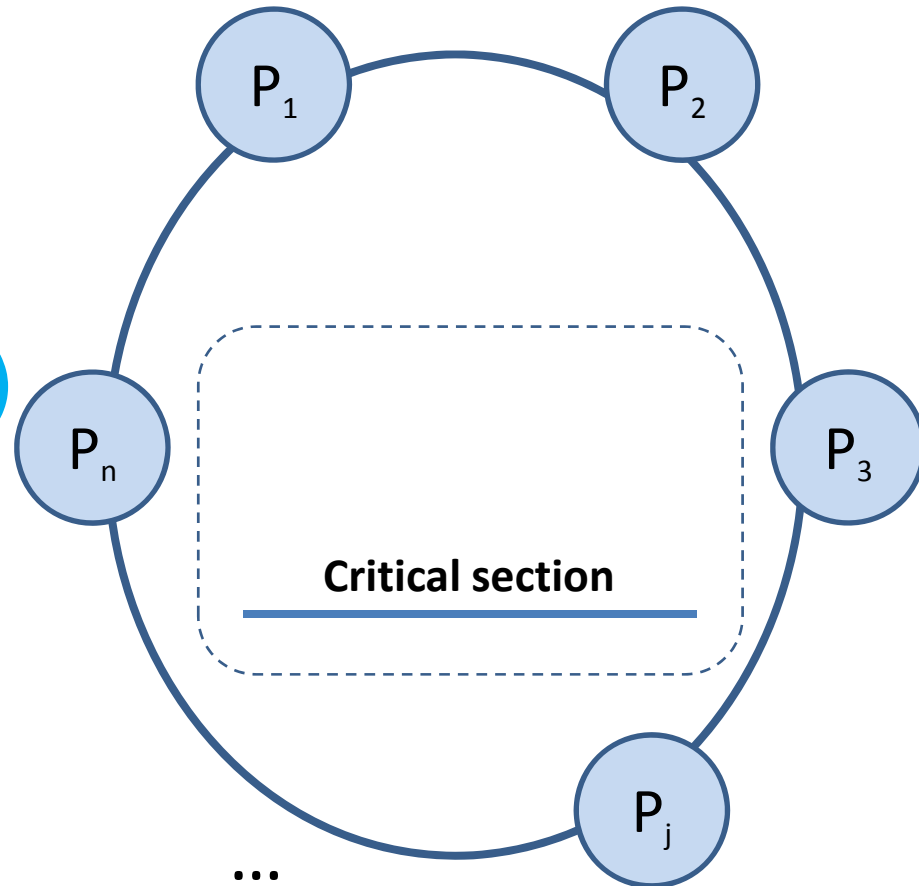
Enter_CS()



Ring-based algorithm

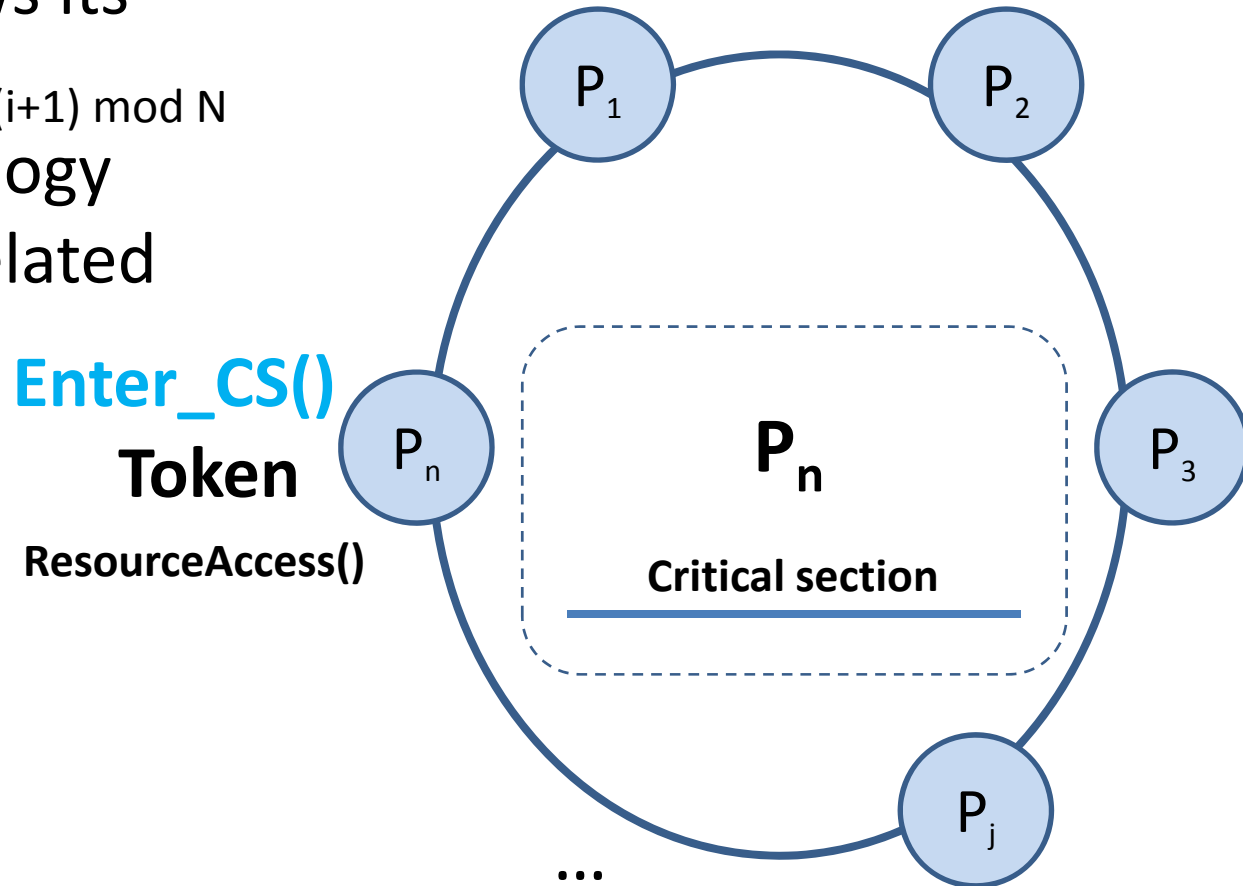
- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology

Enter_CS()
Token



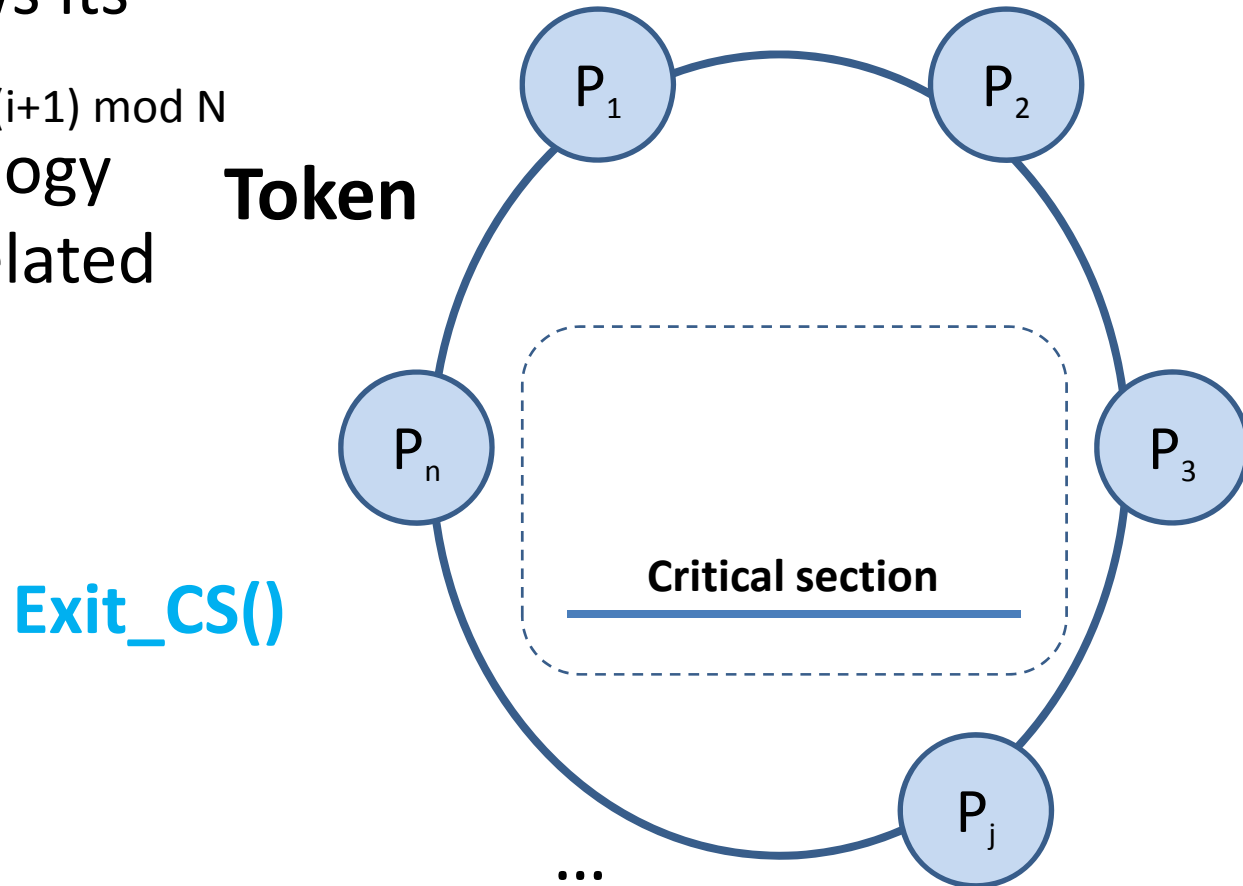
Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology



Ring-based algorithm

- Logical ring of processes
- Each P_i knows its successor, $P_{(i+1) \bmod N}$
- Logical topology a priori unrelated to physical topology



Ring-based algorithm analysis

- **Safe:** Node enters CS only if it holds token
- **Live:** Since finite work is done by each node (can't re-enter), token eventually gets to each node
- **Fair:** Ordering is based on ring topology, no starvation (pass token between accesses)
- **Performance**
 - **Constantly consumes network bandwidth**, even when no processes seek entry, except when inside CS
 - **Synchronization delay:** Between 1 and N messages
 - **Client delay:** 0 to N messages for entry; 0 for exit

Potential problems with ring-based algorithm

(Due to our assumption, do not apply here.)

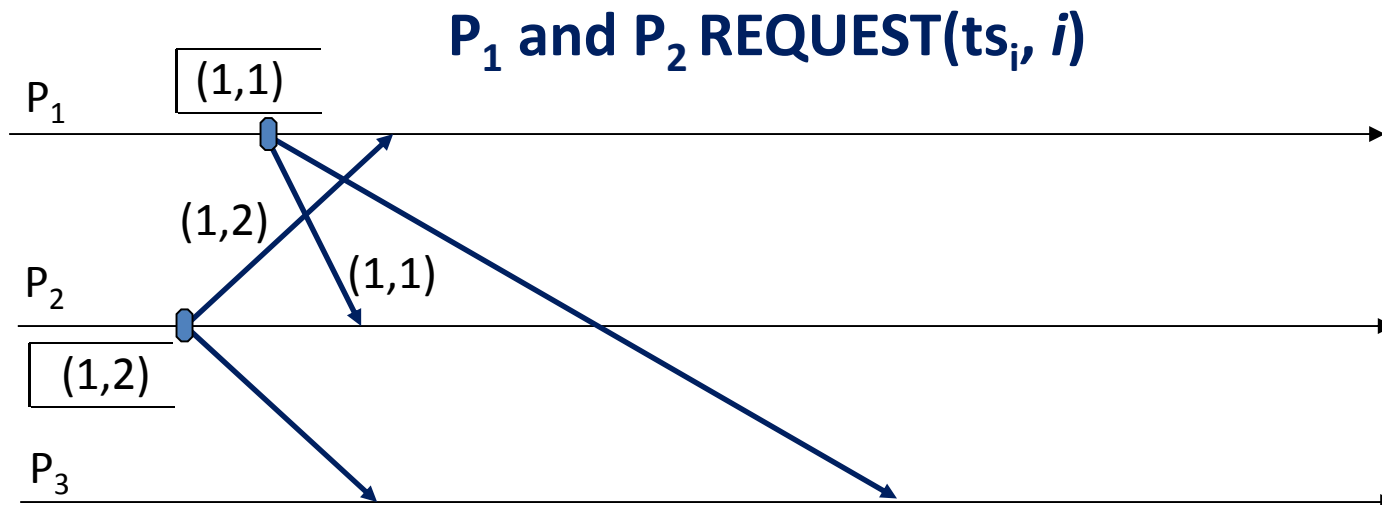
- Node crash
- Lost token
- Duplicate token
- Timeouts on token passing

Lamport's algorithm

- System of n processes
- (ts_i, i) – logical clock timestamp of process P_i
- Non-token based approach uses logical timestamps to order requests for CS
- Smaller timestamps have priority over larger ones
- Request queue maintained at each process ordered by timestamp (**a priority queue!**)
- Assume message delivered in FIFO order

Lamport's algorithm

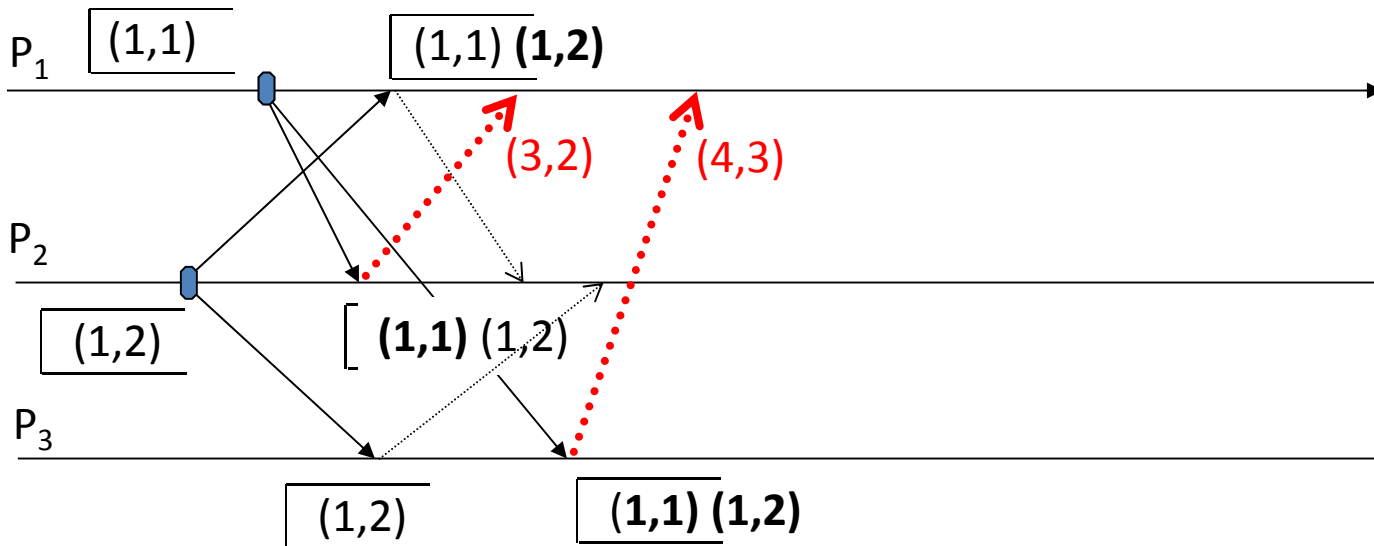
- P_i requesting CS
 - Broadcast **REQUEST**(ts_i, i) message to all processes
 - Place request in *request_queue_i*
 - (ts_i, i) denotes timestamp of request



Lamport's algorithm

- P_k receiving a request to enter CS
 - When P_k receives a REQUEST(ts_i, i) message from P_i ,
 - It places P_i 's request into ***request_queue_k***
 - Sends a timestamped REPLY message to P_i

REPLY(ts_k, k) – only replies to P_i 's request are shown

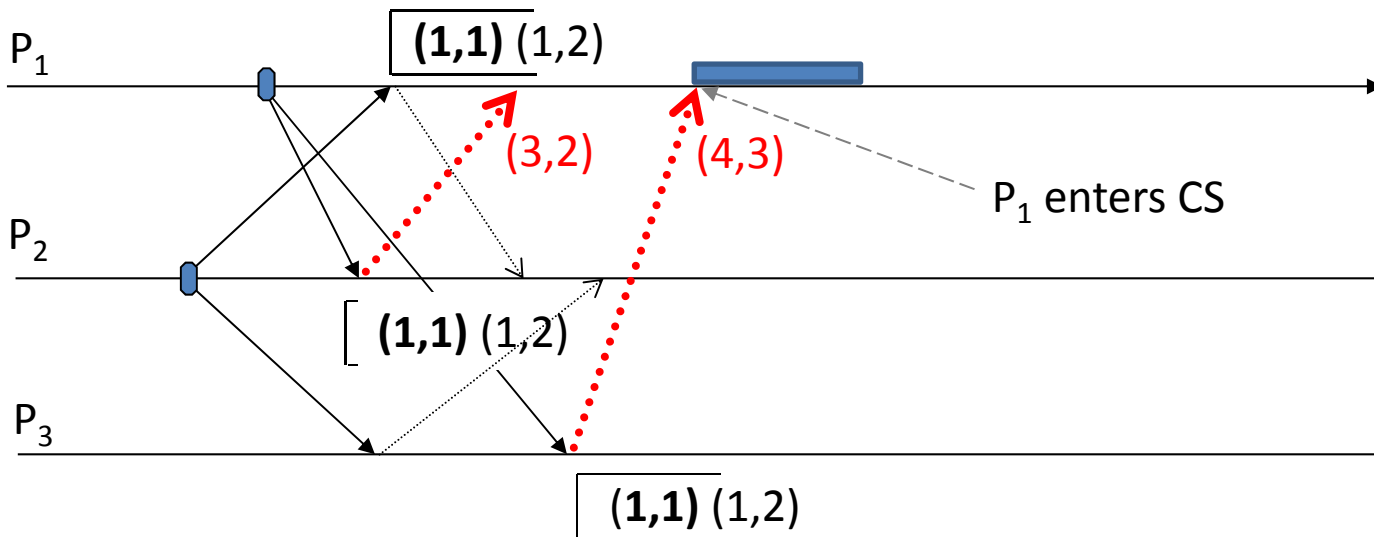


Lamport's algorithm

P_i requesting CS – P_i enters when following conditions hold

1. P_i has received a message with timestamp larger than (ts_i, i) from every other processes
2. P_i 's request is at top of *request_queue* _{i}

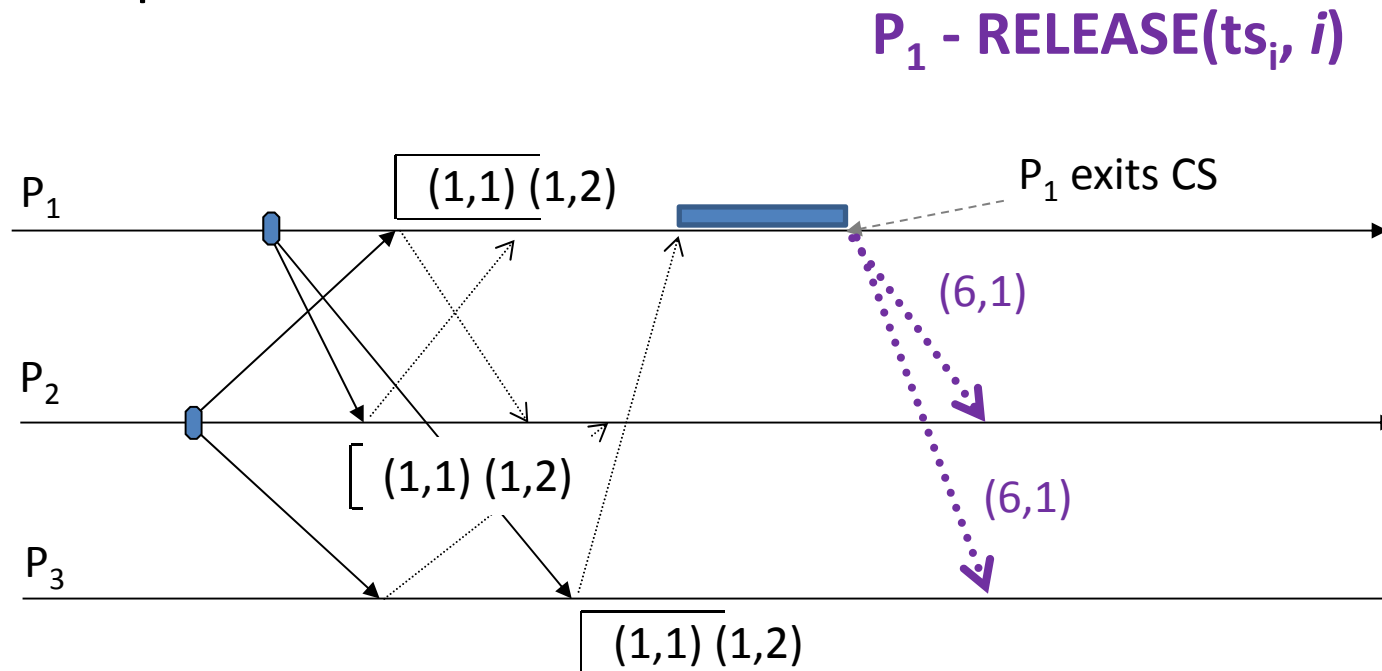
P_1 sent REQ(1, 1) and received REPLY(3,2) and REPLY(4,3)



Lamport's algorithm

P_i releasing CS

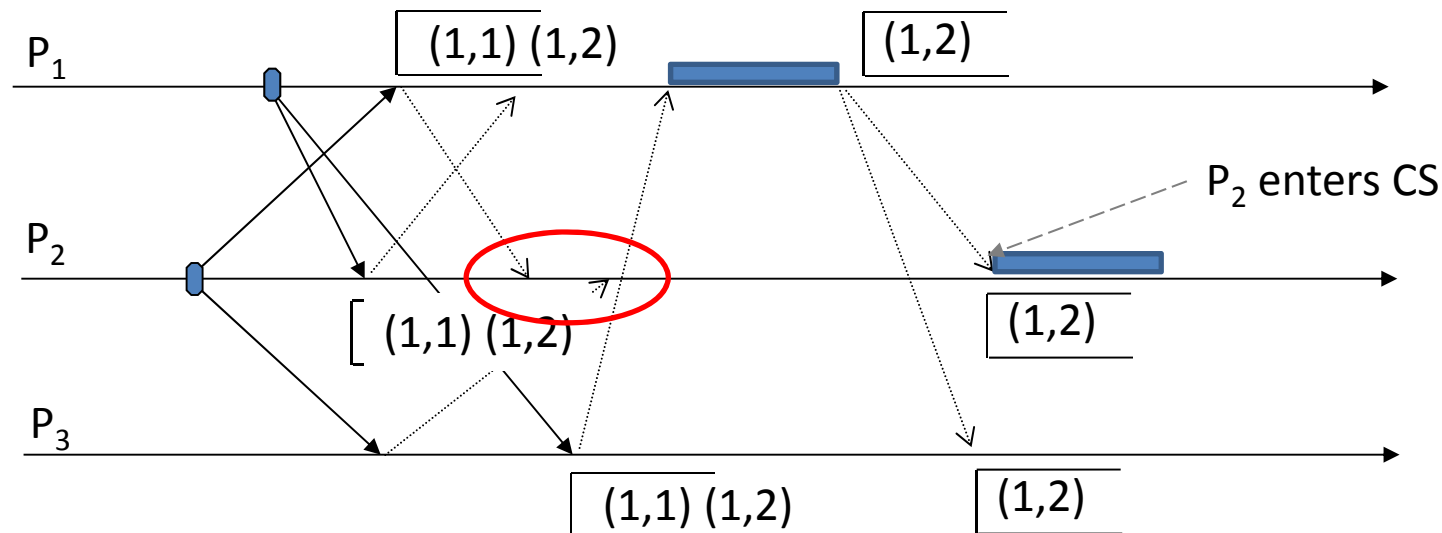
- Removes request from top of the request queue
- Broadcasts a timestamped **RELEASE** message to all processes



Lamport's algorithm

P_k receiving a release message

- When P_k receives a **RELEASE** message from P_i , P_k removes P_i 's request from its request queue



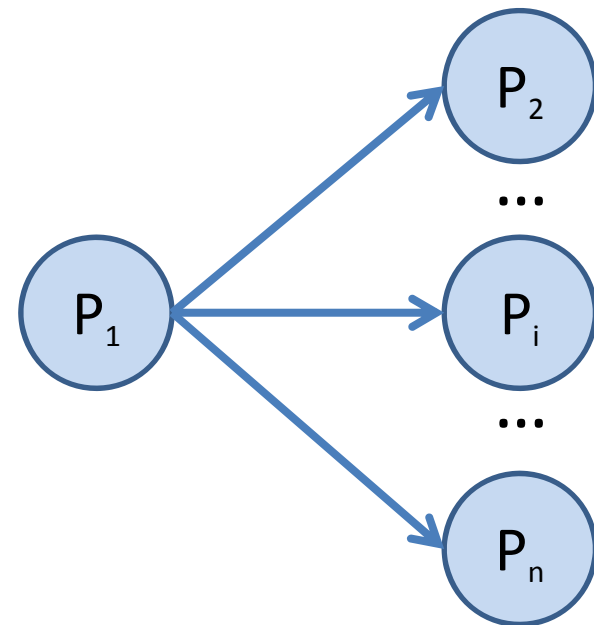
Performance

- $3(N-1)$ messages per CS invocation
 - $(N - 1)$ REQUEST
 - $(N - 1)$ REPLY
 - $(N - 1)$ RELEASE messages

Ricart & Agrawala, 1981, algorithm

(Guarantees mutual exclusion among n processes)

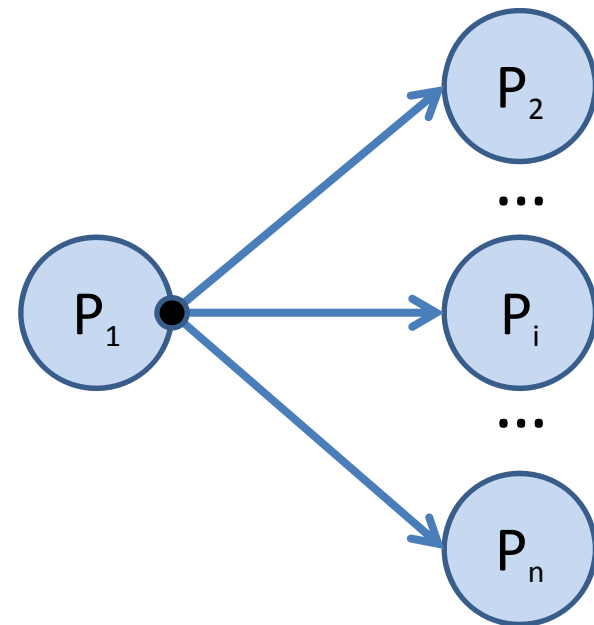
- **Basic idea**
 - Processes wanting to enter CS, **broadcast a request to all processes**
 - Enter CS, once **all** have **granted request**
- Use **Lamport timestamps** to order requests: $\langle T, P_i \rangle$, T the timestamp, P_i the process identifier



Ricart & Agrawala, 1981, algorithm

(Guarantees mutual exclusion among n processes)

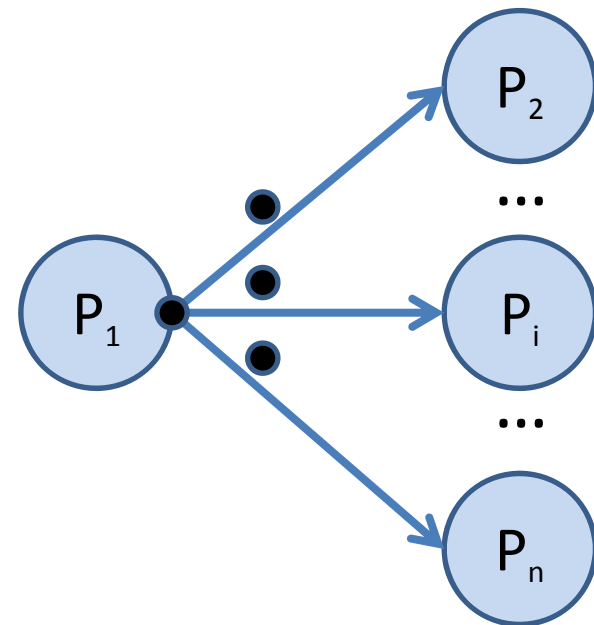
- **Basic idea**
 - Processes wanting to enter CS, **broadcast a request to all processes**
 - Enter CS, once **all** have **granted request**
- Use **Lamport timestamps** to order requests: $\langle T, P_i \rangle$, T the timestamp, P_i the process identifier



Ricart & Agrawala, 1981, algorithm

(Guarantees mutual exclusion among n processes)

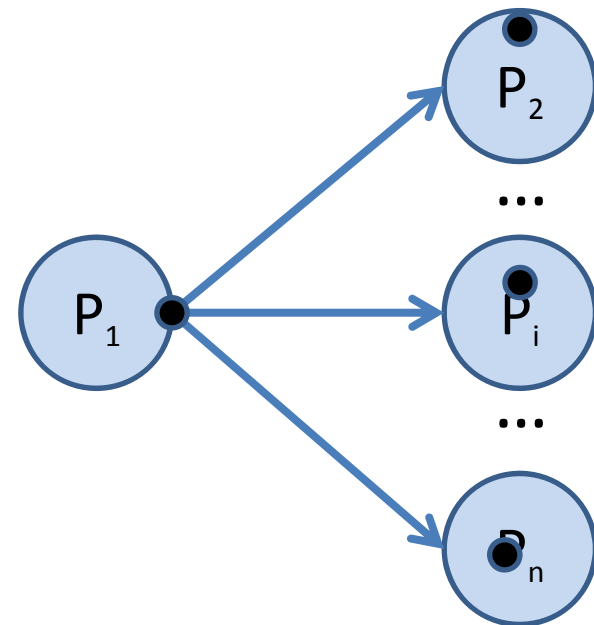
- **Basic idea**
 - Processes wanting to enter CS, **broadcast a request to all processes**
 - Enter CS, once **all** have **granted request**
- Use **Lamport timestamps** to order requests: $\langle T, P_i \rangle$, T the timestamp, P_i the process identifier



Ricart & Agrawala, 1981, algorithm

(Guarantees mutual exclusion among n processes)

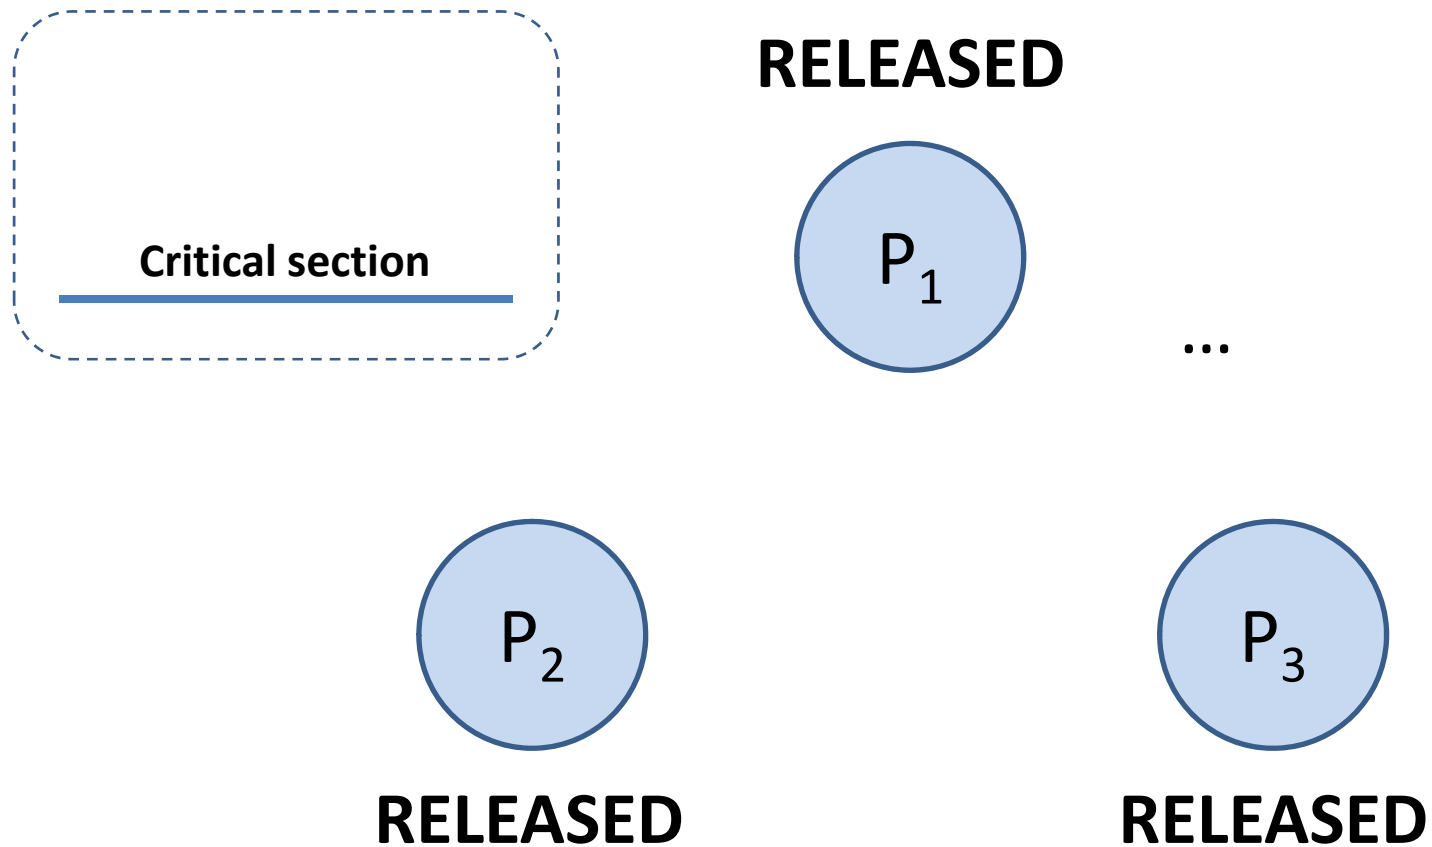
- **Basic idea**
 - Processes wanting to enter CS, **broadcast a request to all processes**
 - Enter CS, once **all** have **granted request**
- Use **Lamport timestamps** to order requests: $\langle T, P_i \rangle$, T the timestamp, P_i the process identifier



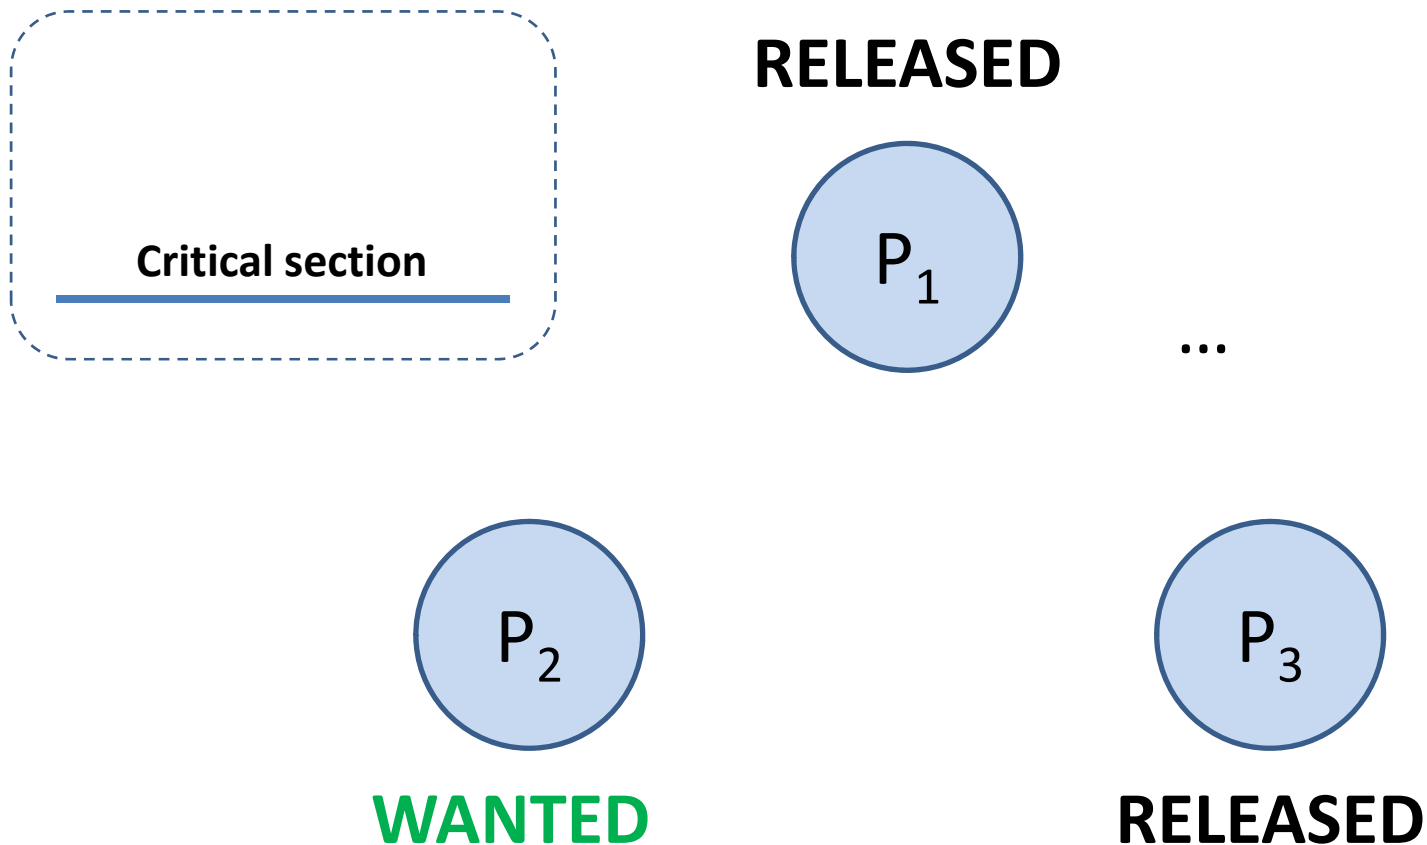
Ricart & Agrawala: Distributed strategy

- Each process is in one of three states
 - **Released** - Outside the CS, Exit_CS()
 - **Wanted** - Wanting to enter CS, Enter_CS()
 - **Held** - Inside CS, RessourceAccess()
- If a process requests to enter CS and **all** other processes are in the ***Released state***, entry is **granted** by each process
- If a process, P_i , requests to enter CS and another process, P_k , is inside the CS (***Held state***), then P_k will not reply, until it is finished with the CS

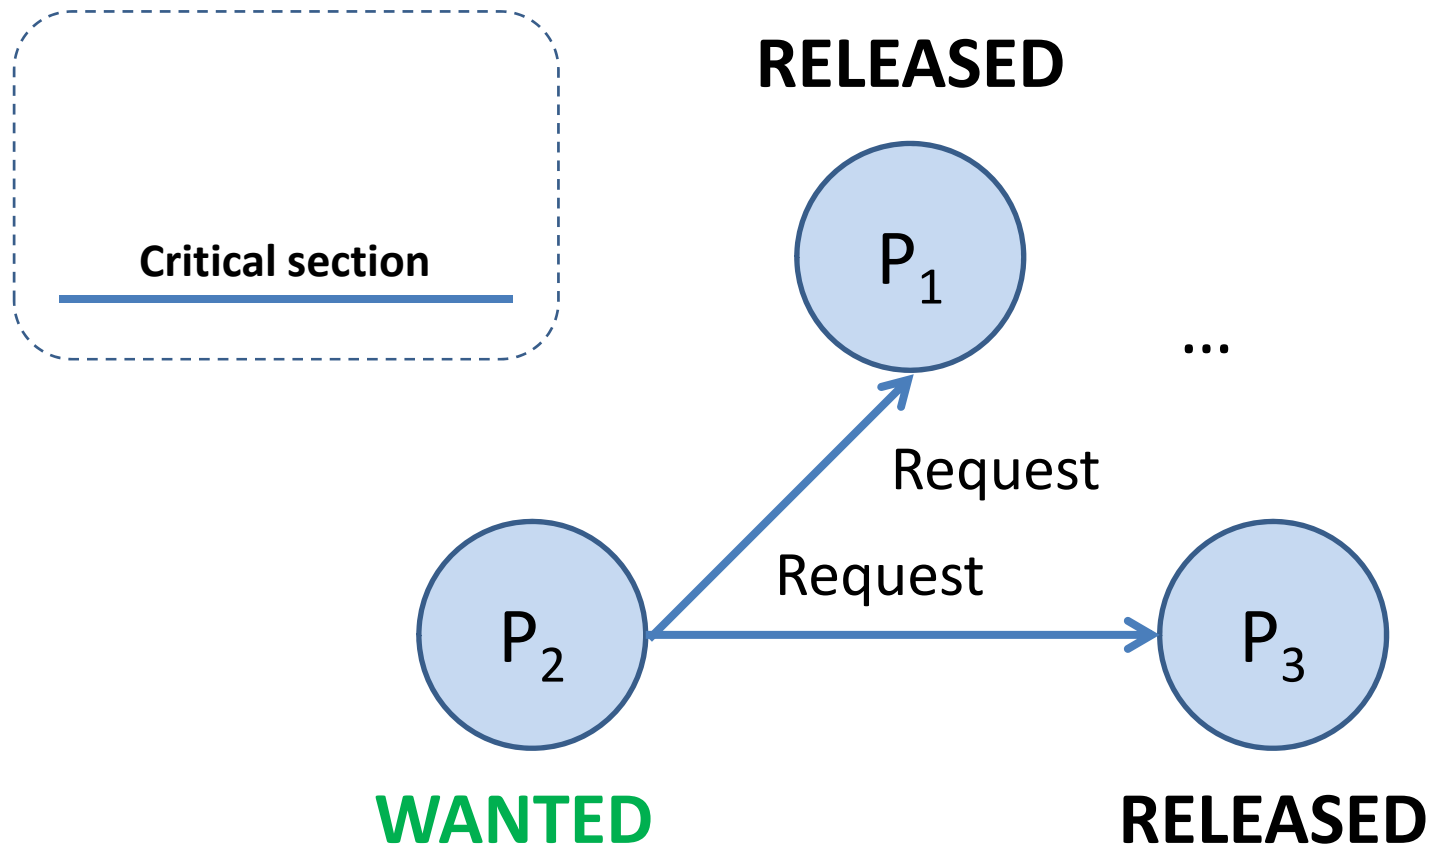
Initialization



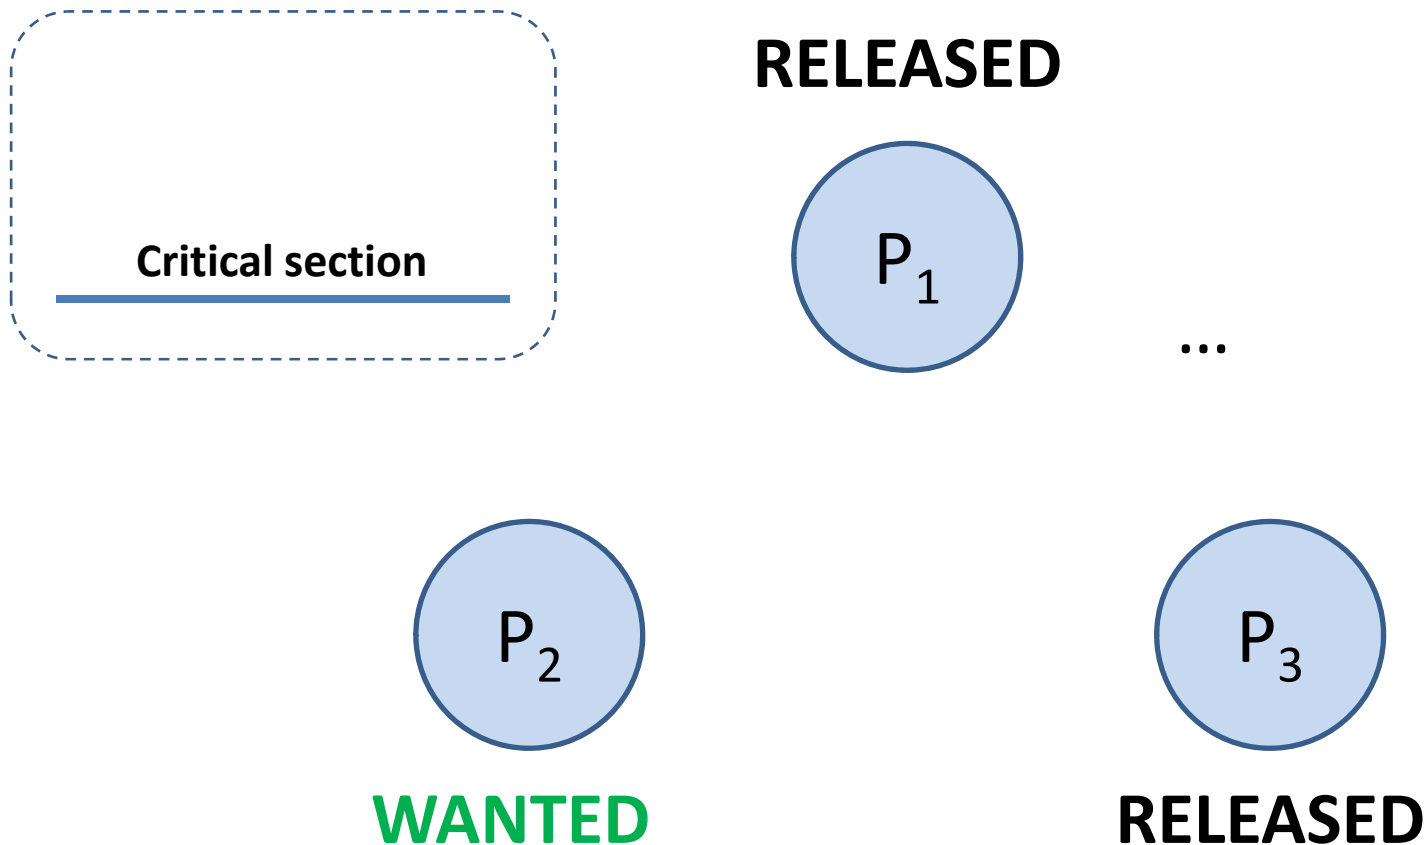
Requesting entry to CS: *All Released*



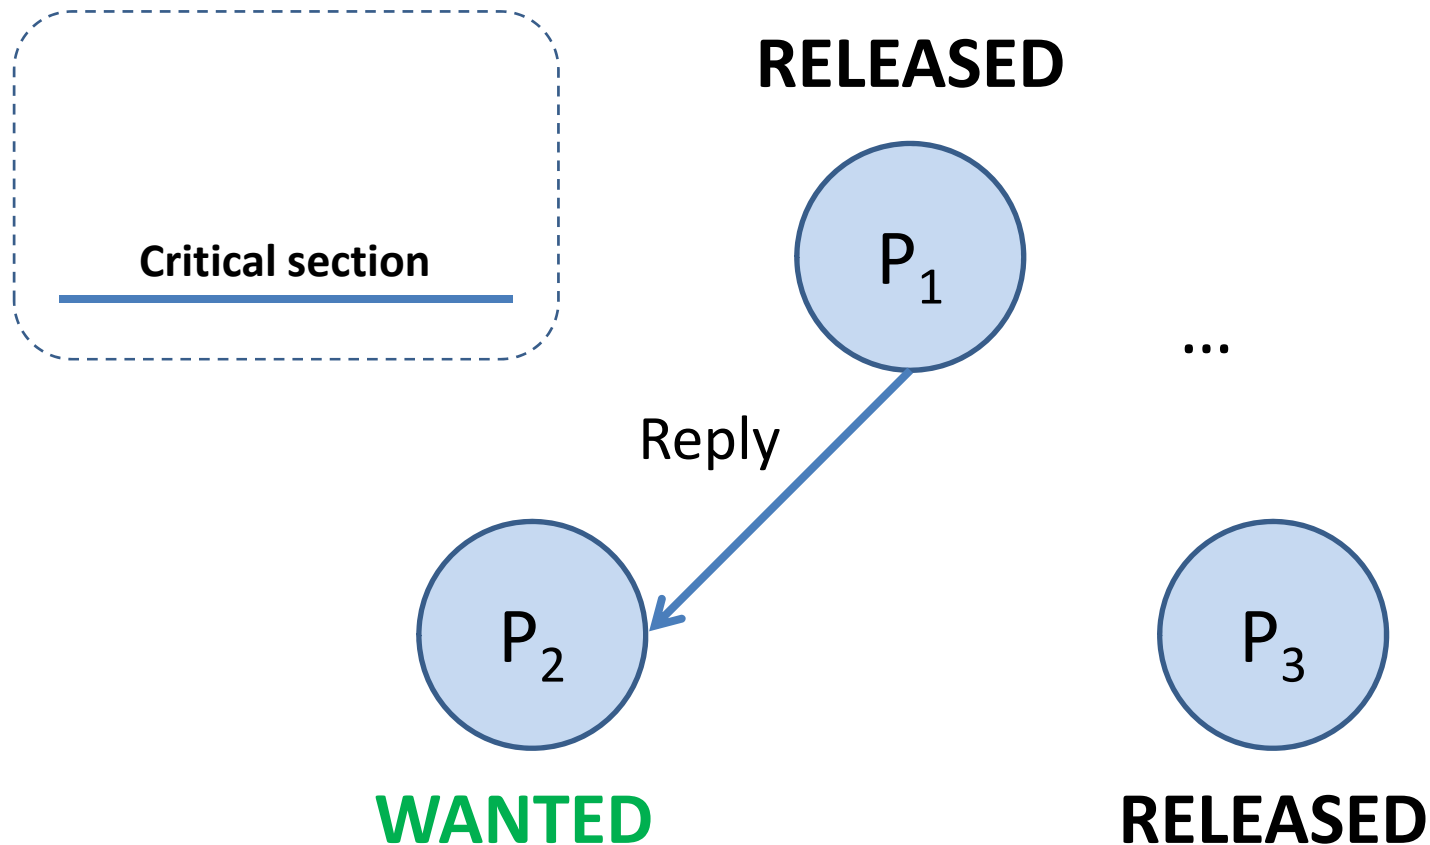
Requesting entry to CS: *All Released*



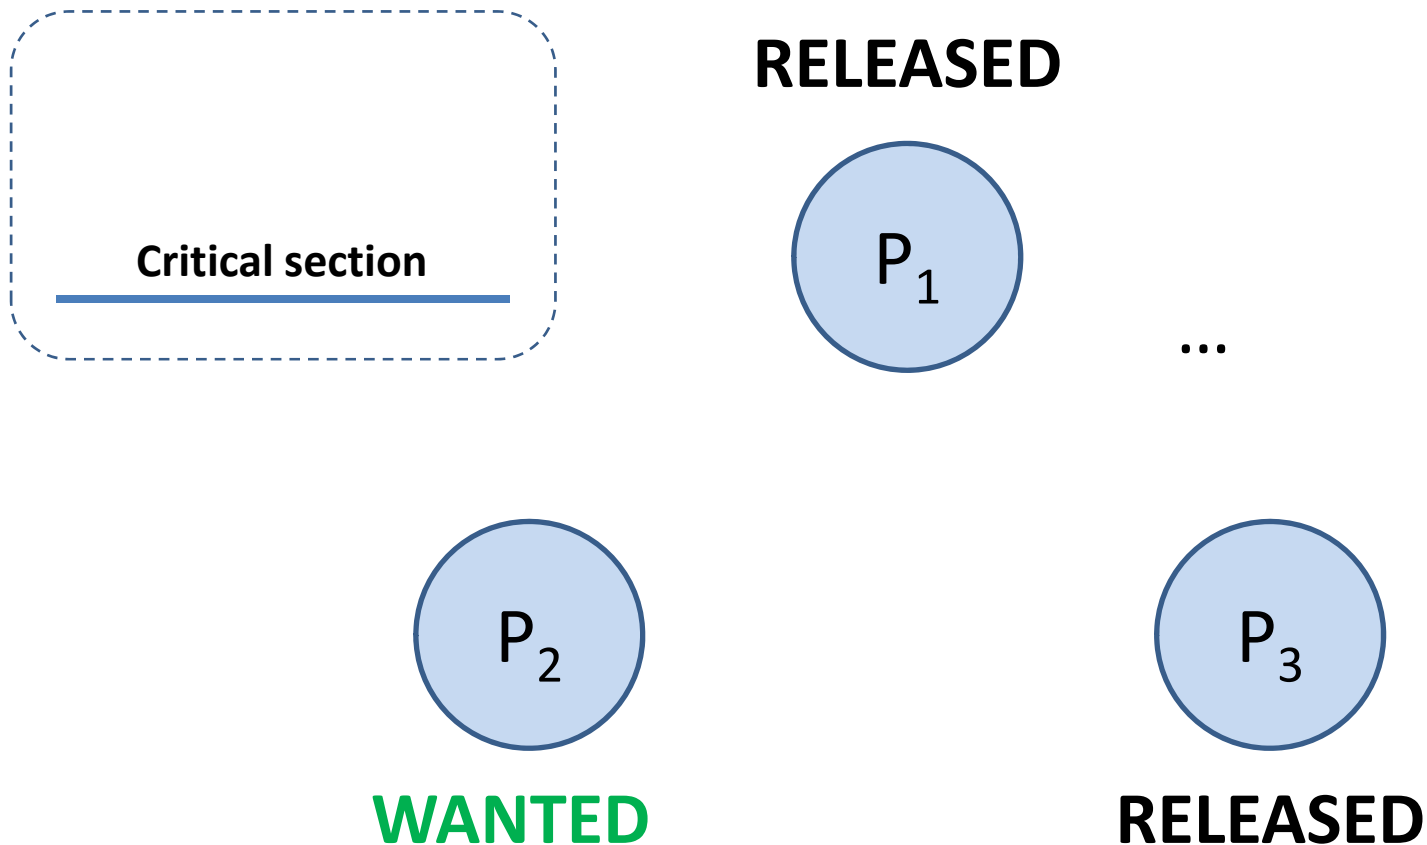
Requesting entry to CS: *All Released*



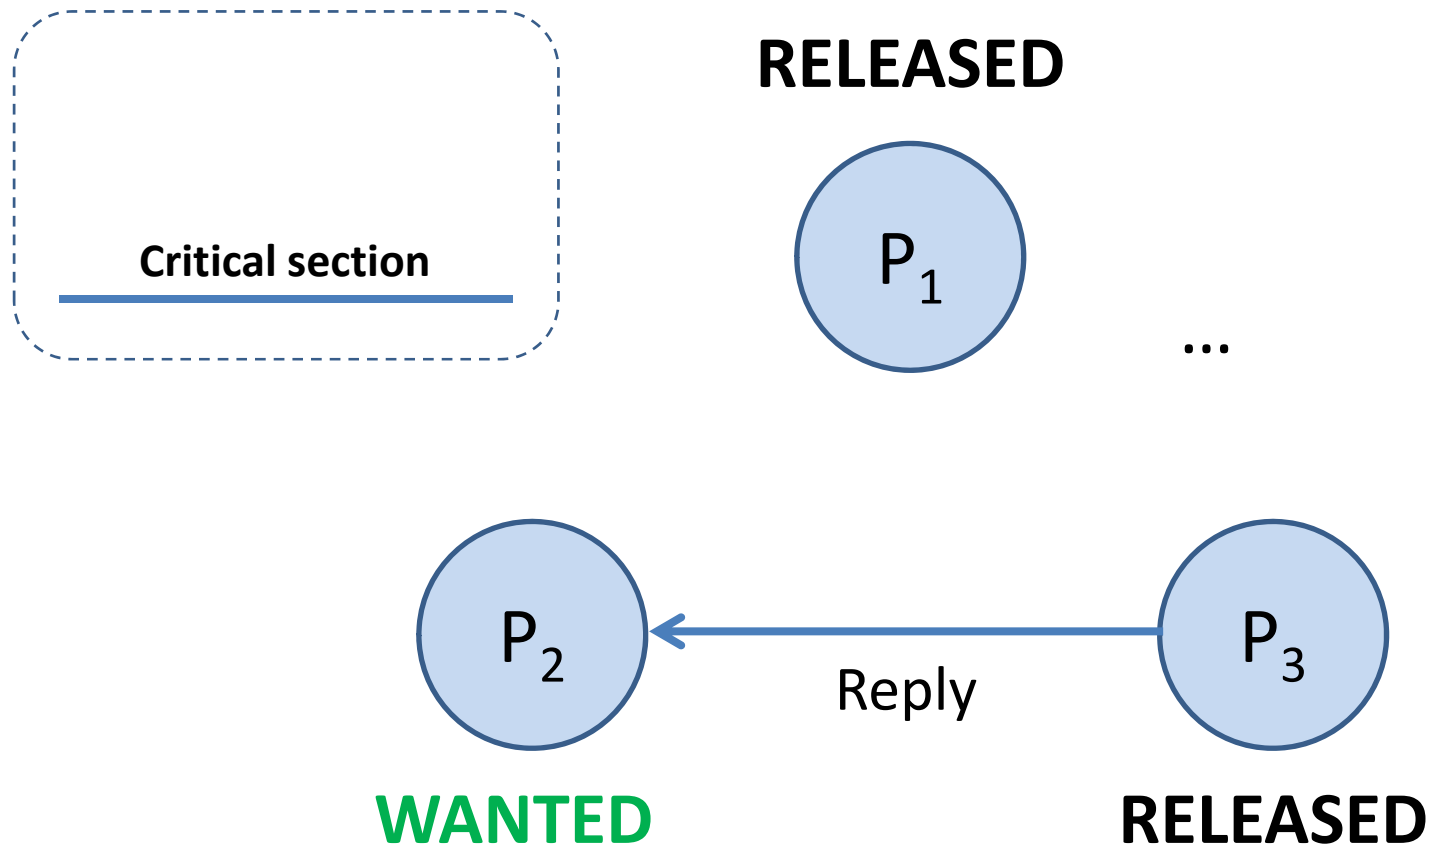
Requesting entry to CS: *All Released*



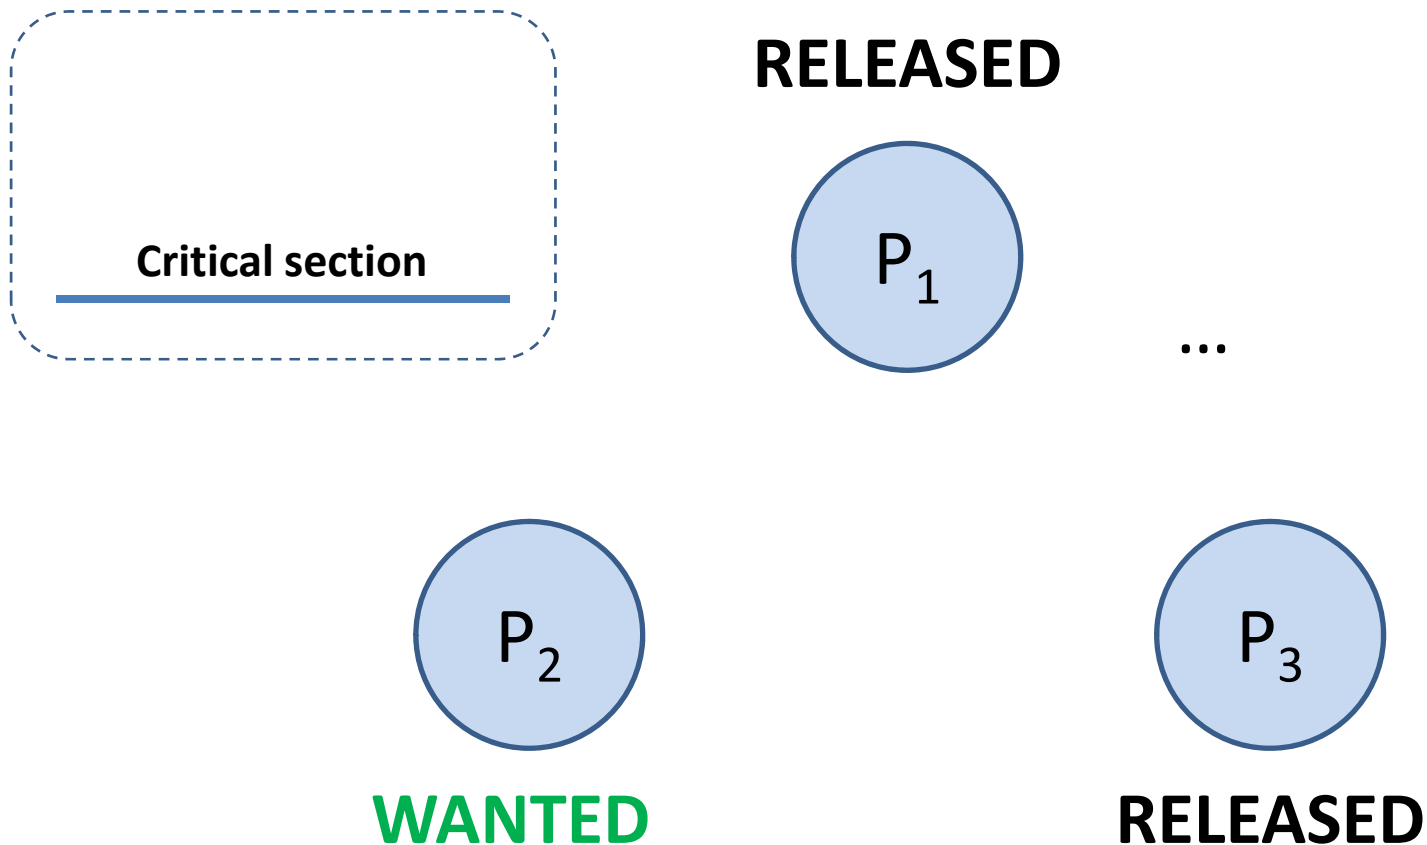
Requesting entry to CS: *All Released*



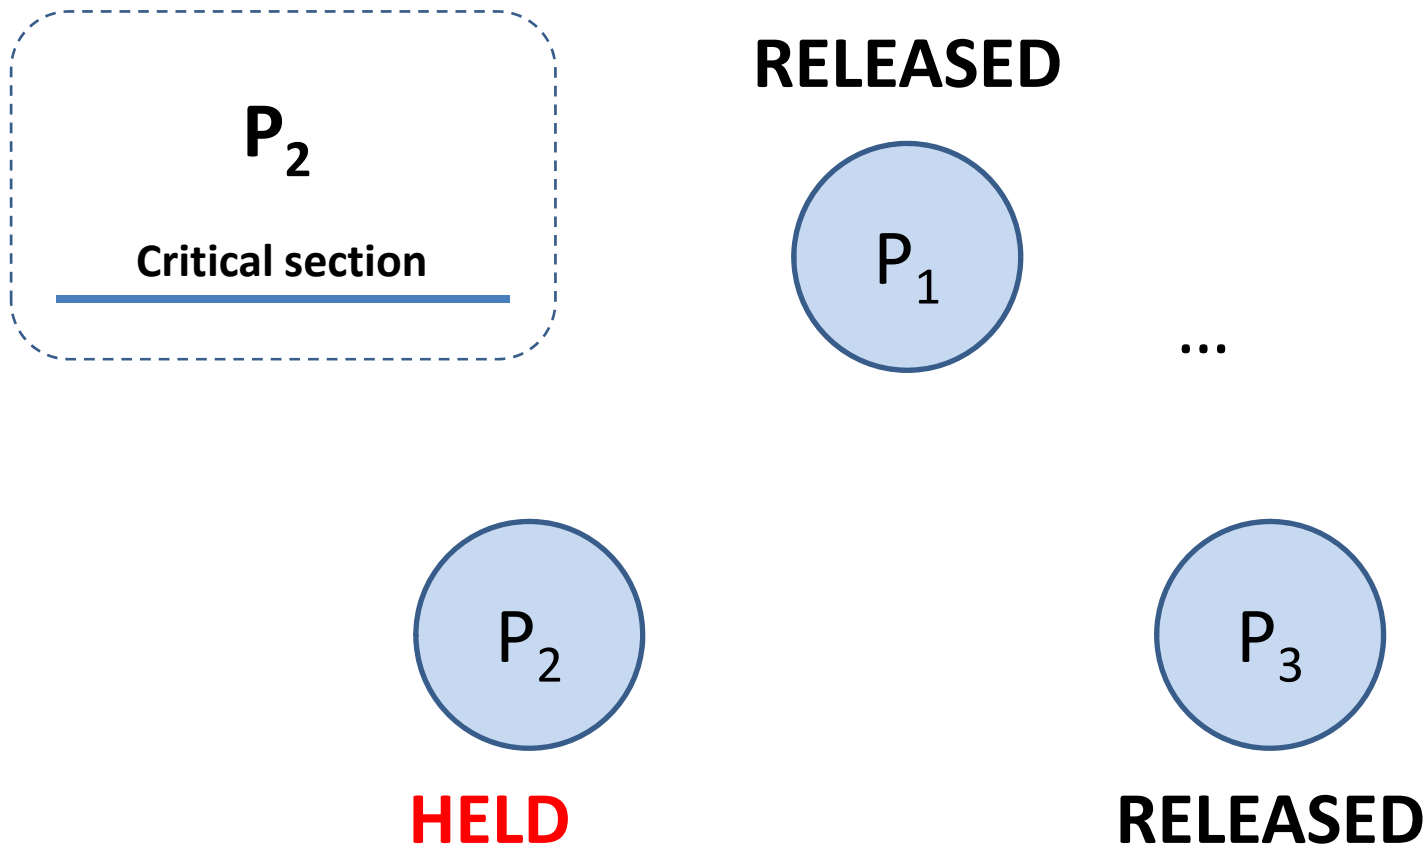
Requesting entry to CS: *All Released*



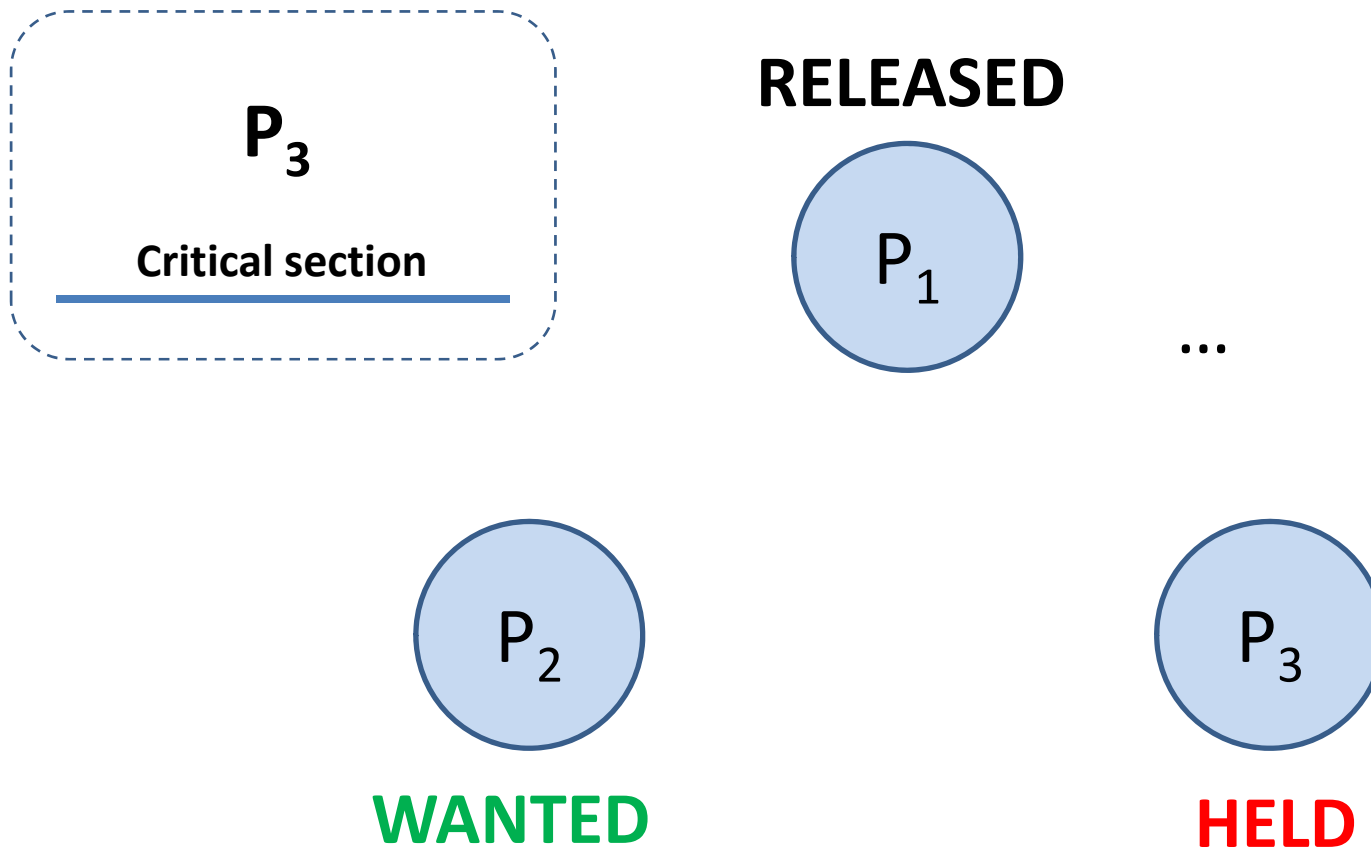
Requesting entry to CS: *All Released*



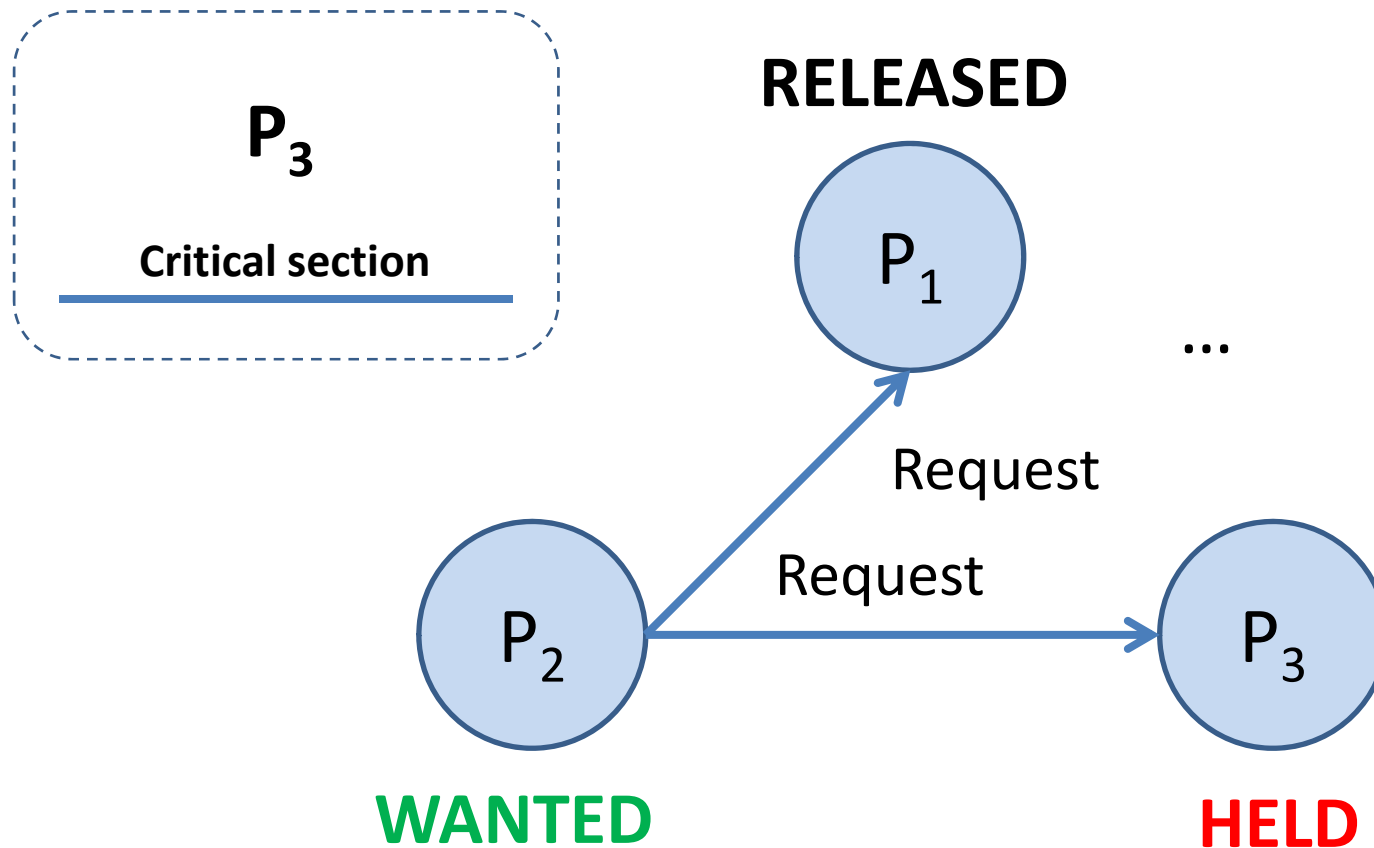
Requesting entry to CS: *All Released*



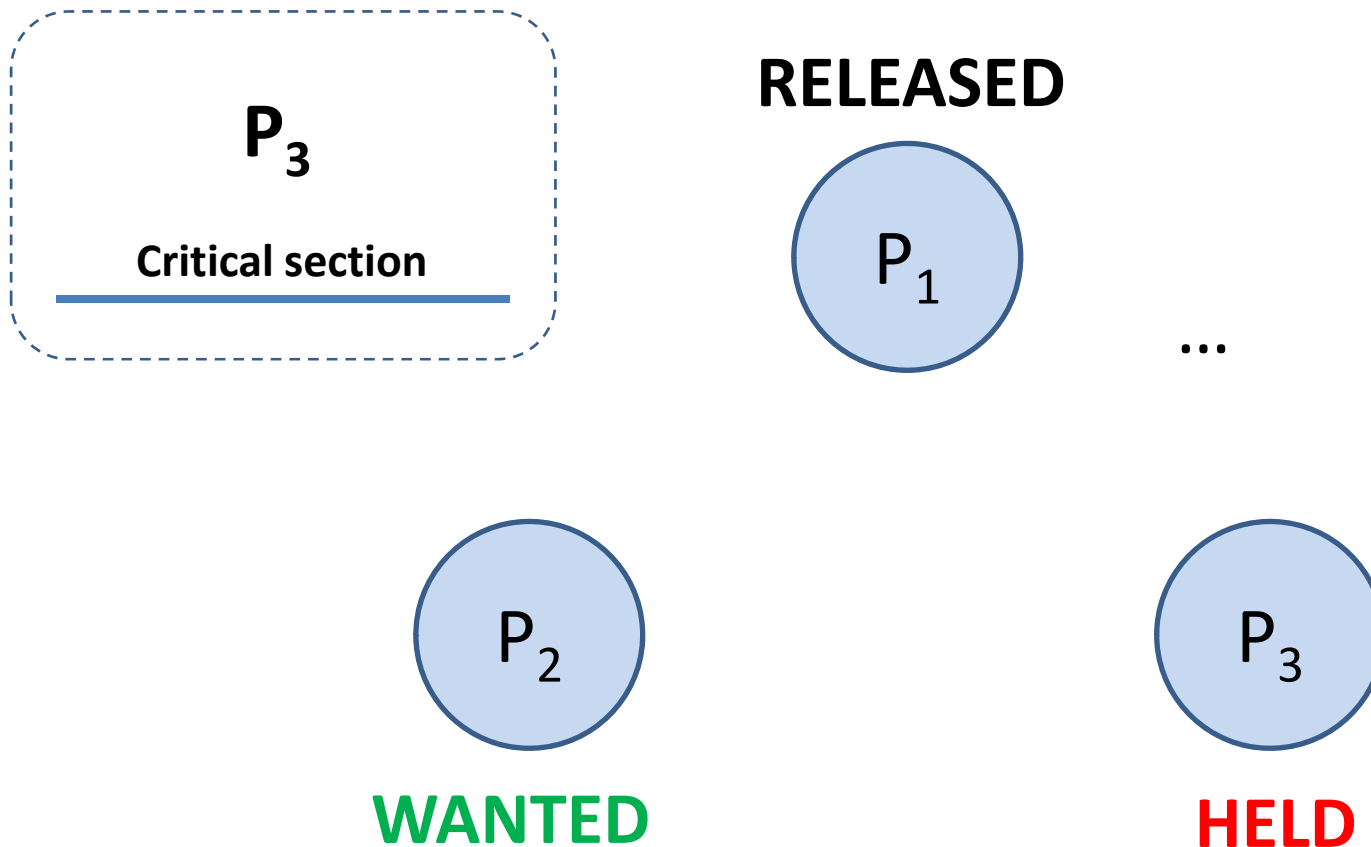
Requesting entry to CS: *Request while Held*



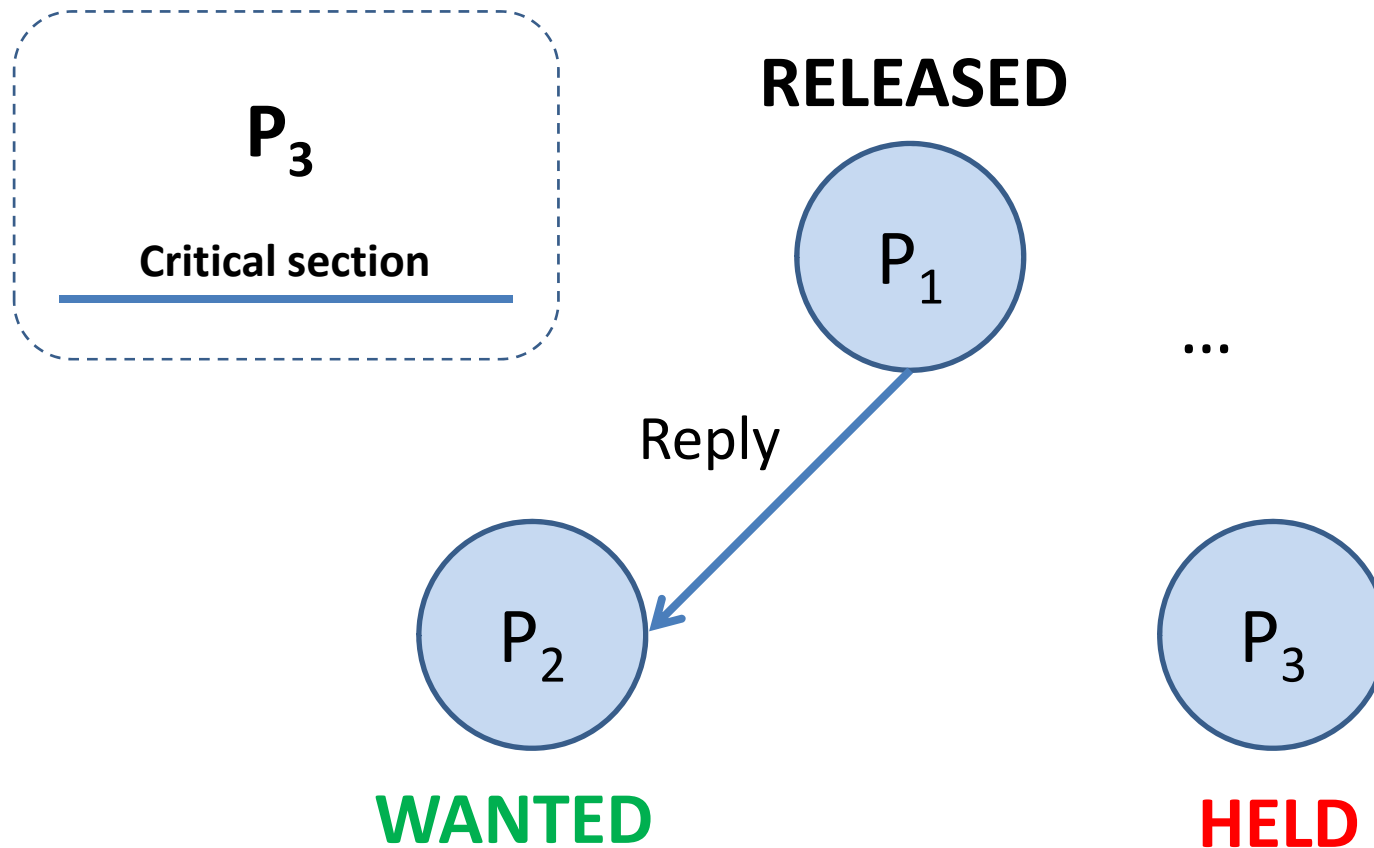
Requesting entry to CS: *Request while Held*



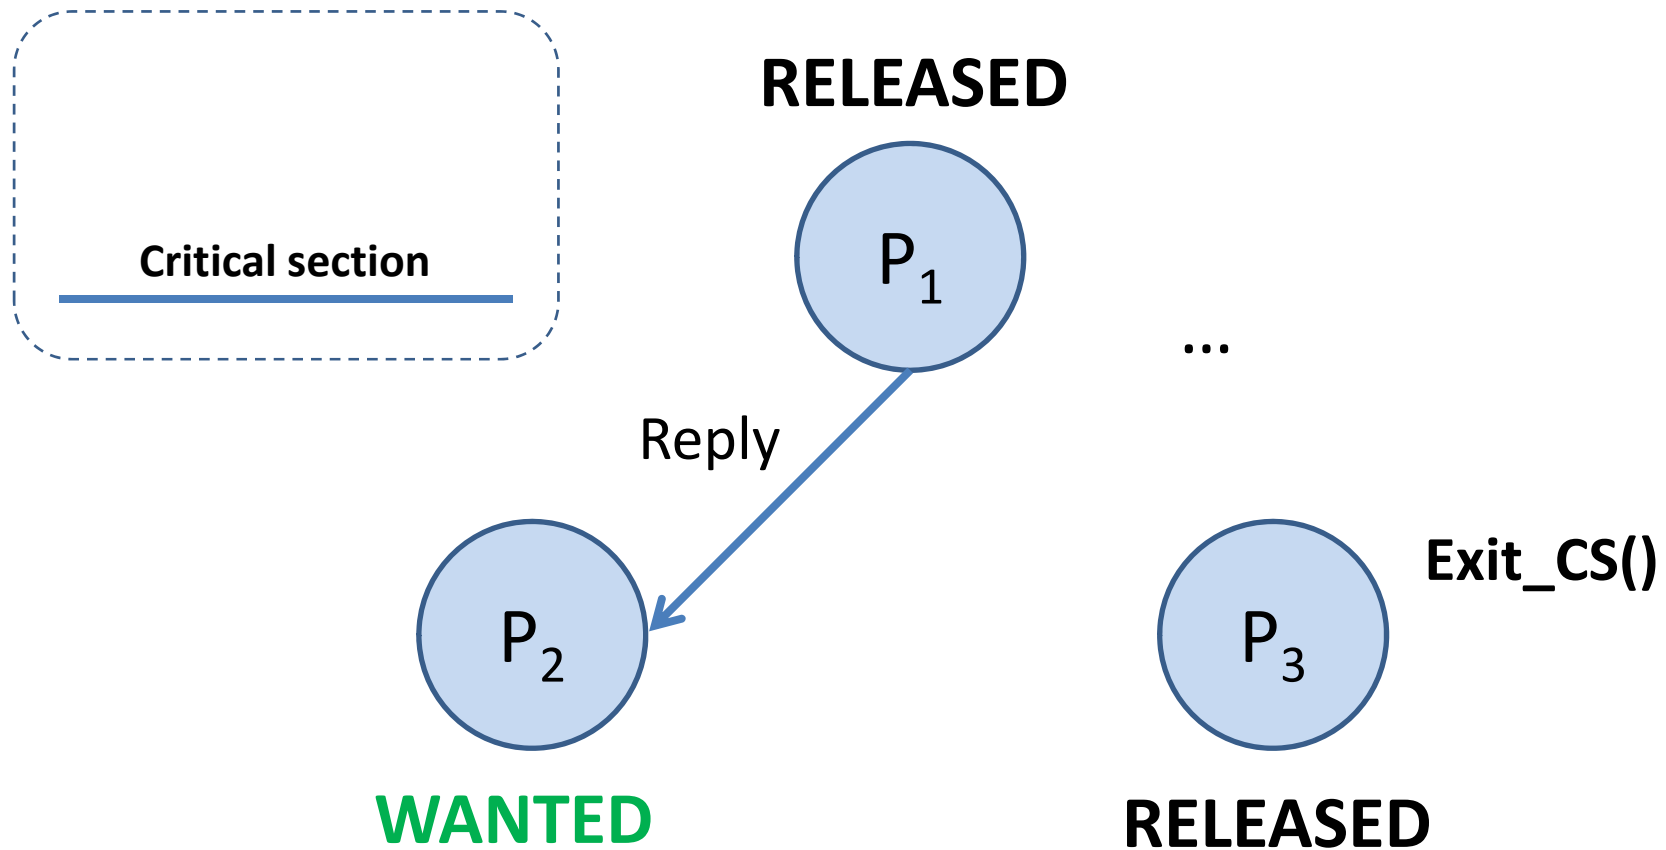
Requesting entry to CS: *Request while Held*



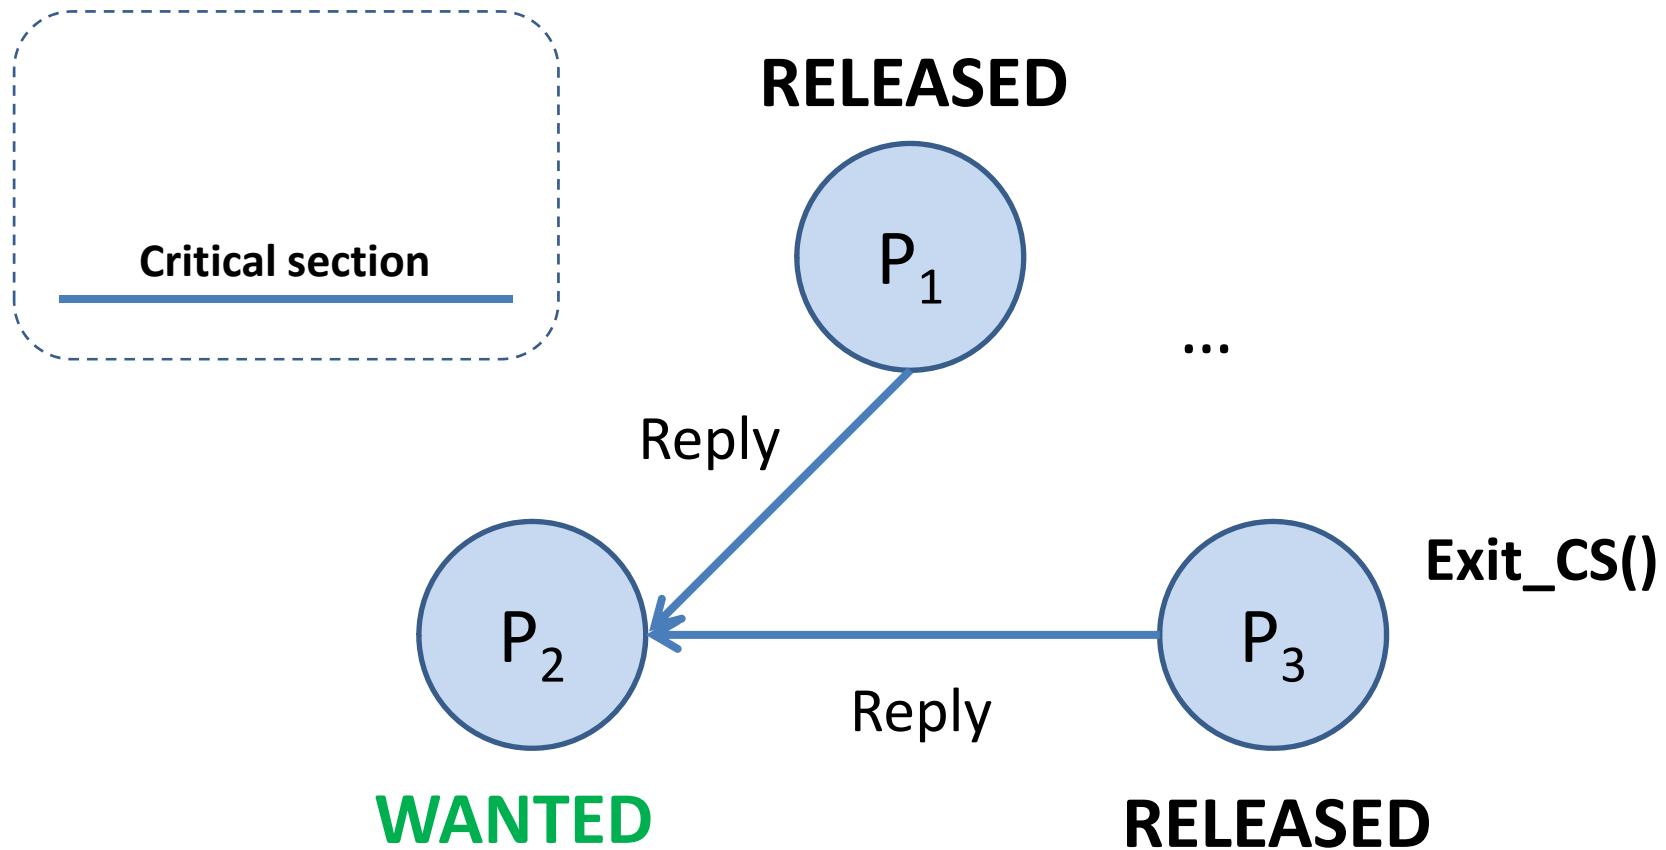
Requesting entry to CS: *Request while Held*



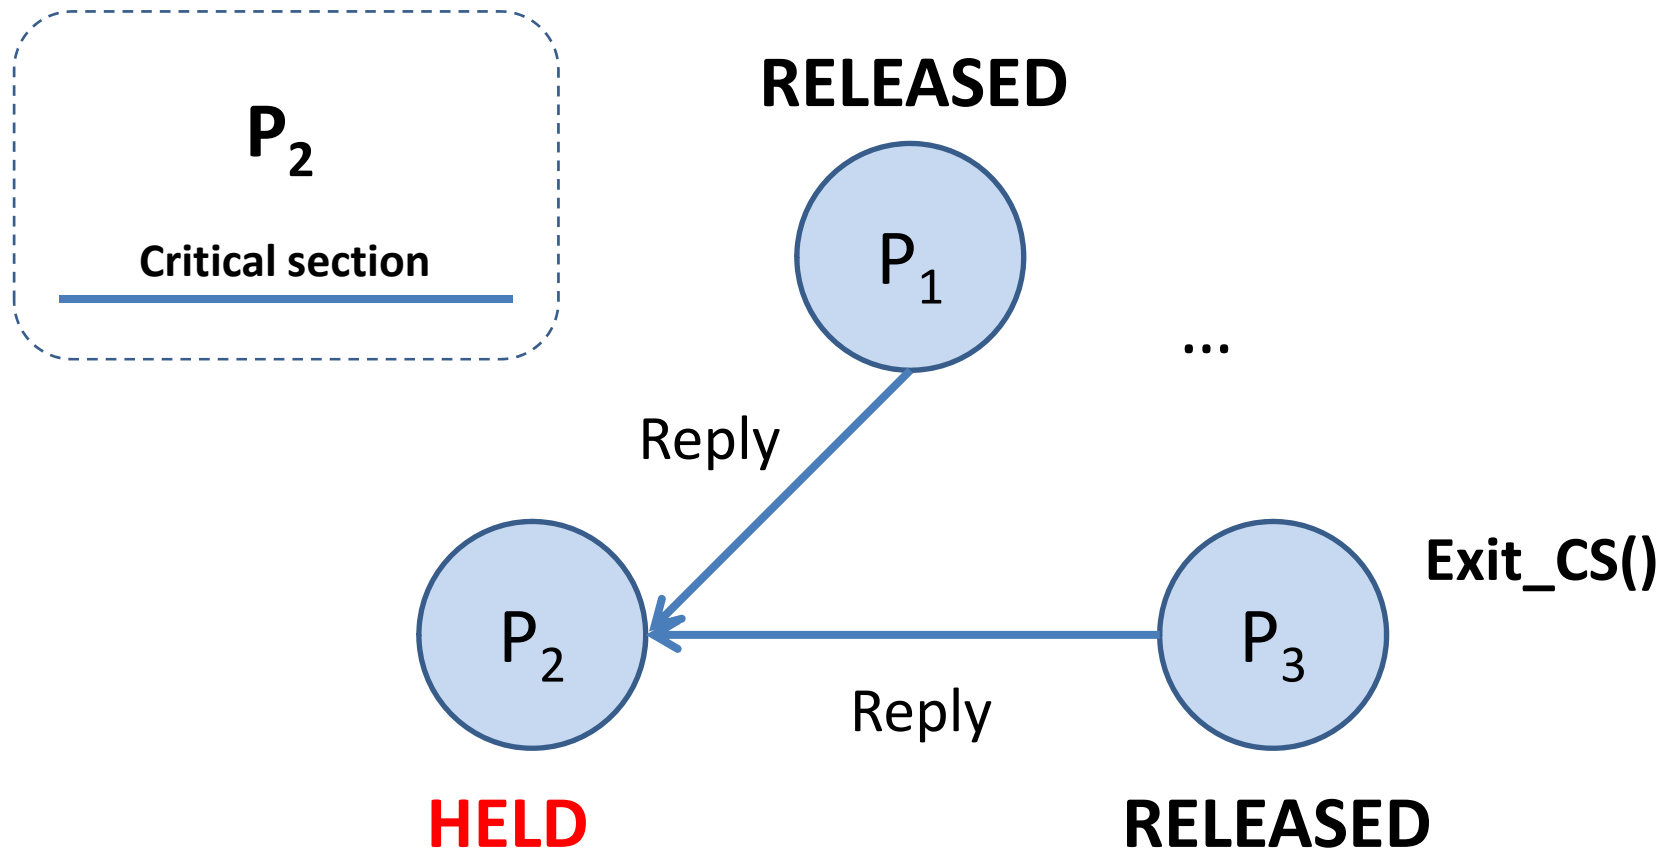
Requesting entry to CS: *Request while Held*



Requesting entry to CS: *Request while Held*



Requesting entry to CS: *Request while Held*

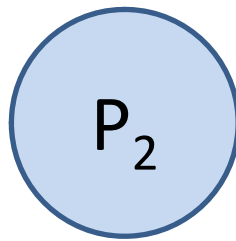
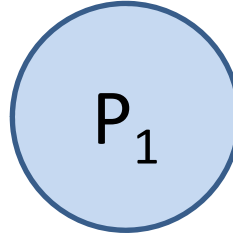


Concurrent entry requests to CS

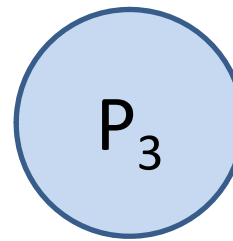
P_2 and P_3 request entry
to CS concurrently

RELEASED

...



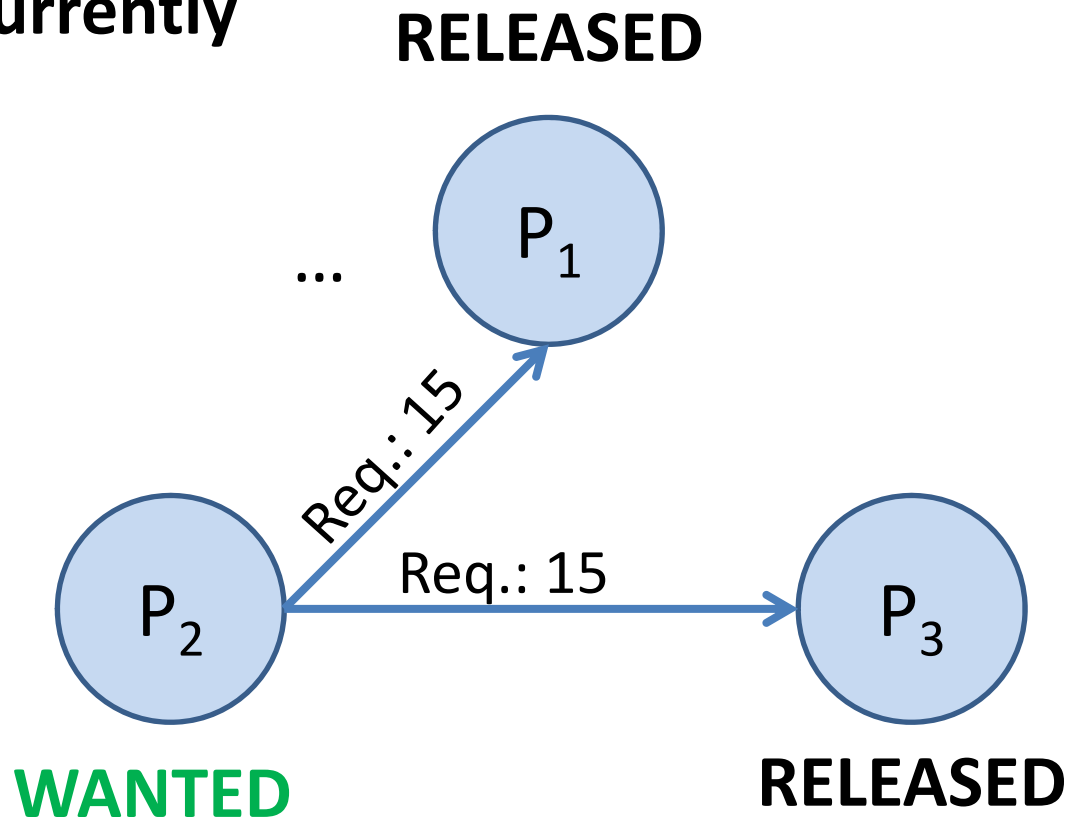
WANTED



RELEASED

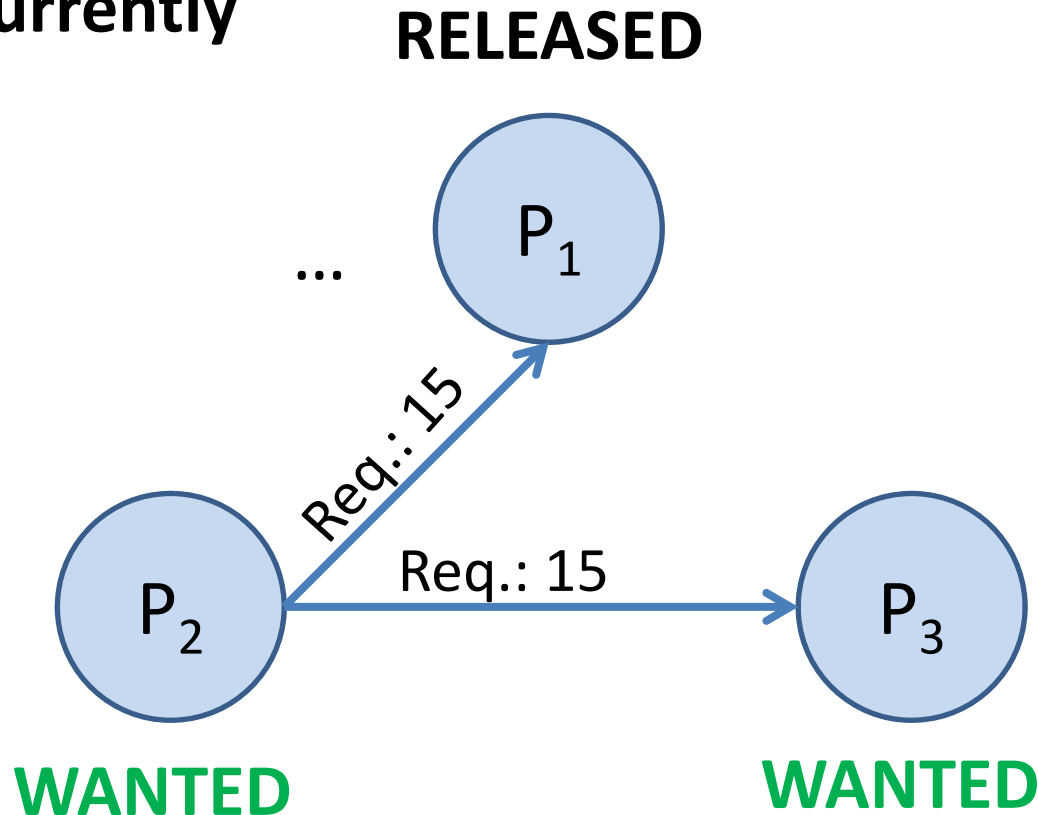
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently



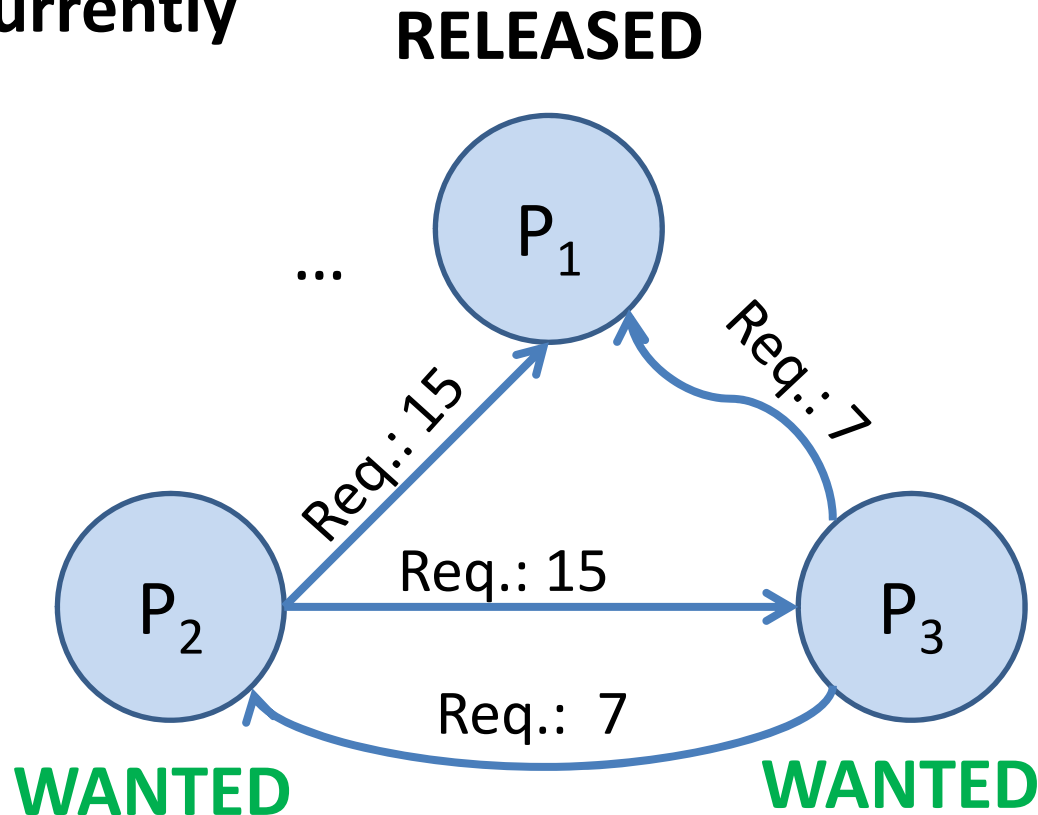
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently



Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

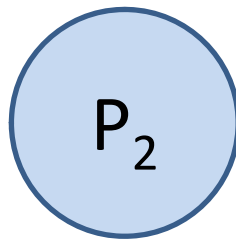
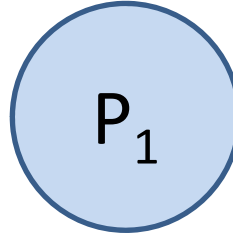


Concurrent entry requests to CS

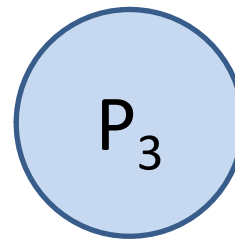
P_2 and P_3 request entry
to CS concurrently

RELEASED

...



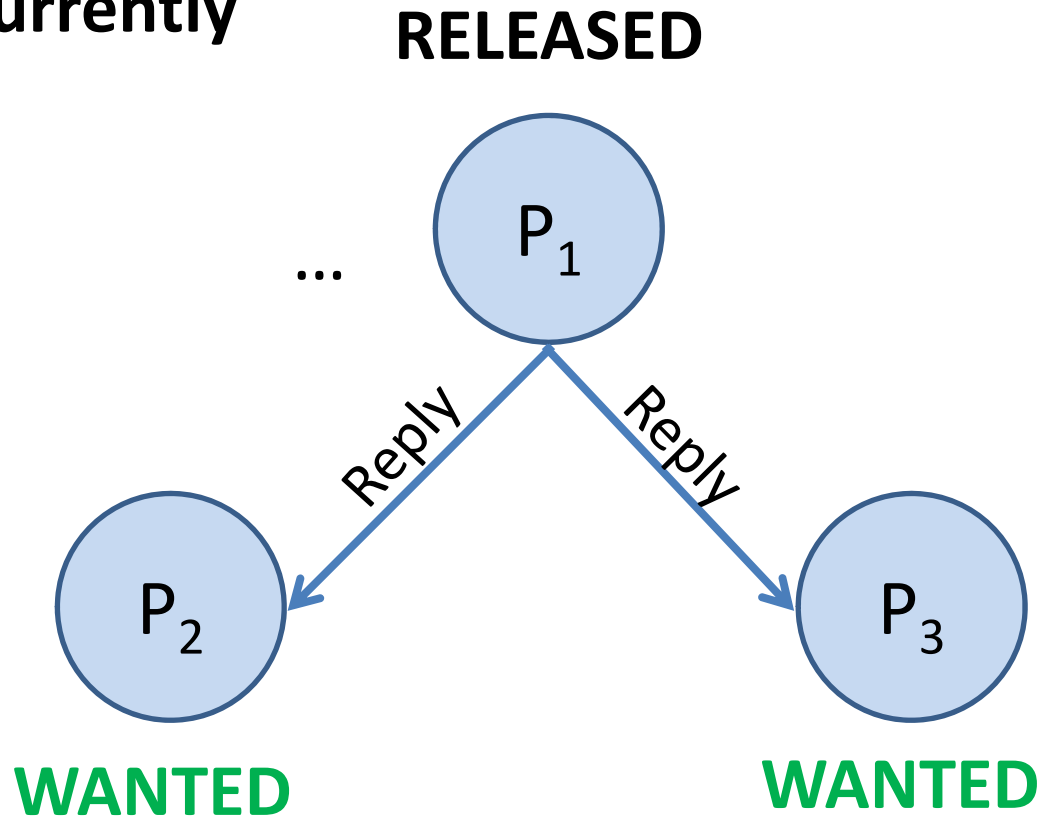
WANTED



WANTED

Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

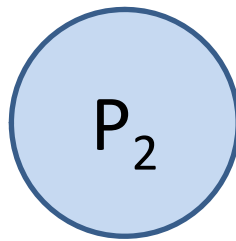
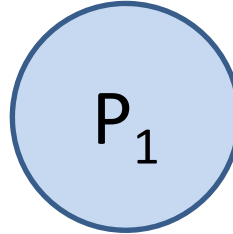


Concurrent entry requests to CS

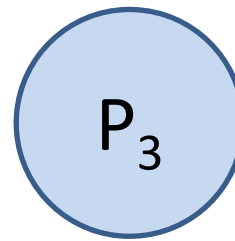
P_2 and P_3 request entry
to CS concurrently

RELEASED

...



WANTED

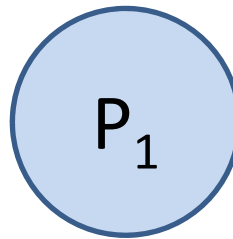


WANTED

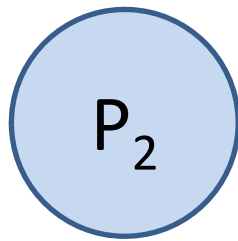
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

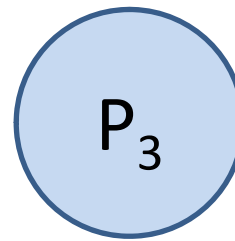
RELEASED



@ P_3 : 7 from $P_3 < 15$ from P_2 ,
Own timestamp lower, therefore,
hold on to reply to P_2



WANTED



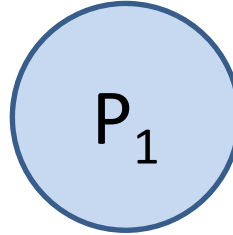
WANTED

Concurrent entry requests to CS

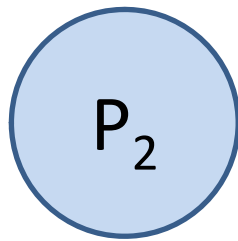
P_2 and P_3 request entry
to CS concurrently

RELEASED

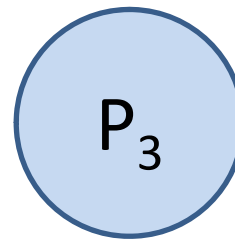
@ P_2 : 7 from $P_3 < 15$ from P_2 ,
therefore grant P_3
access first



@ P_3 : 7 from $P_3 < 15$ from P_2 ,
Own timestamp lower, therefore,
hold on to reply to P_2



WANTED



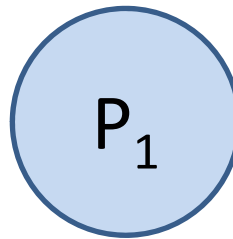
WANTED

Concurrent entry requests to CS

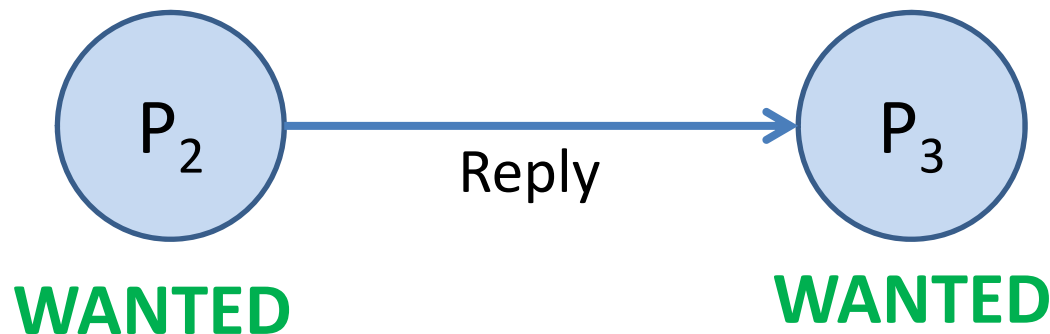
P_2 and P_3 request entry
to CS concurrently

RELEASED

@ P_2 : 7 from $P_3 < 15$ from P_2 ,
therefore grant P_3
access first



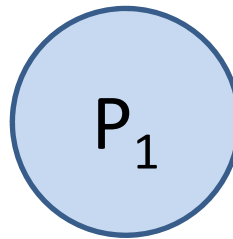
@ P_3 : 7 from $P_3 < 15$ from P_2 ,
Own timestamp lower, therefore,
hold on to reply to P_2



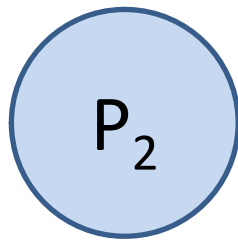
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

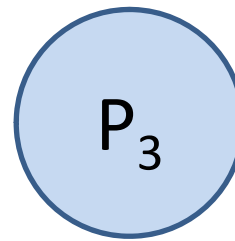
RELEASED



@ P_3 : 7 from $P_3 < 15$ from P_2 ,
Own timestamp lower, therefore,
hold on to reply to P_2



WANTED

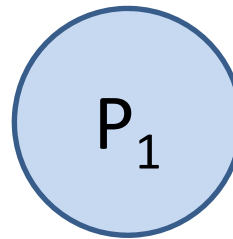
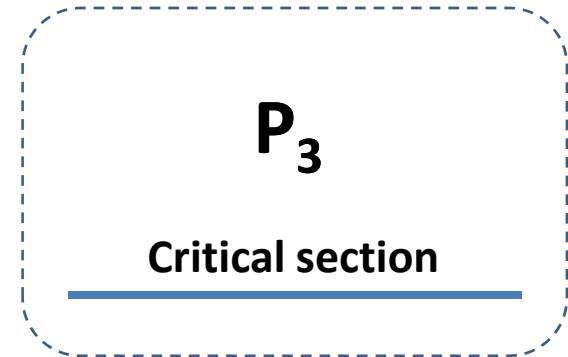


WANTED

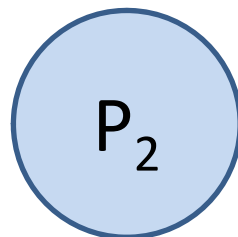
Concurrent entry requests to CS

P_2 and P_3 request entry to CS concurrently

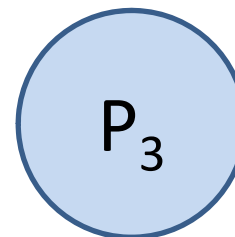
RELEASED



@ P_3 : 7 from P_3 < 15 from P_2 ,
Own timestamp lower, therefore,
hold on to reply to P_2



WANTED

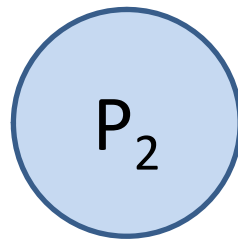
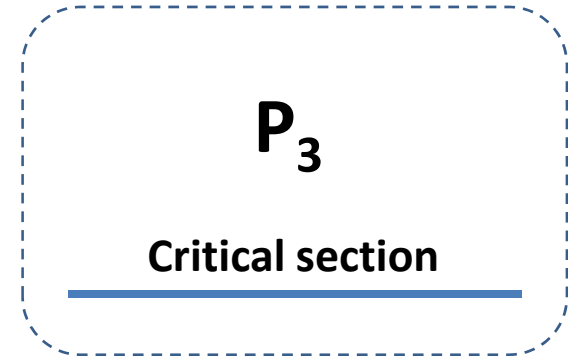
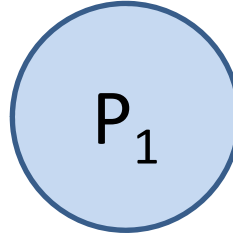


HELD

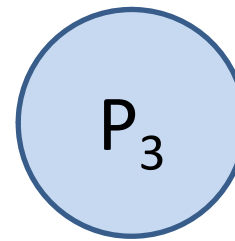
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

RELEASED



WANTED

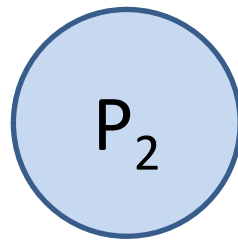
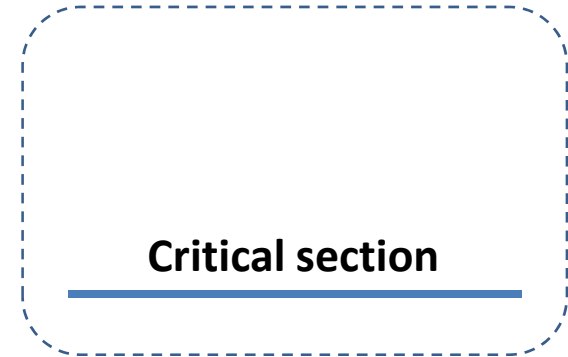
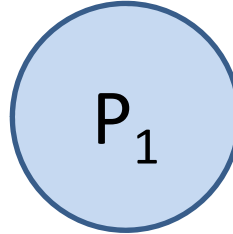


HELD

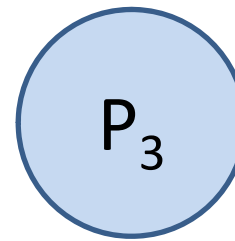
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

RELEASED



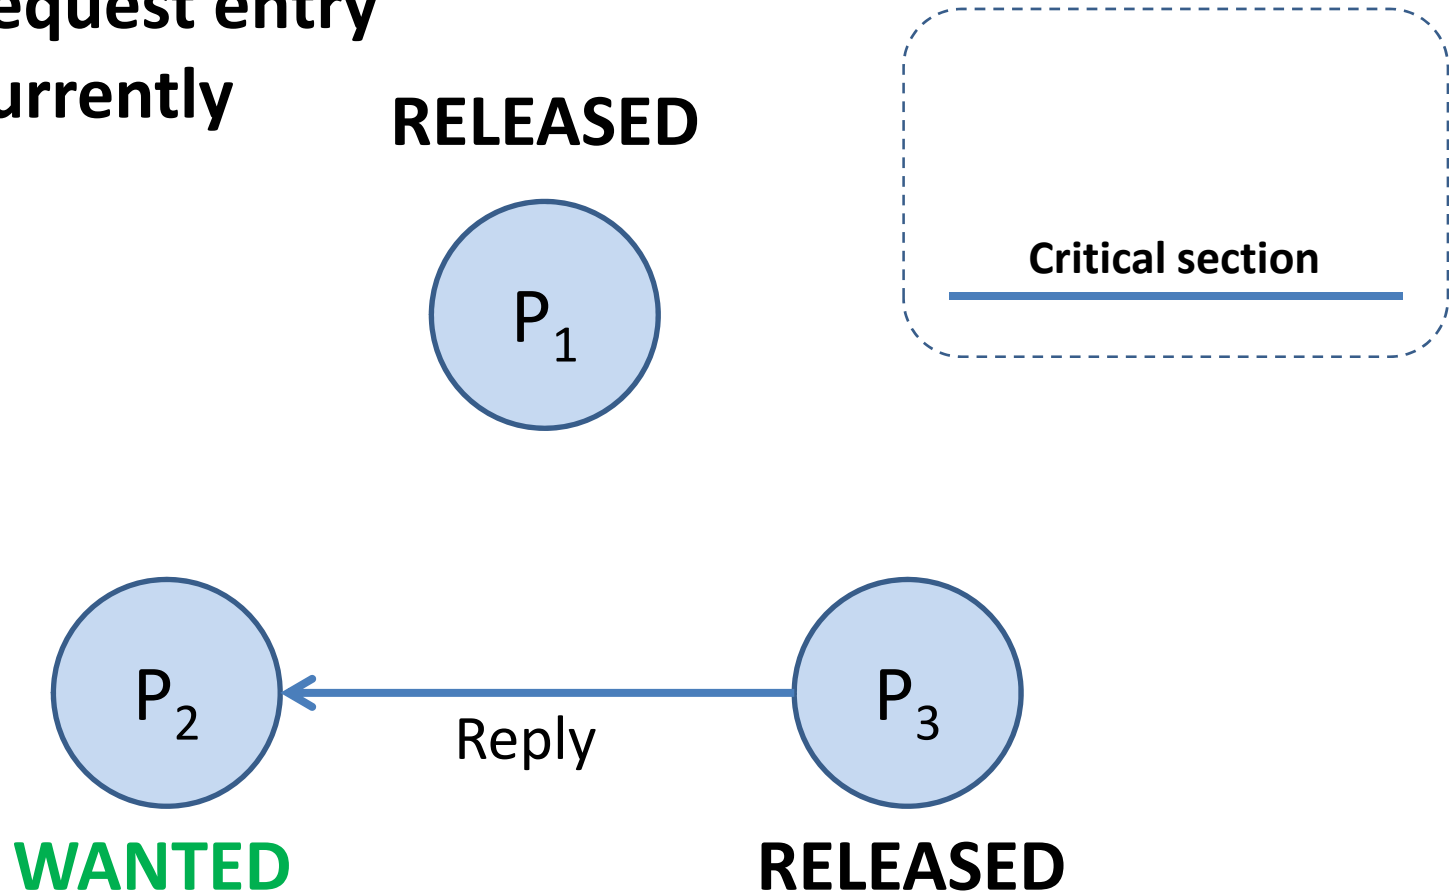
WANTED



RELEASED

Concurrent entry requests to CS

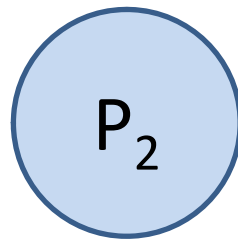
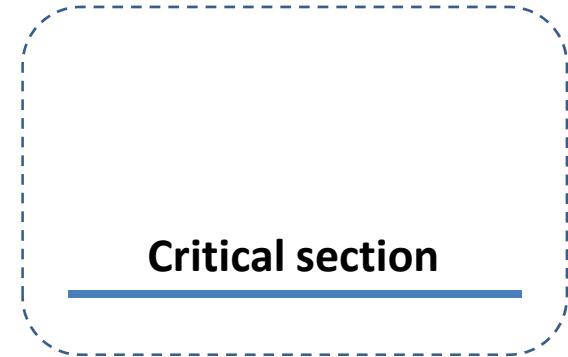
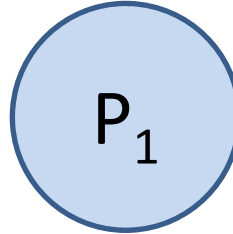
P_2 and P_3 request entry
to CS concurrently



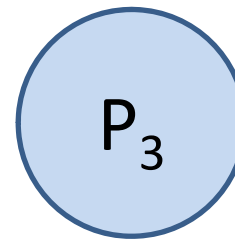
Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently

RELEASED



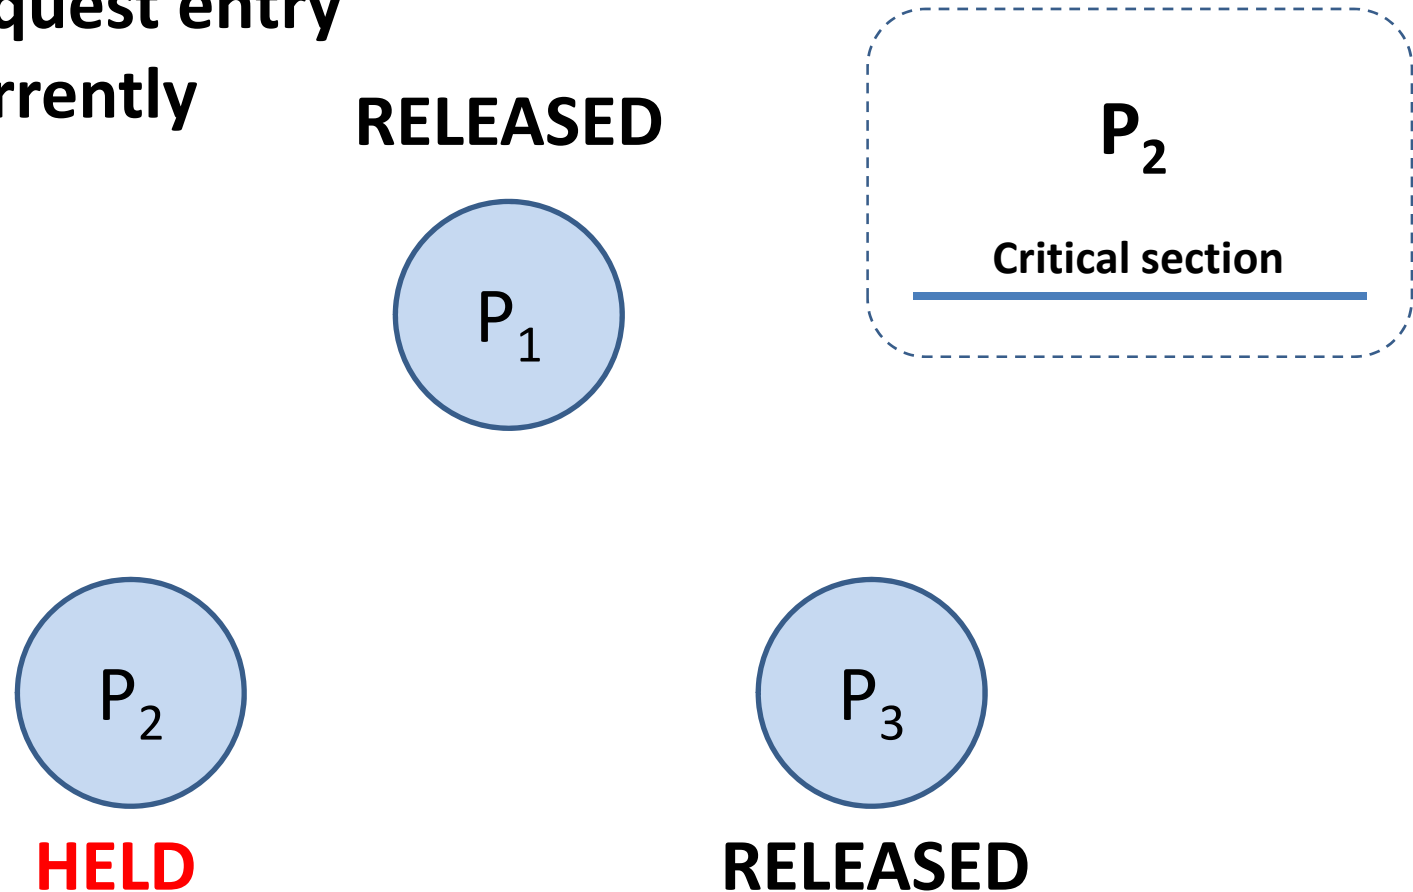
WANTED



RELEASED

Concurrent entry requests to CS

P_2 and P_3 request entry
to CS concurrently



Reminder: Subtlety about timestamps

- Use **Lamport timestamps** to order requests:
 $\langle T, P_i \rangle$, T the timestamp, P_i the process identifier
- If for two timestamps
 - $\langle T, P \rangle$ and $\langle S, Q \rangle$, $T=S$, then break ties by looking at process identifiers
 - Gives rise to a total order over timestamps

Ricart & Agrawala algorithm

Ricart & Agrawala algorithm

On initialization
state = RELEASED

Ricart & Agrawala algorithm

Enter_CS()

state = WANTED

Multicast **request** to all processes

T = request's timestamp

wait until ((n-1) acks received)

state = HELD

On initialization

state = RELEASED

Ricart & Agrawala algorithm

Enter_CS()

state = WANTED

Multicast **request** to all processes

T = request's timestamp

wait until ((n-1) acks received)

state = HELD

On initialization

state = RELEASED

On receiving a request with $\langle T_i, P_i \rangle$ at $P_j (i \neq j)$

if (state==HELD or (state==WANTED and $(T, P_j) < (T_i, P_i)$))

queue request from P_i without replying

else

send a reply to P_i

j

Ricart & Agrawala algorithm

Enter_CS()

state = WANTED

Multicast **request** to all processes

T = request's timestamp

wait until ((n-1) acks received)

state = HELD

On initialization

state = RELEASED

Exit_CS()

state = RELEASED

Reply to all queued requests

On receiving a request with $\langle T_i, P_i \rangle$ at $P_j (i \neq j)$

if (state==HELD or (state==WANTED and $(T, P_j) < (T_i, P_i)$))

queue request from P_i without replying

else

send a reply to P_i

j

Ricart & Agrawala algorithm

Enter_CS()

state = WANTED

Multicast **request** to all processes

T = request's timestamp

wait until ($(n-1)$ acks received)

state = HELD

On initialization

state = RELEASED

Exit_CS()

state = RELEASED

Reply to all queued requests

On receiving a request with $\langle T_i, P_i \rangle$ at $P_j (i \neq j)$

if (state==HELD or (state==WANTED and $(T, P_j) < (T_i, P_i)$))

queue request from P_i without replying

else

send a reply to P_i

j

Ricart & Agrawala algorithm

Enter_CS()

state = WANTED

Multicast **request** to all processes

T = request's timestamp

wait until ((n-1) acks received)

state = HELD

On initialization

state = RELEASED

Exit_CS()

state = RELEASED

Reply to all queued requests

On receiving a request with $\langle T_i, P_i \rangle$ at $P_j (i \neq j)$

if (state==HELD or (state==WANTED and $(T, P_j) < (T_i, P_i)$))

queue request from P_i without replying

else

send a reply to P_i

j

Potential fault-tolerance issues

(Our assumptions bracketed them out.)

- None of the algorithms tolerate message loss
- Ring-based algorithm cannot tolerate crash of single process
- Centralized algorithm can tolerate crash of processes that are neither in the CS, nor have requested entry
- Lamport, R&A can not tolerate faults

LEADER ELECTION

Leader election

- **Problem:** A group of processes, P_1, \dots, P_n , **must agree** on some **unique P_k** to be the “**leader**”
- Often, the leader then **coordinates another activity**
- Election runs when leader (a.k.a., coordinator) failed
- Any process who hasn't heard from the leader in some predefined time interval **may call for an election**
- **False alarm** is a possibility (new elections initiated, while current leader still alive)
- **Several processes** may initiate **elections concurrently**
- Algorithm should allow for process crash during election

Applications of leader election

- Berkeley clock synchronization algorithm
- Centralized mutual exclusion algorithm
- Leader election for choosing the master in Hbase / Bigtable (using ZooKeeper/Chubby)
- Choosing the master among the 5 servers of Chubby or ZooKeeper cell
- *Primary-backup replication algorithms*
- *Two-phase commit protocol*

As compared to mutual exclusion

- Losers return to what they were doing ...
... **instead of waiting**
- **Fast election** is important ...
... not starvation avoidance
- **All processes** must know result ...
... not just the winner

As compared to mutual exclusion

- Losers return to what they were doing ...
... **instead of waiting**
- **Fast election** is important ...
... not starvation avoidance
- **All processes** must know result ...
... not just the winner

ME can be reduced to LE!
(e.g., Hbase wants LE,
ZooKeeper provides ME)

Process identifier

- Elected **leader** must be **unique**
- Active process with **largest identifier** wins
- Identifier could be any “useful value”
 - I.e., **unique & totally ordered value**
- E.g., based on OS process identifiers, IP adr., port
- E.g., based on least computational load
 - $\langle 1/\text{load}, i \rangle$, $\text{load} > 0$, i is process ID to break ties
- Each process, P_i , has a variable **elected_i** that holds the value of the leader or is “ \perp ” (undefined)

Election algorithm requirement

- **Safety**

- A participating process, P_i , has variable **electd_i** = “ \perp ” or **electd_i** = P , where P is chosen as the non-crashed process at the end of the election run with the **largest identifier**.
- **Only one leader at a time!**

- **Liveness**

- **All processes participate** in the election and **eventually** either set **electd_i** \neq “ \perp ” or **crash**.

Chang & Roberts Ring-based algorithm, 1978

Essentially three phases

1. **Initialization**

2. **Election phase** (concurrent calls for election)

- Determine election winner (voting phase)
- Reach point of message originator

3. **Announce leader phase (victory announcement phase)**

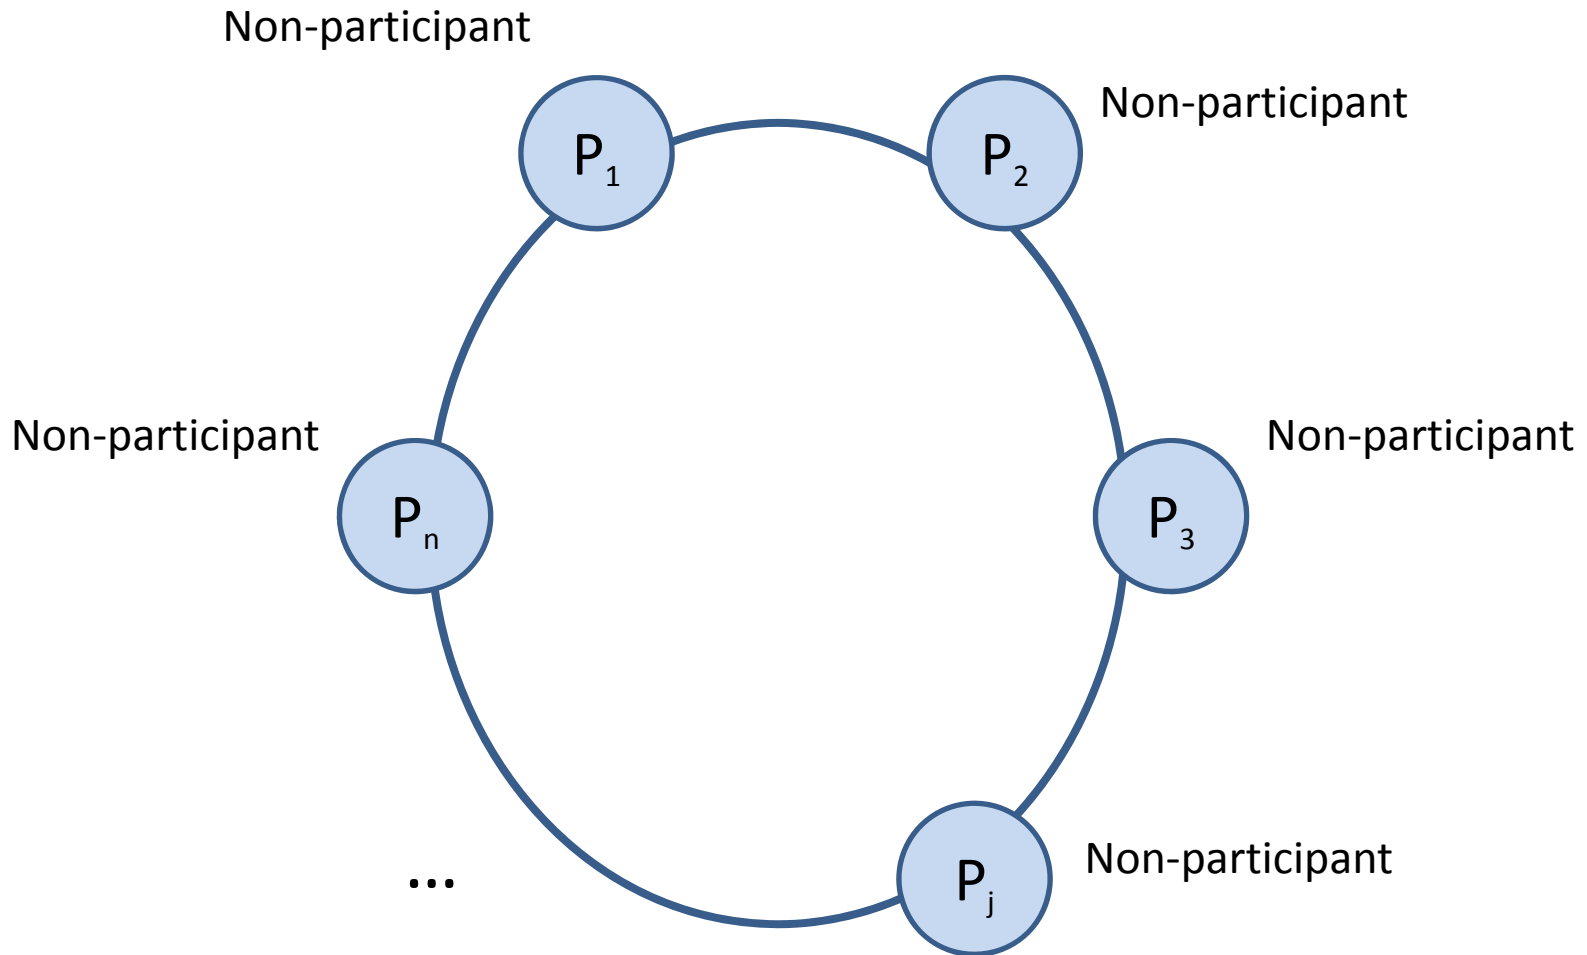
Ring-based election algorithm I

- Construct a ring (cf. ring-based mutual exclusion)
- Assume, each P has **unique ID (e.g., unique integer)**
- Assume, no failures and asynchronous system, **but failures can happen before the election!**
- Any P_i may begin an election by sending an **election message** to its successor (i.e., after detecting leader failure)
- Election message holds P_i 's ID

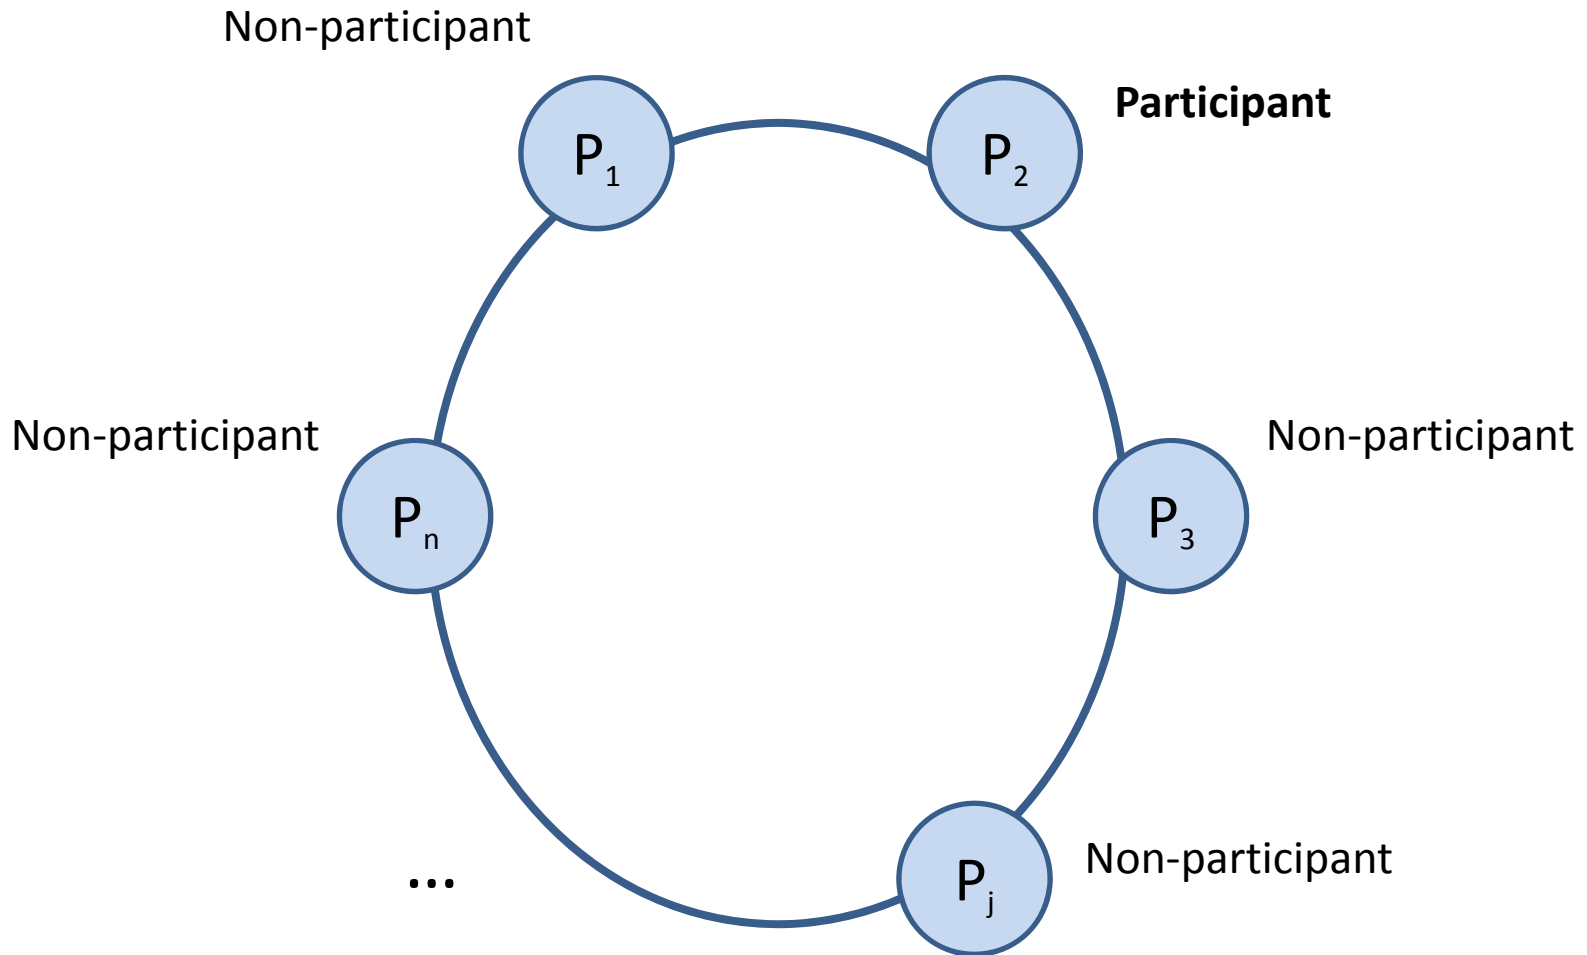
Ring-based election algorithm II

- Election message holds P_i 's ID
- Upon receiving an election message, P_k **compares its own ID to ID it received in election message**
 - If received message ID is **greater**: **Forward** election message
 - If ... **smaller**: **Forward** election message **with P_k 's ID**, unless P_k has already sent a message, i.e., has participated in current election run
 - If ... **equal**: **P_k is now leader**. Forward victory message to notify all other processes

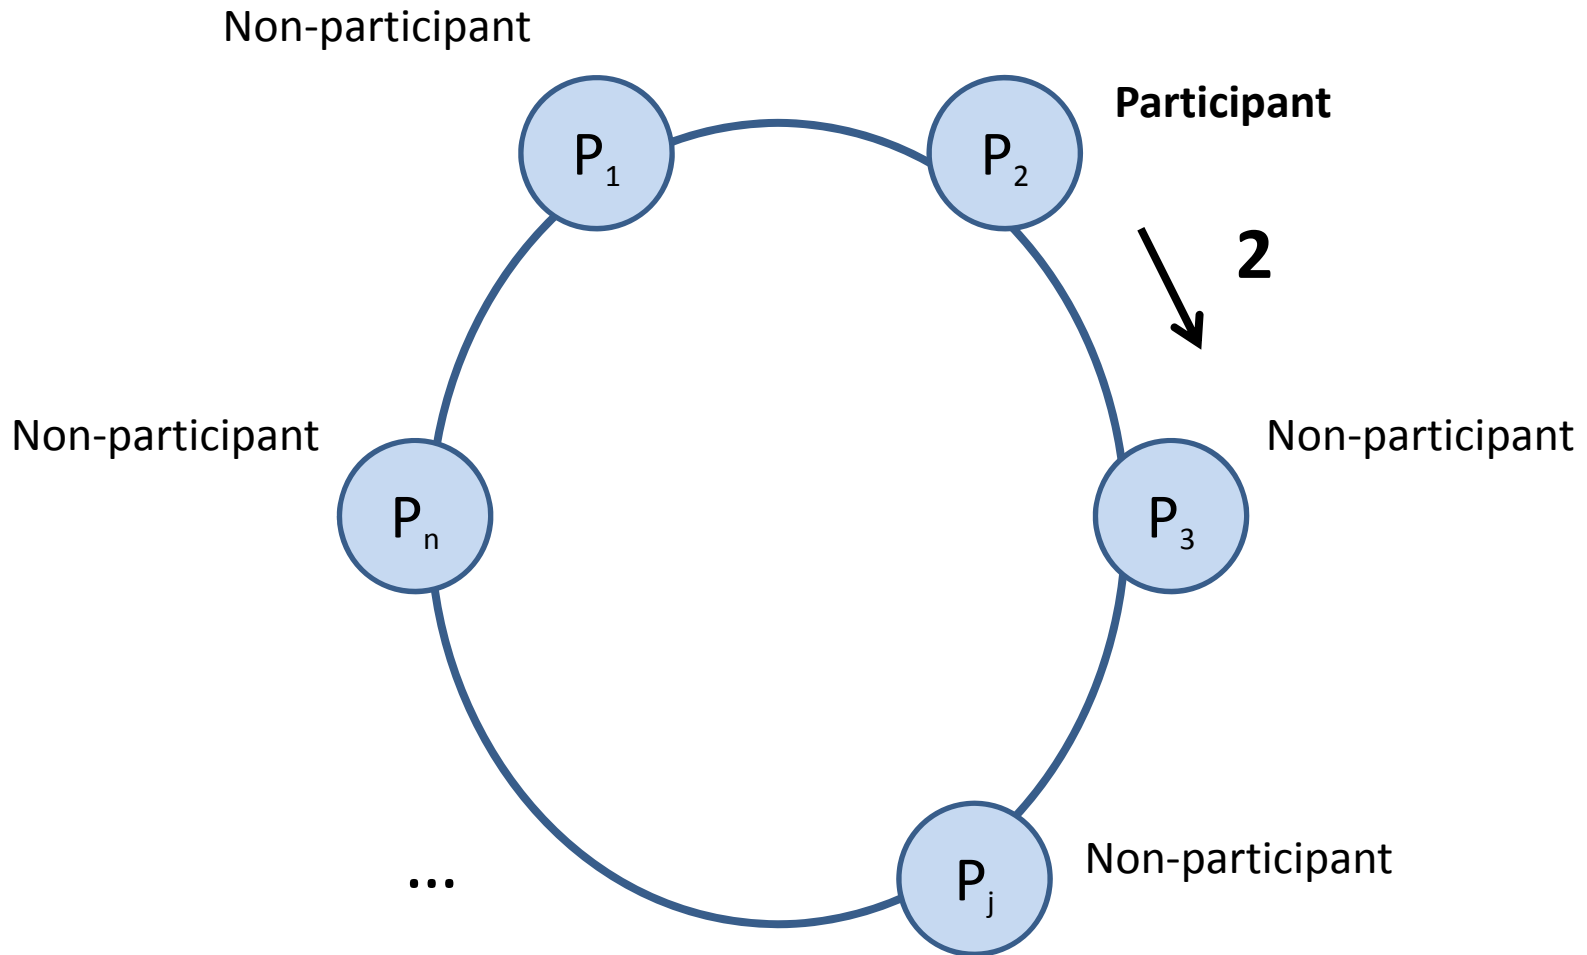
Ring-based algorithm: Calling an Election (determine winner)



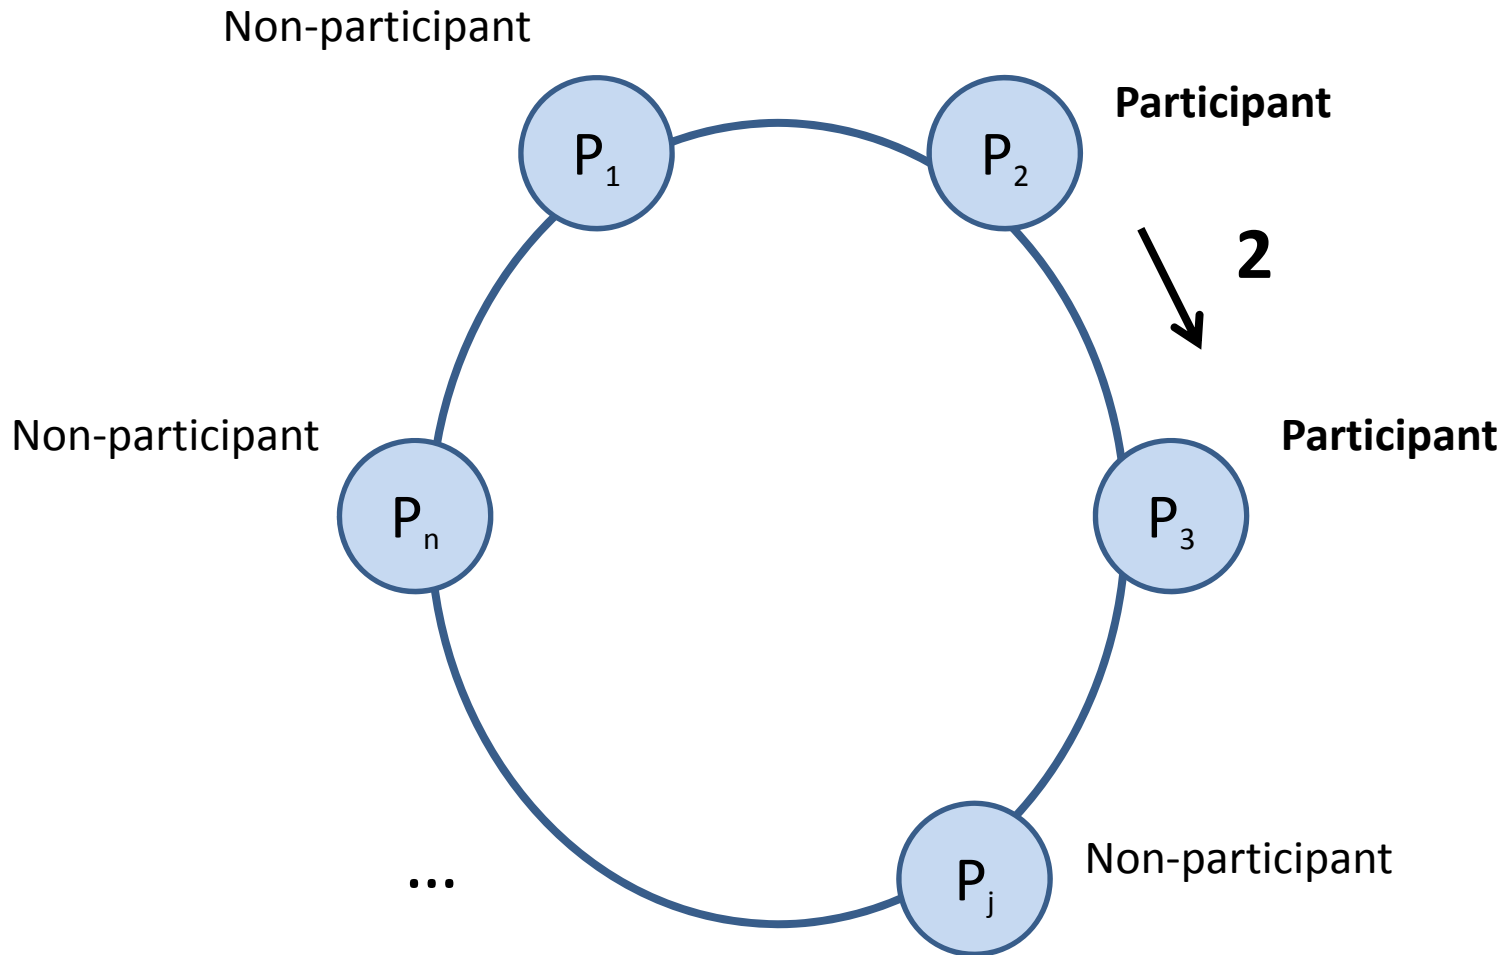
Ring-based algorithm: Calling an Election (determine winner)



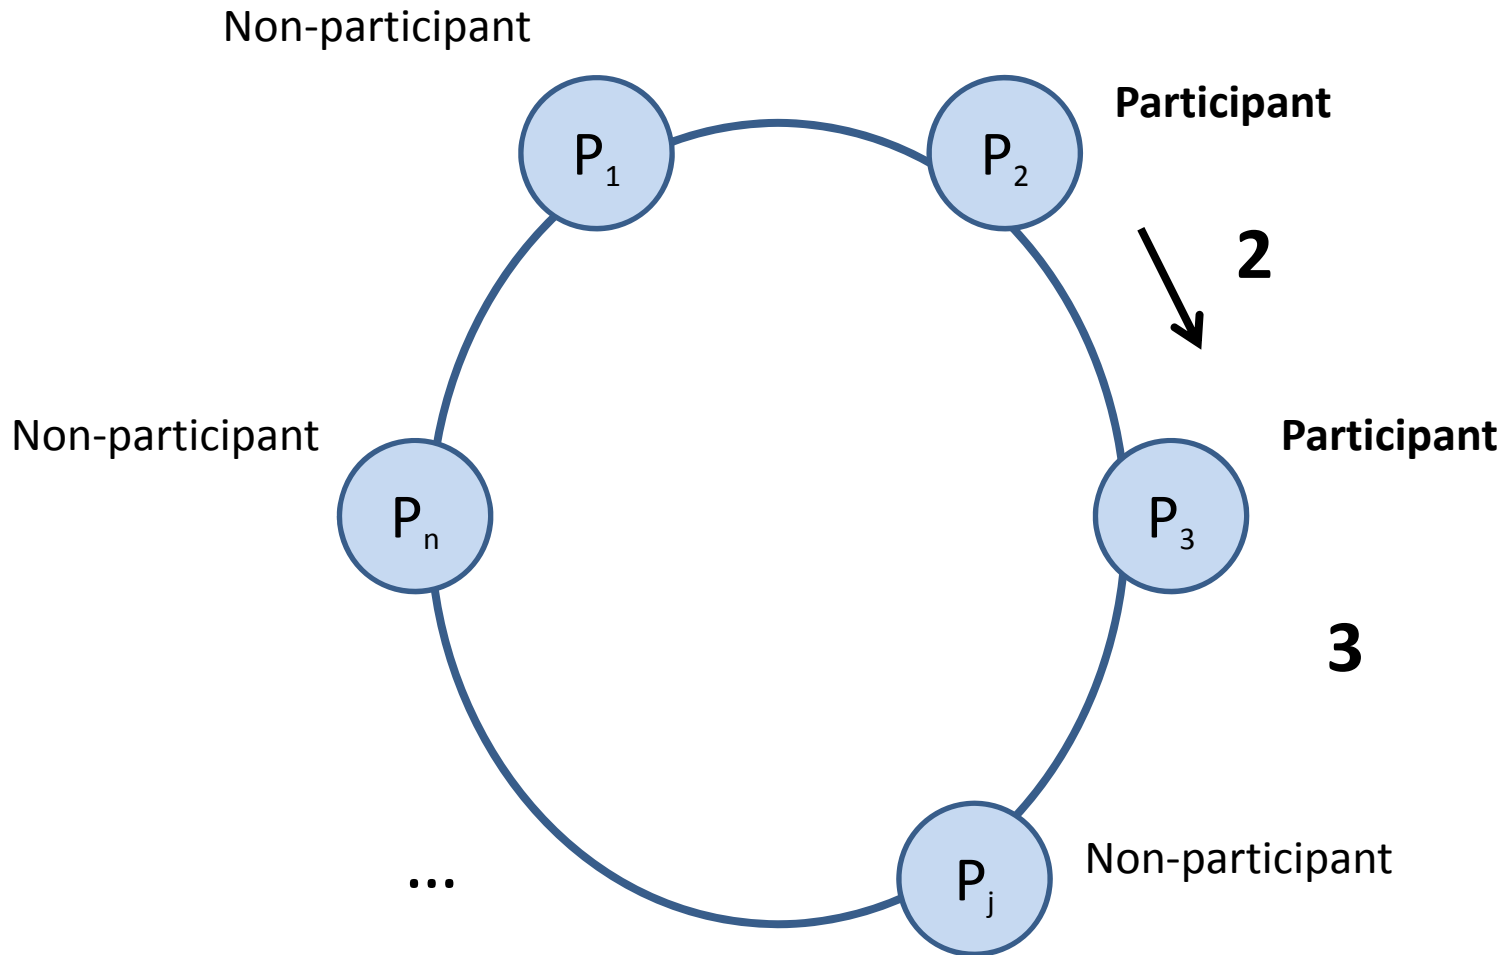
Ring-based algorithm: Calling an Election (determine winner)



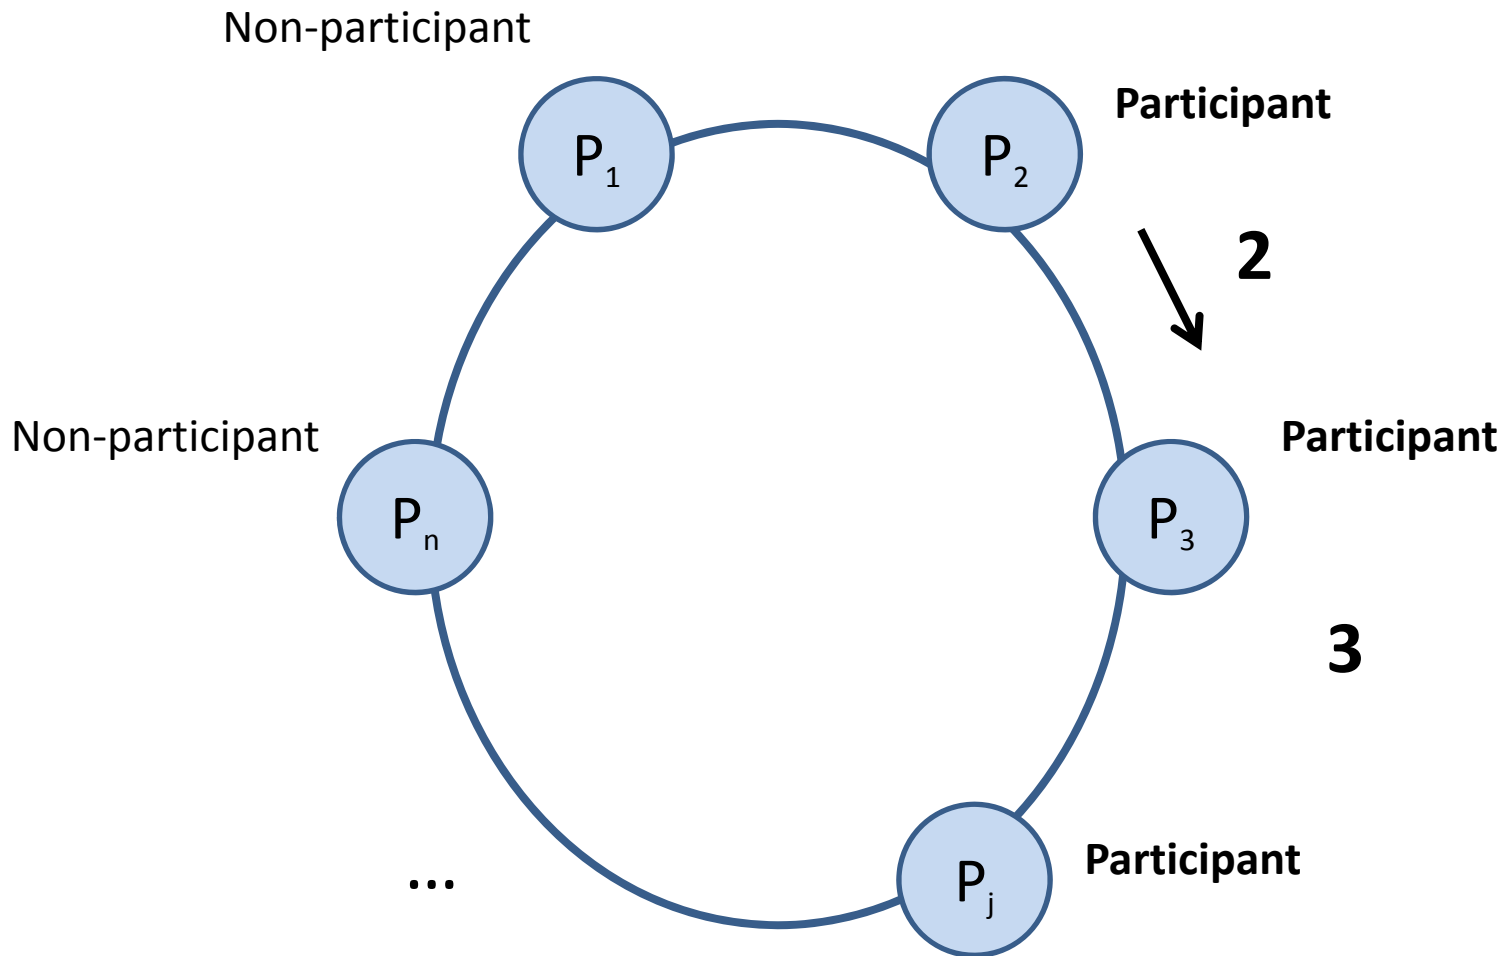
Ring-based algorithm: Calling an Election (determine winner)



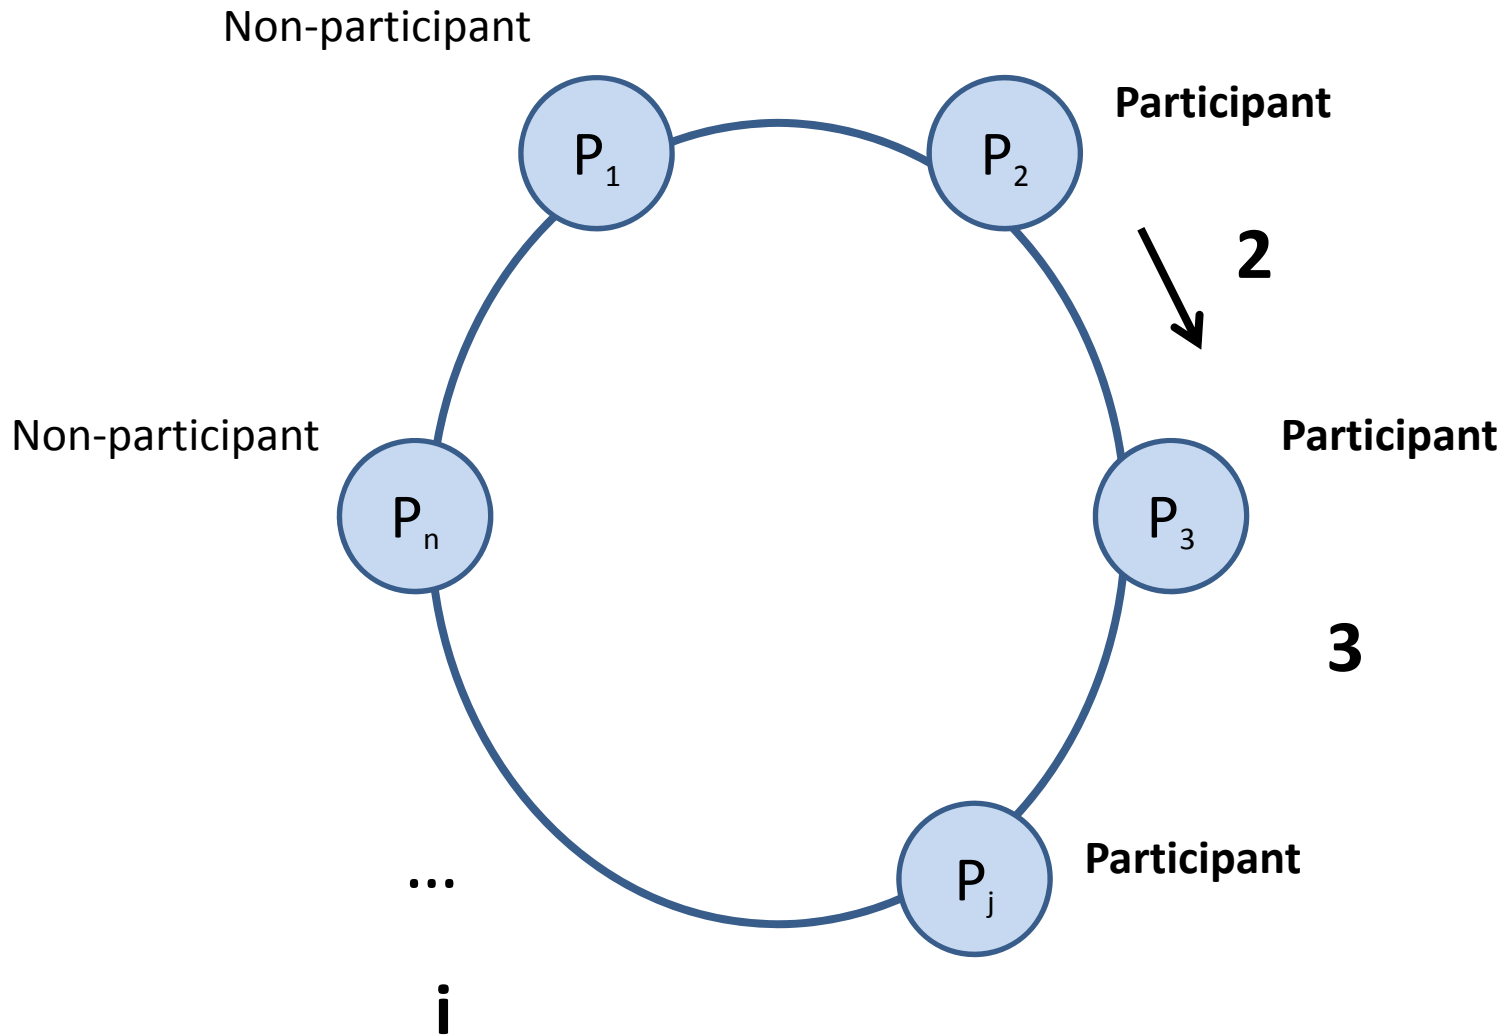
Ring-based algorithm: Calling an Election (determine winner)



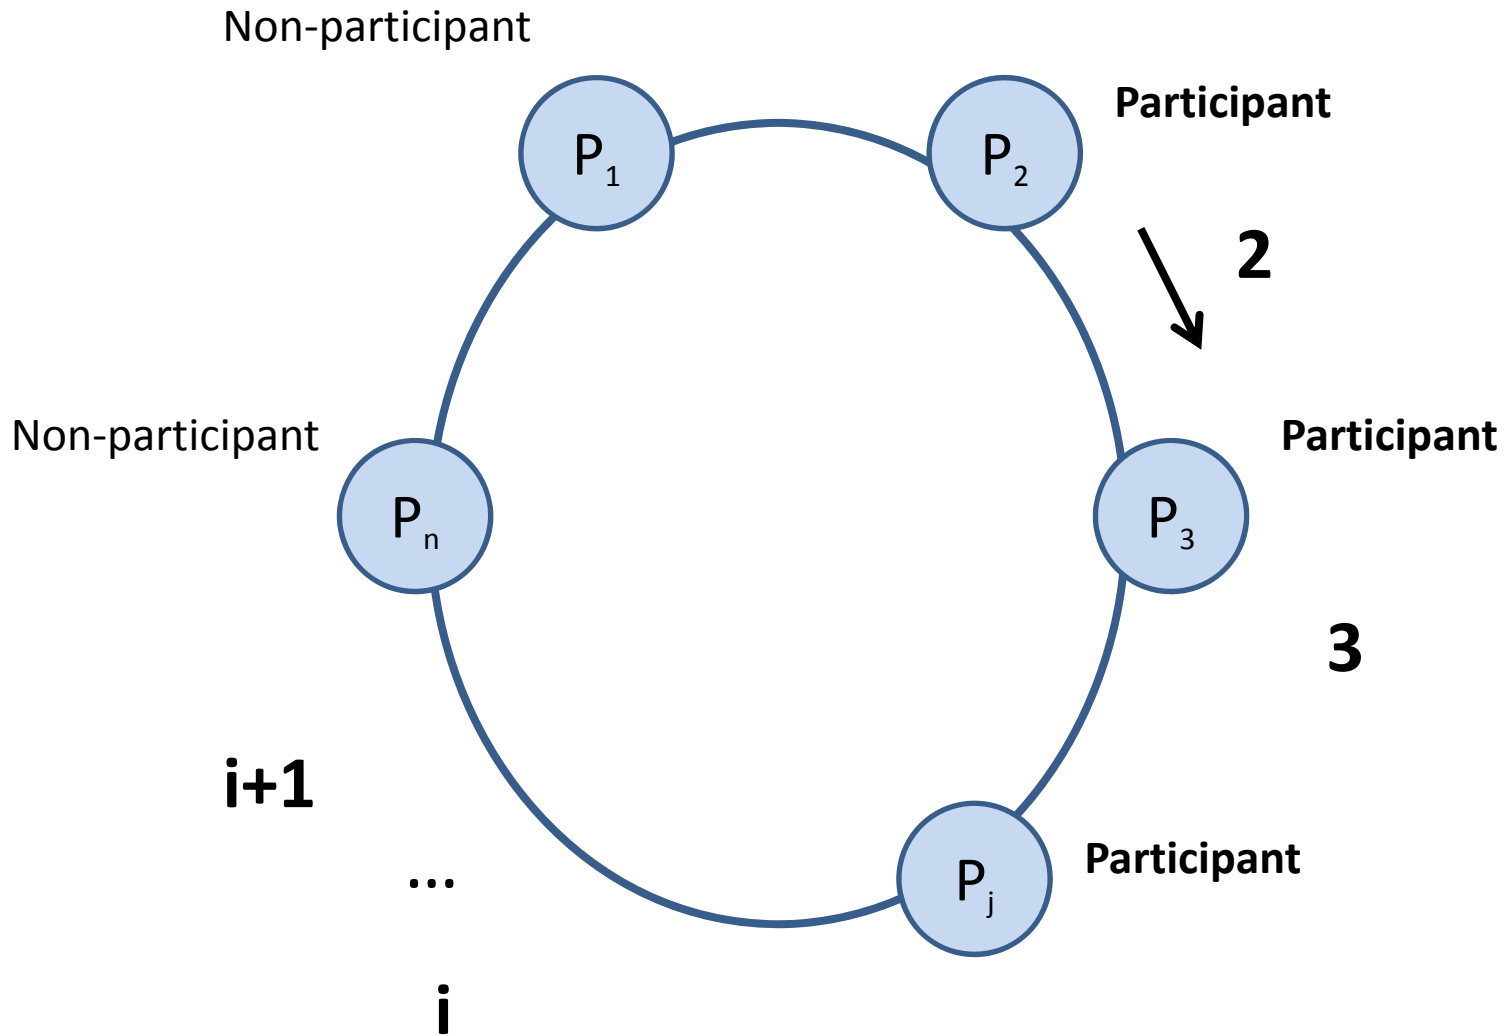
Ring-based algorithm: Calling an Election (determine winner)



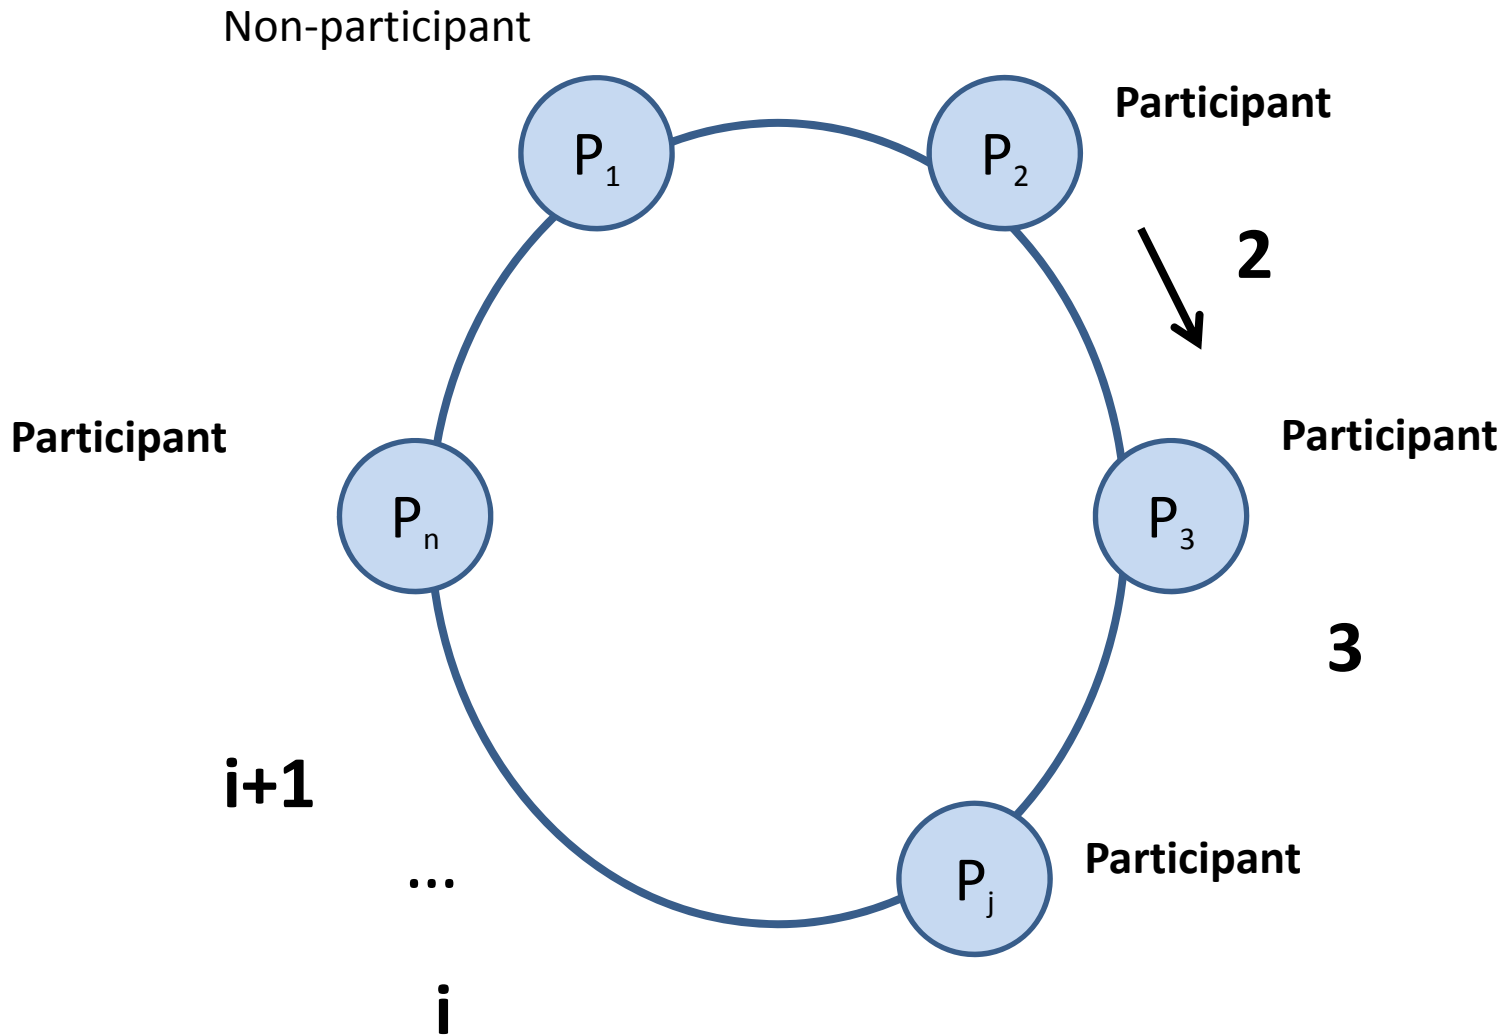
Ring-based algorithm: Calling an Election (determine winner)



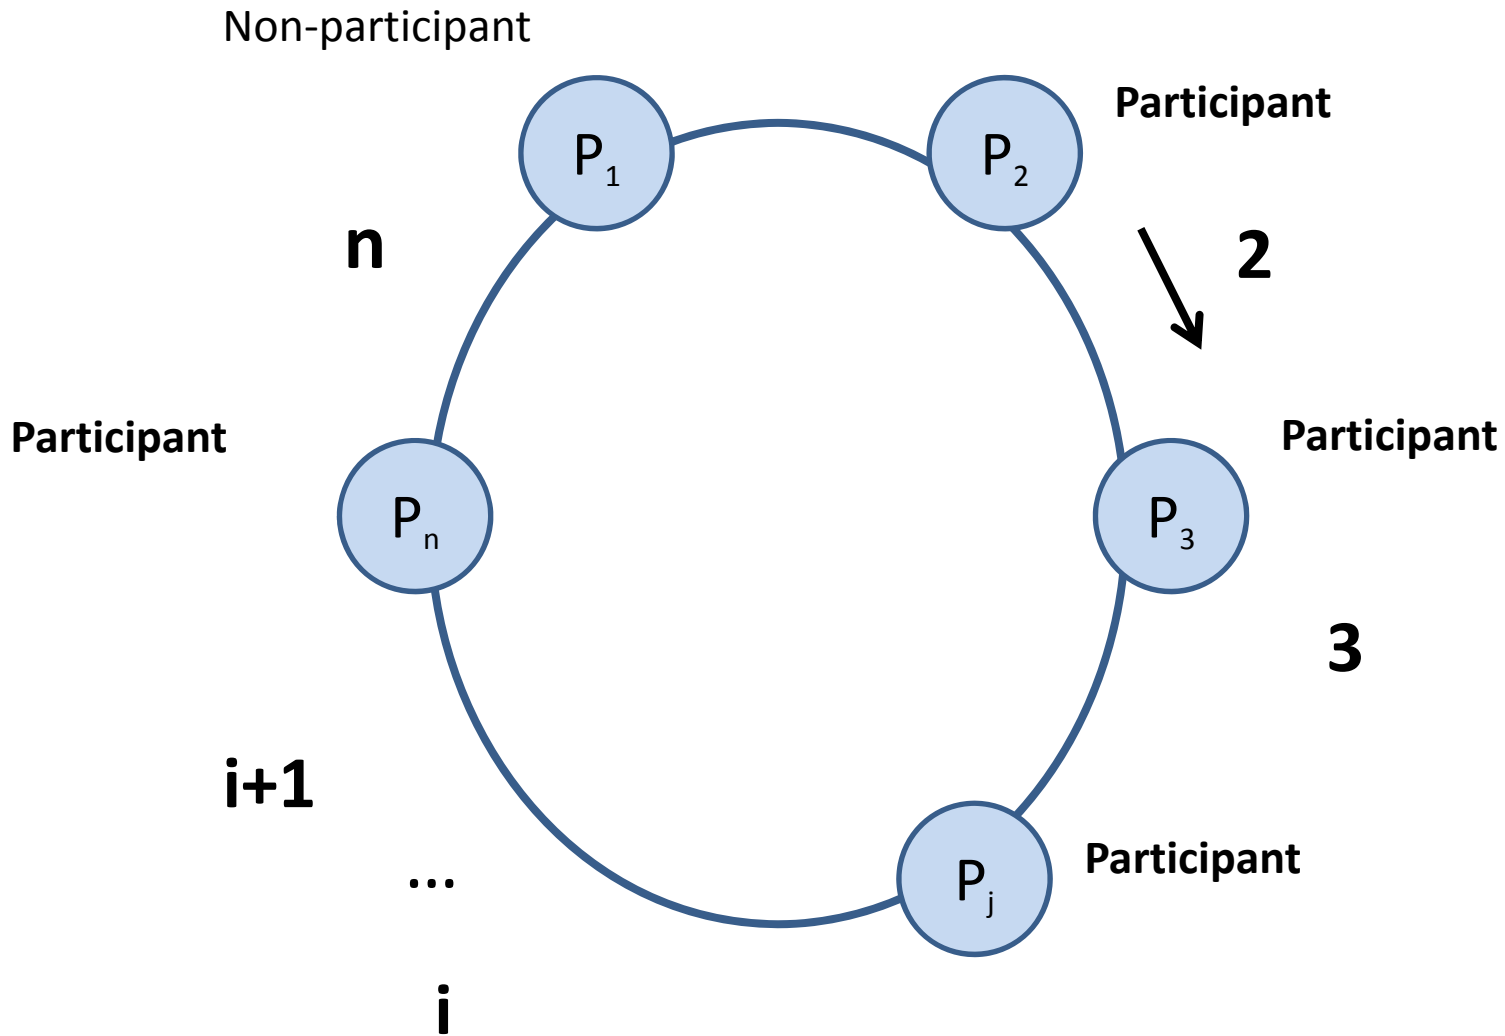
Ring-based algorithm: Calling an Election (determine winner)



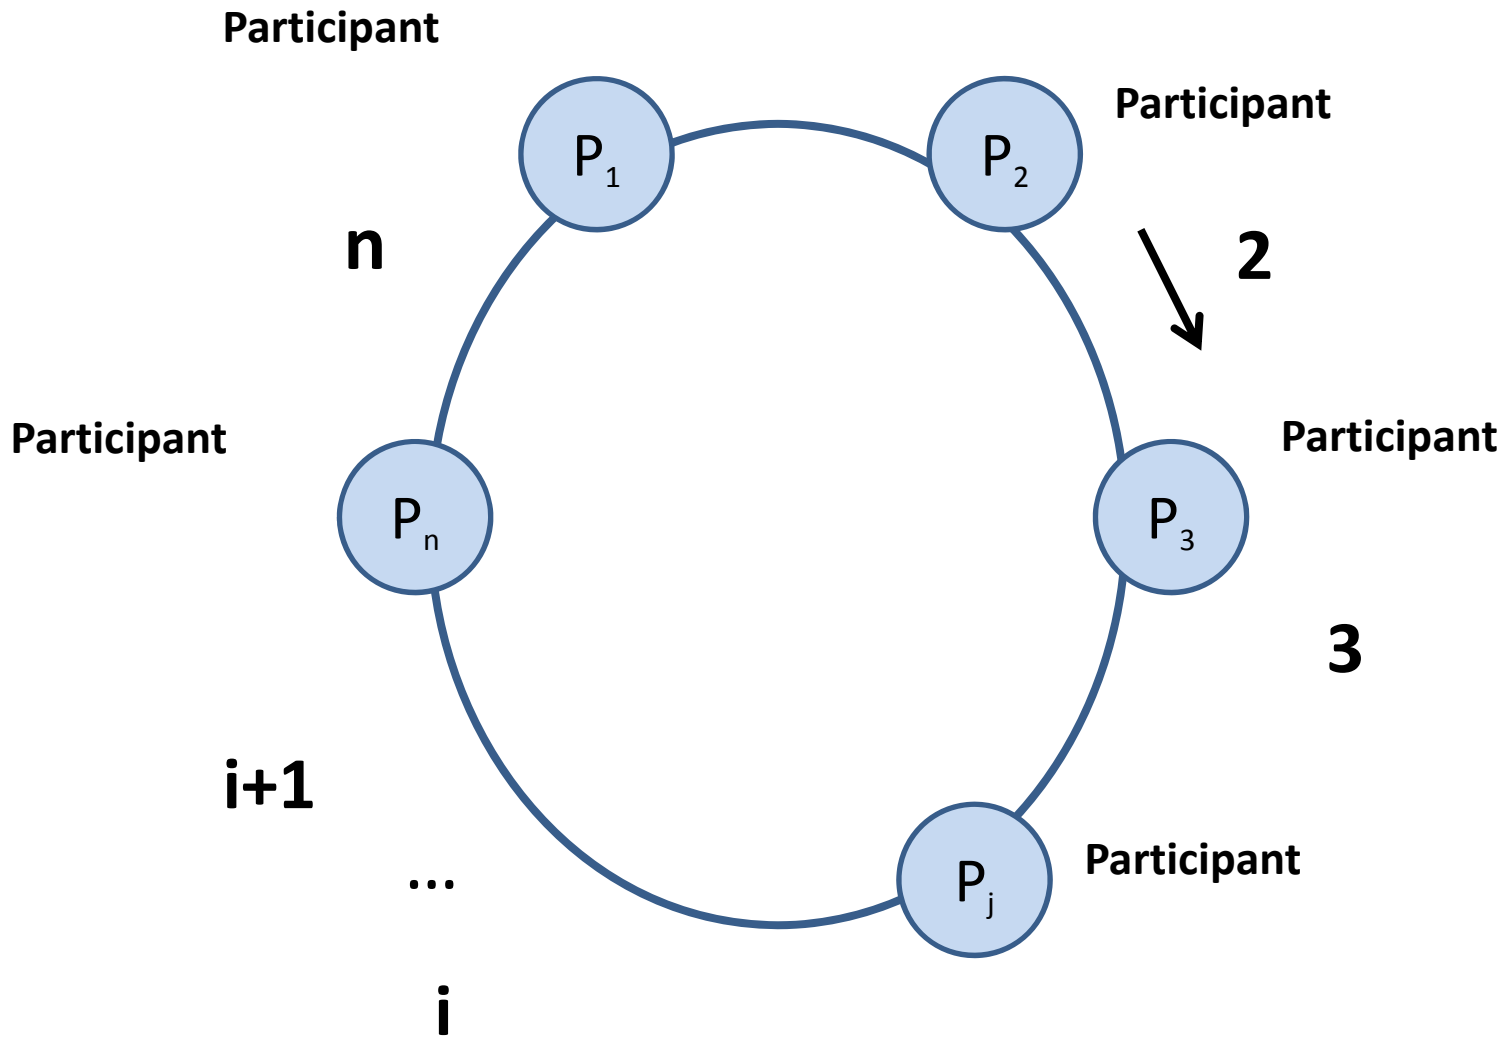
Ring-based algorithm: Calling an Election (determine winner)



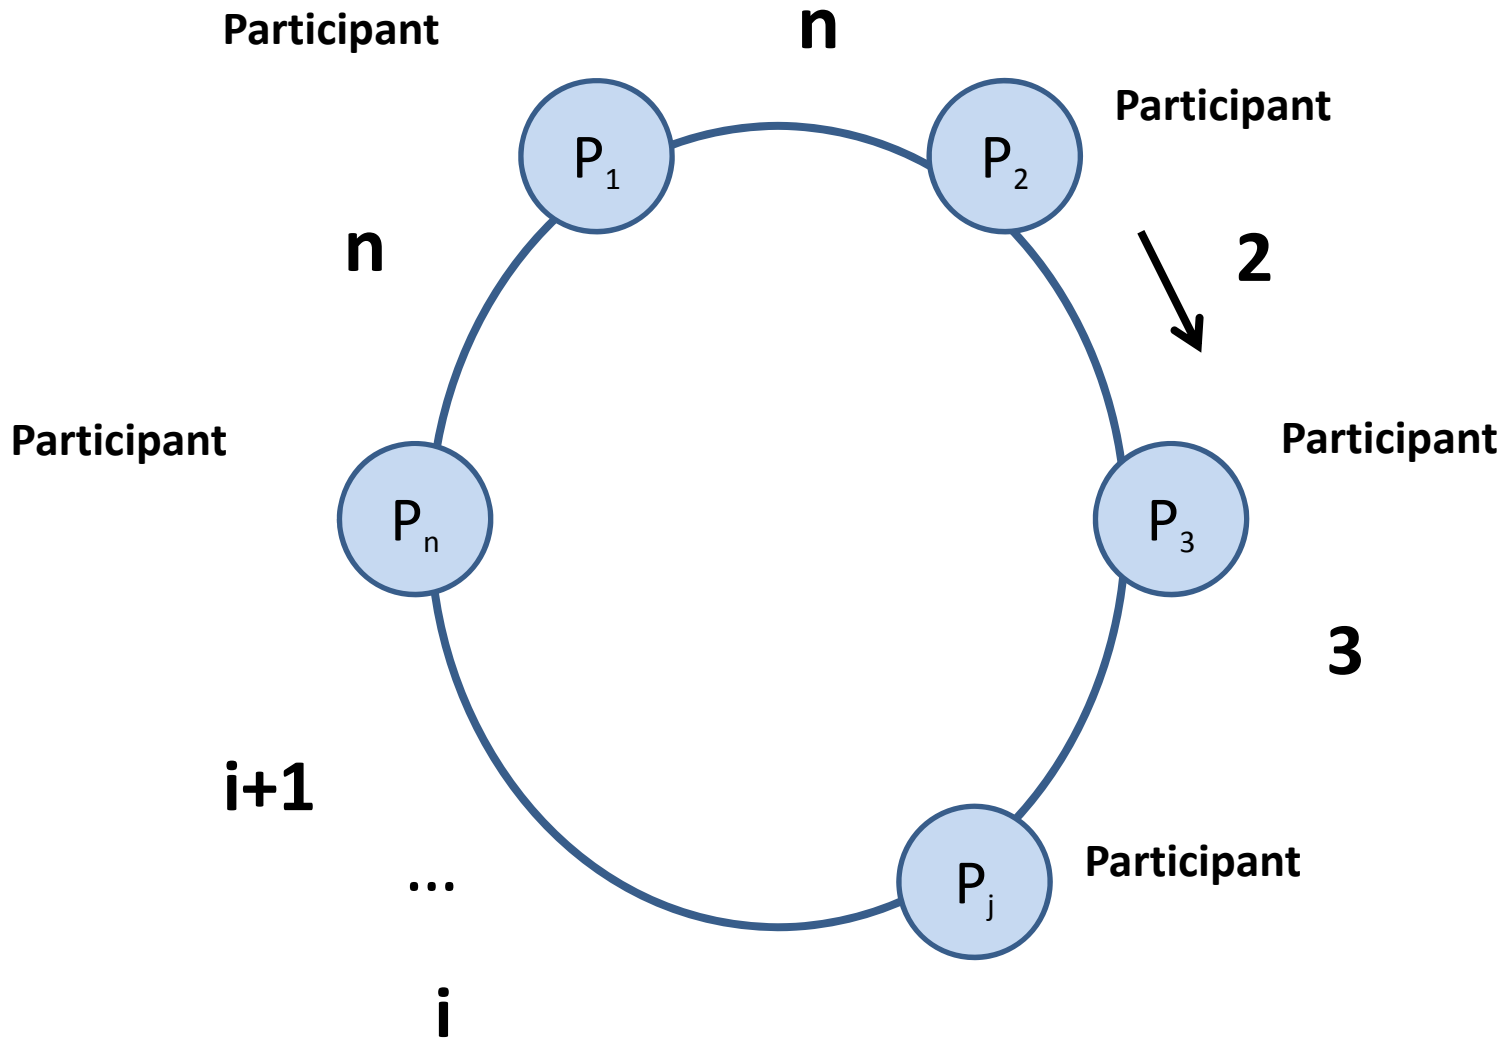
Ring-based algorithm: Calling an Election (determine winner)



Ring-based algorithm: Calling an Election (determine winner)

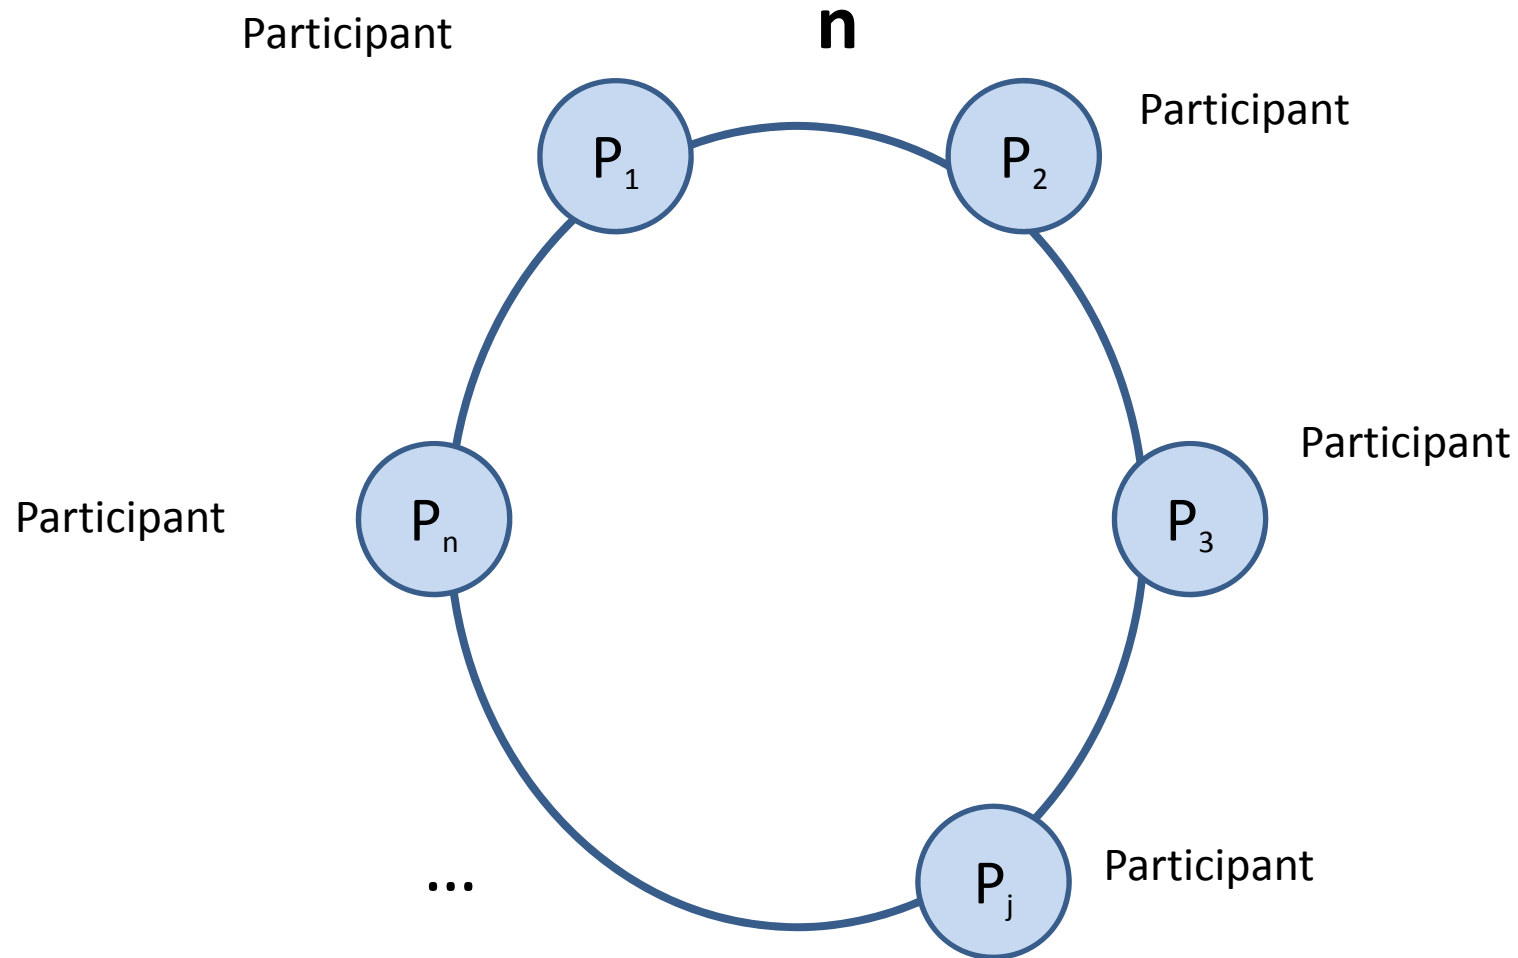


Ring-based algorithm: Calling an Election (determine winner)

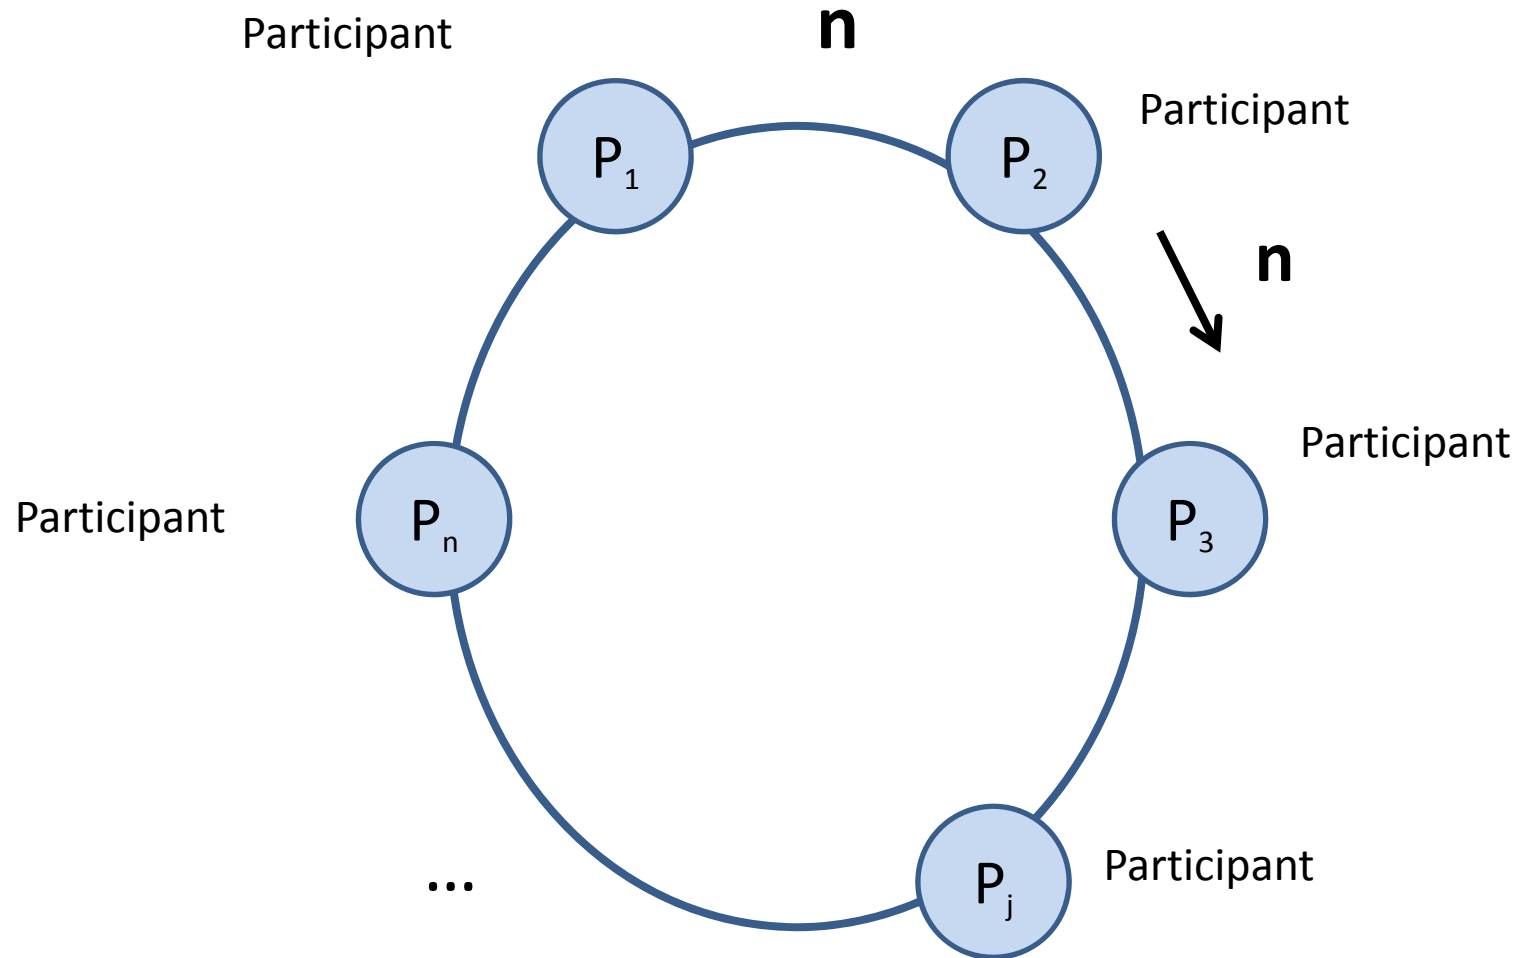


Ring-based algorithm:

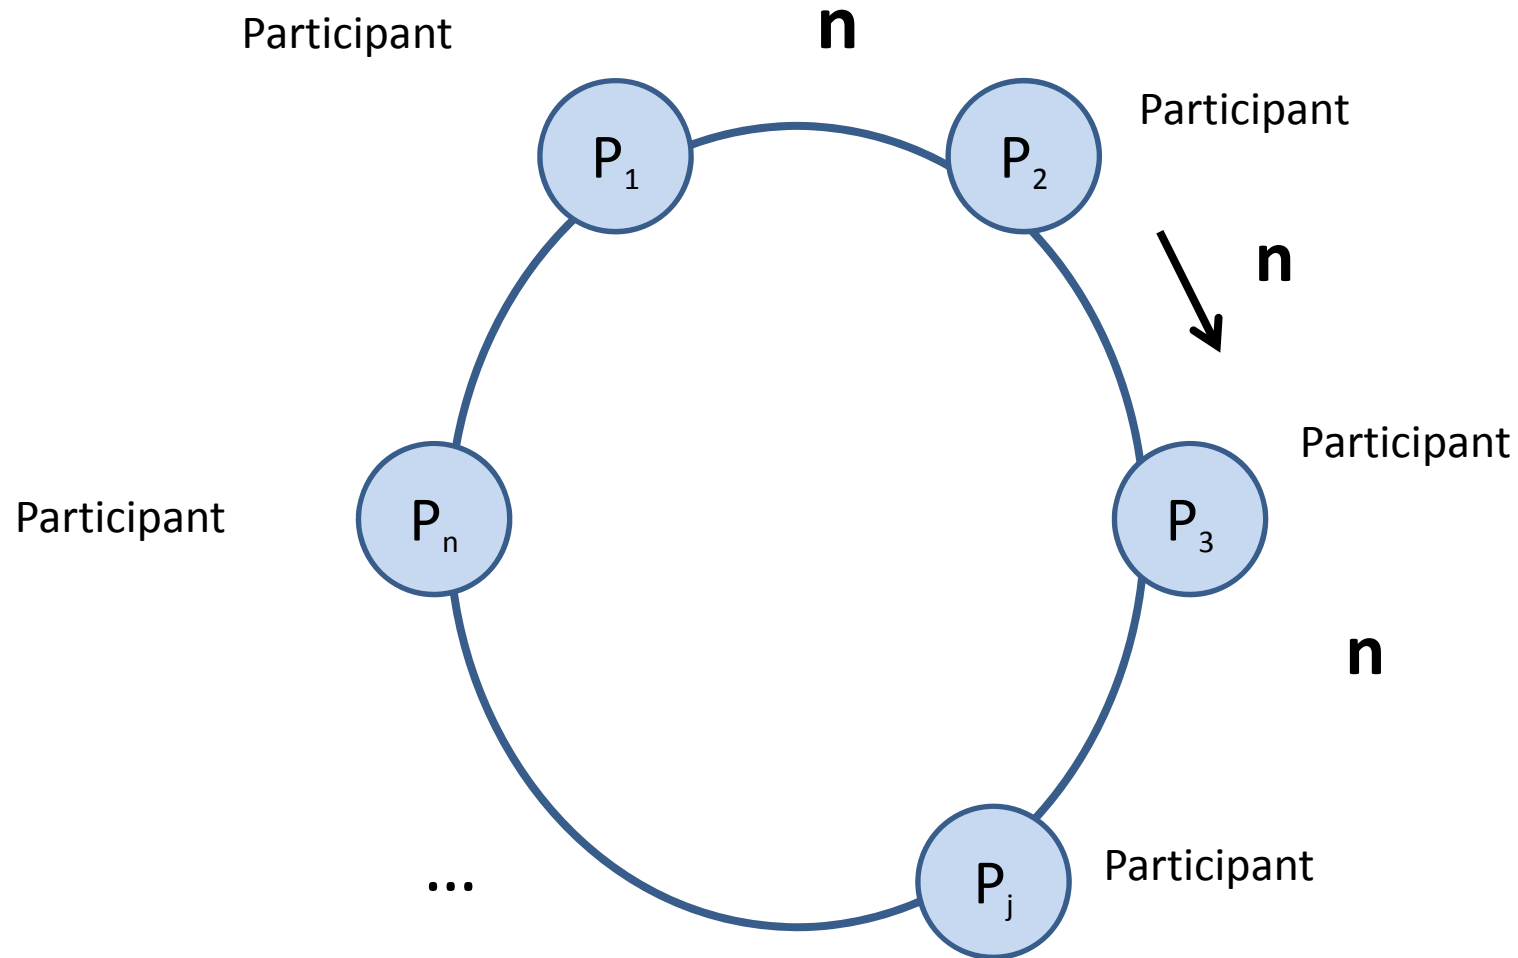
Calling an Election (**origin & victory**)



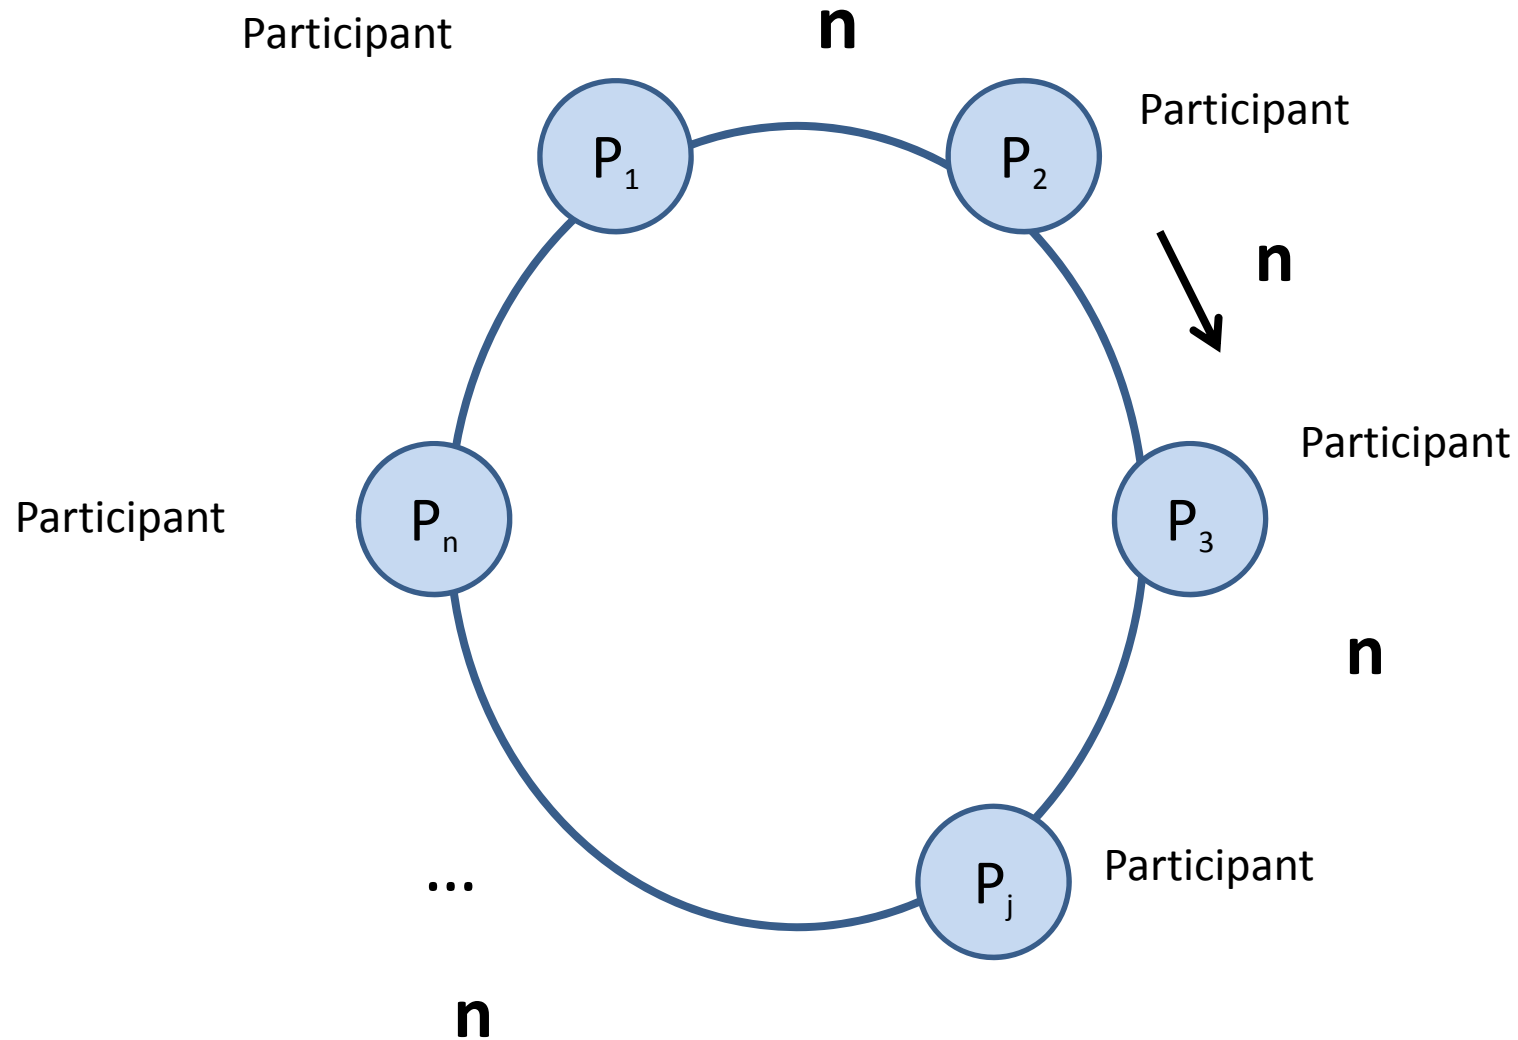
Ring-based algorithm: Calling an Election (**origin & victory**)



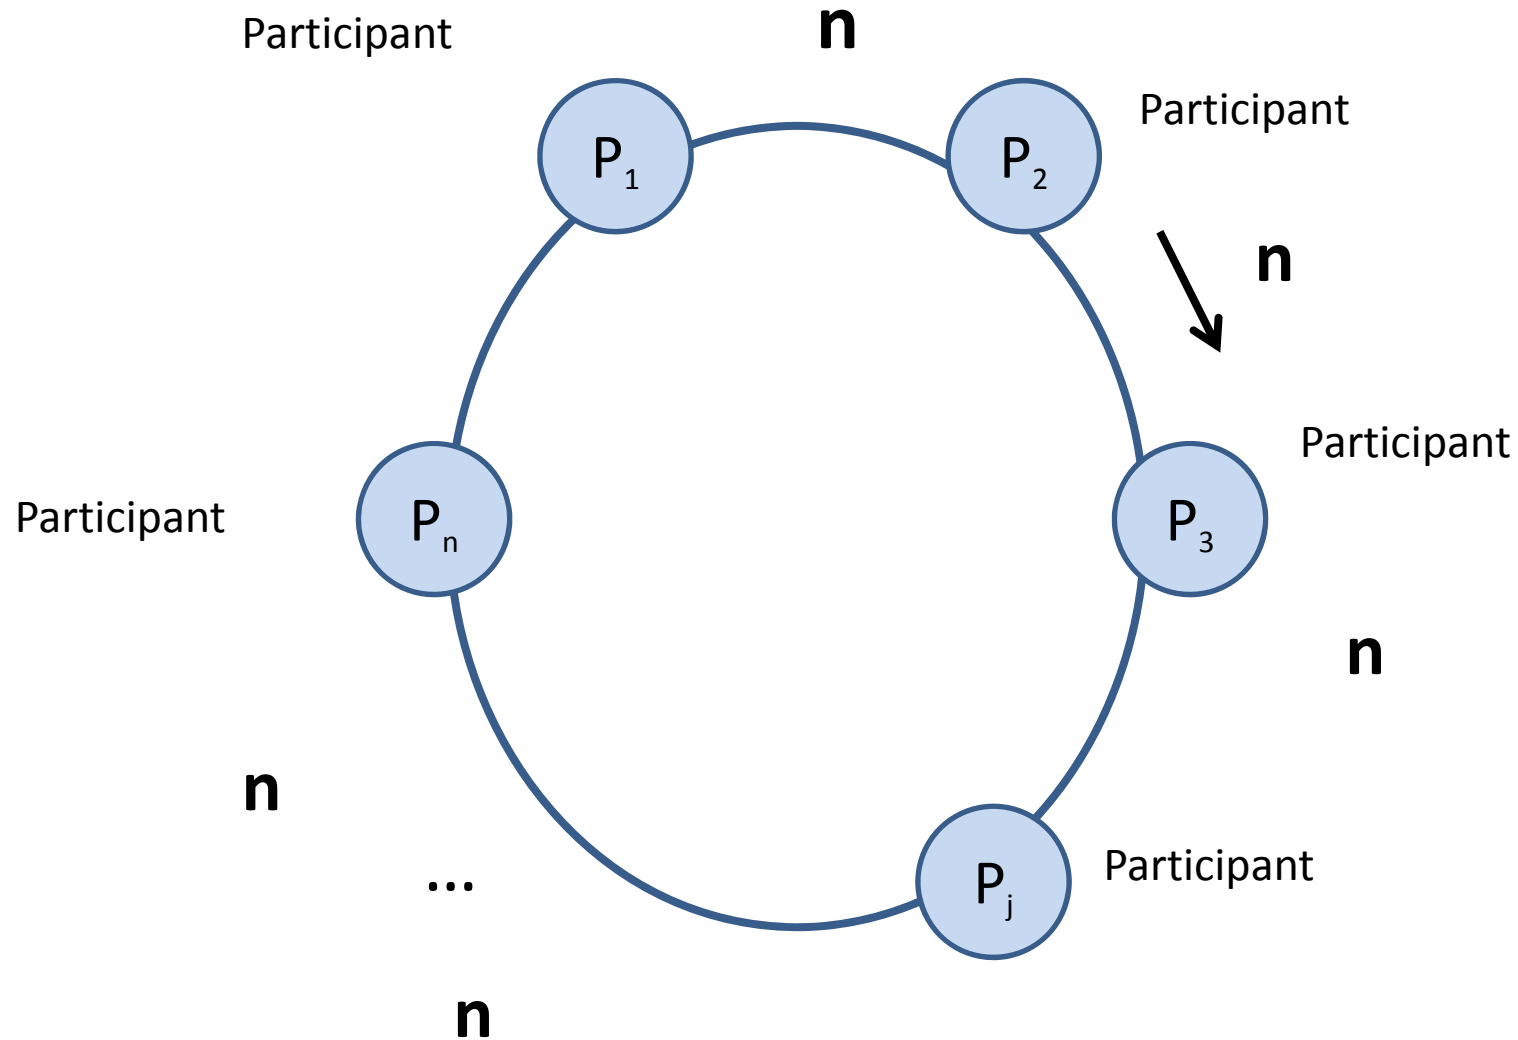
Ring-based algorithm: Calling an Election (**origin & victory**)



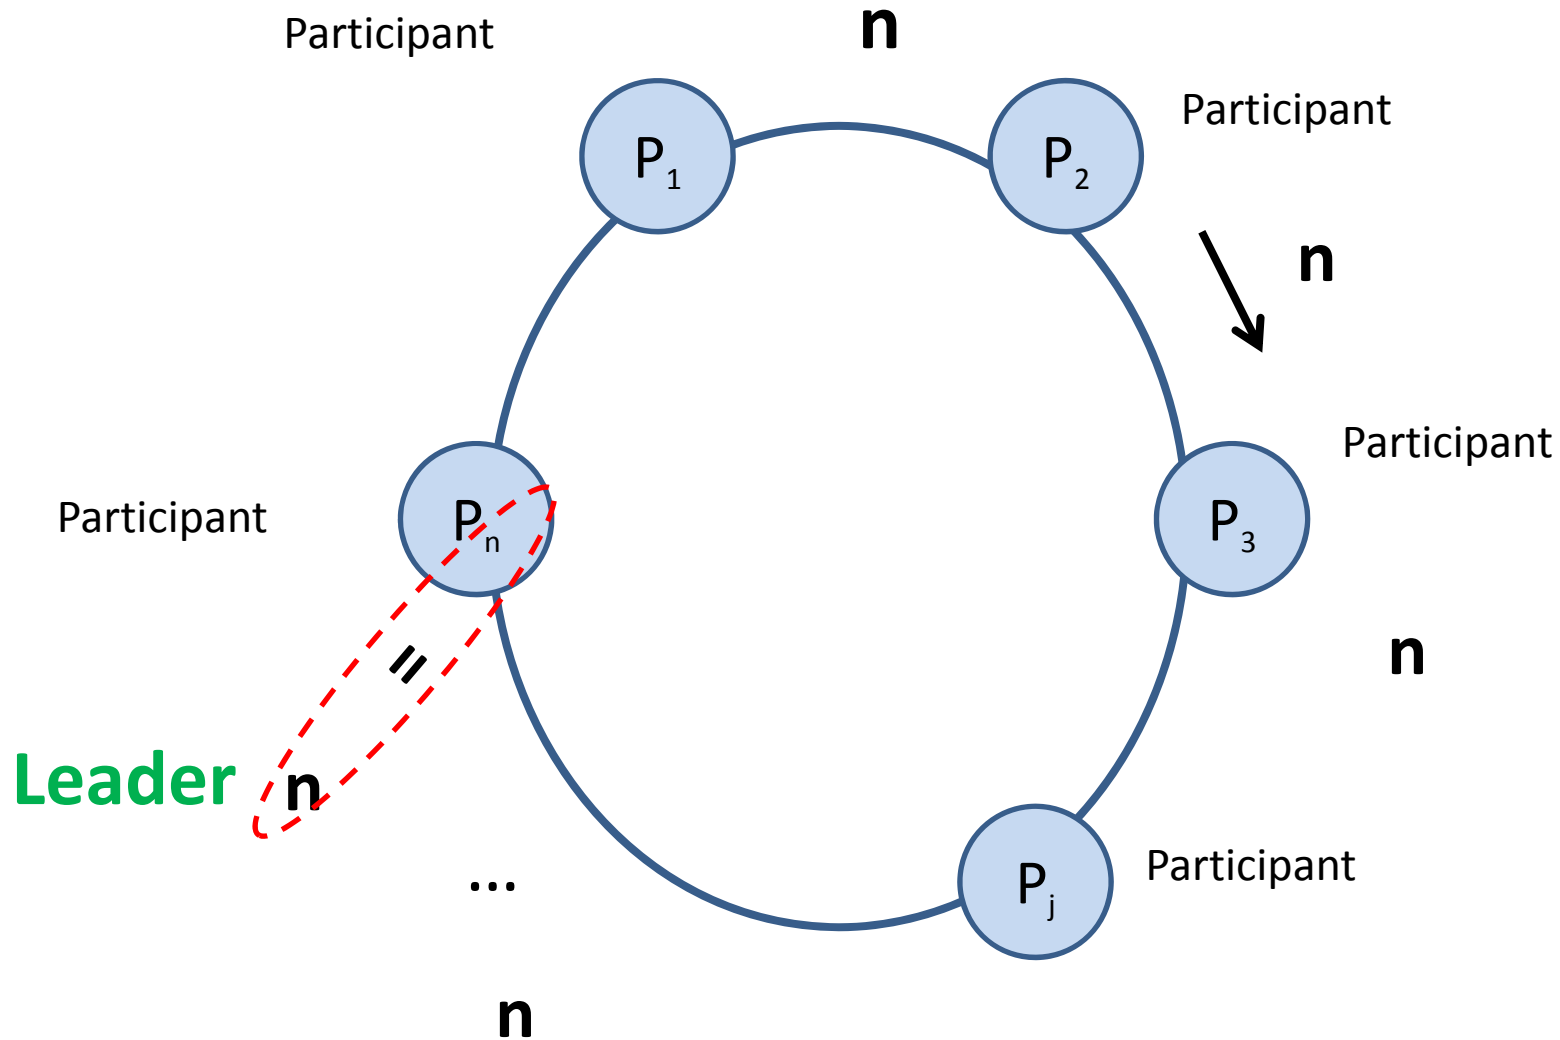
Ring-based algorithm: Calling an Election (**origin & victory**)



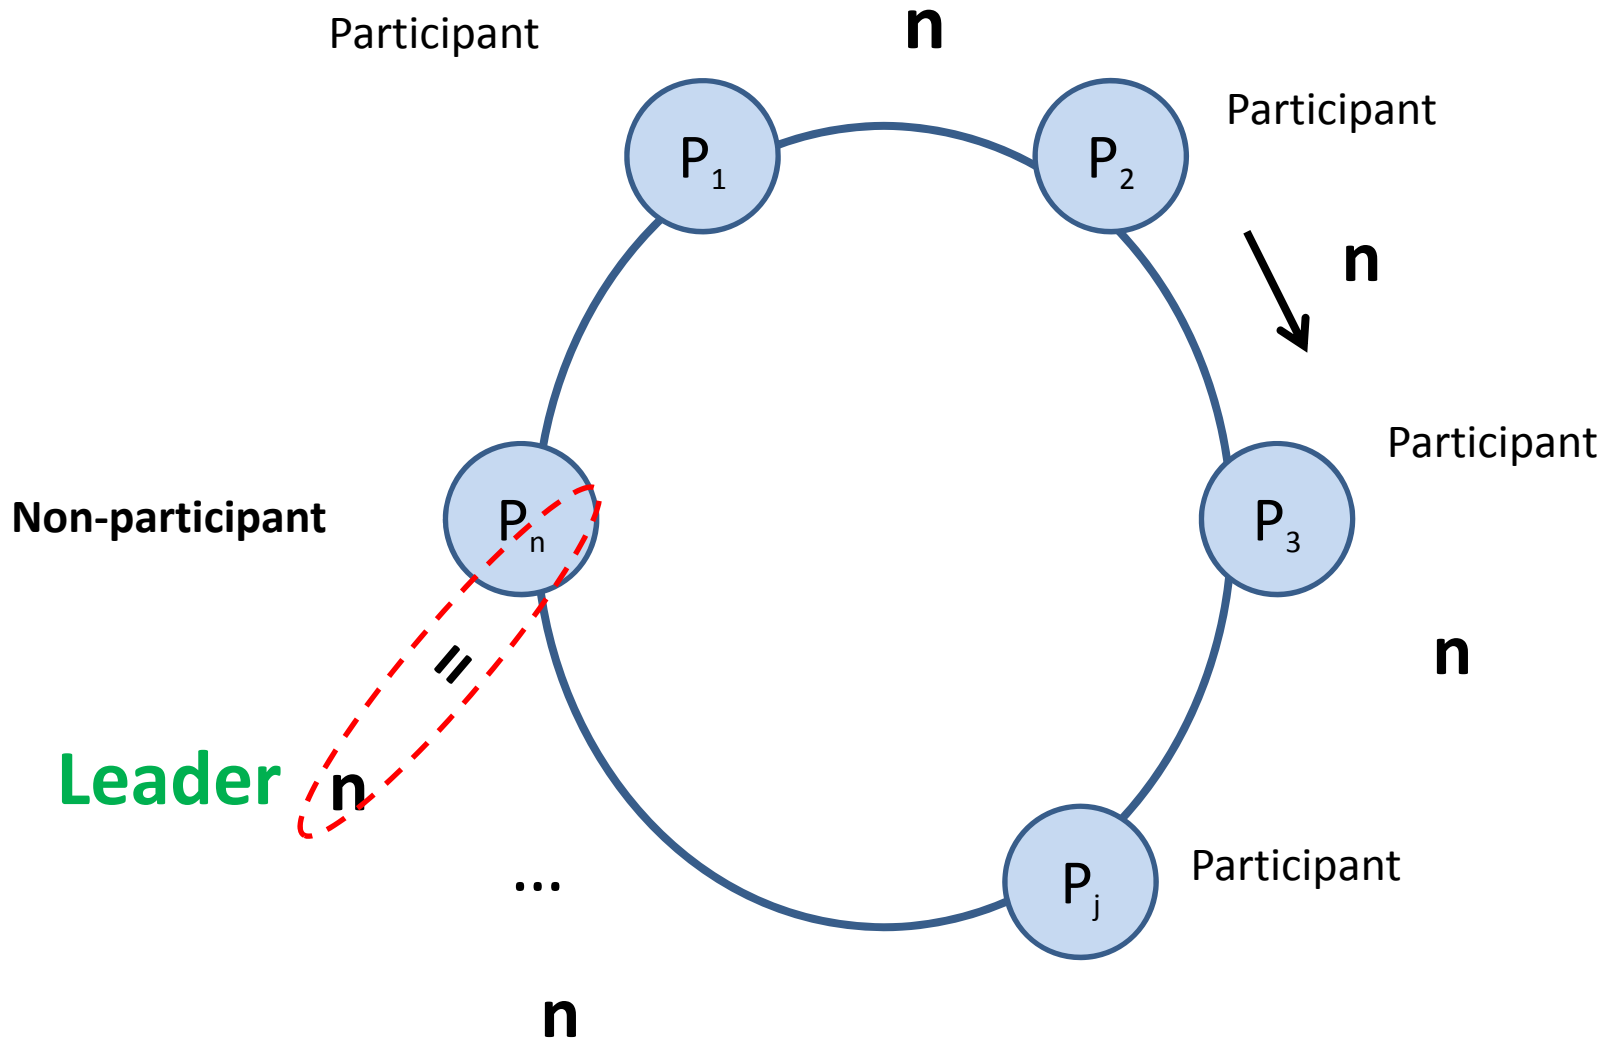
Ring-based algorithm: Calling an Election (**origin & victory**)



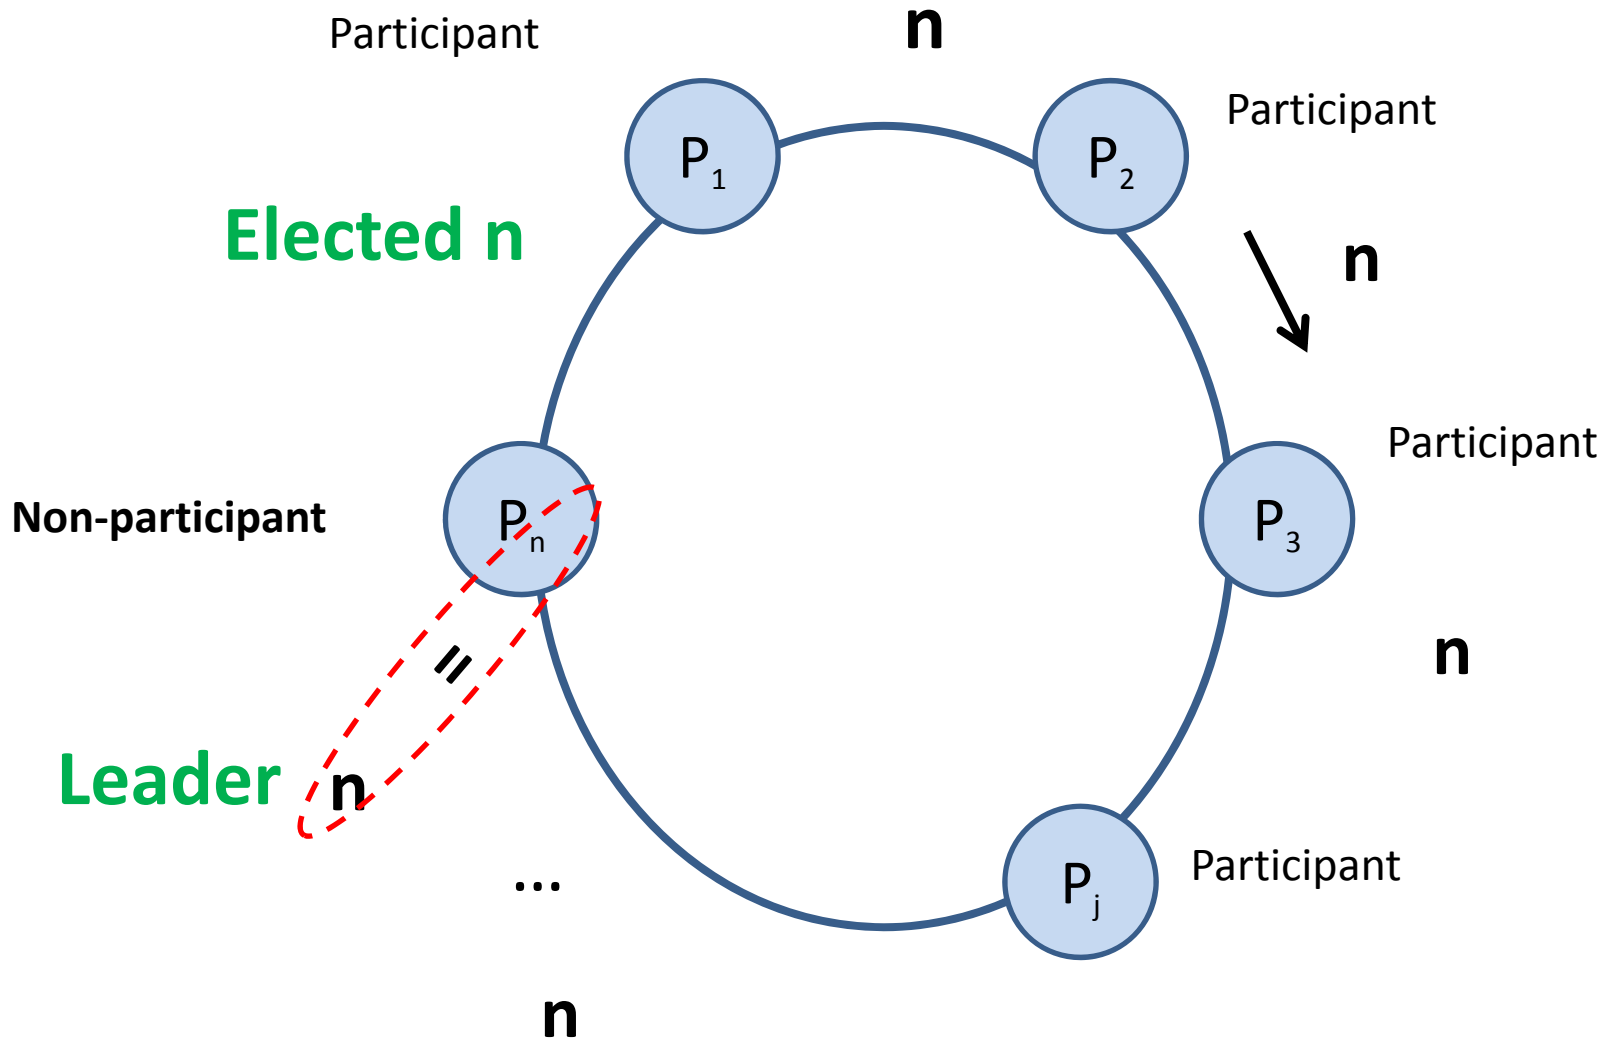
Ring-based algorithm: Calling an Election (**origin & victory**)



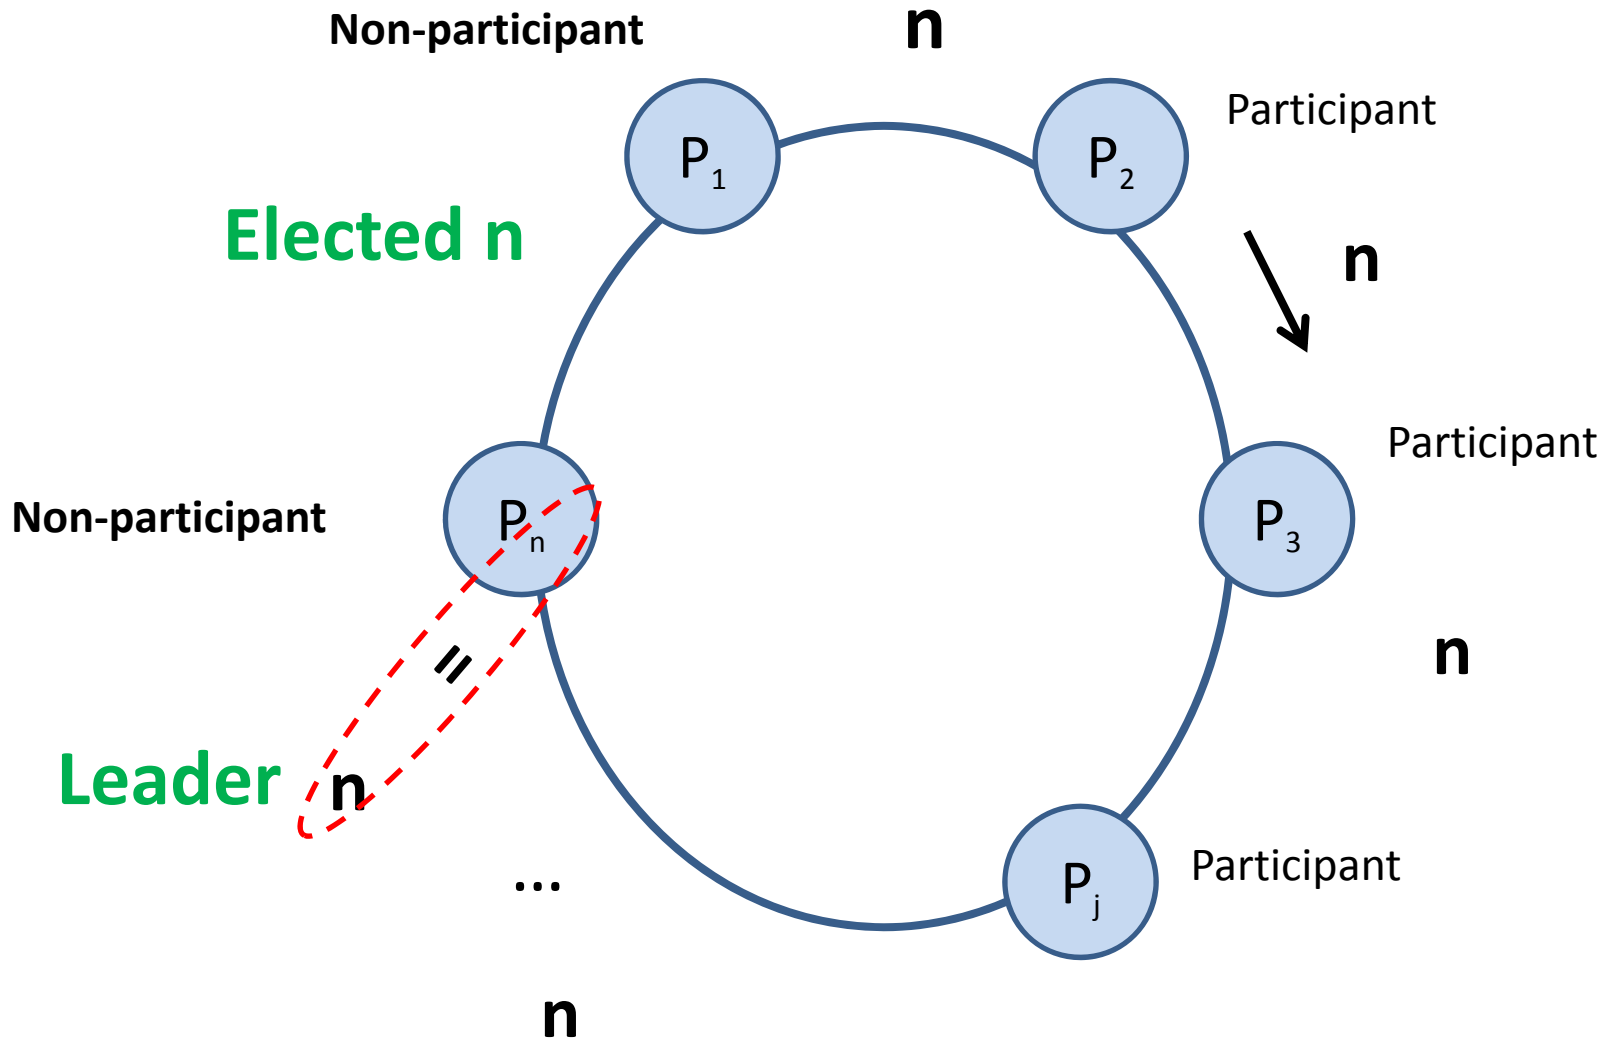
Ring-based algorithm: Calling an Election (**origin & victory**)



Ring-based algorithm: Calling an Election (**origin & victory**)

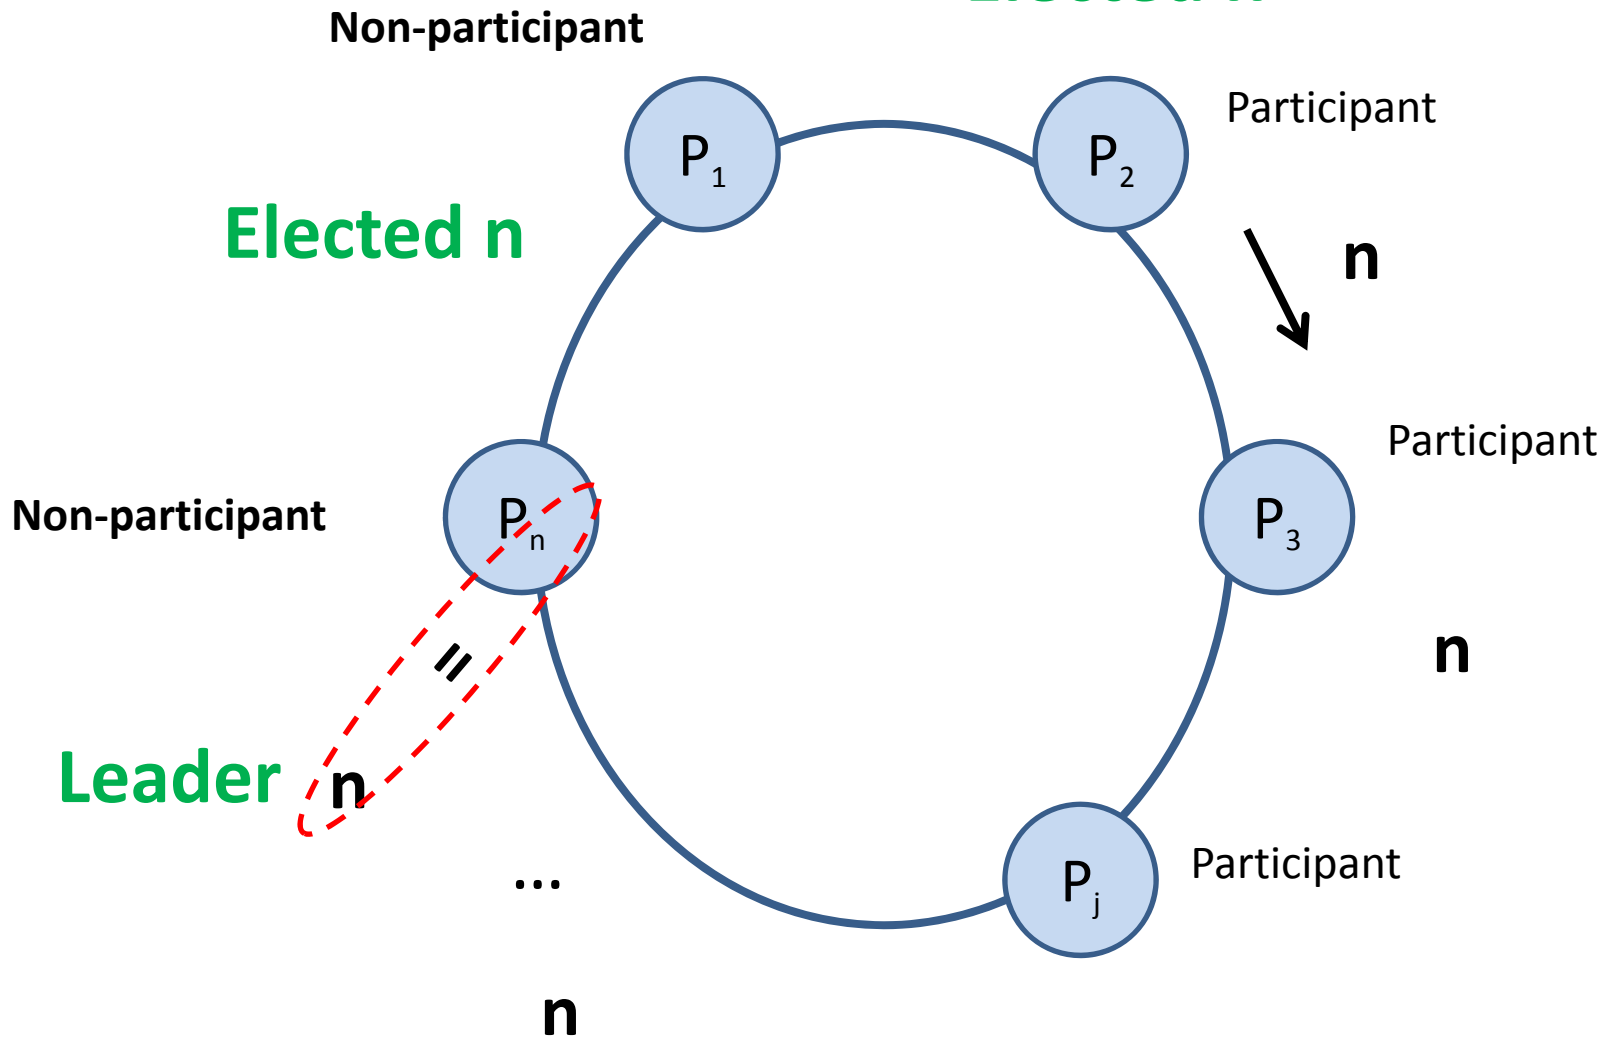


Ring-based algorithm: Calling an Election (**origin & victory**)



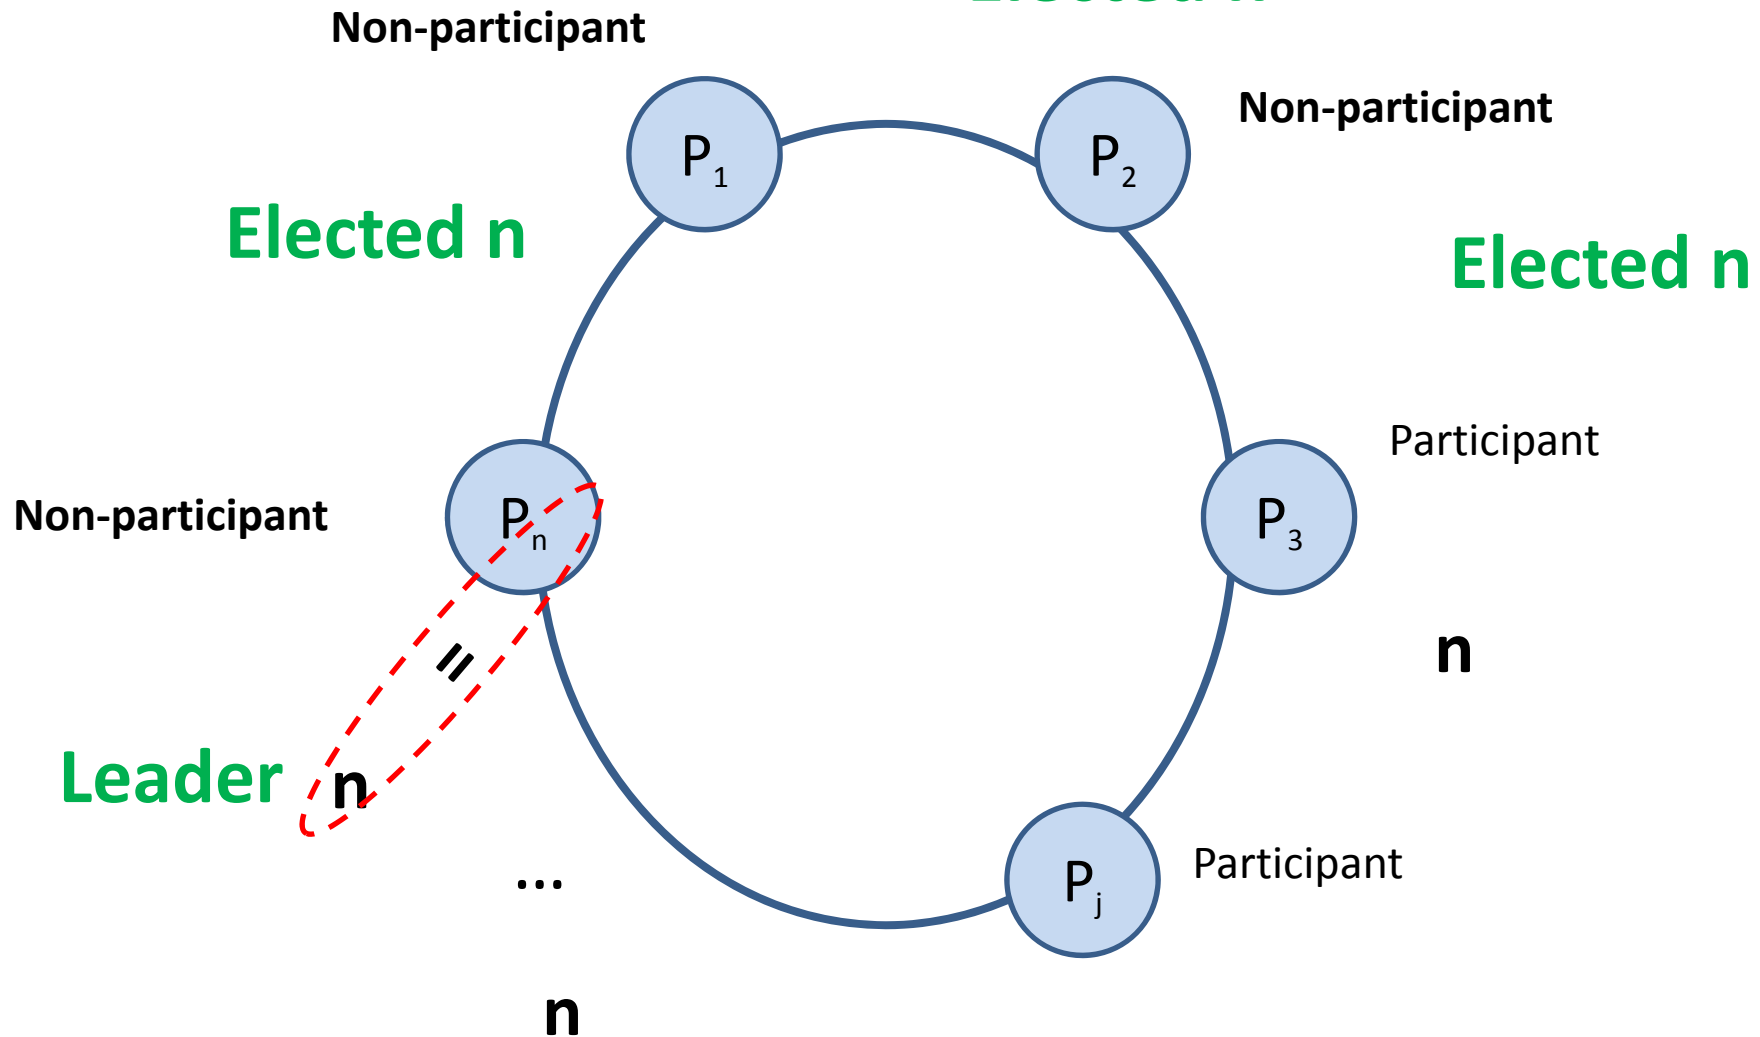
Ring-based algorithm: Calling an Election (**origin & victory**)

Elected n



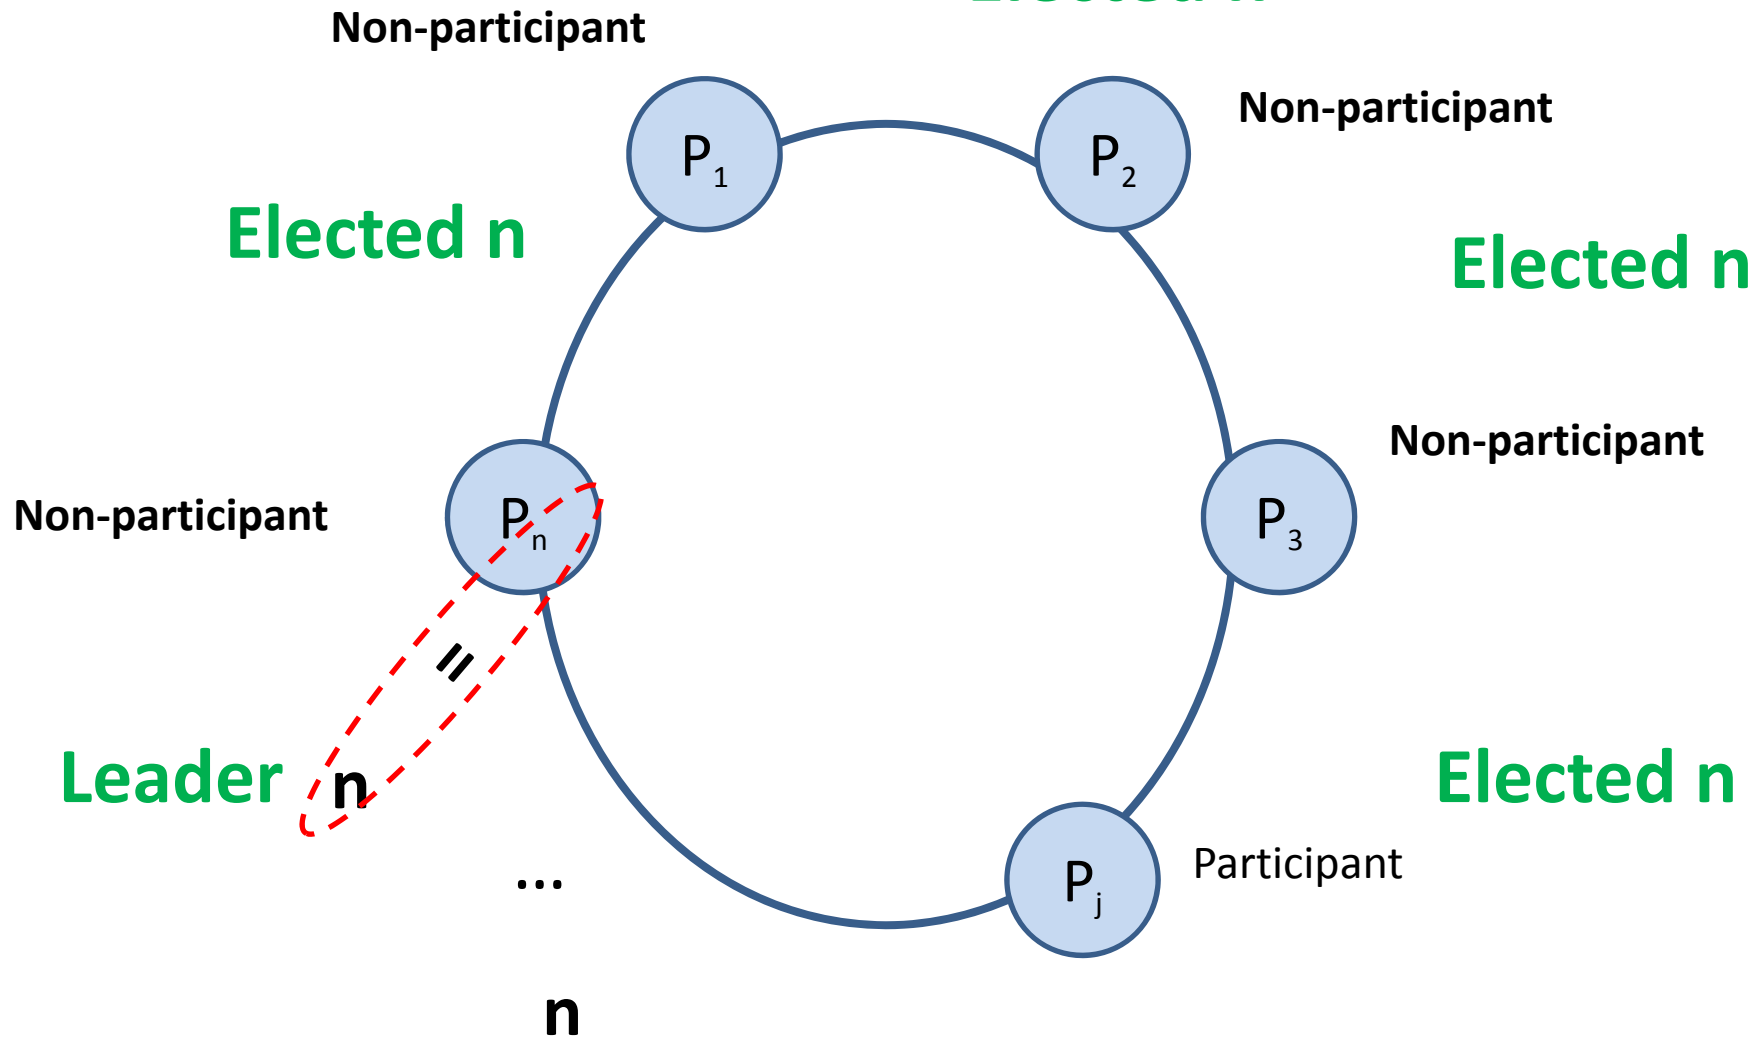
Ring-based algorithm: Calling an Election (**origin & victory**)

Elected n

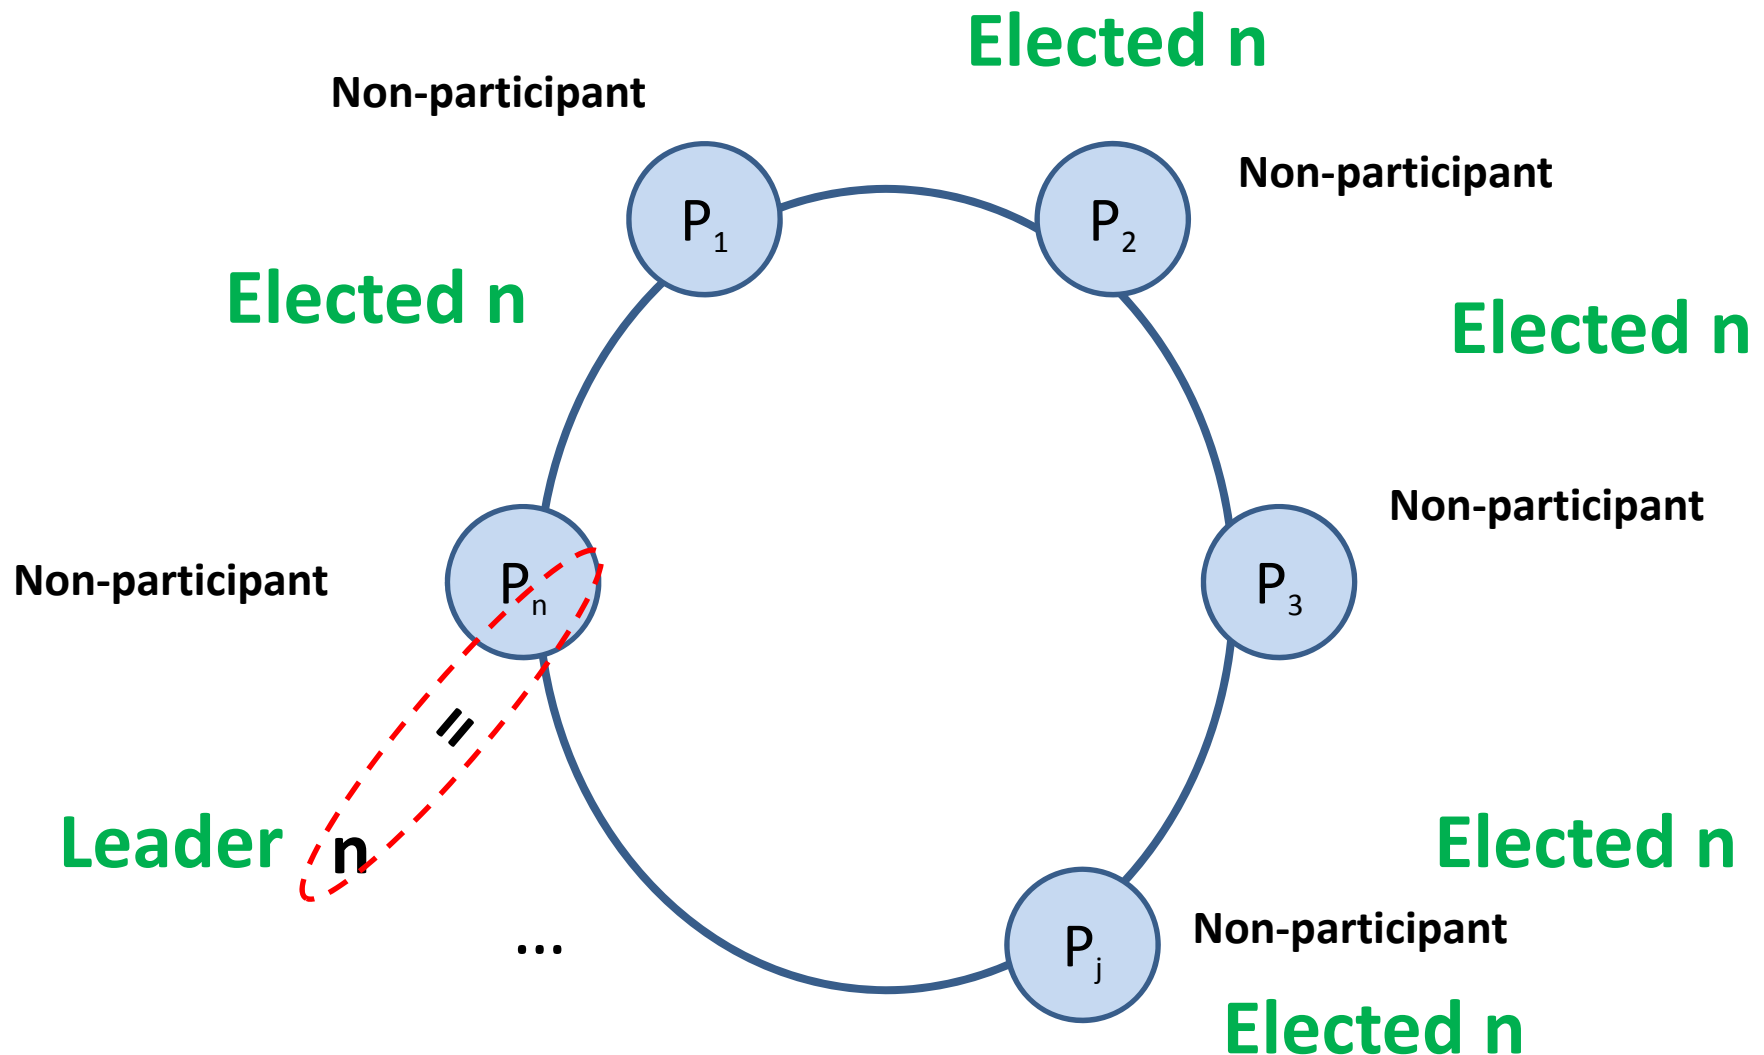


Ring-based algorithm: Calling an Election (**origin & victory**)

Elected n



Ring-based algorithm: Calling an Election (**origin & victory**)



Different cases

S (sender)

R (receiver)

$$S_{ID} > R_{ID}$$

Participant



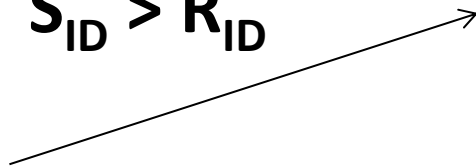
Different cases

S (sender)

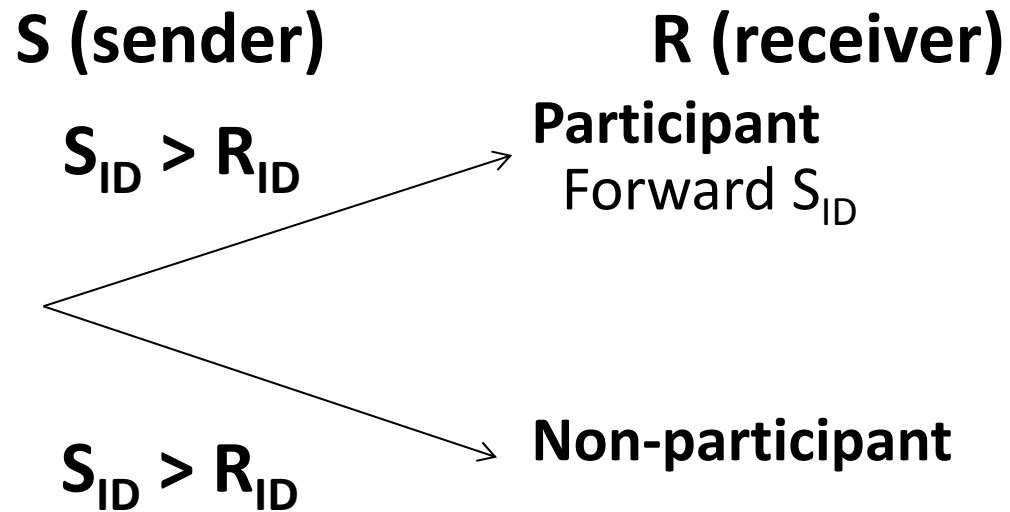
$S_{ID} > R_{ID}$

R (receiver)

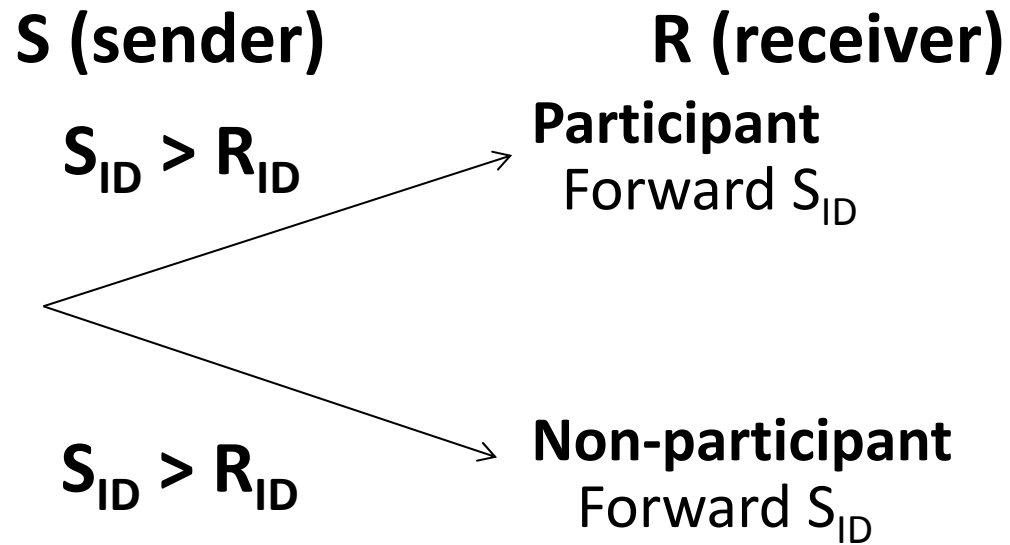
Participant
Forward S_{ID}



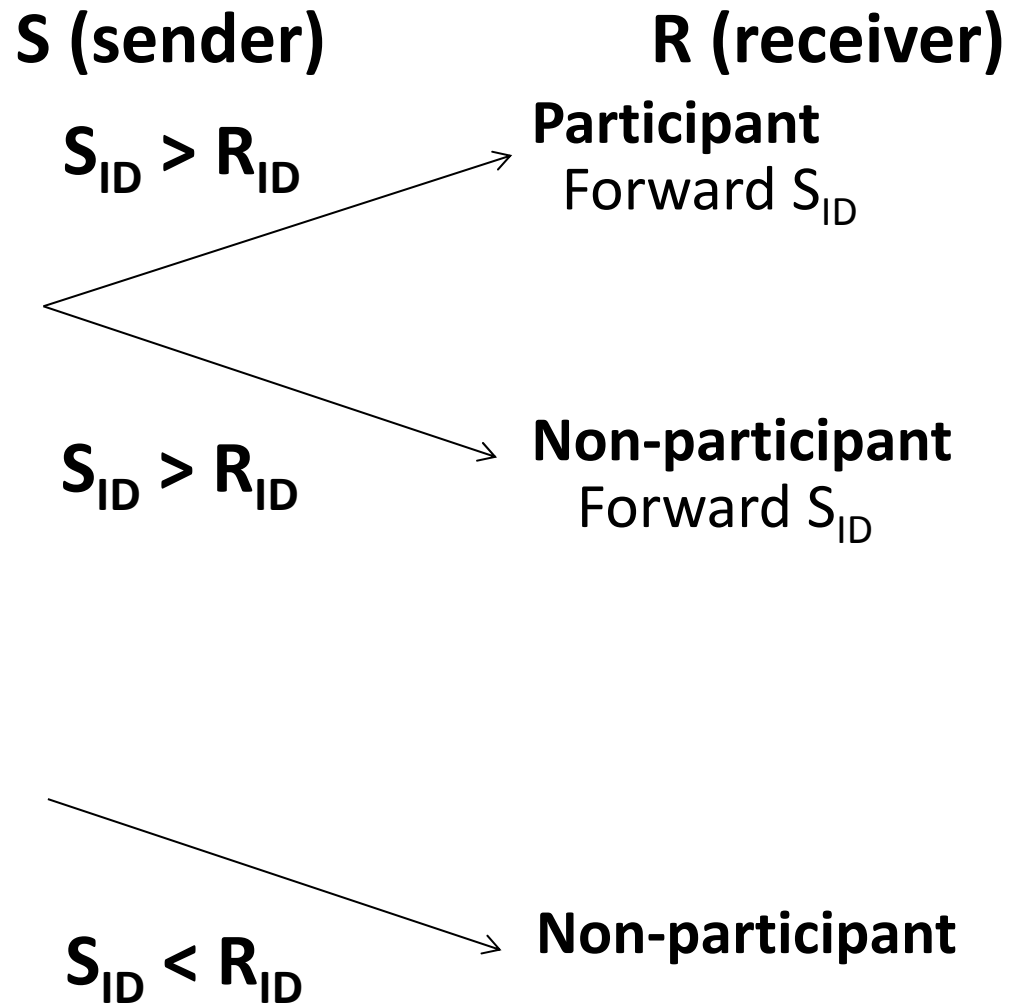
Different cases



Different cases



Different cases



Different cases

S (sender)

R (receiver)

$S_{ID} > R_{ID}$

Participant
Forward S_{ID}

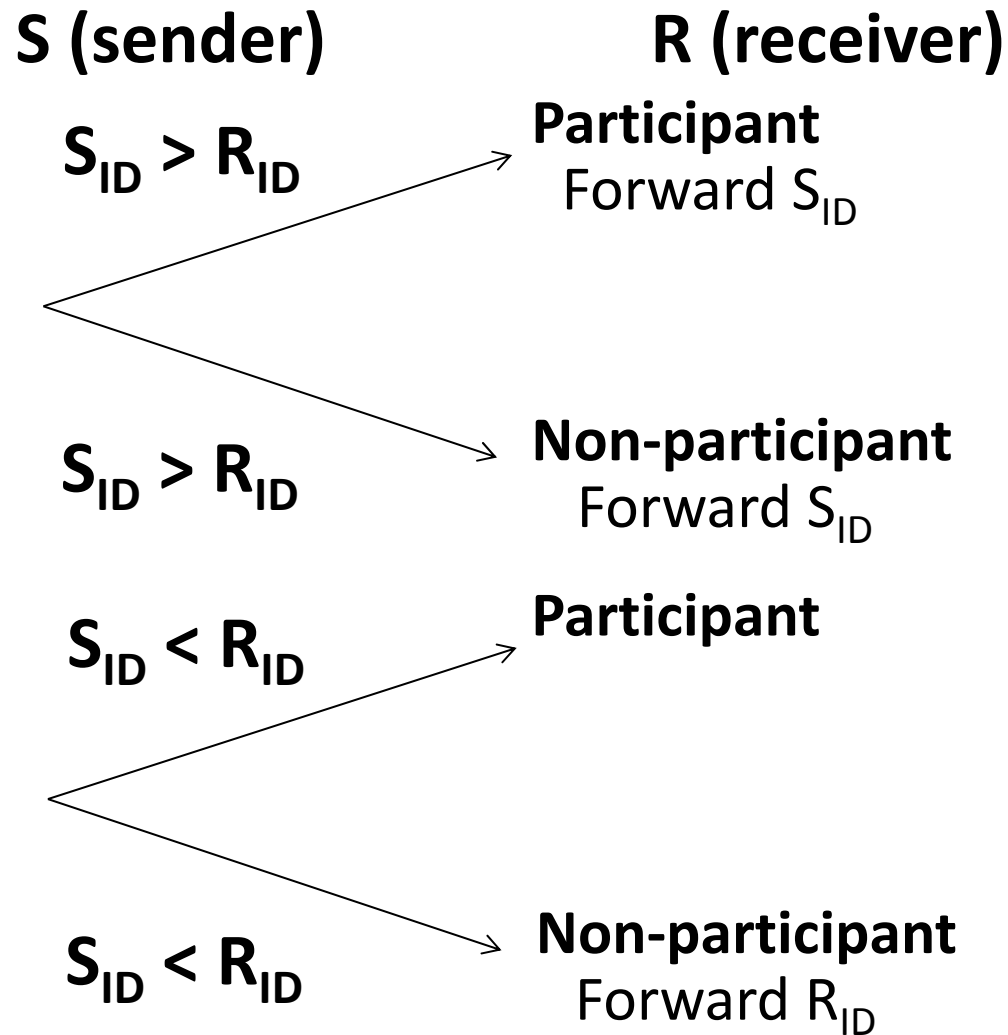
$S_{ID} > R_{ID}$

Non-participant
Forward S_{ID}

$S_{ID} < R_{ID}$

Non-participant
Forward R_{ID}

Different cases



Different cases

S (sender)

R (receiver)

$S_{ID} > R_{ID}$

Participant
Forward S_{ID}

$S_{ID} > R_{ID}$

Non-participant
Forward S_{ID}

$S_{ID} < R_{ID}$

Participant
No forwarding (own ID already sent)

$S_{ID} < R_{ID}$

Non-participant
Forward R_{ID}

Different cases

S (sender)

R (receiver)

$S_{ID} > R_{ID}$

Participant
Forward S_{ID}

If $S_{ID} = R_{ID}$, it follows
R elected as leader

$S_{ID} > R_{ID}$

Non-participant
Forward S_{ID}

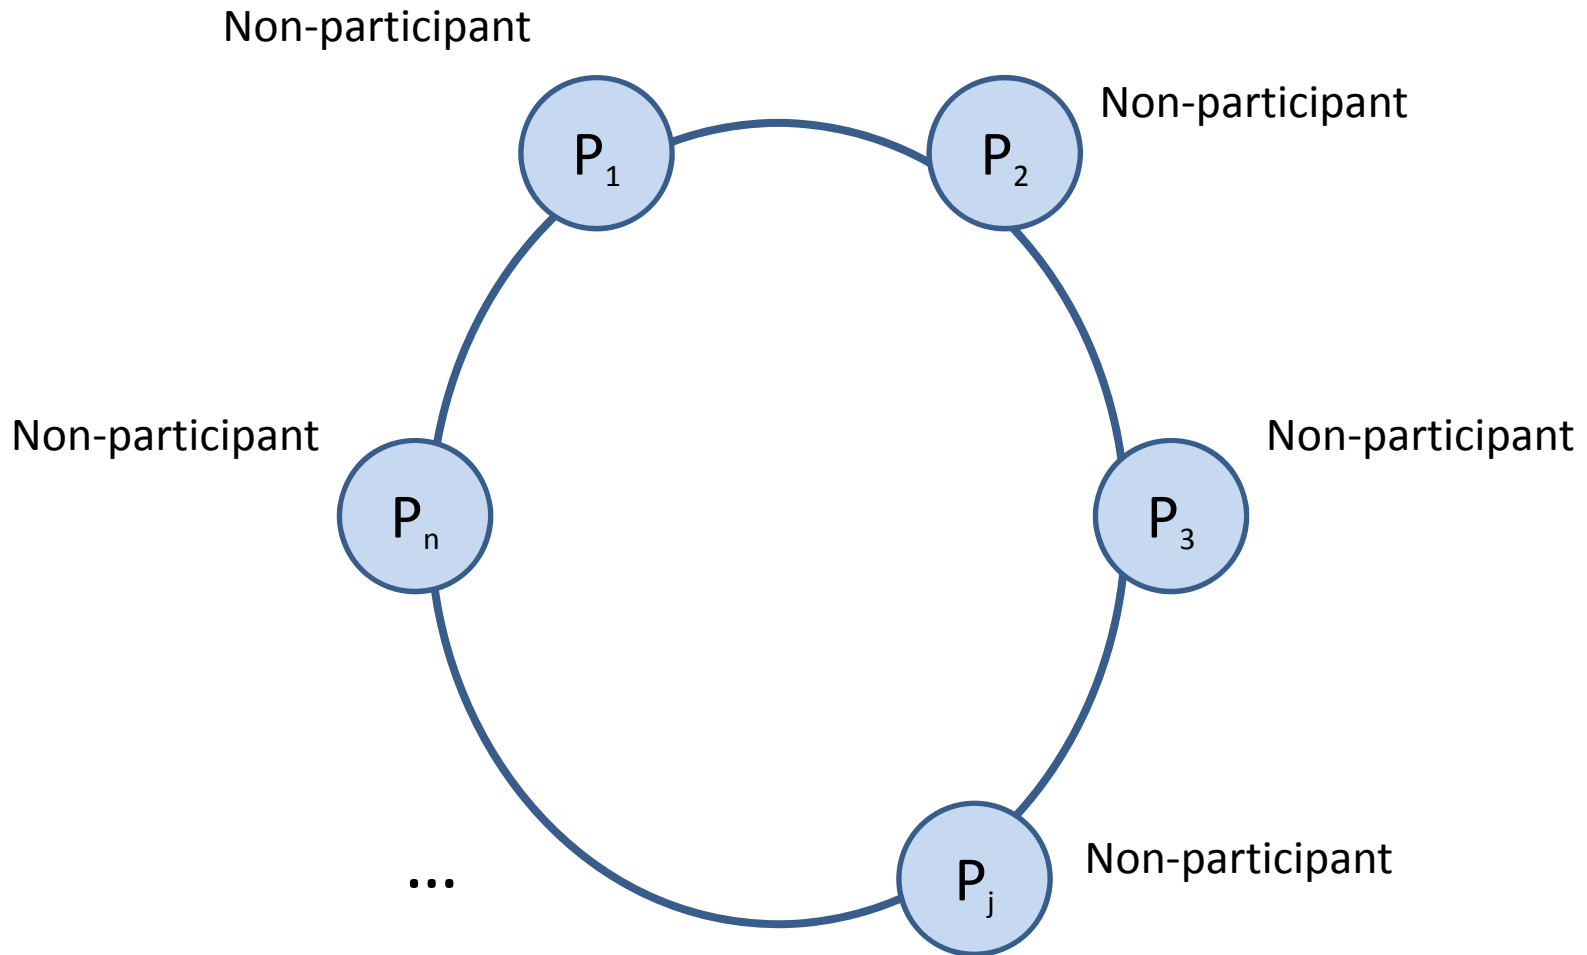
$S_{ID} < R_{ID}$

Participant
No forwarding (own ID already sent)

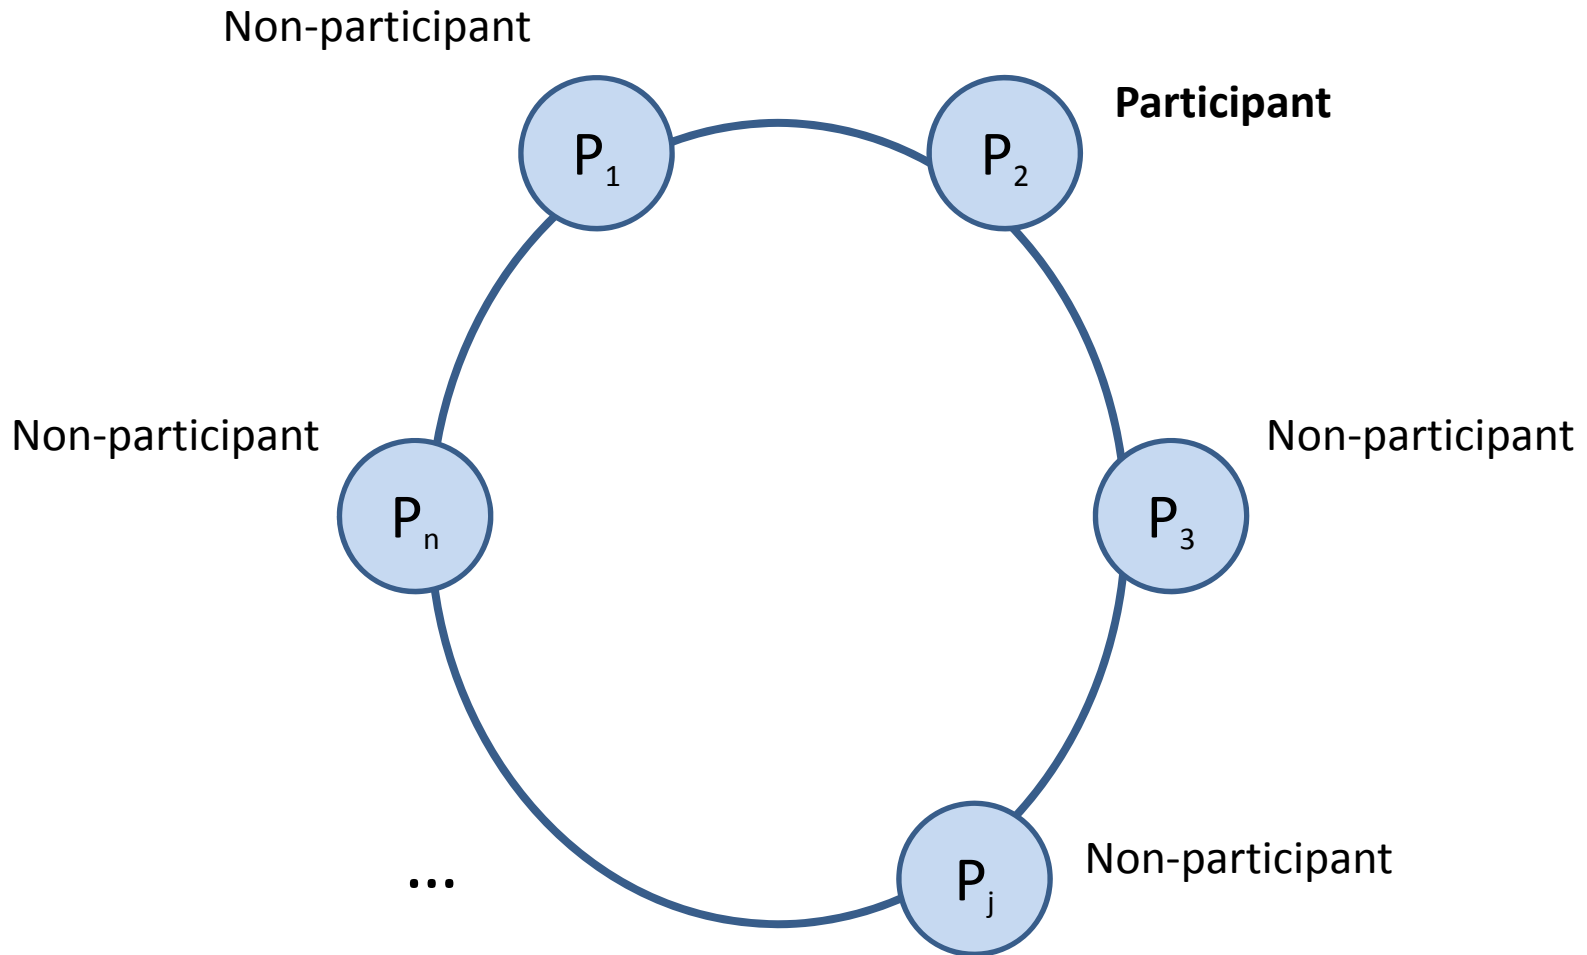
$S_{ID} < R_{ID}$

Non-participant
Forward R_{ID}

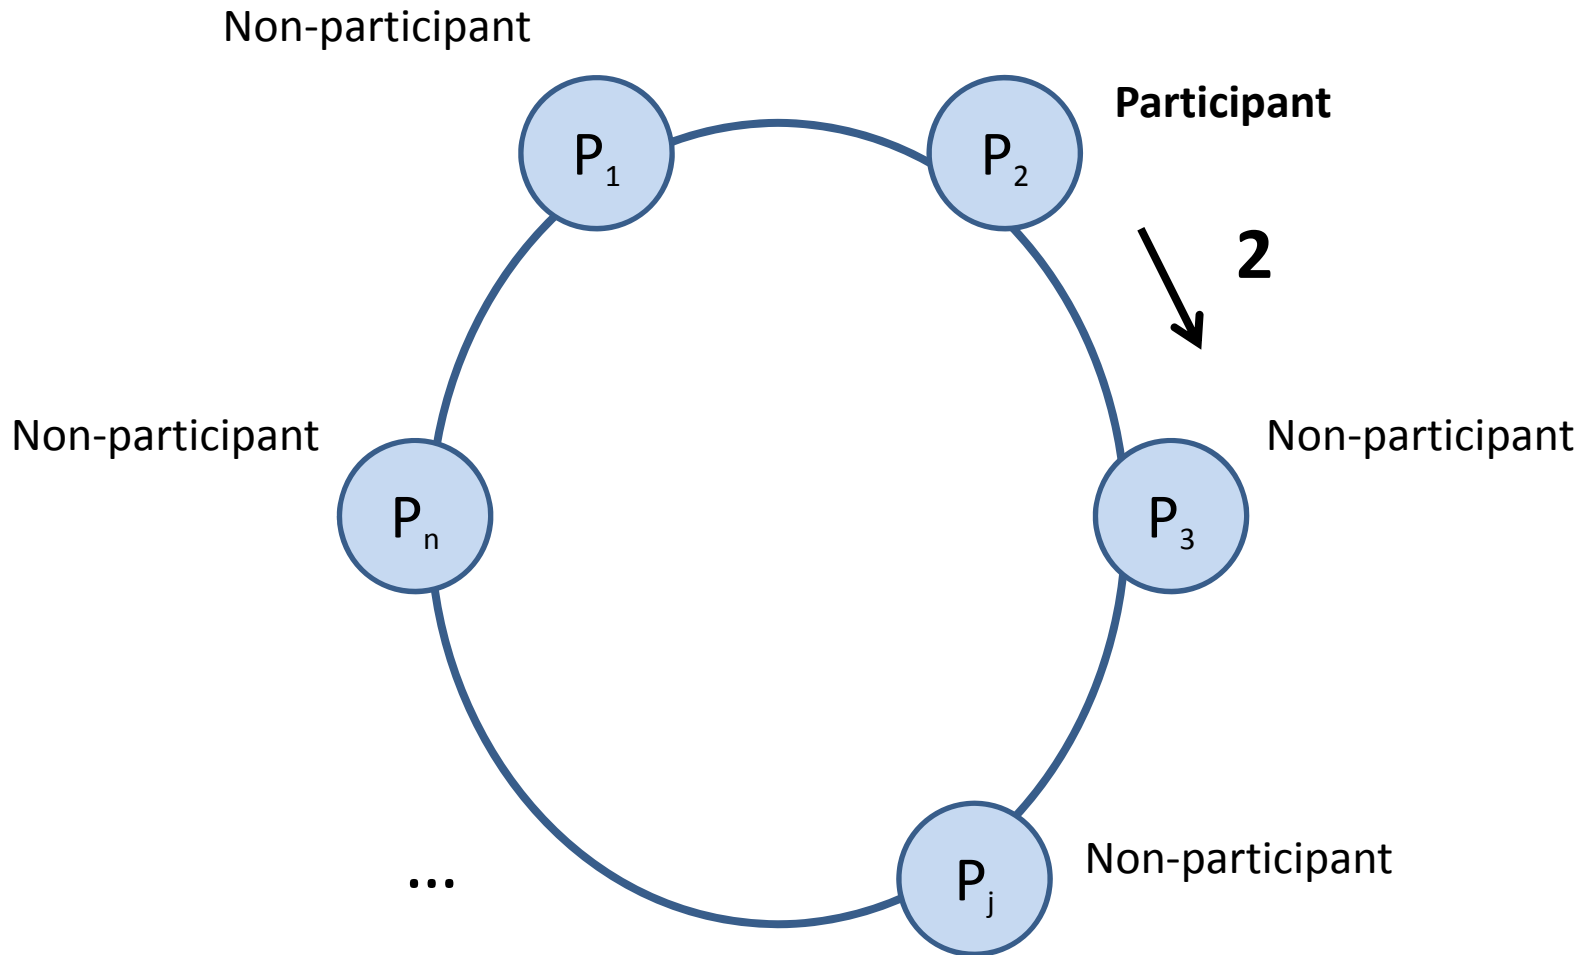
Ring-based algorithm: Concurrent election start



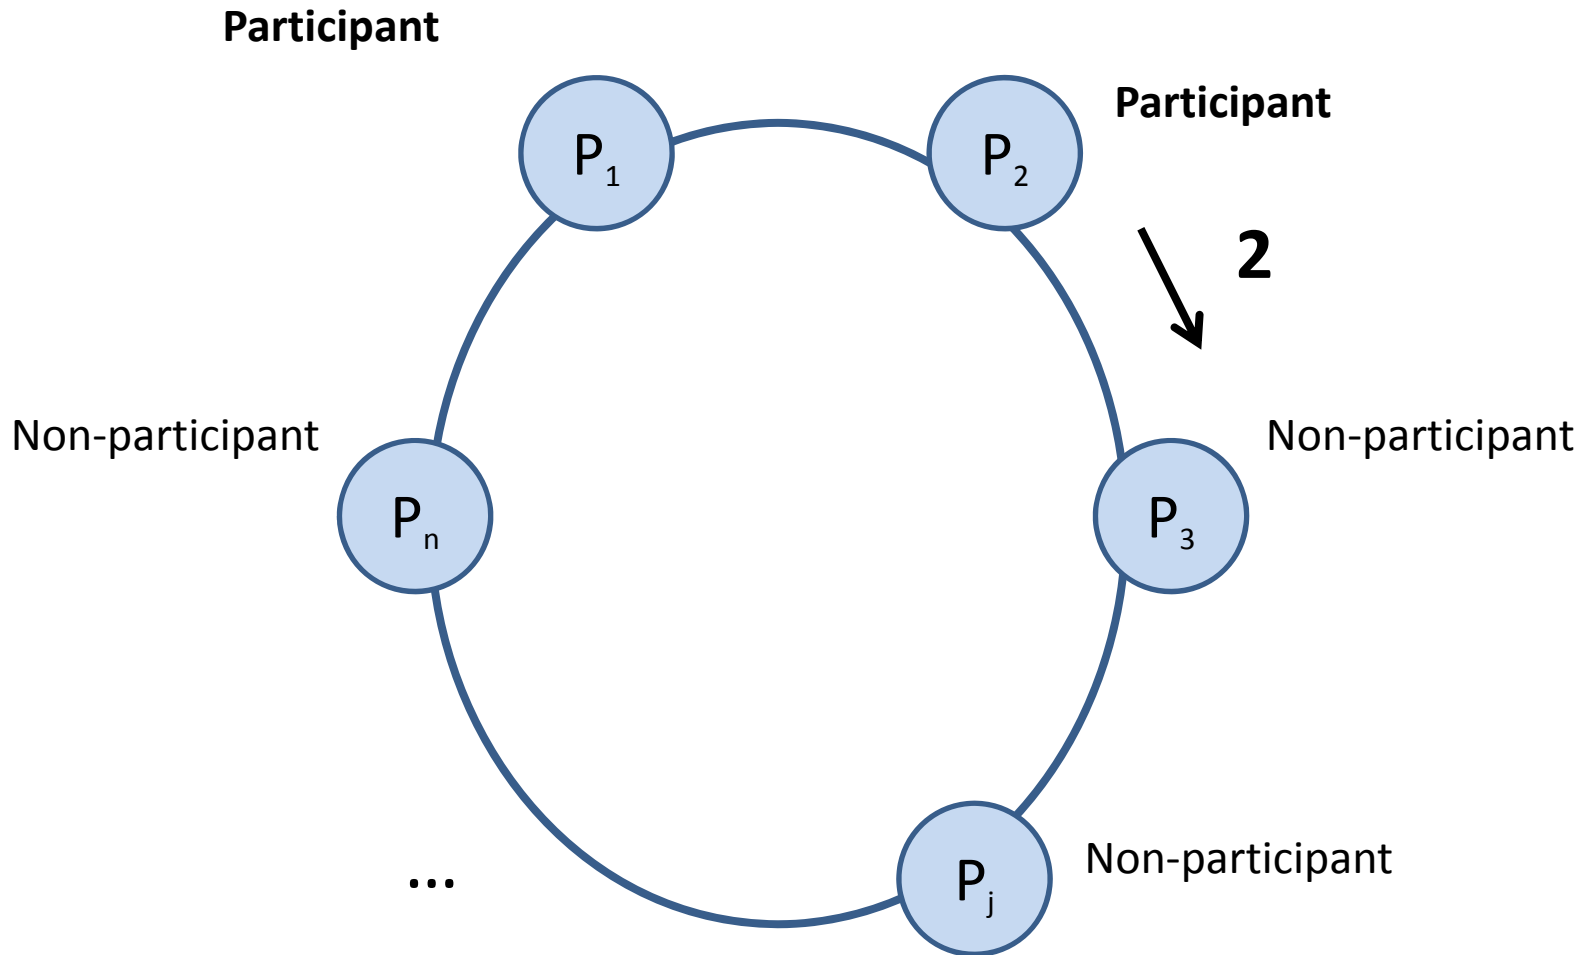
Ring-based algorithm: Concurrent election start



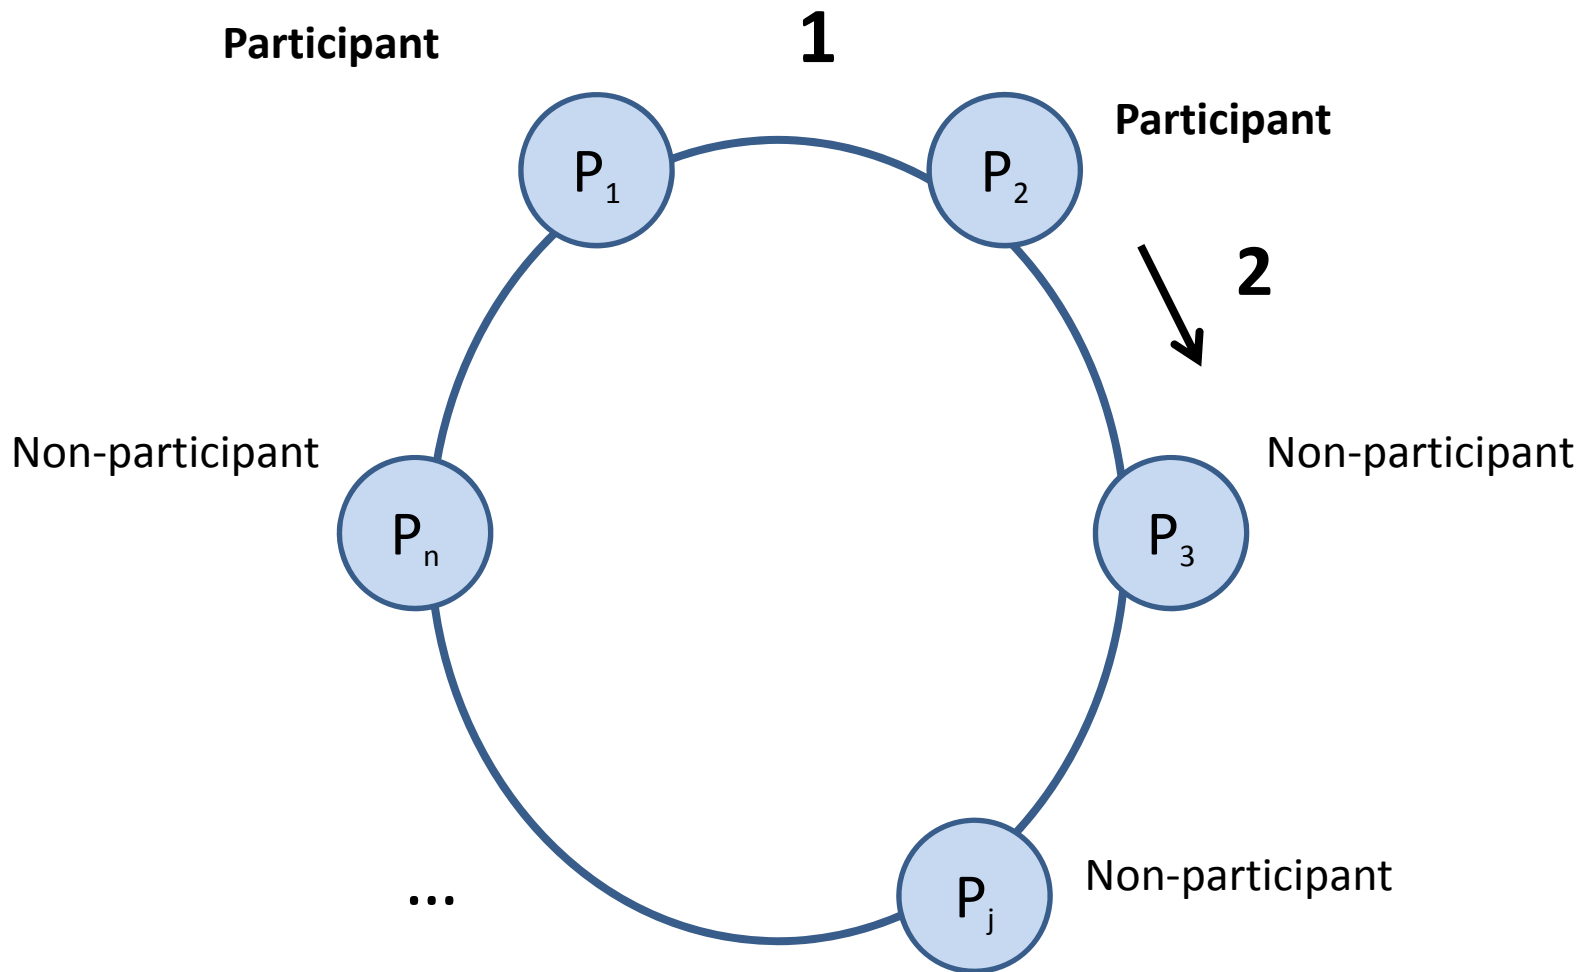
Ring-based algorithm: Concurrent election start



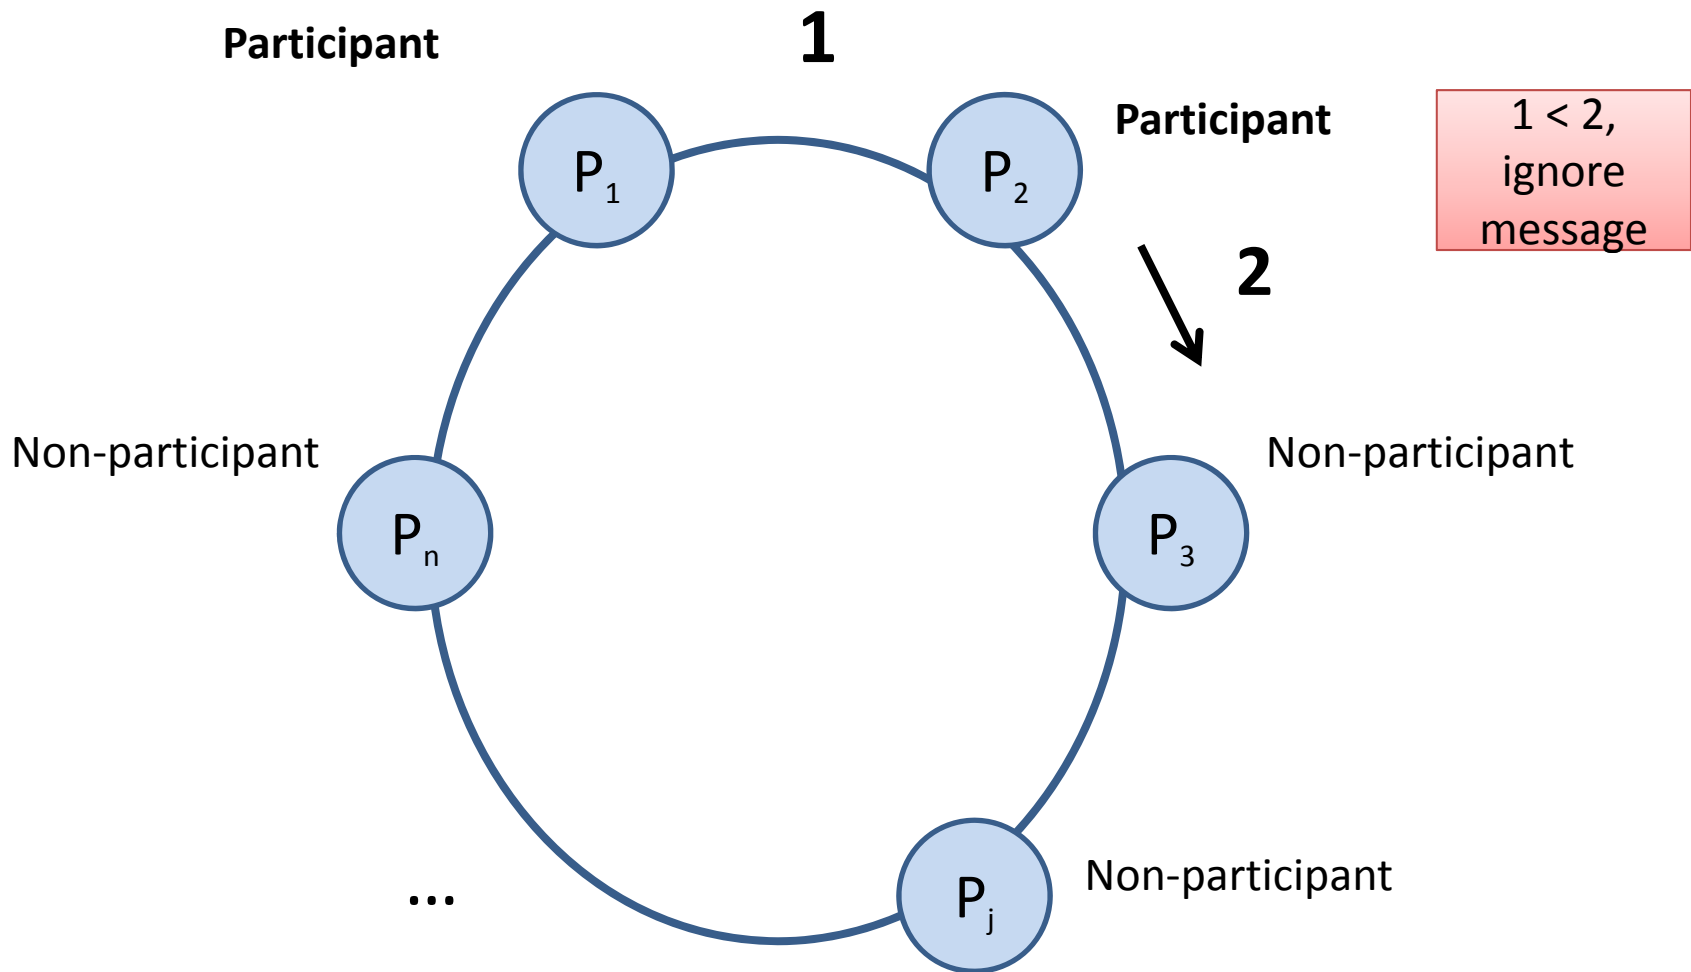
Ring-based algorithm: Concurrent election start



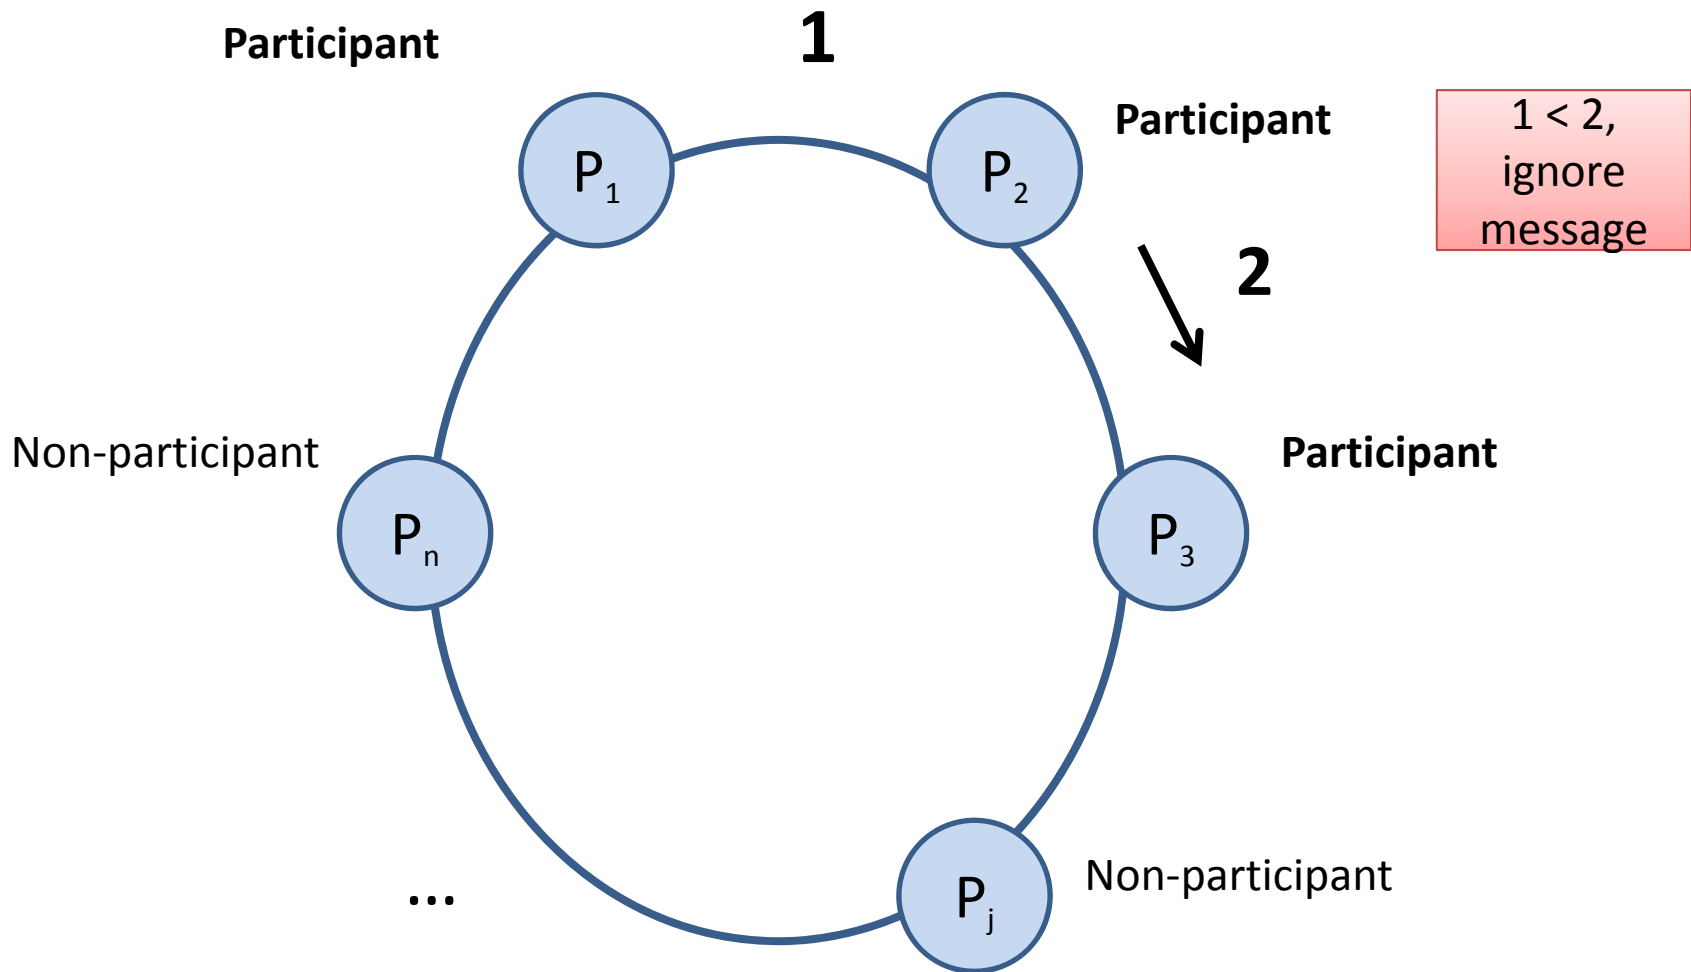
Ring-based algorithm: Concurrent election start



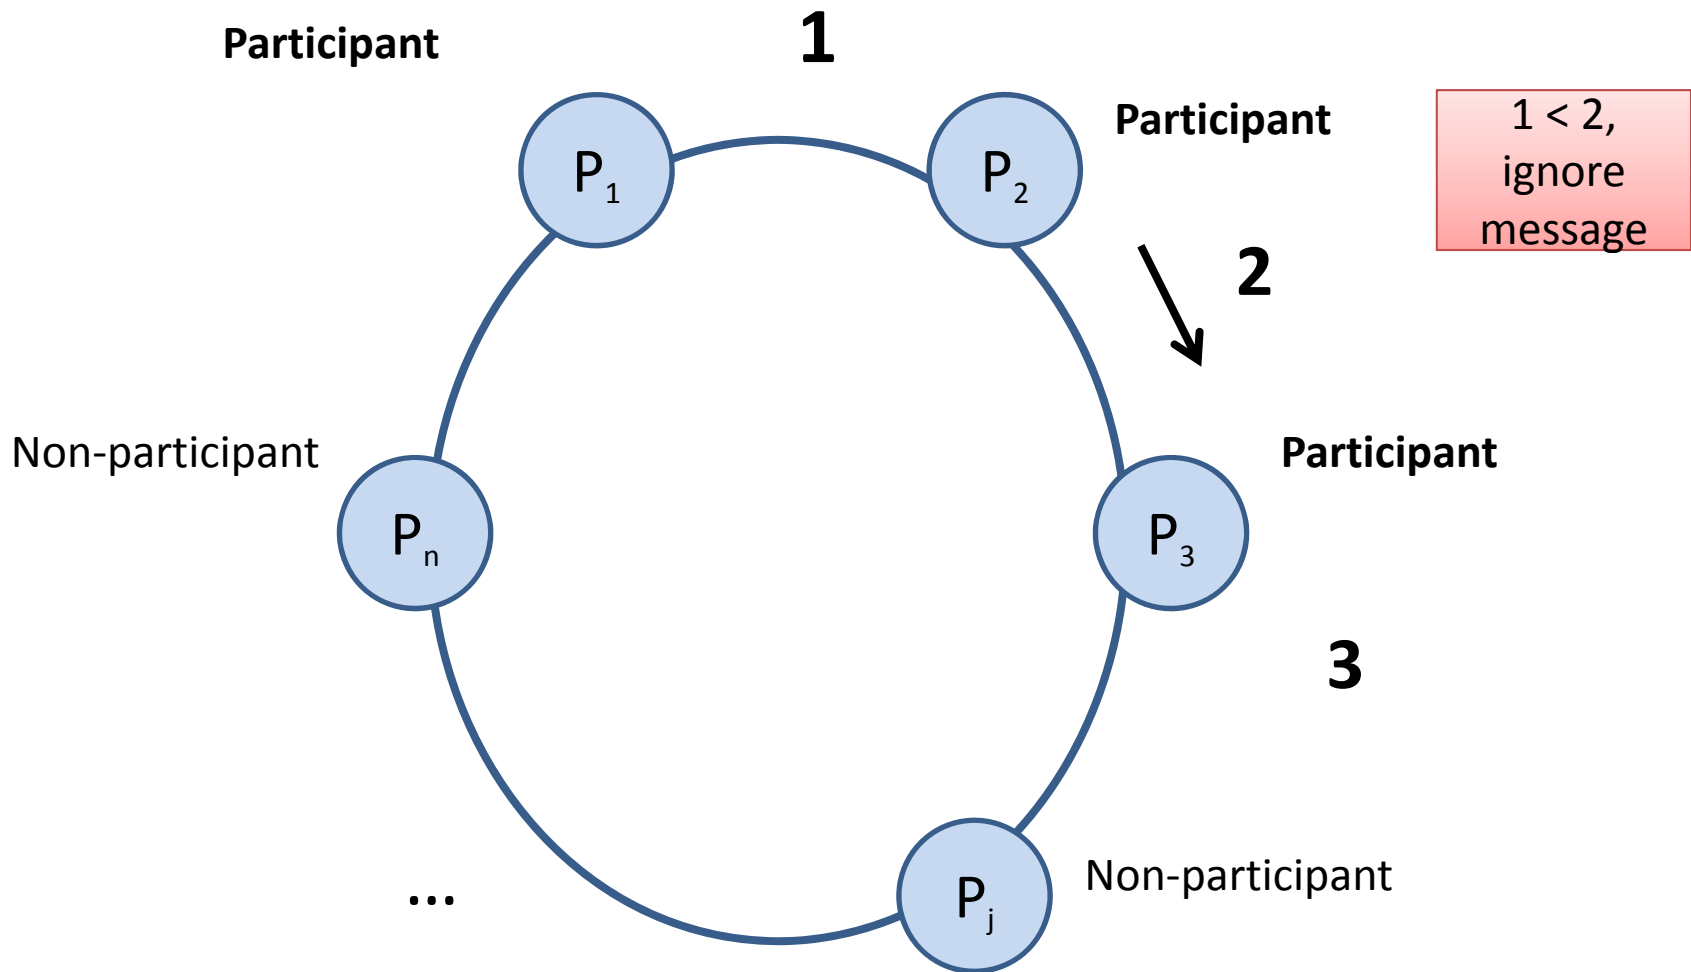
Ring-based algorithm: Concurrent election start



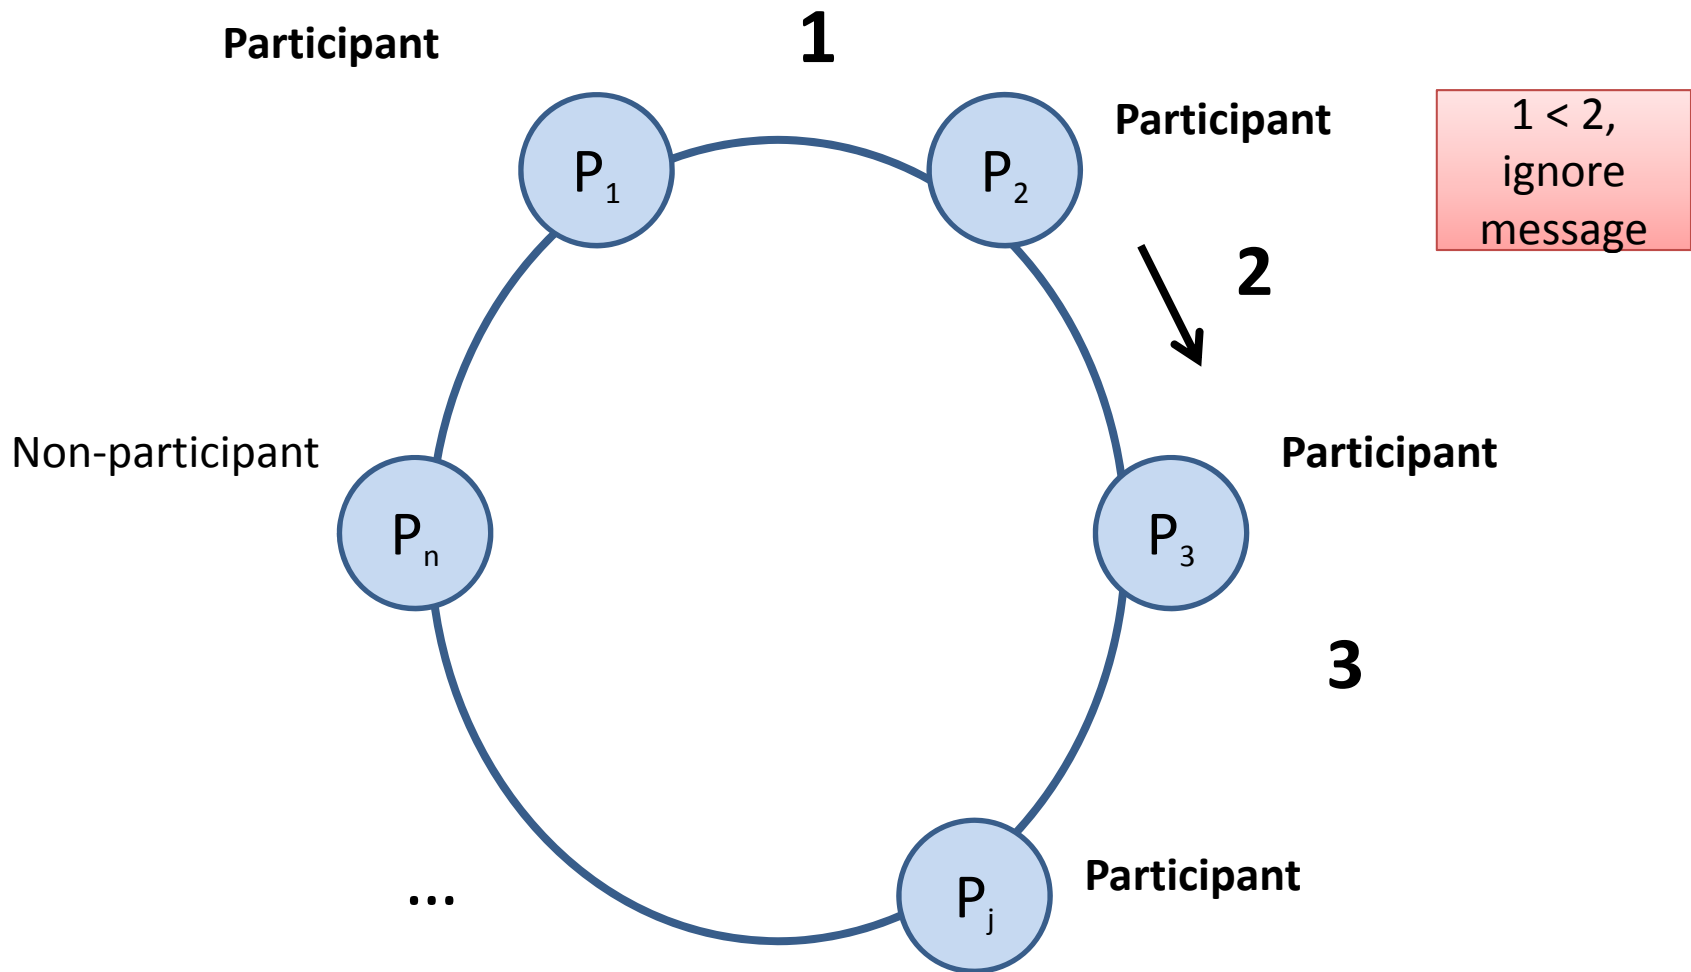
Ring-based algorithm: Concurrent election start



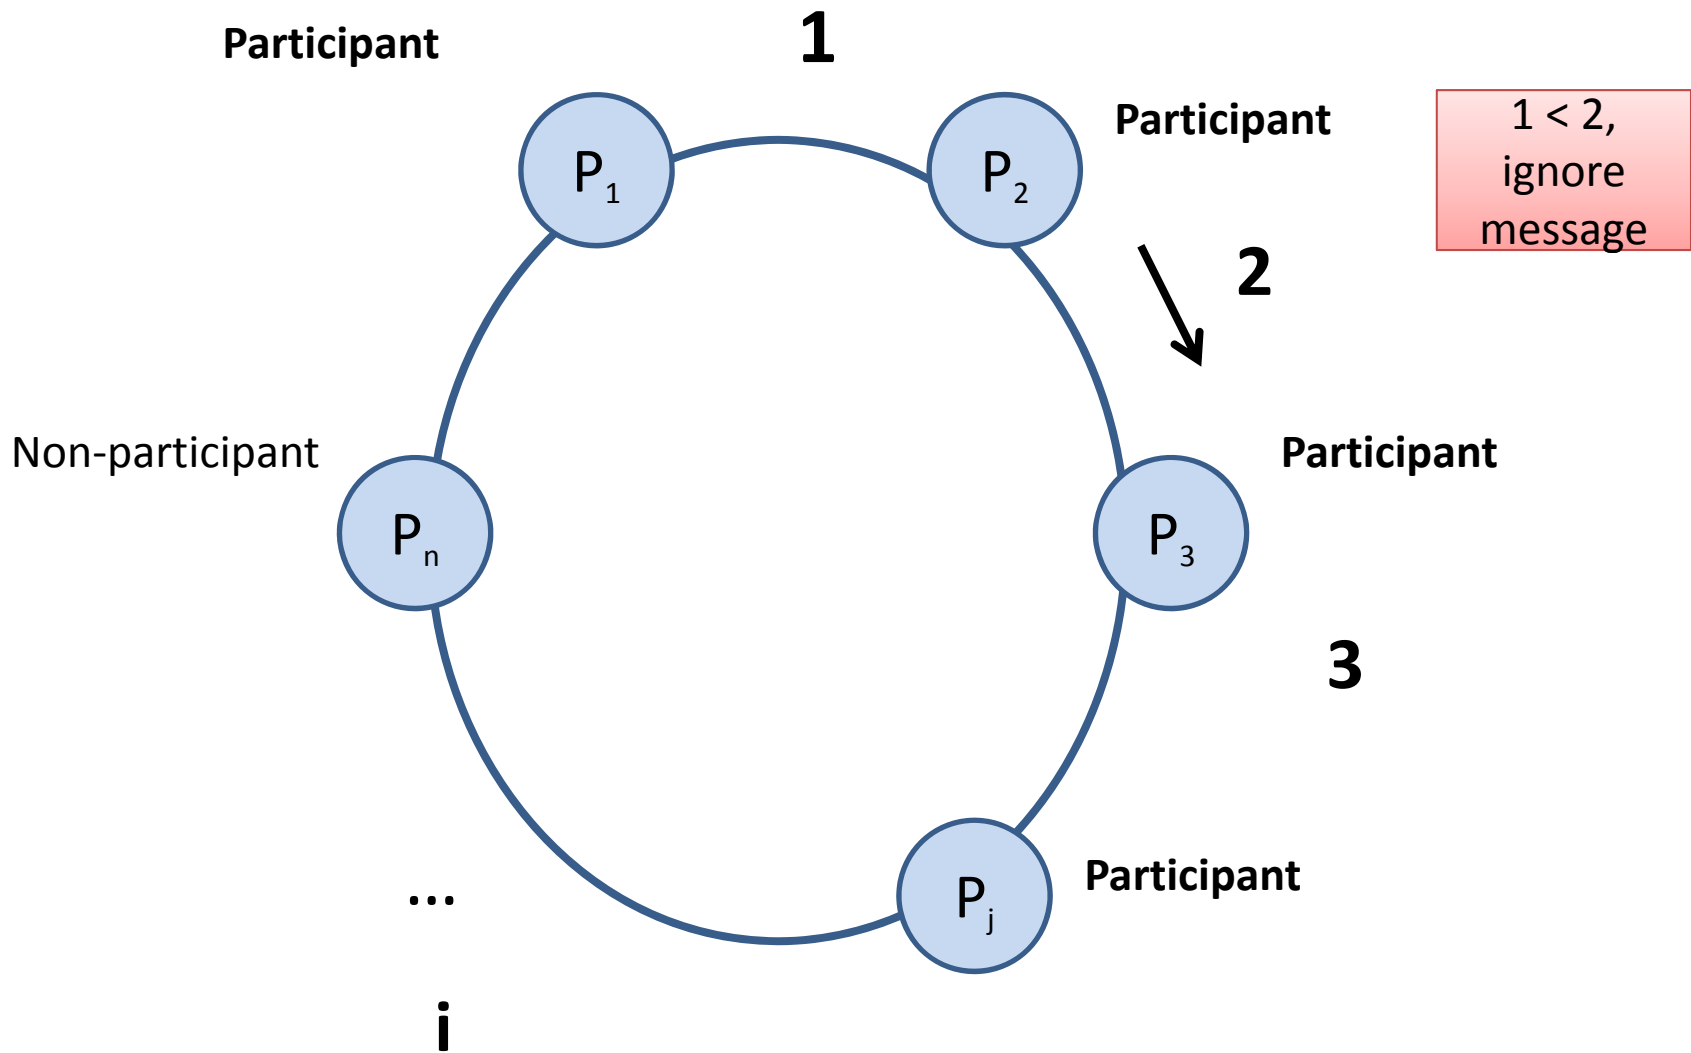
Ring-based algorithm: Concurrent election start



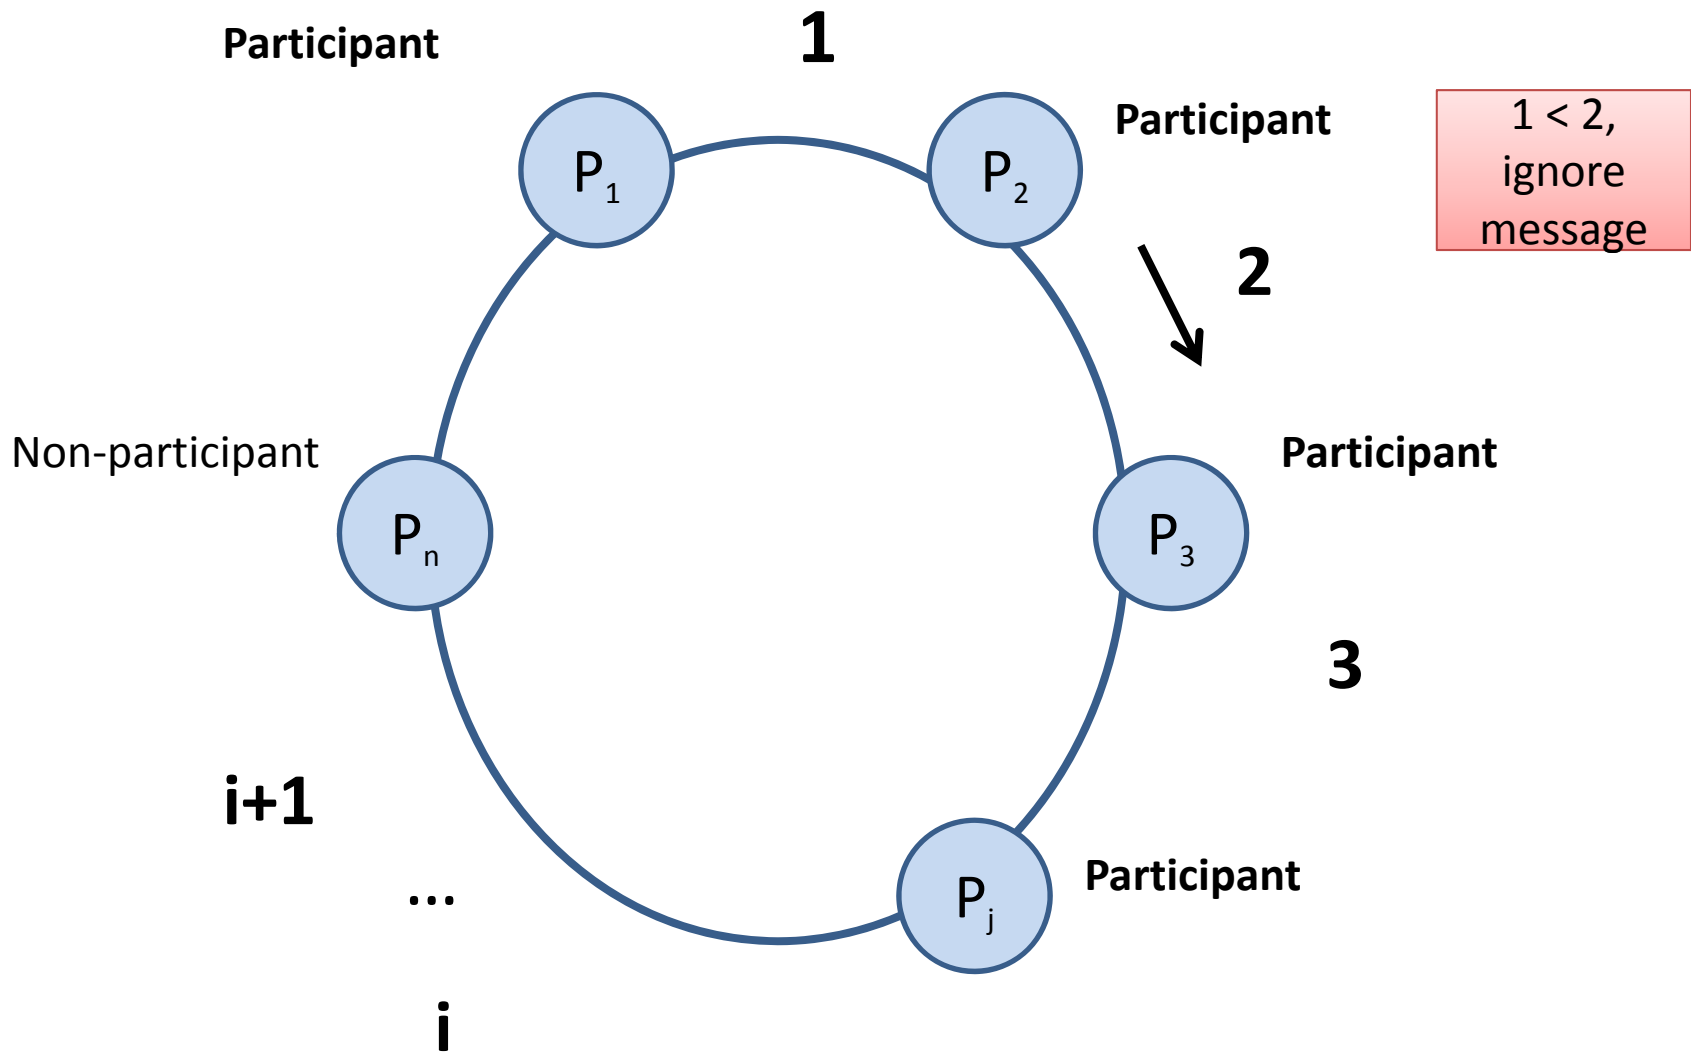
Ring-based algorithm: Concurrent election start



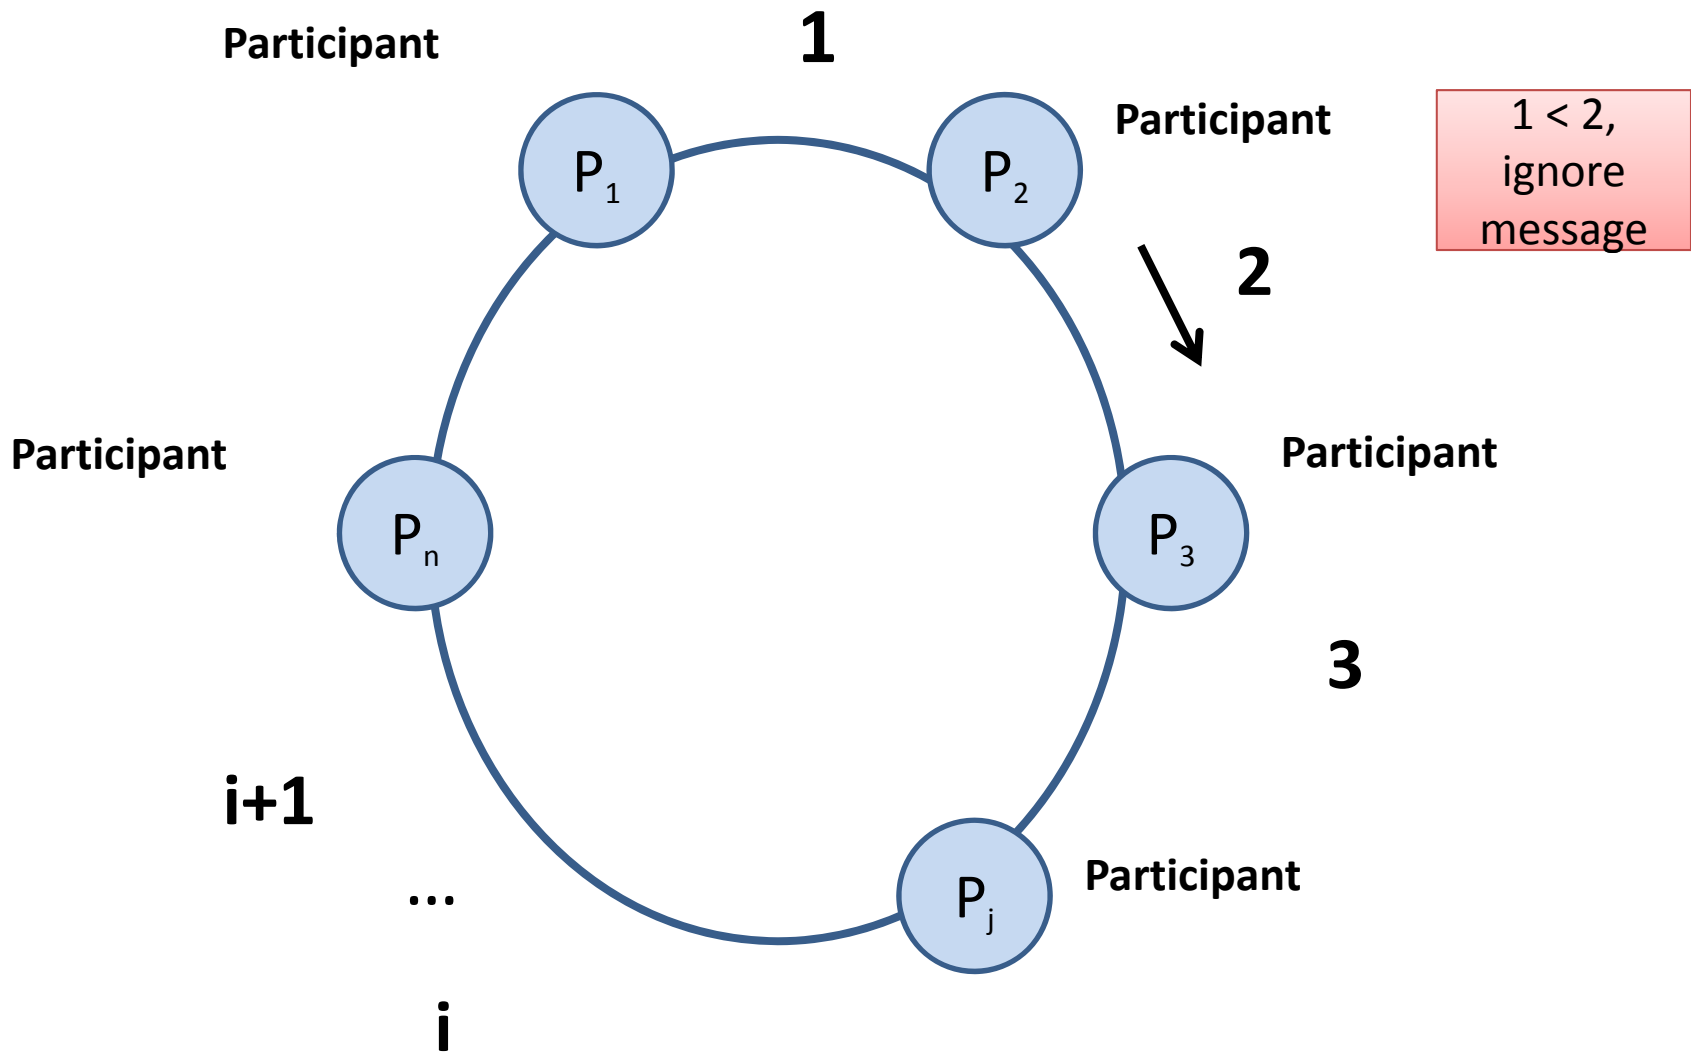
Ring-based algorithm: Concurrent election start



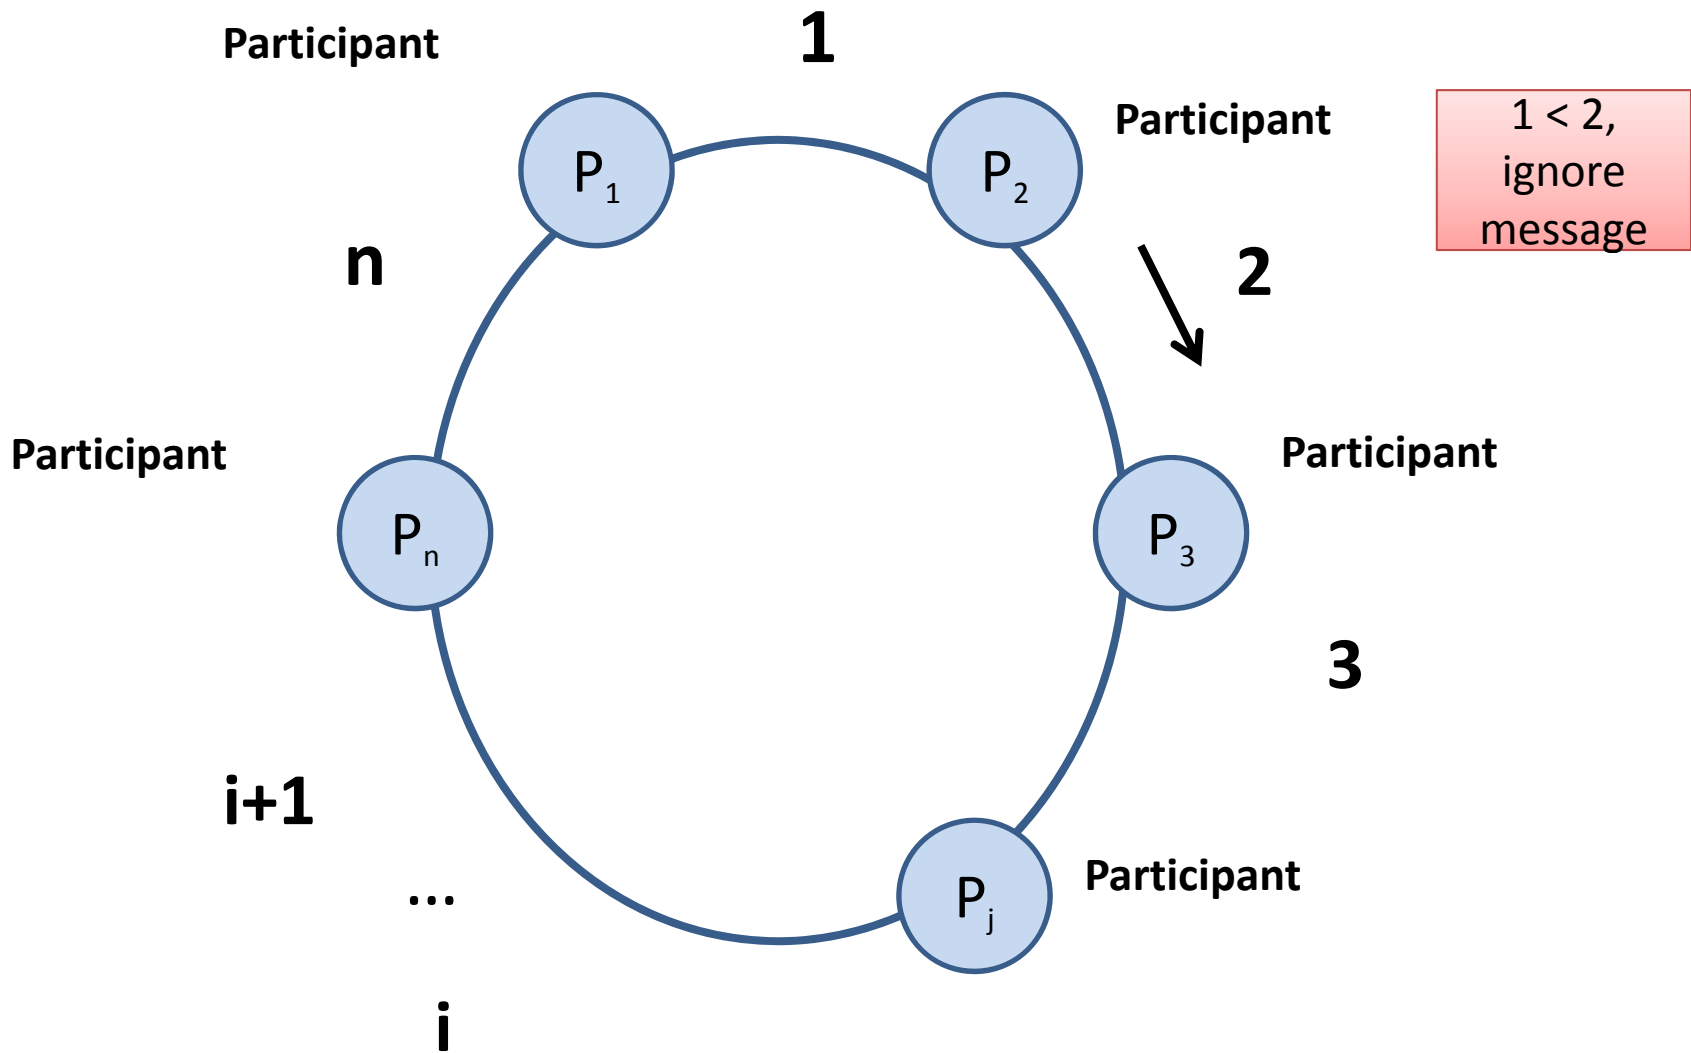
Ring-based algorithm: Concurrent election start



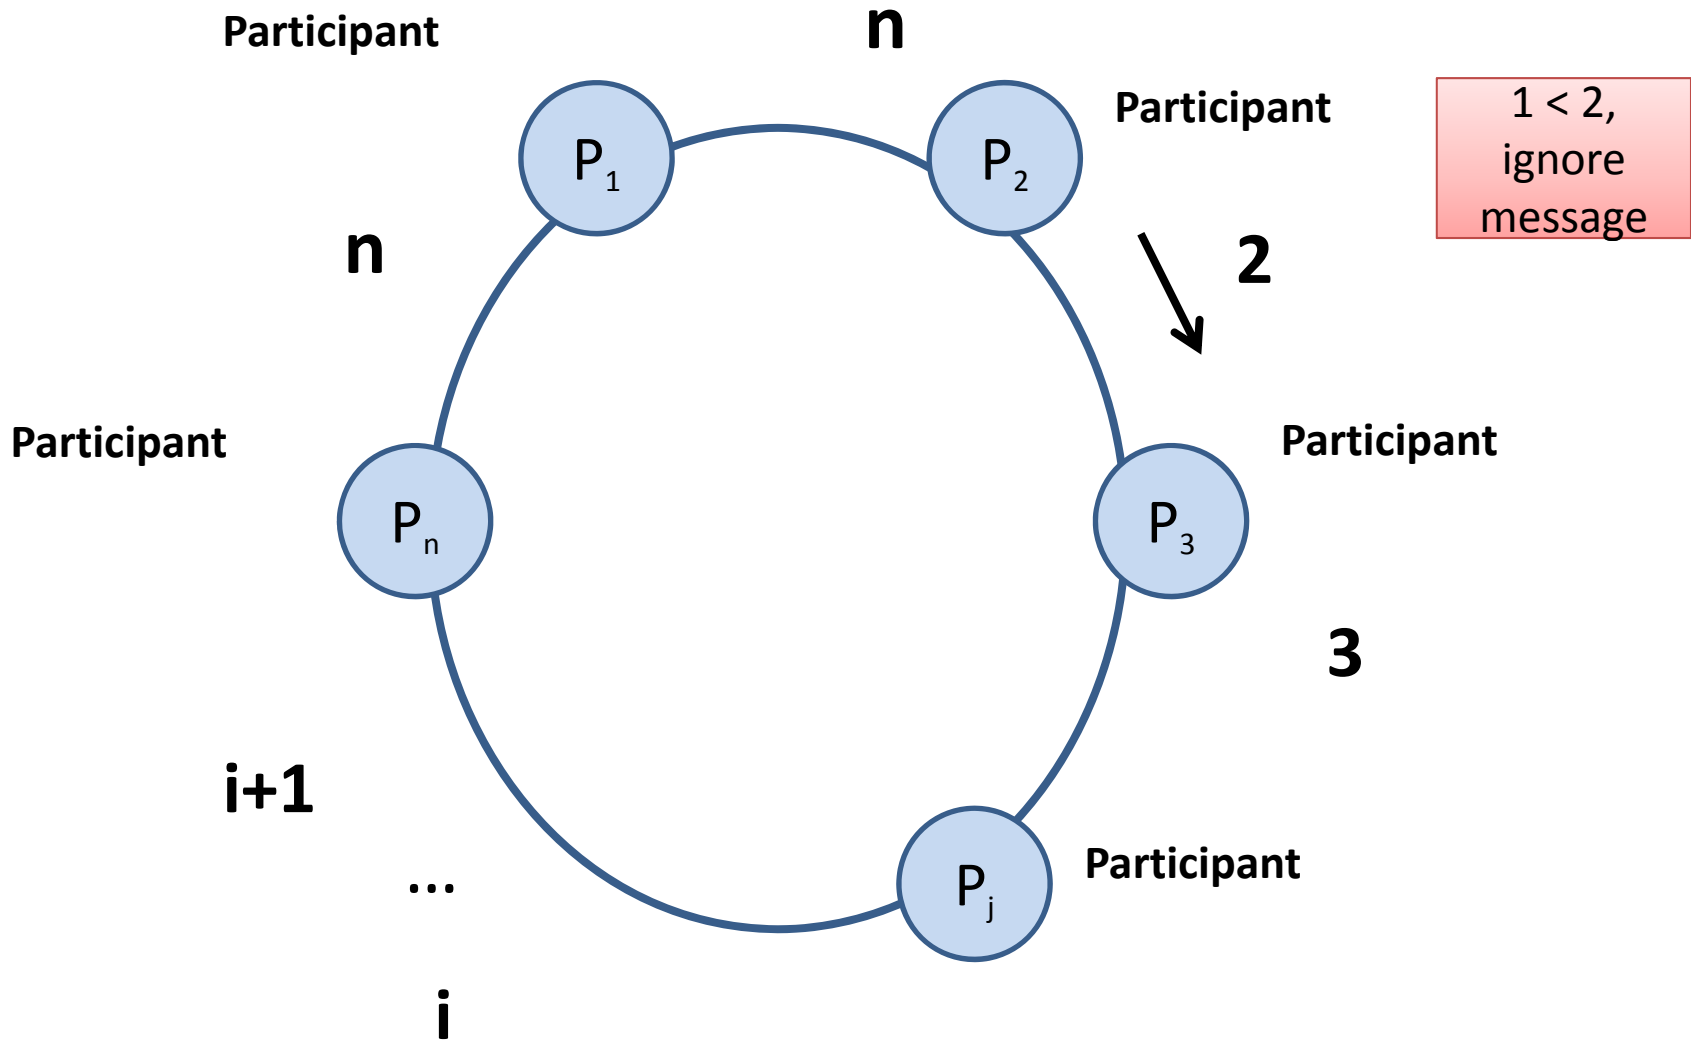
Ring-based algorithm: Concurrent election start



Ring-based algorithm: Concurrent election start



Ring-based algorithm: Concurrent election start



Ring-based election algorithm analysis

- Worst case: $3N - 1$ messages
 - $N - 1$ messages to reach highest ID from lowest ID
 - N messages to reach point of origin
 - Leader announcement takes another N messages
- Safety: even with multiple processes starting election
- Liveness: guaranteed progress *if no failures during the election occur*

Bully algorithm, 1982

- Assumes each process has a unique ID, reliable message delivery, and synchronous system
- Assumes processes know each others' IDs and can communicate with one another
 - Higher IDs have priority
 - Can “bully” lower numbered processes
- Initiated by any process that **suspects leader failure**
- Tolerates processes crashing during elections

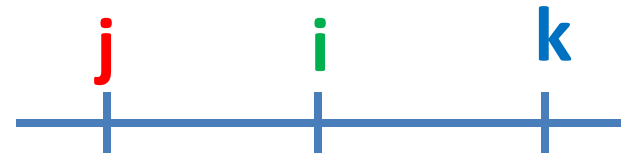


Bully algorithm messages

- Operates with three types of **messages**
 - ***Election*** announces an election
 - ***Answer*** responds to an election message
 - ***Coordination*** announces identity of leader
- Algorithm is triggered by any process that detects leader to have crashed (synchronous system)

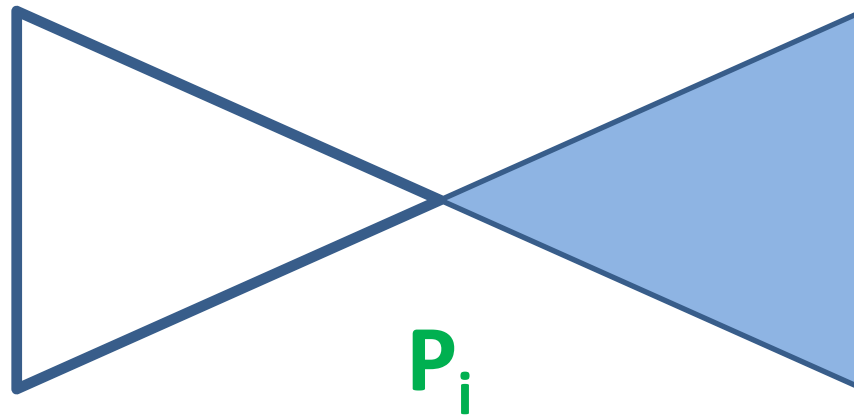
@ P_i P_i detects failure of leader

For any $j < i$ and any $i < k$

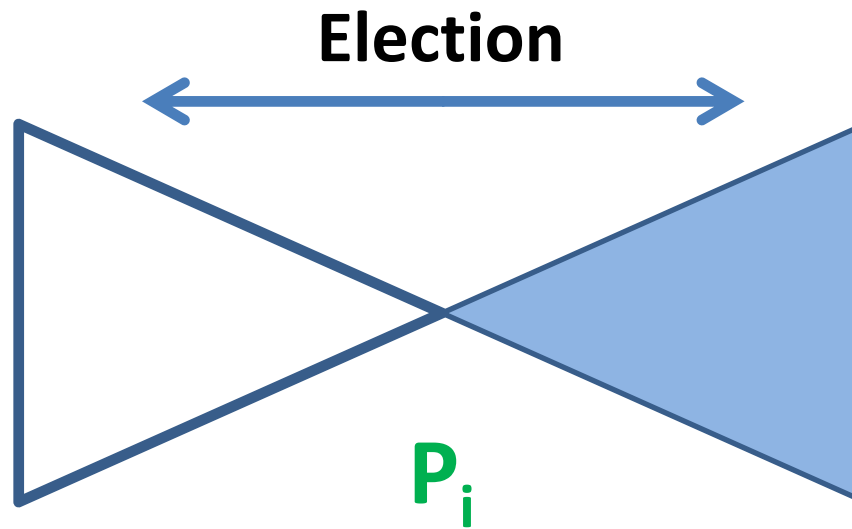


1. Broadcasts **election message**
2. Any P_k receiving election message **responds with answer message** and starts another election
3. Any P_j receiving election message **does not respond**
4. If P_i **does not receive any answer message (timeout)** then it **broadcasts victory** via **coordination message**
5. If P_i **does receive answer message(s)** then waits to receive **coordination message**
6. Restarts election, if no coordination message received (failure happened)

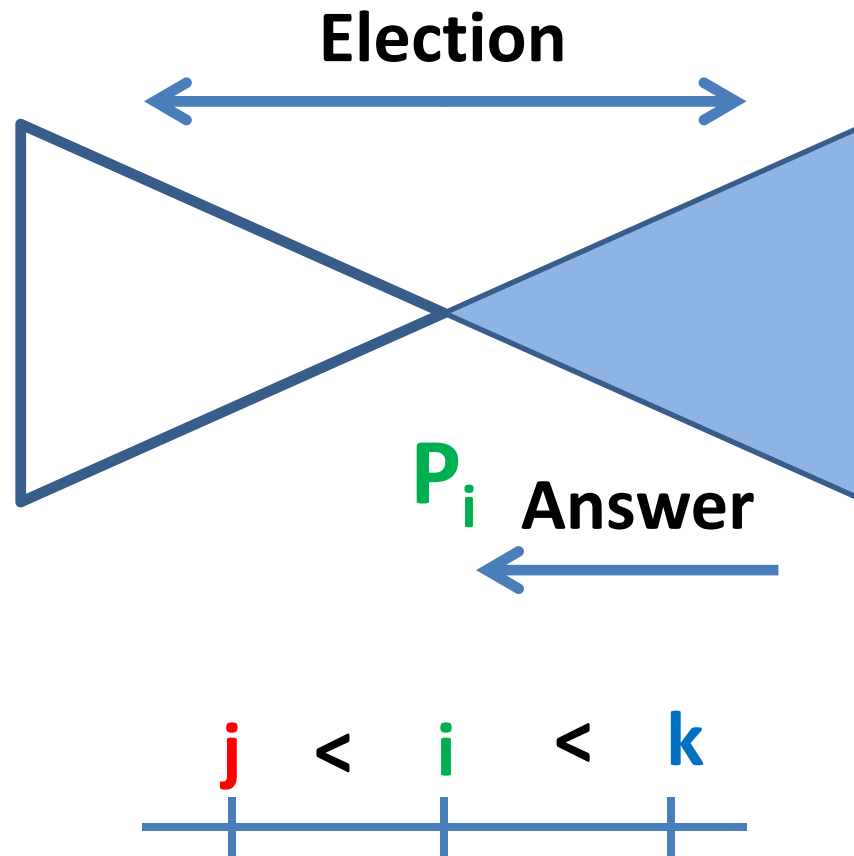
Election & answer



Election & answer



Election & answer



Upon crash of a process

- Suppose process eventually recovers (*no problem if it stays down, why?*)
- Process may determine that it has the highest identifier, thus, pronounce itself as leader
 - Even though system may have an existing leader (elected during crash)
- New process “**bullies**” current leader



