



# Consistent Hashing

*Organization-Based Analysis  
of Web-Object Sharing and  
Caching*  
Alec Wolma et al.



# Caching

- Five cache nodes
- Client determines node to contact



# Problem: Mapping objects to caches

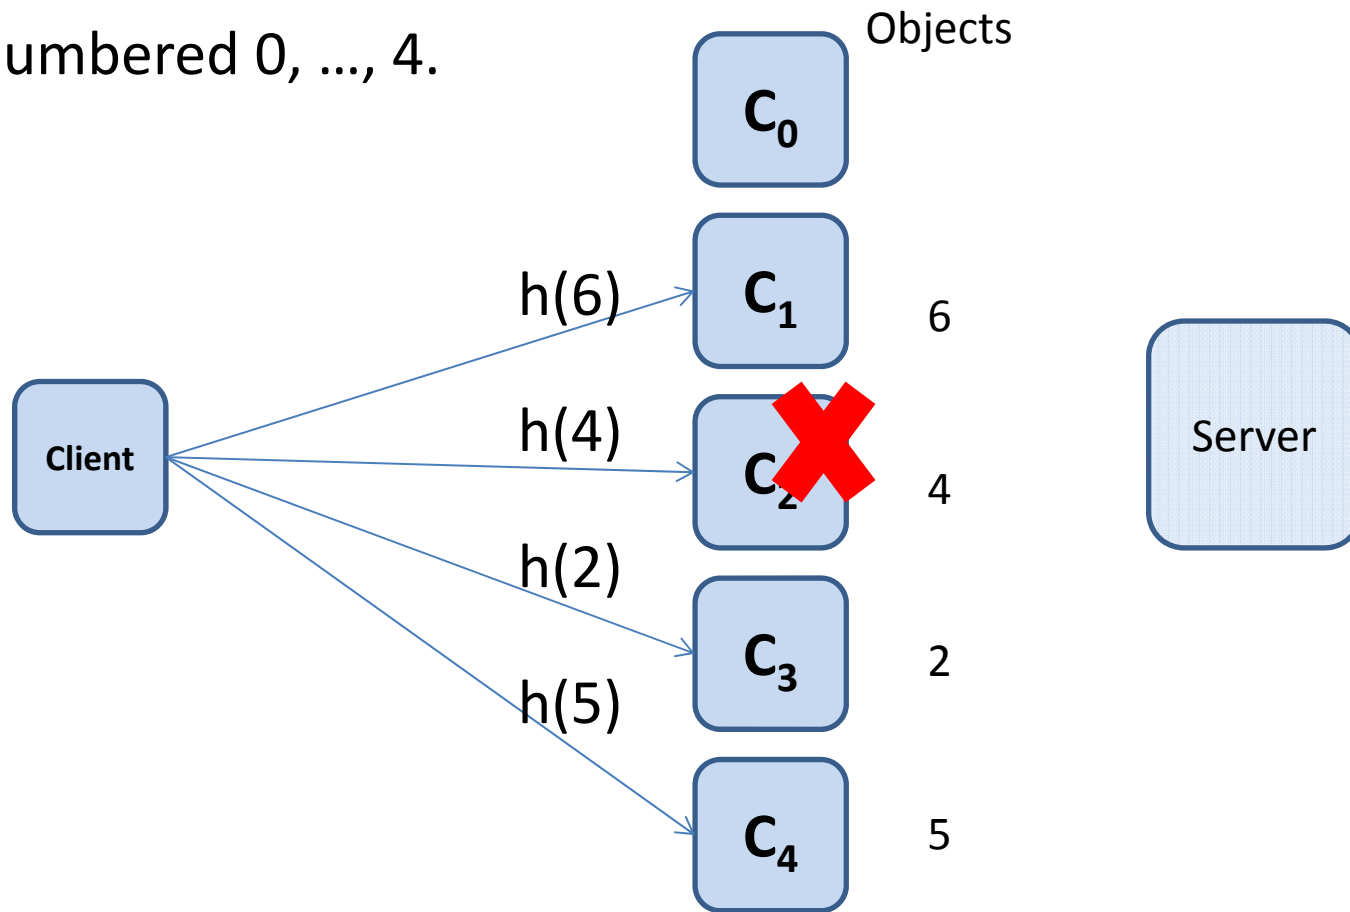
- Given a number of caches (e.g., cooperative caching, CDNs, etc.)
- Each cache should carry an **equal share of objects**
- **Clients** need to **know what cache to query** for a given object
- **Horizontally partition** (shard) object ID space
  - **Doesn't work** with **skewed distributions**: e.g., 10 servers, each handles 100 IDs, but all objects have IDs between 1-100 or 900-1000
- **Caches** should be able to **come and go** without disrupting the whole operation (i.e., non-effected caches)

# Solution attempt: Use hashing

- **Map object ID** (e.g., URL  $u$ ) **into one of the caches**
- Use a hash function that **maps  $u$  to node  $h(u)$** 
  - For example,  $h(x) = (ax + b) \bmod p$ , where  $p$  is range of  $h(x)$ , i.e., the number of caches
  - Interpret  $u$  as a number based on bit pattern of object ID (or URL)
- Hashing tends to **distribute input uniformly** across range of hash function
  - Objects (URLs) are **equally balanced** across caches, even if object IDs are skewed (i.e., highly clustered in ID space)
- No one cache responsible for an **uneven share of objects/URLs**
- No disproportionately loaded node (potential bottleneck)

$$h(u) = (7u + 4) \bmod 5$$

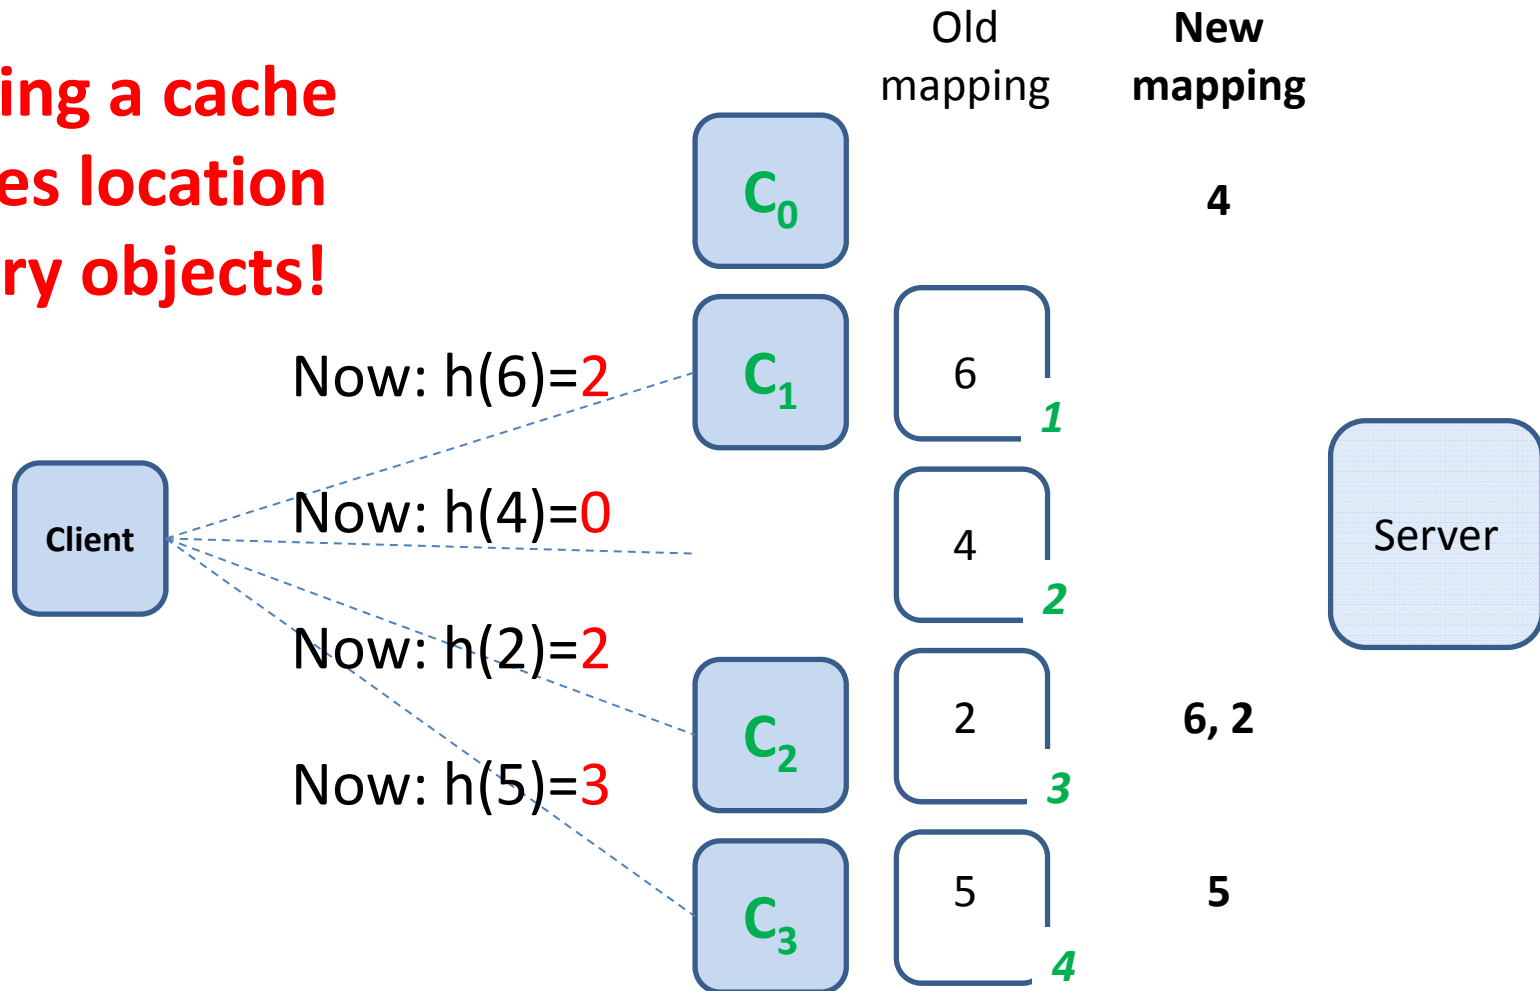
Assume, we have **five**  
**caches**, numbered 0, ..., 4.



$$h(u) = (7u + 4) \bmod 4$$

(now have to map across 4 caches)


Removing a cache  
changes location  
of every objects!



$$h(u) = 7u + 4 \bmod 4$$

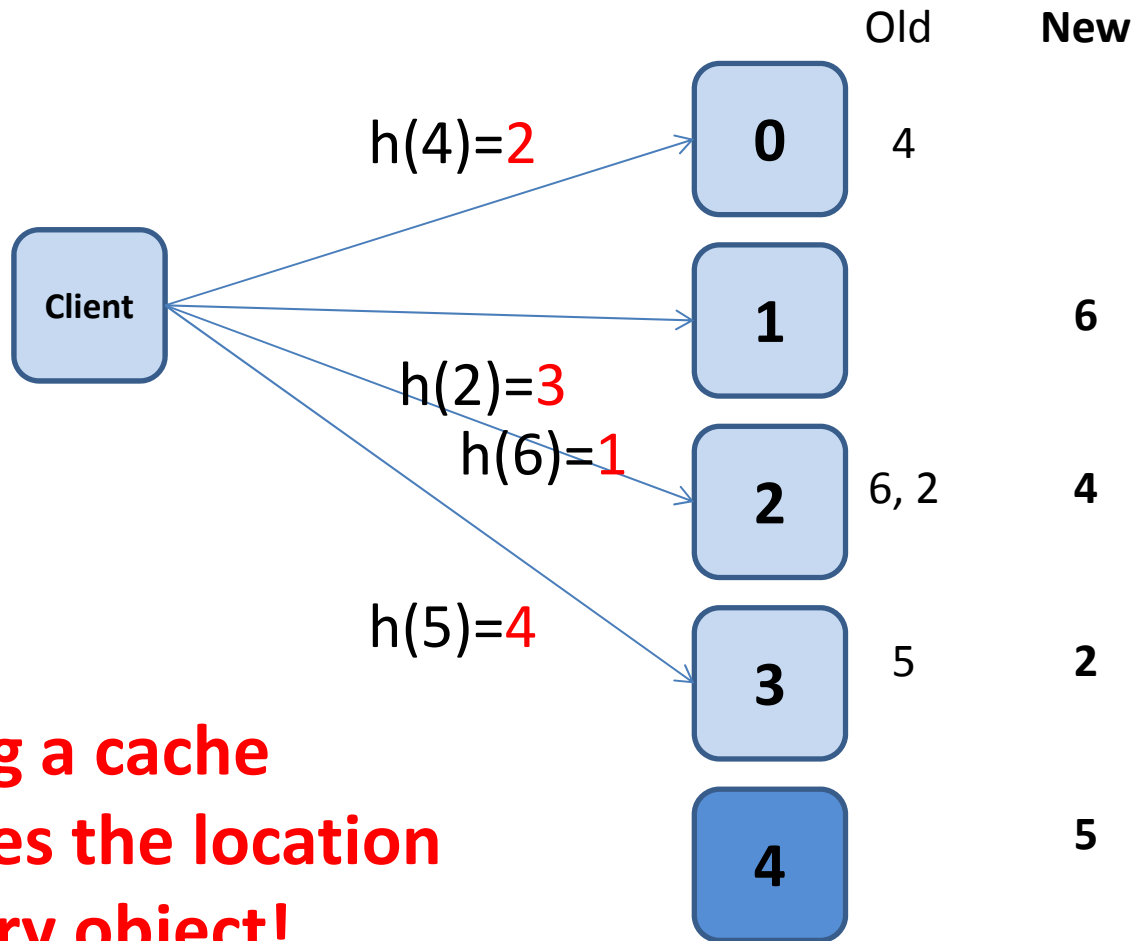
(mapped across 4 nodes)

Objects



$$h(u) = (7u + 4) \bmod 5$$

(adding a cache again)



**Adding a cache  
changes the location  
of every object!**




# Consistent hashing

- **Goals**
  - Uniform distribution of objects across nodes
  - Easily find objects
  - Let any client perform a local computation mapping a URL to node that contains referenced object
  - **Allow for nodes to be added/removed without much disruption**
- D. Karger *et al.*, MIT, 1997
- Basis for Akamai
  - CDN company (content distribution network)
  - Web cache as a service

# Consistent hashing

## Key idea intuition


- Select a **base hash function** that maps input identifier to the number range  $[0, \dots, m-1]$
- *E.g.*,  $h(x) = (ax + b) \bmod m$
- Interpret range of  $h(..)$  as array that wraps around (i.e., a circle)
- $h(..)$  gives slot in array (circle) and wraps around at  $m-1$  to 0
- Each **object** is mapped to a **slot** via  $h(..)$
- Each **cache** is mapped to a **slot** via  $h(..)$
- Assign **each object** to the **closest cache slot** in **clockwise direction** on the circle




# Consistent hashing

## Original interpretation

- Select a **base hash function** that maps input identifier to the number range  $[0, \dots, M]$
- Divide by  $M$ , re-mapping  $[0, \dots, M]$  to  $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via  $h(..)$
- **Each cache** is mapped to a **point** on unit circle via  $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle




# Mapping items to caches



Items	2, 3	mapped to <b>B</b>
Items	4, 5, 6	mapped to <b>C</b>
Items	7, 1	mapped to <b>A</b>


● Cache  
● Object

# Removing a cache



Items	2, 3, 7, 1	mapped to	<b>B</b>
Items	4, 5, 6	mapped to	<b>C</b>
Items	7, 1	mapped to	<b>A</b>

# Adding a cache

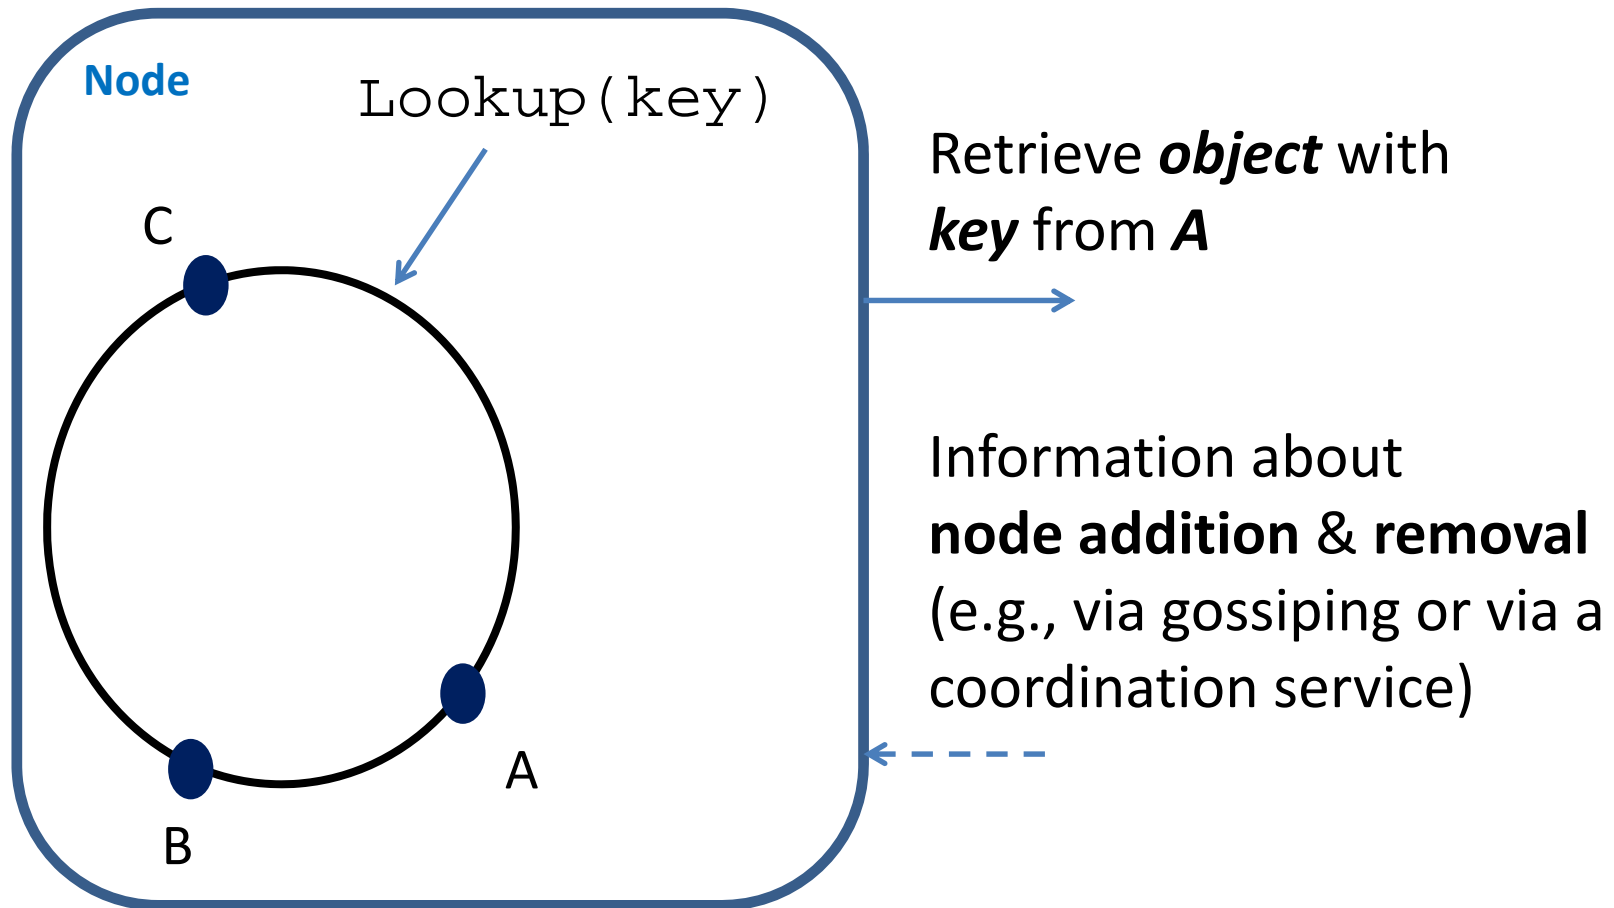


Items ~~3~~, 1, 2, 3 mapped to **B**

Items 4, 5, 6 mapped to **C**

Items **7**, **1**, **2** mapped to **A**


# Processing a Lookup (key)



# Cache lookup data structure at each node

- Store **cache points** in a **binary tree**
- Find **clockwise successor** of a **URL point** by single search in tree (takes  **$O(\log n)$**  time)
- For a constant time technique, cf. Karger et al., 1997

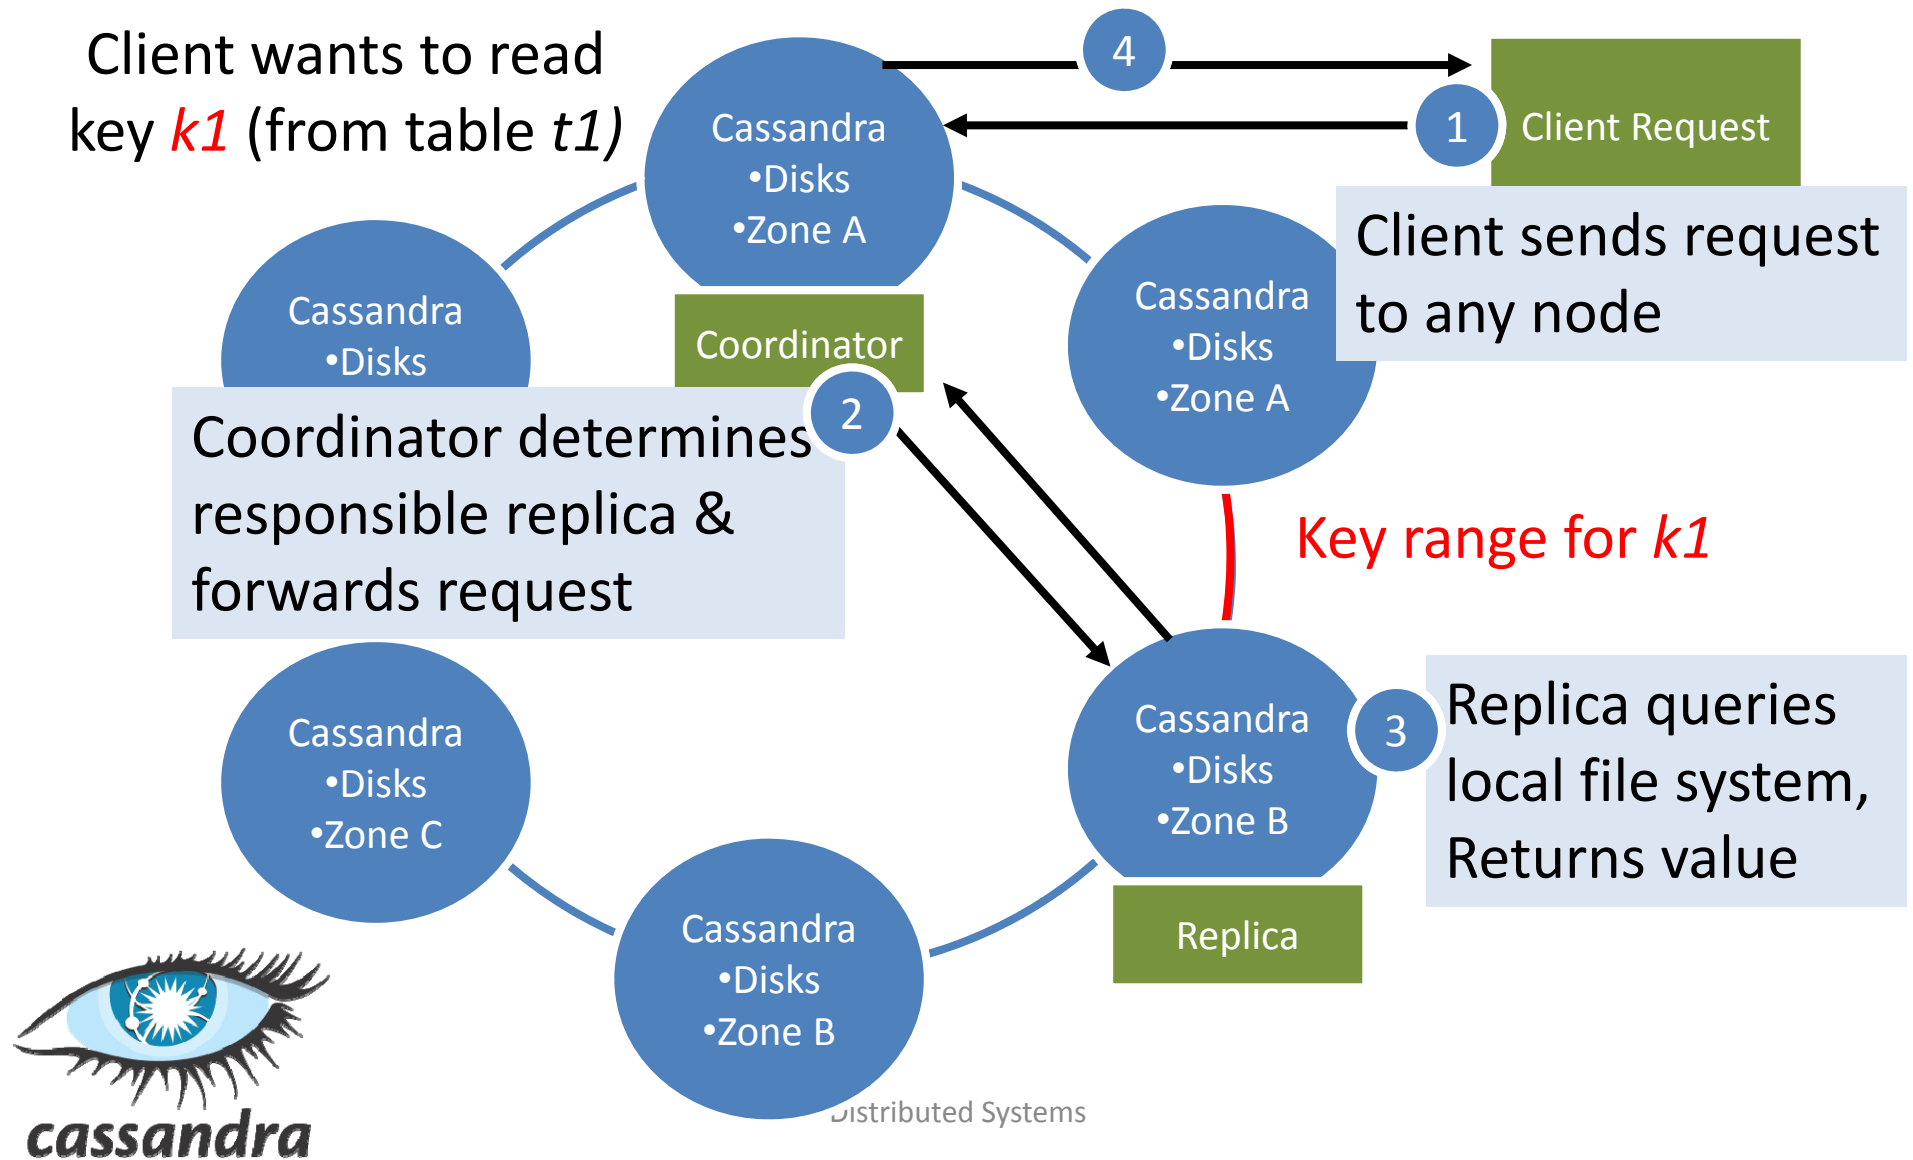
$\text{key}[\text{left}(x)] \leq \text{key}[x] \leq \text{key}[\text{right}(x)]$





# Cassandra global read-path

Client wants to read  
key *k1* (from table *t1*)



# Base hash function: MD5

- **Message Digest 5 (MD5)**, R. Rivest, **1992** (MD1, ..., MD6)
- **Hash function** that produces a **128-bit** (16-byte) **hash value**
- Maps variable-length message into a **fixed-length output**
- MD5 hash is typically expressed as a hex number (32 digits)
- It's been shown that **MD5 is not collision resistant**
- US-CERT about MD5 “*should be considered cryptographically broken and unsuitable for further use*” (for security, not for caching)
- SHA-2 is a more appropriate cryptographic hash function
- For consistent hashing, MD5 is sufficient

# MD5 examples

- MD5(*"The quick brown fox jumps over the lazy dog"*) = 9e107d9d372bb6826bd81d3542a419d6
- MD5(*"The quick brown fox jumps over the lazy dog."*) = e4d909c290d0fb1ca068ffaddf22cbd0
- MD5(*"*) = d41d8cd98f00b204e9800998ecf8427e