

ZooKeeper

*“Because Coordinating
Distributed Systems is a Zoo”*



Motivation

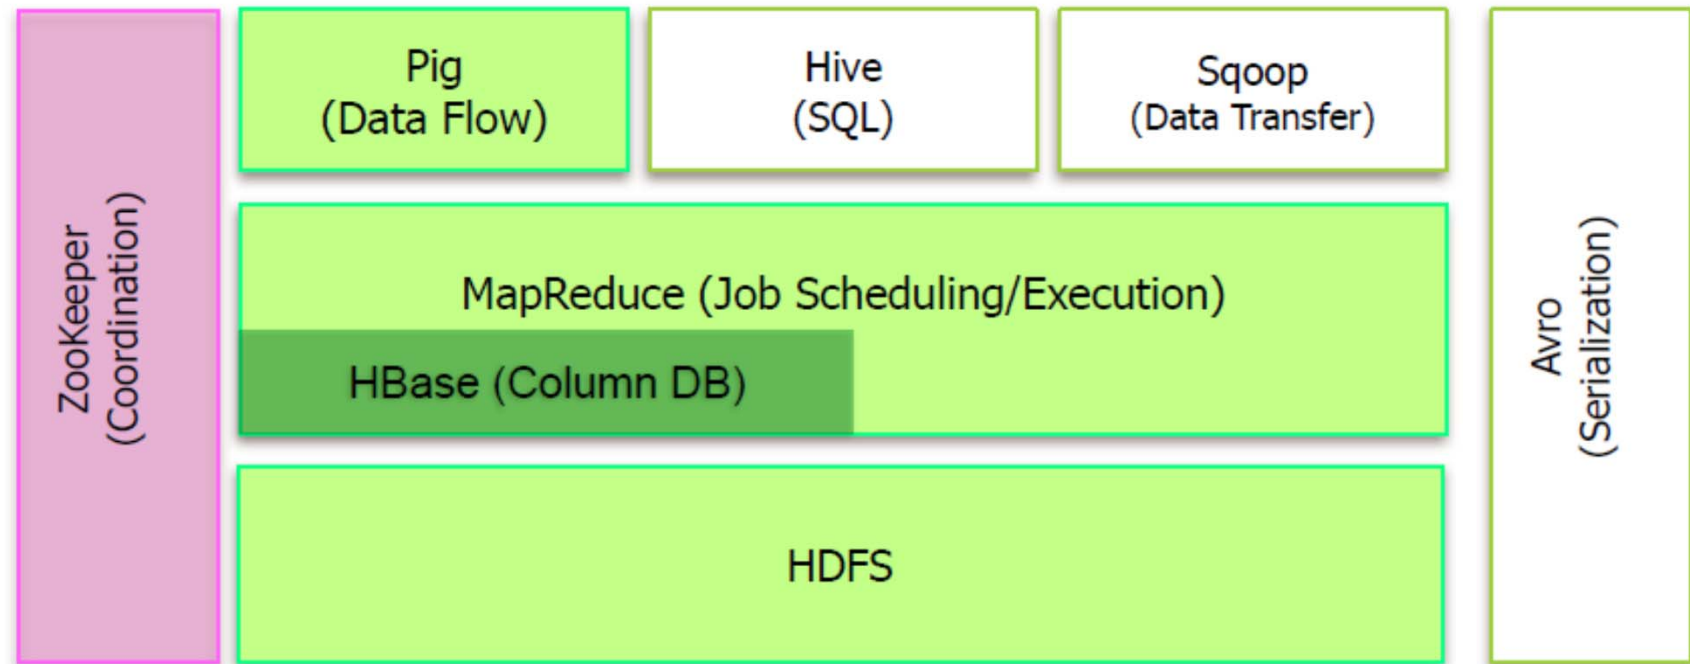
Why do we need a coordination service?

- Formerly, a **single** process running on a **single** node with a **single** CPU – no coordination required
- Today, applications, so called **services**, consist of **independent** processes running on a **changing** set of nodes
- Difficulty: **coordination** of those independent processes
- Developers have to deal with **coordination logic** and **application logic** at the same time

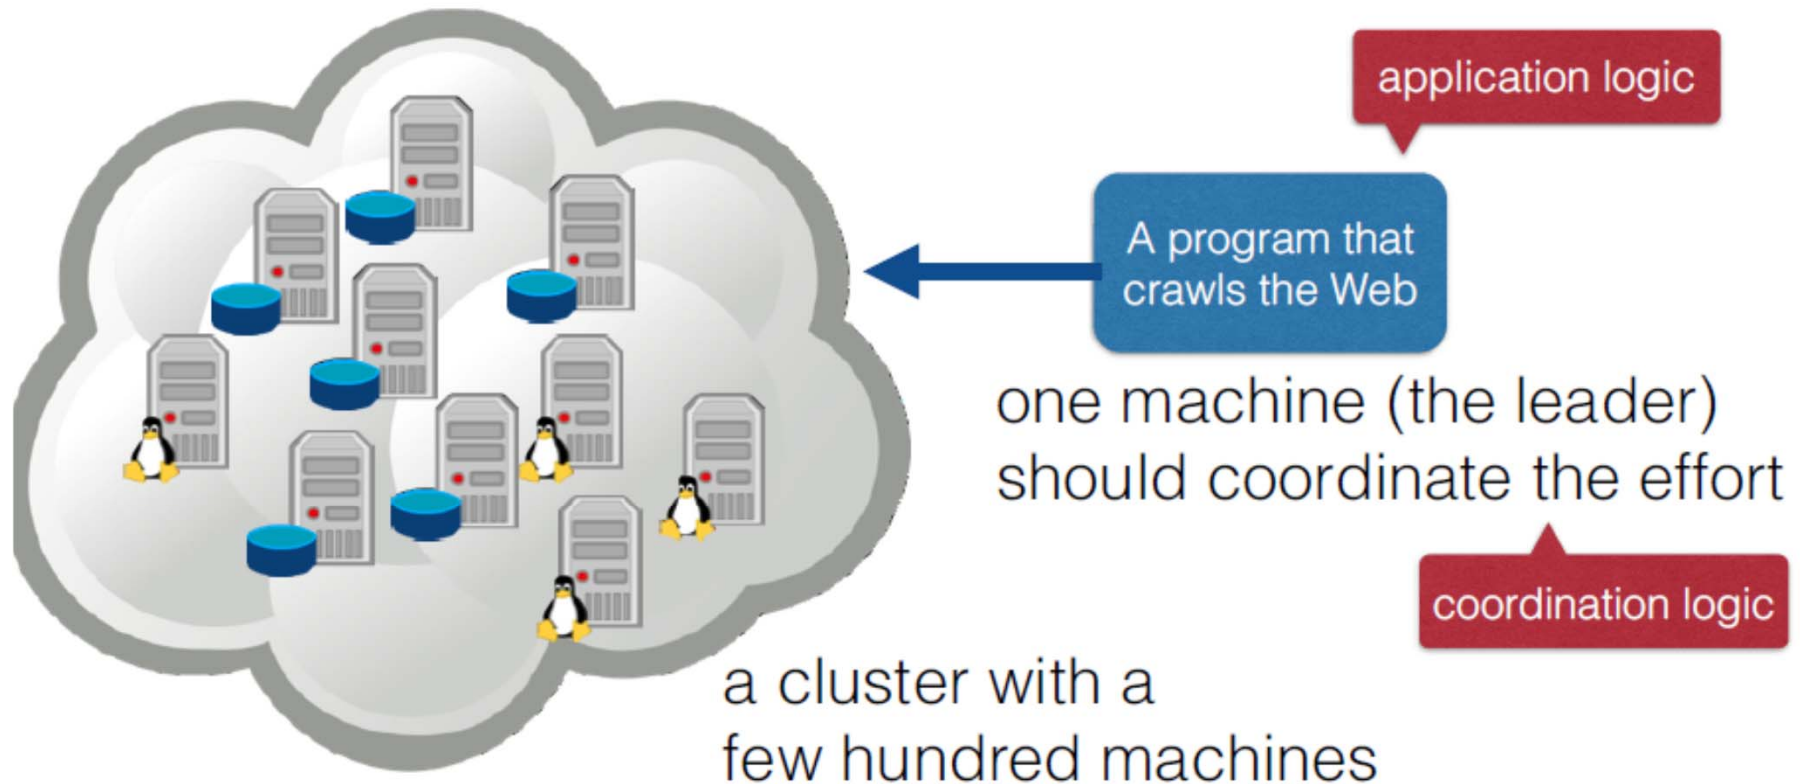
ZooKeeper

- ZooKeeper is designed to relieve developers from writing coordination logic code
- ZooKeeper is a highly-available service for coordinating processes of distributed applications

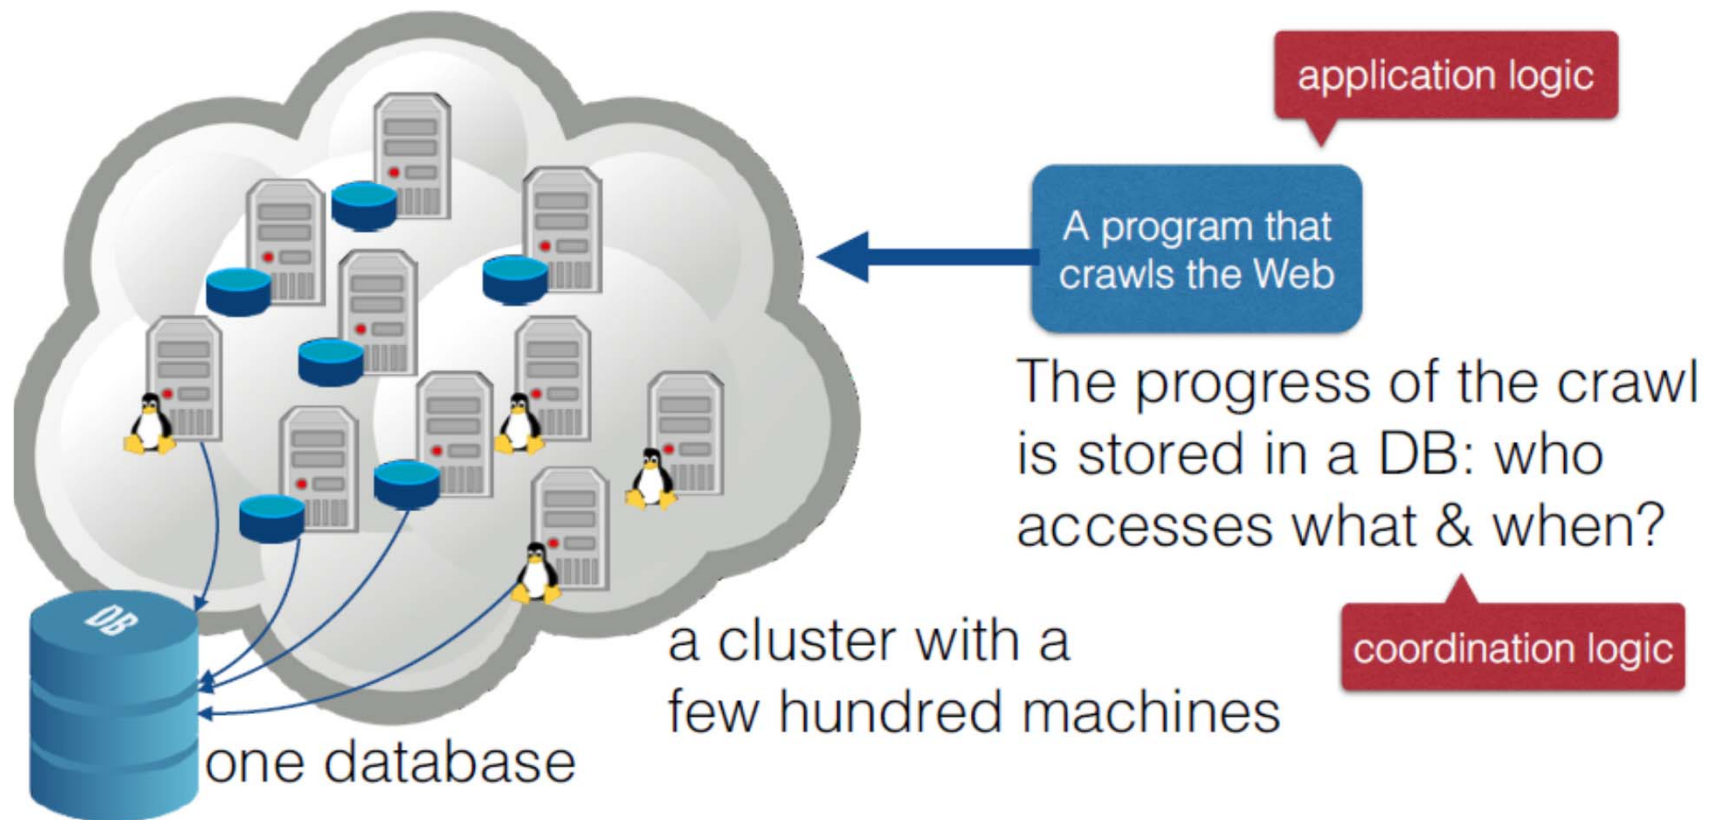
ZooKeeper in the Hadoop ecosystem



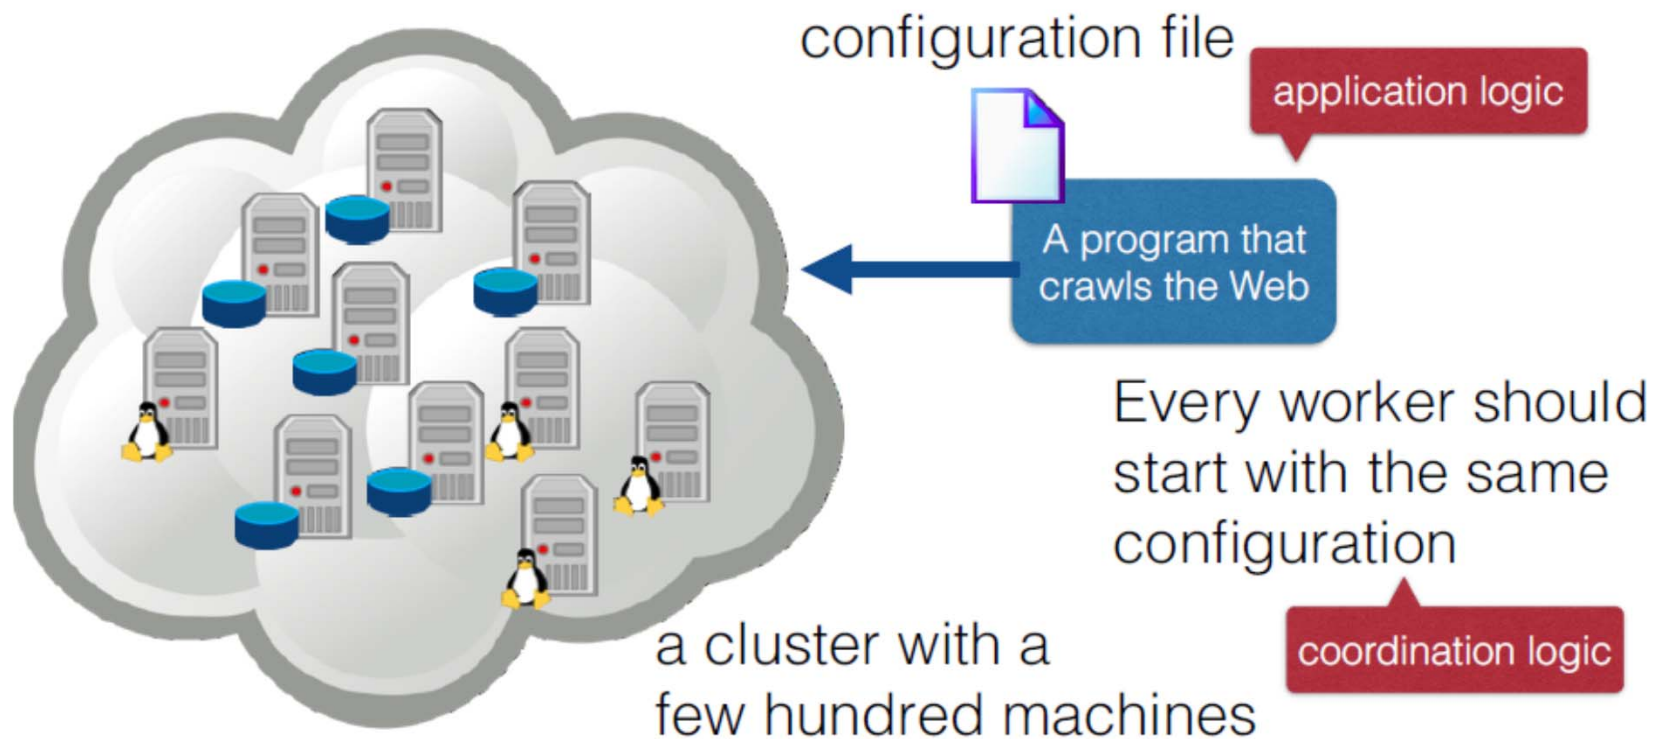
How do you elect the leader?



How do you lock a service?



How do you distribute the configuration?



ZooKeeper philosophy

- Be **specific** and develop a particular service for each coordination task
 - Locking service
 - Leader election service
 - Distribute coordination information
- Be **general** and provide an API to enable other services
- ZooKeeper offers API for developers to build their own primitives

Typical coordination problems

- **Static configuration:** list of operational parameters for system processes
- **Dynamic configuration:** parameter changes on the fly
- **Group membership:** who is alive?
- **Leader election:** who is in charge, who is the backup?
- **Mutually exclusive access** to critical resources (locks)
- **Barriers** (e.g., supersteps in computational workflows)

ZooKeeper API enables all these coordination tasks.

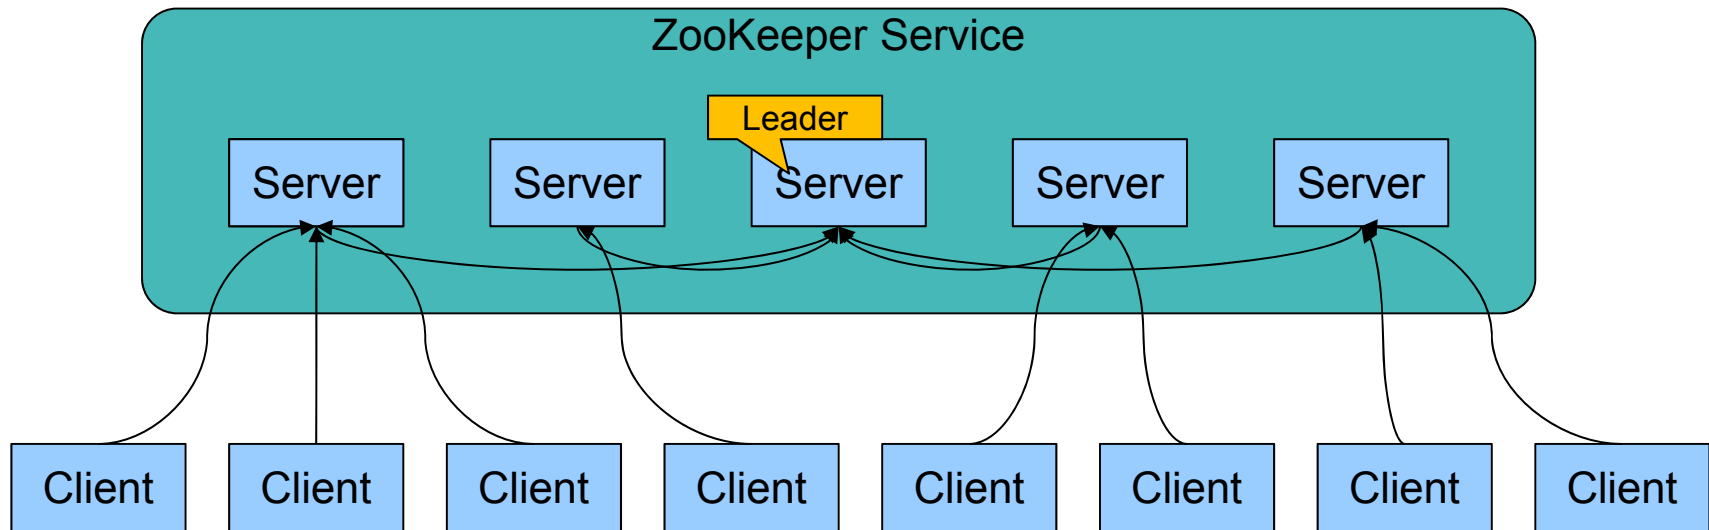
ZooKeeper (ZK) principles

- Design principles
 - API is **wait-free**
 - No blocking primitives in API
 - No deadlocks
 - Blocking can be implemented at client (deadlock possible)
- Guarantees
 - **Writes** to ZooKeeper **are linearizable** (appear atomic)
 - Client requests are processed in **FIFO order**
- Clients receive notifications of changes before the changed data becomes visible

ZK terminology

- **Client** is a user of ZK service
- **Server** is a process providing ZK service
- **Znode** is an **in-memory** data node in ZK, organised in a **hierarchical namespace** (data tree)
- **Write** (update) is any operation which modifies the state of data tree
- Clients establish a **session** when connecting to ZK

ZooKeeper servers & service I

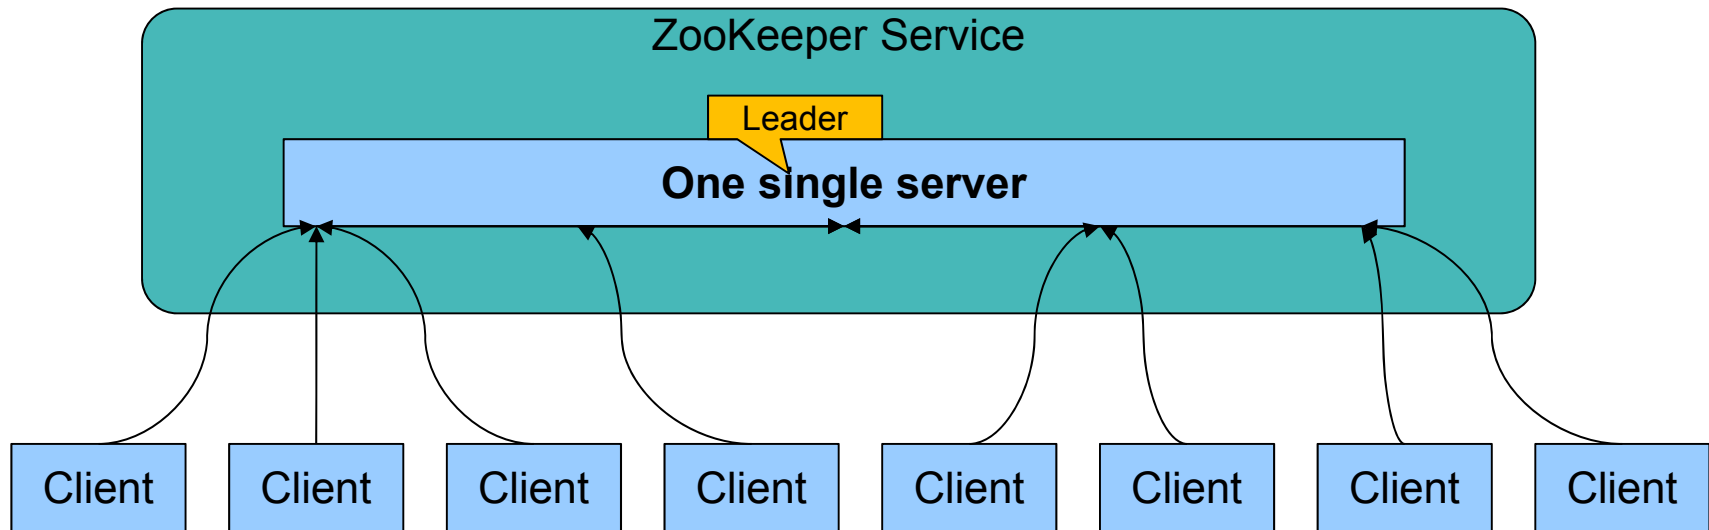


- ZooKeeper service comprised of **multiple servers**
- A designated **leader server** processing client **write requests** and a set of **follower servers** processing client **read requests**
- For development, ZooKeeper can be run in **standalone mode** (a **single server** which is “leader and follower” at the same time; **our modus operandi in M2-M4**)

ZooKeeper servers & service II

- All servers store a copy of data tree (in memory)
- **Leader** is elected at startup (or upon leader failure)
- **Followers** service client read requests, all updates go through leader
- **Write acknowledgements** are sent when a **majority of servers** persisted a write

ZK standalone mode



- Provides no guarantees
- Is not fault-tolerant
- Meant to ease development
- **Our expectation and modus operandi in M2-M4**

Client sessions

- Clients establish a **session** when connecting to ZK
- Session is with the ZK service
- Clients may connect to different ZK servers within the same session

ZooKeeper request processing

- ZooKeeper server services clients
- Clients connect to exactly one server to submit requests
 - **Read requests** served **from local replica**
 - **Write requests** processed by an **agreement protocol** (a server, elected leader, initiates processing of write)


Data model

(Similar to a file system)

- znodes are organised in a hierarchical namespace
- znodes can be manipulated by clients through ZK API
- znodes are referred to by UNIX-style file system paths (always absolute paths)



znodes can store data (**file like**; KBs up to 1 MB) & can have children (**directory like**).

/app1/p_2



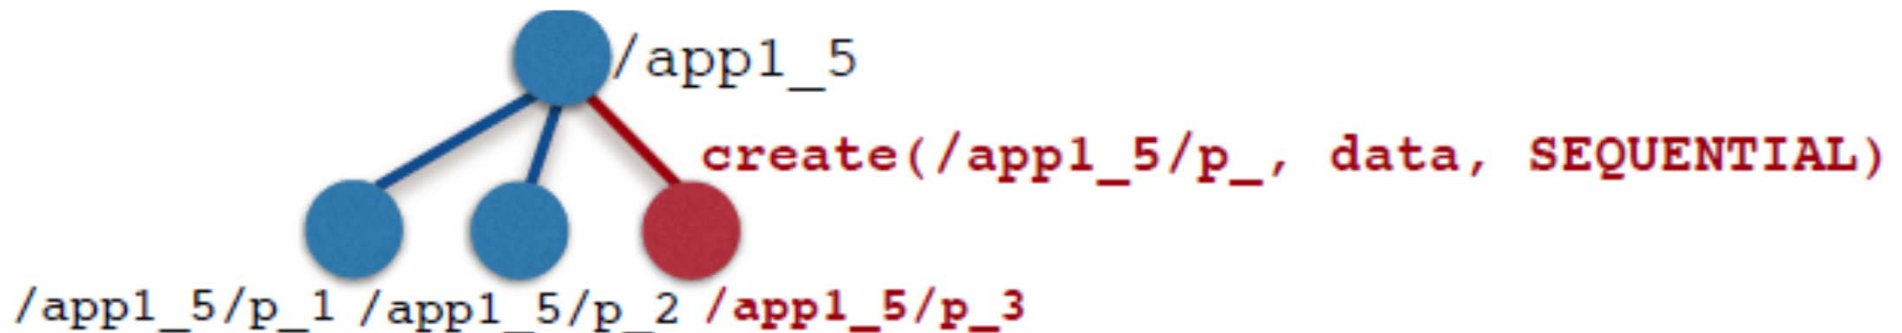
znodes

- znodes are not meant for general data storage
- znodes map to abstractions of the client (i.e., the application)
- **Group membership** protocol:
 - Client process p_i creates znode p_i under **/app1**
 - **/app1/ p_i** persists as long as p_i is running (via session)



znode flags

- Clients manipulate znodes by creating and deleting them
- **EPHEMERAL** flag (*passing, short-lived*)
 - Clients create znodes which are deleted at the end of the client's session
- **SEQUENTIAL** flag
 - Monotonically increasing counter appended to a znode's path
 - Counter value of a new znode under a parent is always larger than value of existing children



znodes & watch flag

- Clients can issue read operations on znodes with a **watch flag** set
- Server **notifies** client when data on znode changes
- Watches are **one-time triggers** associated with a session (unregistered once triggered or session closes)
- Watch notifications **indicate change**, not the new data (client has to retrieve data, post notification)

ZooKeeper API I

(simplified, cf. ZK API documentation)

- Create a znode with path name *path*, store *data* in it and set *flags* (ephemeral, sequential)
 - string **create**(*path*, *data*, *flags*)
- Delete the node *path*, if it is at expected *version*
 - void **delete**(*path*, *version*)
- Let client set a *watch* on znode
 - stat **exists**(*path*, *watch*)

ZooKeeper API II

(simplified, cf. ZK API documentation)

- Return data and meta-data of znode
 - (*data*, *stat*) **getData**(*path*, *watch*)
- Write data if version number is current version of znode
 - *stat* **setData**(*path*, *data*, *version*)
- Return all children of node
 - *string*[] **getChildren**(*path*, *watch*)

ZooKeeper API Notes

- No partial read/writes
- No **open**, **seek** or **close** methods
- No create-lock, lock, unlock

Configuration management example

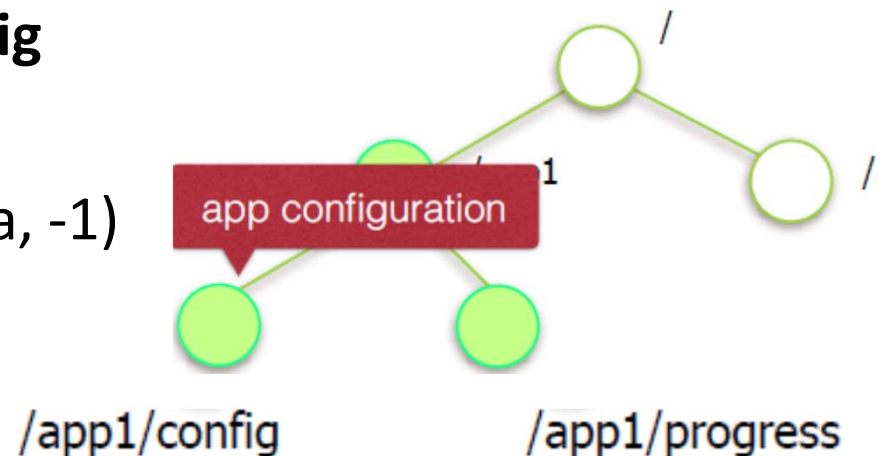
1. How does a **new** node query ZK for a configuration?
2. How does a node **change** the configuration?
3. How do nodes read the **new** configuration?

Configuration stored in /app1/config

1. `getData(/app1/config, true)`
2. `setData(/app1/config/config_data, -1)`

Watch for configuration changes

3. `getData(/app1/config, true)`



Group membership example

1. How can all nodes of an application **register themselves** at ZK?
2. How can a node find out about all **active** nodes of an application?

Create znode to store nodes

1. `create(/app1/workers/worker_, data, SEQUENTIAL)`
2. `getChildren(/app1/workers, true)`

Simple locking

- Client creates a lock file

`create(/app1/zlock, data, EPHEMERAL)`

If successful, client holds lock, otherwise not

- Client releases lock if it **explicitly deletes** it or if it **crashes**

`delete(/app1/zlock)`

- Could be held by another client, watch for status changes

`getData(/app1/zlock, TRUE)`

Caller is notified once status of lock changes

- Simple lock may lead to “**herd effect**”

Simple lock without herd effect

Intuition: Line up clients for locks; client with lowest node name wins (obtains lock).


Lock

```
1 n = create(l + "/lock-", EPHMERAL|SEQUENTIAL)
2 C = getChildren(l, false)
3 if n is lowest znode in C, exit
4 p = znode in C ordered just before n
5 if exists(p, true) wait for watch event
6 goto 2
```

Holds
lock

Unlock

```
1 delete(n)
```



Read/Write lock

(multiple readers, single writer)

Write Lock

```
1  n = create(l + "/write-", EPHMERAL|SEQUENTIAL)
2  C = getChildren(l, false)
3  if n is lowest znode in C, exit
4  p = znode in C ordered just before n
5  if exists(p, true) wait for event
6  goto 2
```

Holds
write lock

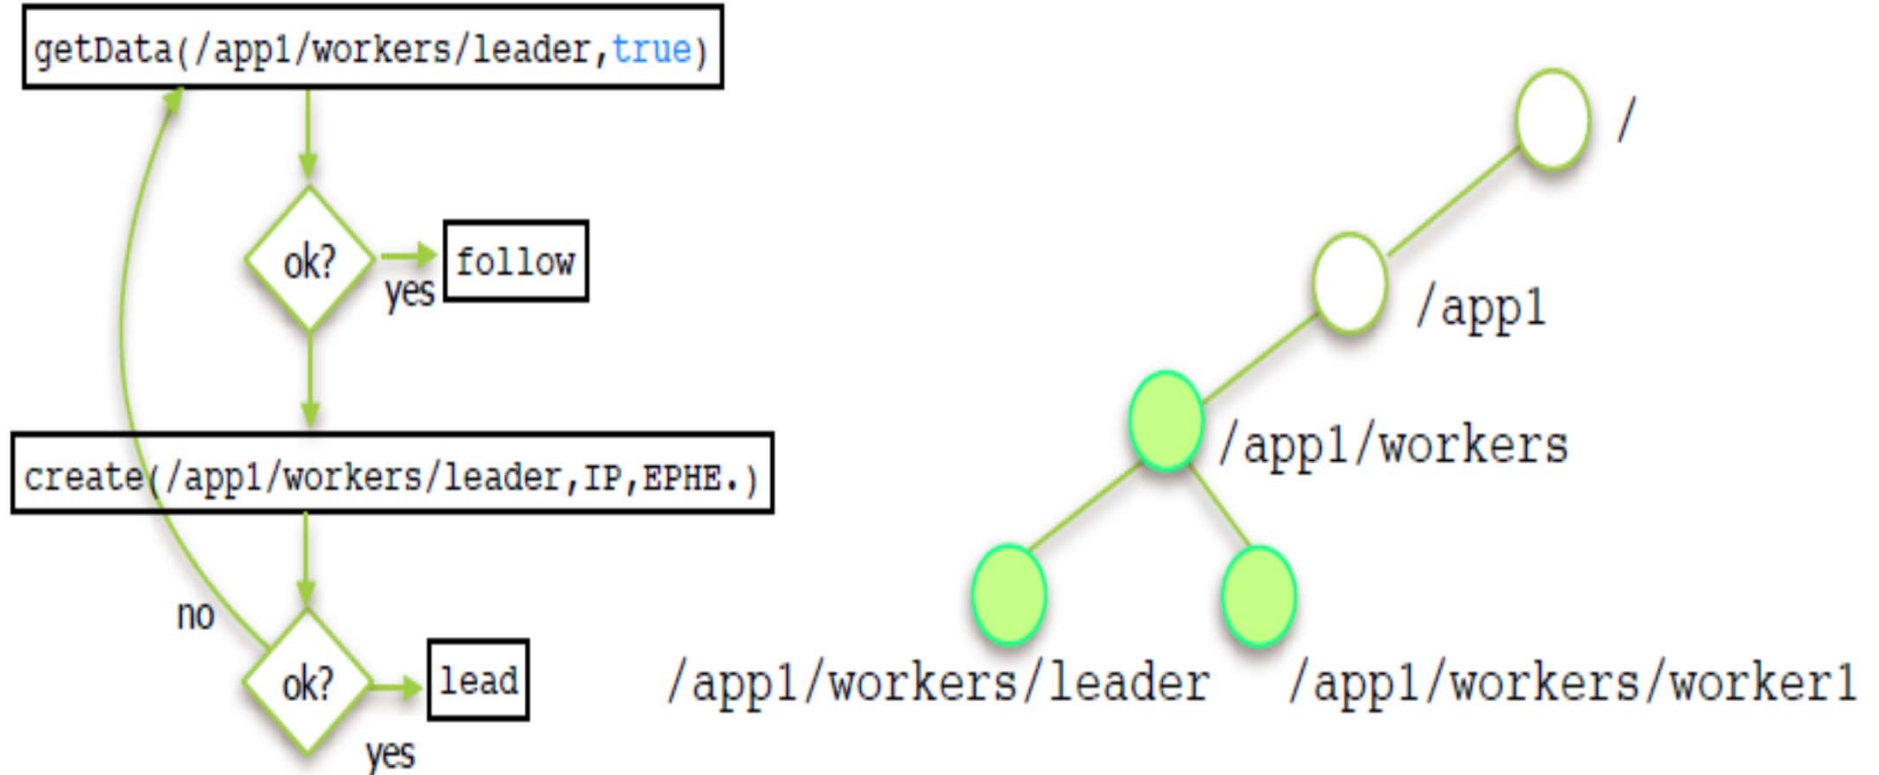
Read Lock

```
1  n = create(l + "/read-", EPHMERAL|SEQUENTIAL)
2  C = getChildren(l, false)
3  if no write znodes lower than n in C, exit
4  p = write znode in C ordered just before n
5  if exists(p, true) wait for event
6  goto 2
```

Holds
read lock

Leader election example

- How can all nodes elect a leader?



ZooKeeper internals

