

# ERASURE CODES

## INTRODUCTION

# Erasures

- Byte errors where we ***know*** the **position** of the **dropped** or corrupted **bytes** is called an ***erasure***
- As opposed to byte errors where we don't know the position of the dropped or corrupted bytes

# Types of Erasure Codes

- **Linear block code**
  - Reed Solomon Code
  - Can sustain lost bytes of known position
  - Distributed file systems
- **Fountain code**
  - LT Codes (Luby Transform)
  - Can sustain lost bytes of unknown position
  - P2P systems, torrents, video streaming ...

# Reed Solomon Code

- Error-correcting code
  - Used in QR codes
- Block encoding
  - **Data blocks**
  - Error-correcting blocks (**parity blocks**)
  - Read data blocks first
  - Else, decode with parity blocks

# General Problem

## $(n, m)$ -code

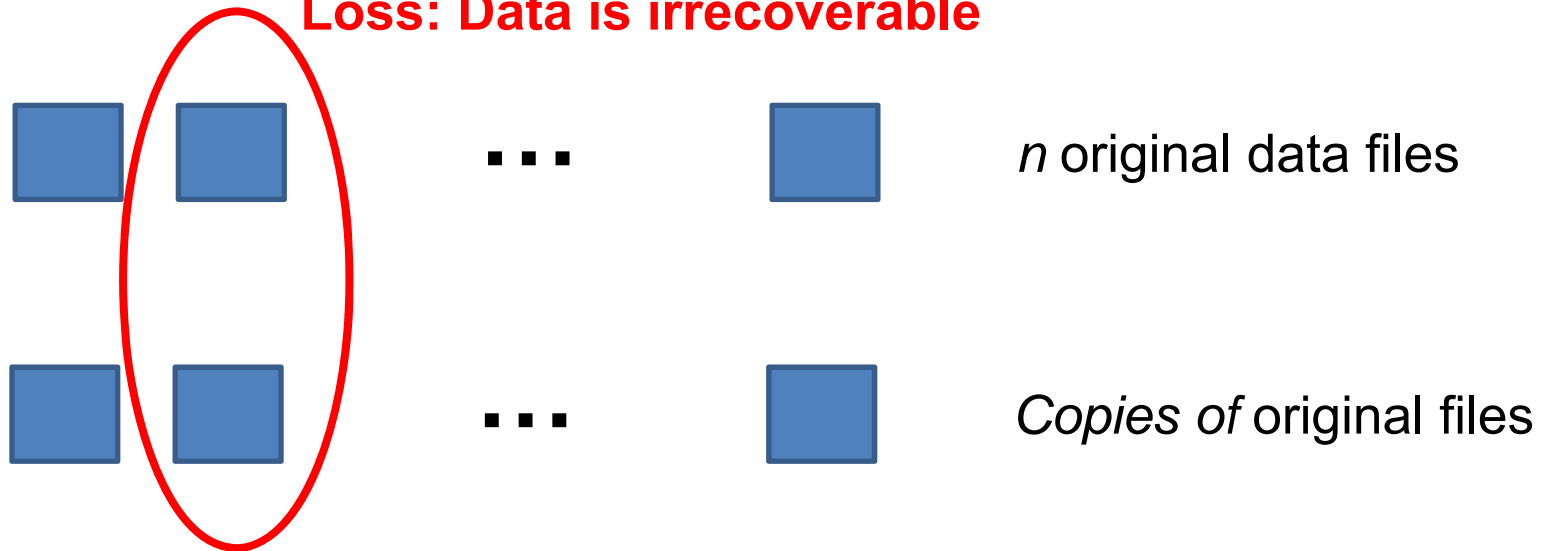


- We have  **$n$  data files**, guard against **losing  $m$**  of them
- Generate  **$m$  parity files**
- Lose up to  $m$  data files, can use equal number of parity files to recover data files
- Also works if some parity files are lost, as long as there are  **$n$  files left** (parity or data), can recover original data files
- Compare creating  $n$  parity files to making *a* copy of  $n$  data files (a.k.a. common backup strategy)

# Recoverability

## Via backup strategy

Loss: Data is irrecoverable




- With  $n$  original files and  $n$  copies (one each), if we loose two files (original and its copy), we can't recover loss
- Total storage requirement is  $2n$

# Recoverability

## Via an $(n, m)$ -code

Parity files take **same amount of space** but provide **superior (recovery capabilities!**



- With an  $(n, m)$ -code, we can protect  $n$  data files against the loss of  $m$  of them by generating  $m$  parity files
- Say, we use an  $(n, n)$ -code with total storage requirement  $2n$
- Could lose up to  $n$  files (any combination of data and parity files)

# Optimal Erasure Code

- Above sketched  $(n, m)$ -code is an erasure code because it guards against byte erasures
- It does not guard against more general errors where we don't know which data bytes have been corrupted
- Called optimal because in general, we **need at least  $n$  known bytes to recover  $n$  data bytes**, bound achieved here



# Erasure Code Computations

## Parity & Data Reconstruction Computation

- Given a pair  $(n, m)$  and  $n$  data bytes
- **Compute parity data:** compute  $m$  parity bytes, given  $n$  data bytes
- **Reconstruct data:** given a partial list of at least  $n$  data and parity bytes
  - Return full list of data bytes, i.e., there are no more than  $m$  omitted data or parity bytes
  - Error otherwise
- Generally, operates on byte level

# Erasure Code

$n, m = 1$

- For  $d_i$  a data byte, compute parity byte  $p$ 
  - $p = d_0 + d_1 + \dots + d_{n-1}$
- For  $m = 1$ , guard against loss of a single  $d_i$
- Reconstruct data
  - $d_i = p - (d_0 + d_1 + d_{i-1} + d_{i+1} + \dots + d_{n-1})$

# Nota Bene

- In practise, above and below computations are done modulo 256
- Generally speaking, on finite fields, a.k.a. Galois fields
- For simplicity and illustration, we'll use decimal arithmetic
- Including in assignments, *etc.*
- Unless explicitly stated otherwise

# Erasure Code

$$n = 3, m = 2$$

- (3, 2)-code
- Here,  $p_i$  **parity bytes**,  $d_i$  **data bytes** (i.e., numbers)
- Parity byte equations must be “*sufficiently different*”

$$p_0 = d_0 + d_1 + d_2$$

$$p_1 = 1 * d_0 + 2 * d_1 + 3 * d_2$$

- For example, the following is not “sufficiently different”

$$p_1 = 2 * d_0 + 2 * d_1 + 2 * d_2$$

- Here,  $p_1 = 2 p_0$  (**avoid linear combinations**)

For two **missing data bytes**,  $d_i, d_j$   $i < j$ , and **given parity bytes**  $p_0, p_1$ , we rearrange parity equations to move unknown (i.e., missing data bytes) to left hand side:

Given parity  
equations:

$$p_0 = d_0 + d_1 + d_2$$

$$p_1 = 1 * d_0 + 2 * d_1 + 3 * d_2$$

Rearranged (to solve for  
missing data bytes):

$$d_i + d_j = X = p_0 - d_k$$

$$(i + 1) * d_i + (j + 1) * d_j = Y = p_1 - (k + 1) * d_k$$

Equations kept generic since  
we don't know upfront which  
bytes get lost

where  $d_k$  is the **known (not missing) data byte**

For two **missing data bytes**,  $d_i, d_j$   $i < j$ , and **given parity bytes**  $p_0, p_1$ , we rearrange parity equations to move unknown (i.e., missing data bytes) to left hand side:

Given parity  
equations:

$$p_0 = d_0 + d_1 + d_2$$

$$p_1 = 1 * d_0 + 2 * d_1 + 3 * d_2$$

Rearranged (to solve for  
missing data bytes):

$$d_i + d_j = X = p_0 - d_k$$

$$(i + 1) * d_i + (j + 1) * d_j = Y = p_1 - (k + 1) * d_k$$

where  $d_k$  is the **known (not missing) data byte**

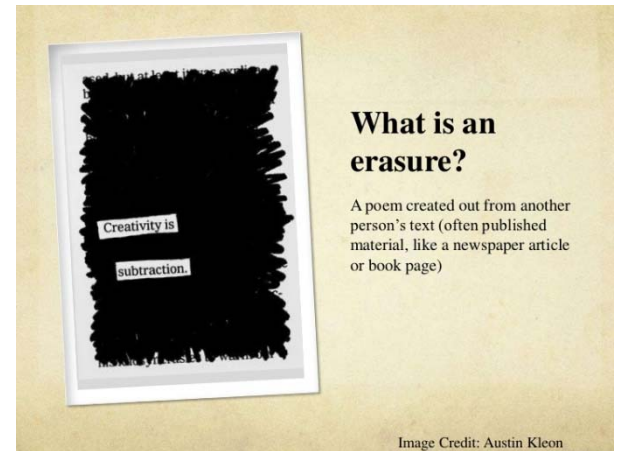
Multiply first equation by  $(i+1)$  and subtract it from second one to cancel the  $d_i$  term; then, use first equation to solve for  $d_j$

$$\begin{aligned} d_i + d_j &= X = p_0 - d_k \\ (i + 1) * d_i + (j + 1) * d_j &= Y = p_1 - (k + 1) * d_k \end{aligned}$$

$$d_j = (Y - (i + 1) * X) / (j - i)$$

$$d_i = X - d_j = ((j + 1) * X - Y) / (j - i)$$

We now have equations for  $d_i, d_j, i < j$  to reconstruct the missing data from known parity bytes.



# ERASURE CODES

## BASIC LINEAR ALGEBRA



# Basic Linear Algebra

## Digression

Equations of parity numbers from above example are of the form:

$$p = a_0 * d_0 + a_1 * d_1 + a_2 * d_2$$

where  $a_i$  are constants

These are **linear combinations** of  $d_i$ s and written as

$$p = (a_0 \quad a_1 \quad a_2) \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

# Correspondence

## Linear Functions and Matrices: Taking $n$ inputs to $m$ outputs

Two parity numbers, each a linear combination of  $d_i$ s:

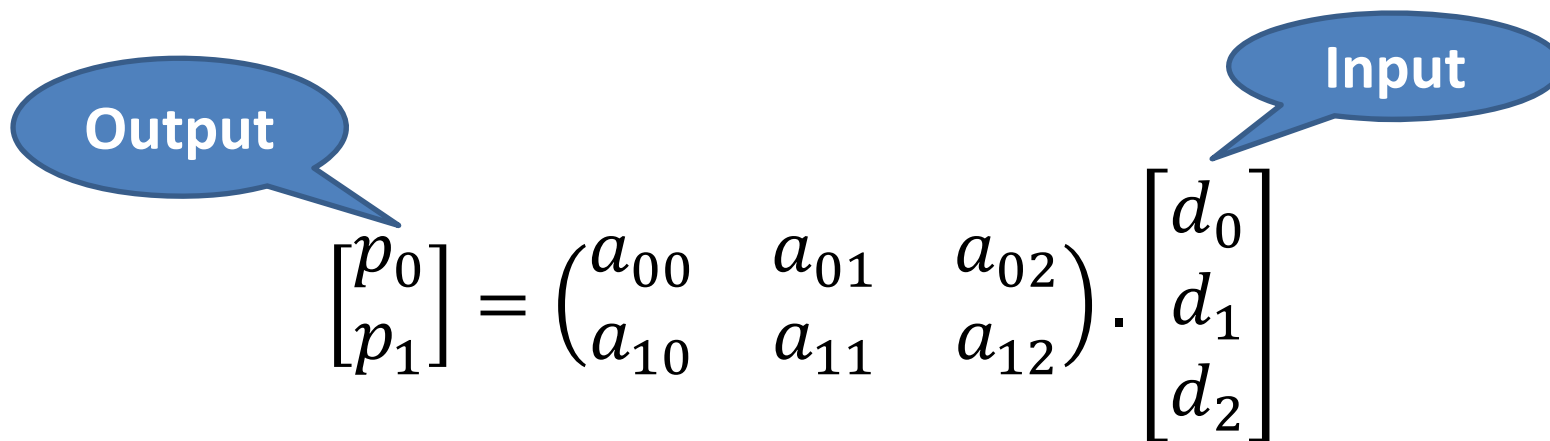
$$p_0 = a_{00} * d_0 + a_{01} * d_1 + a_{02} * d_2$$

$$p_1 = a_{10} * d_0 + a_{11} * d_1 + a_{12} * d_2$$

$$\begin{bmatrix} p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

# Linear Function vs. Matrix

- *Deleting a row of a matrix corresponds to deleting an output of a linear function*


$$\begin{bmatrix} p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$


# Linear Function vs. Matrix

- *Deleting a row* of a matrix corresponds to *deleting an output* of a linear function
- *Swapping two rows* of a matrix corresponds to *swapping two outputs* of a linear function
- *Appending a row* to a matrix corresponds to *adding an output* to a linear function

# Matrix Multiplication

- Multiply matrices  $A$ ,  $B$ , if they are **compatible**
- Number of columns of  $A$  must equal number of rows of  $B$ 
  - $A$  is  $p \times q$  matrix (#rows x #columns)
  - $B$  is a  $q \times r$  matrix
  - Resulting  $C$  is an  $r \times p$  matrix

# Matrix Multiplication



# Matrix Multiplication

MATRIX-MULTIPLY (*A*, *B*)

1 **if** *A.columns*  $\neq$  *B.rows*

2     **error** “incomplete dimensions”

3 **else** let *C* be a new *A.rows*  $\times$  *B.columns* matrix

4     **for** *i* = 1 **to** *A.rows*

5         **for** *j* = 1 **to** *B.columns*

6              $c_{ij} = 0$

7             **for** *k* = 1 **to** *A.columns*

8                  $c_{ij} = c_{ij} + a_{ik} * b_{kj}$

9 **return** *C*

# Identity Matrix

**Identity function** returns its  $n$  inputs as its outputs, corresponding matrix is the **identity matrix**

$$I_n = \begin{pmatrix} 1 & 0 & \cdots & 0 & 0 \\ 0 & 1 & \cdots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & 0 \\ 0 & 0 & \cdots & 0 & 1 \end{pmatrix}$$



# The Inverse

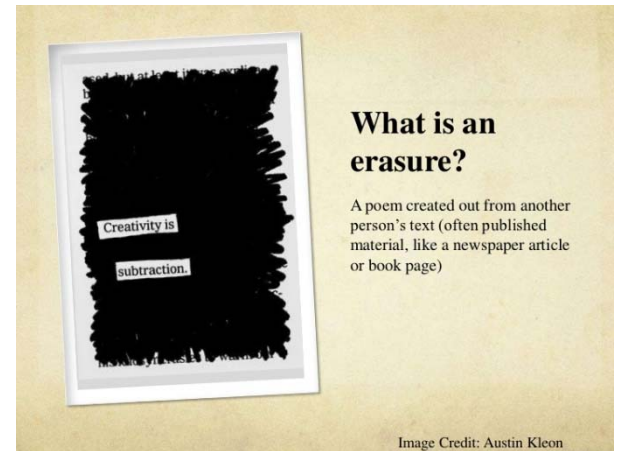
**$M$**  a square matrix ( $n \times n$ ) and  
 **$M^{-1}$**  its inverse, if it exists

$$M * M^{-1} = M^{-1} * M = I_n$$

## ***M* Invertible**

$$M = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix} \quad M^{-1} = \begin{pmatrix} -2 & 1 \\ 3/2 & -1/2 \end{pmatrix}$$

$$M * M^{-1} = M^{-1} * M = I_n$$



# ERASURE CODES

## (3, 2)-CODE EXAMPLE CONTINUED

# Erasure Codes

$n = 3, m = 2$  from above example (3, 2)-code

$$\begin{aligned} p_0 &= d_0 + d_1 + d_2 \\ p_1 &= 1 * d_0 + 2 * d_1 + 3 * d_2 \end{aligned} \quad P = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix}$$

$$\begin{bmatrix} p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

# Data Reconstruction Matrix

A.k.a. encoding matrix

$$\begin{bmatrix} d_0 \\ d_1 \\ d_2 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} * \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

*Identity matrix*

*Parity matrix / encoding matrix fragment*

Linear function that returns input data and parity bytes

# Data Reconstruction Matrix

A.k.a. encoding matrix

$$\begin{bmatrix} d_0 \\ d_1 \\ d_2 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} * \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

Linear function that returns input data and parity bytes

# Data Loss Example

$$\begin{bmatrix} \text{red X} \\ d_1 \\ \text{red X} \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} \text{red X} & \text{red X} & \text{red X} \\ 0 & 1 & 0 \\ \text{red X} & \text{red X} & \text{red X} \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} * \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

Linear function mapping data bytes to non-lost data and parity bytes:

$$\begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

*On the wrong side of equation ☹*

# Solve for $d_i$ s

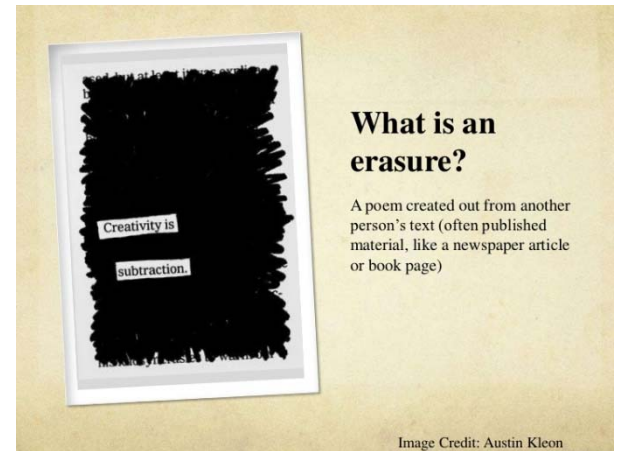
Need Inverse of Reconstruction Matrix

$$\begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix} \quad M^{-1} = \begin{pmatrix} -1/2 & 3/2 & -1/2 \\ 1 & 0 & 0 \\ -1/2 & -1/2 & 1/2 \end{pmatrix}$$

Gives us **original data bytes** in terms of **known data** and **parity bytes**!

$$\begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix} = \begin{pmatrix} -1/2 & 3/2 & -1/2 \\ 1 & 0 & 0 \\ -1/2 & -1/2 & 1/2 \end{pmatrix} \begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix}$$





# ERASURE CODES

## GENERALIZING TO (N, M)-CODES

# Generalizing

Arbitrary  $n, m$  - compute parity matrix

Generate an  $m \times n$  **parity matrix  $P$**  (rows need to be “*sufficiently different*”)

$$\begin{bmatrix} p_0 \\ \vdots \\ p_{m-1} \end{bmatrix} = \begin{pmatrix} \mathbf{p}_0 \\ \vdots \\ \mathbf{p}_{m-1} \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ \vdots \\ d_{n-1} \end{bmatrix}$$

$\mathbf{p}_i$  are rows of  $P$  – an  $m \times n$  parity matrix

# Compute Data Reconstruction Matrix

Append rows of  $P$  to  $I_n$ , denoted as  $e_i$

$$\begin{bmatrix} d_0 \\ \vdots \\ d_{n-1} \\ p_0 \\ \vdots \\ p_{m-1} \end{bmatrix} = \begin{pmatrix} e_0 \\ \vdots \\ e_{n-1} \\ \mathbf{p}_0 \\ \vdots \\ \mathbf{p}_{m-1} \end{pmatrix} * \begin{bmatrix} d_0 \\ \vdots \\ d_{n-1} \end{bmatrix}$$

# Data Loss with Resulting Matrix

- Indices  $k \leq m$  of **missing data bytes** are  $i_0, \dots, i_{k-1}$
- Remove rows corresponding to missing data bytes
- Keep  $k$  **parity rows**  $p_0, \dots, p_{k-1}$
- $j_0, \dots, j_{n-k-1}$  indices of present  $n-k$  data bytes

$$\begin{bmatrix} d_{j_0} \\ \vdots \\ d_{j_{n-k-1}} \\ p_0 \\ \vdots \\ p_{k-1} \end{bmatrix} = \begin{pmatrix} e_{j_0} \\ \vdots \\ e_{j_{n-k-1}} \\ \mathbf{p}_0 \\ \vdots \\ \mathbf{p}_{k-1} \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ \vdots \\ d_{n-1} \end{bmatrix}$$

# Compute $M$ 's Inverse $M^{-1}$

Reconstruct data by multiplying with  $M^{-1}$

If  $P$  was chosen correctly,  $M^{-1}$  exists.

$$\begin{bmatrix} d_0 \\ \vdots \\ d_{n-1} \end{bmatrix} = M^{-1} * \begin{bmatrix} d_{j_0} \\ \vdots \\ d_{j_{n-k-1}} \\ \mathbf{p}_0 \\ \vdots \\ \mathbf{p}_{k-1} \end{bmatrix}$$

# Loose Ends

- *How do we compute matrix inverses?*
- *How do we generate “optimal” parity matrices  $P$  s.t.  $M^{-1}$  always exists?*
- *How do we compute parity bytes instead of parity numbers (informally referred to as bytes, above)?*

# Facts: Non-invertible Matrix

More a negative result

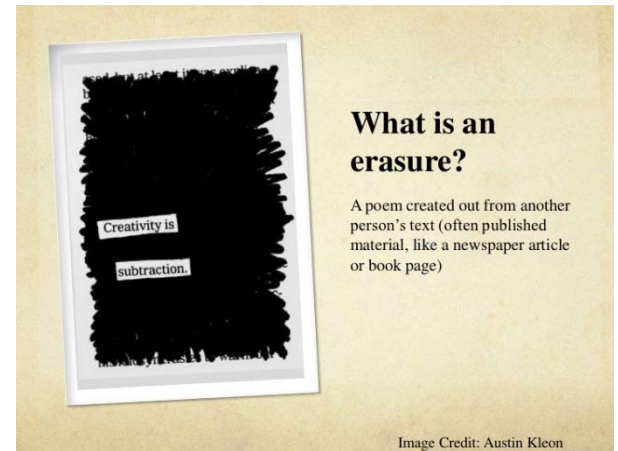
- $M$  has a row that can be expressed as a linear combination of other rows of  $M$  then  $M$  is non-invertible
- Said differently:
  - Linear function corresponding to  $M$  has one of its outputs redundant with the other outputs, so it is essentially a linear function taking  $n$  inputs to fewer than  $m$  outputs

# “Not Sufficiently Different”

- If one parity function is a linear combination of other parity functions then it is redundant, i.e., not sufficiently different
- Therefore, choose  $P$  wisely
- **No row of  $P$  should be expressed as linear combination of other rows of  $P$**
- ...







# ERASURE CODES

## ANOTHER EXAMPLE

# Erasure Coding Example

- Storage: “My private key”
- Pad with spaces to obtain right length if needed
  - “My private key\_\_” (added two spaces to obtain length of 16 characters)
- Build data matrix, here, a 4x4 matrix (ASCII code)

```
My p
riva
te k
ey__
```

# Encoding Matrix

Protect against loss of 2 bytes

- Identity matrix appended with
- Parity matrix
  - Derived from parity equations, i.e., one per parity byte (a.k.a. encoding matrix fragment)
- Results in full encoding matrix

# Encoding Matrix

(For a (3, 2)-code)

The diagram illustrates the encoding matrix for a (3, 2)-code. It shows the equation:

$$\begin{bmatrix} d_0 \\ d_1 \\ d_2 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

Two callouts provide additional information:

- Identity matrix:** A callout points to the top three rows of the matrix, which form the identity matrix.
- Encoding matrix fragment – derived from parity equations:** A callout points to the bottom two rows of the matrix, which are derived from parity equations.

Parity equations need to be “sufficiently different,” e.g., not be linear combinations of each other.

# Parity Matrix

- Results from **multiplying encoding matrix with original data matrix**

$$\begin{bmatrix} \text{---} \\ \text{---} \\ \text{---} \end{bmatrix} \begin{matrix} \text{\# columns} \\ p \times q \end{matrix} * \begin{bmatrix} \text{---} \\ \text{---} \\ \text{---} \end{bmatrix} \begin{matrix} \text{\# rows} \\ q \times r \end{matrix} = \begin{bmatrix} \text{---} \\ \text{---} \\ \text{---} \end{bmatrix} \begin{matrix} \text{\# result} \\ p \times r \end{matrix}$$

# Data Loss Example

$$\begin{bmatrix} \cancel{d_0} \\ d_1 \\ \cancel{d_2} \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} \cancel{1} & \cancel{0} & \cancel{0} \\ 0 & 1 & 0 \\ \cancel{0} & \cancel{0} & \cancel{1} \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

Data reconstruction matrix:

$$\begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix}$$

# Recovering Lost Data

- Find **Inverse** to data reconstruction matrix (from previous slide; it is a **square matrix**)
- For large matrix, use online tools to determine matrix inverse
- Recover data by multiplying **Inverse with relevant rows of parity matrix** (rows not affected by data loss)



# Solve for $d_i$ s

Need Inverse of Data Reconstruction Matrix

$$\begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix} \cdot \begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix} M^{-1} = \begin{pmatrix} -1/2 & 3/2 & -1/2 \\ 1 & 0 & 0 \\ -1/2 & -1/2 & 1/2 \end{pmatrix}$$

Gives us original data bytes in terms of  
known data and parity bytes!

$$\begin{bmatrix} d_0 \\ d_1 \\ d_2 \end{bmatrix} = \begin{pmatrix} -1/2 & 3/2 & -1/2 \\ 1 & 0 & 0 \\ -1/2 & -1/2 & 1/2 \end{pmatrix} \begin{bmatrix} d_1 \\ p_0 \\ p_1 \end{bmatrix}$$